SYBASE®

Open Client DB-Library™/C Reference Manual

Open Client

12.5

DOCUMENT ID: 32600-01-1250-01

LAST REVISED: May 2001

Copyright © 1989-2001 by Sybase, Inc. All rights reserved.

This publication pertains to Sybase database management software and to any subsequent release until otherwise indicated in new editions or technical notes. Information in this document is subject to change without notice. The software described herein is furnished under a license agreement, and it may be used or copied only in accordance with the terms of that agreement.

To order additional documents, U.S. and Canadian customers should call Customer Fulfillment at (800) 685-8225, fax (617) 229-9845.

Customers in other countries with a U.S. license agreement may contact Customer Fulfillment via the above fax number. All other international customers should contact their Sybase subsidiary or local distributor. Upgrades are provided only at regularly scheduled software release dates. No part of this publication may be reproduced, transmitted, or translated in any form or by any means, electronic, mechanical, manual, optical, or otherwise, without the prior written permission of Sybase, Inc.

Sybase, the Sybase logo, ADA Workbench, Adaptable Windowing Environment, Adaptive Component Architecture, Adaptive Server, Adaptive Server Anywhere, Adaptive Server Enterprise, Adaptive Server Enterprise Monitor, Adaptive Server Enterprise Replication, Adaptive Server Everywhere, Adaptive Server IQ, Adaptive Warehouse, AnswerBase, Anywhere Studio, Application Manager, AppModeler, APT Workbench, APT-Build, APT-Edit, APT-Execute, APT-FORMS, APT-Translator, APT-Library, Backup Server, ClearConnect, Client-Library, Client Services, Data Pipeline, Data Workbench, DataArchitect, Database Analyzer, DataExpress, DataServer, DataWindow, DB-Library, dbQueue, Developers Workbench, Direct Connect Anywhere, DirectConnect, Distribution Director, E-Anywhere, E-Whatever, Embedded SQL, EMS, Enterprise Application Studio, Enterprise Client/Server, Enterprise Connect, Enterprise Data Studio, Enterprise Manager, Enterprise SQL Server Manager, Enterprise Work Architecture, Enterprise Work Designer, Enterprise Work Modeler, EWA, Financial Fusion, Financial Fusion Server, Gateway Manager, ImpactNow, InfoMaker, Information Anywhere, Information Everywhere, InformationConnect, InternetBuilder, iScript, Jaguar CTS, jConnect for JDBC, KnowledgeBase, MainframeConnect, Maintenance Express, MAP, MDI Access Server, MDI Database Gateway, media.splash, MetaWorks, MySupport, Net-Gateway, Net-Library, ObjectConnect, ObjectCycle, OmniConnect, OmniSQL Access Module, OmniSQL Toolkit, Open Client, Open ClientConnect, Open Client/Server, Open Client/Server Interfaces, Open Gateway, Open Server, Open ServerConnect, Open Solutions, Optima++, PB-Gen, PC APT Execute, PC DB-Net, PC Net Library, Power++, power.stop, PowerAMC, PowerBuilder, PowerBuilder Foundation Class Library, PowerDesigner, PowerDimensions, PowerDynamo, PowerJ, PowerScript, PowerSite, PowerSocket, Powersoft, PowerStage, PowerStudio, PowerTips, Powersoft Portfolio, Powersoft Professional, PowerWare Desktop, PowerWare Enterprise, ProcessAnalyst, Report Workbench, Report-Execute, Replication Agent, Replication Driver, Replication Server, Replication Server Manager, Replication Toolkit, Resource Manager, RW-DisplayLib, RW-Library, S-Designor, SDF, Secure SQL Server, Secure SQL Toolset, Security Guardian, SKILS, smart.partners, smart.parts, smart.script, SQL Advantage, SQL Anywhere, SQL Anywhere Studio, SQL Code Checker, SQL Debug, SQL Edit, SQL Edit/TPU, SQL Everywhere, SQL Modeler, SQL Remote, SQL Server, SQL Server Manager, SQL SMART, SQL Toolset, SQL Server/CFT, SQL Server/DBM, SQL Server SNMP SubAgent, SQL Station, SQLJ, STEP, SupportNow, Sybase Central, Sybase Client/Server Interfaces, Sybase Financial Server, Sybase Gateways, Sybase MPP, Sybase SQL Desktop, Sybase SQL Lifecycle, Sybase SQL Workgroup, Sybase User Workbench, SybaseWare, Syber Financial, SyberAssist, SyBooks, System 10, System 11, System XI (logo), SystemTools, Tabular Data Stream, Transact-SQL, Translation Toolkit, UNIBOM, Unilib, Uninull, Unisep, Unistring, URK Runtime Kit for UniCode, Viewer, Visual Components, VisualSpeller, VisualWriter, VQL, WarehouseArchitect, Warehouse Control Center, Warehouse Studio, Warehouse WORKS, Watcom, Watcom SQL, Watcom SQL Server, Web Deployment Kit, Web.PB, Web.SQL, WebSights, WebViewer, WorkGroup SQL Server, XA-Library, XA-Server and XP Server are trademarks of Sybase, Inc. 3/01

Unicode and the Unicode Logo are registered trademarks of Unicode, Inc.

All other company and product names used herein may be trademarks or registered trademarks of their respective companies.

Use, duplication, or disclosure by the government is subject to the restrictions set forth in subparagraph (c)(1)(ii) of DFARS 52.227-7013 for the DOD and as set forth in FAR 52.227-19(a)-(d) for civilian agencies.

Sybase, Inc., 6475 Christie Avenue, Emeryville, CA 94608.

Contents

CHAPTER 1 **Introducing DB-Library 1**

Client/server architecture 1
Types of clients 2
Types of servers 2
The Open Client and Open Server products 3
Open Client 3
Open Server 4
Open Client libraries 4
What is in DB-Library/C? 4
Comparing the library approach to Embedded SQL 5
Other Sybase development tools 6
Data structures for communicating with servers 6
DB-Library/C programming 7
DB-Library/C datatypes 12
DB-Library/C routines 12
Initialization 13
Command processing 15
Results processing 17
Message and error handling 22
Information retrieval 24
Browse mode 26
text and image handling 28
Datatype conversion 29
Process control flow 30
Remote procedure call processing 30
Registered procedure call processing 31
Gateway passthrough routines 33
Datetime and money 34
Cleanup 35
Secure support 35
Miscellaneous routines 35
Bulk copy special library 36
Two-phase commit service special library 36
Summary of changes for version 10.0 36
DB-Library routines added to version 10.0 37

DB-Library 12.0 and later 38
Recompiling DB-Library 4.x, 10.0.x, and 11.x applications 38
Sample programs and sample databases 39

CHAPTER 2 Routines 41
db12hour 49
dbadata 51
dbadlen 54
dbaltbind 55
dbaltbind_ps 62
dbaltcolid 68
dbaltlen 69
dbaltop 70
dbalttype 71
dbaltutype 72
dbanullbind 73
dbbind 74
dbbind_ps 80
dbbufsize 84
dbbylist 85
dbcancel 86
dbcanquery 88
dbchange 89
dbcharsetconv 90
dbclose 90
dbclrbuf 91
dbclropt 92
dbcmd 94
DBCMDROW 95
dbcolbrowse 96
dbcollen 97
dbcolname 98
dbcolsource 99
dbcoltype 100
dbcoltypeinfo 102
dbcolutype 103
dbconvert 104
dbconvert_ps 109
DBCOUNT 115
DBCURCMD 116
DBCURROW 116
dbcursor 117
dbcursorbind 119
dbcursorclose 121

dbcursorcolinfo 122
dbcursorfetch 123
dbcursorinfo 125
dbcursoropen 125
dbdata 129
dbdate4cmp 130
dbdate4zero 131
dbdatechar 132
dbdatecmp 133
dbdatecrack 134
dbdatename 136
dbdateorder 139
dbdatepart 140
dbdatezero 141
dbdatlen 142
dbdayname 143
DBDEAD 144
dberrhandle 145
dbexit 149
dbfcmd 150
DBFIRSTROW 153
dbfree_xlate 154
dbfreebuf 155
dbfreequal 156
dbfreesort 156
dbgetchar 158
dbgetcharset 158
dbgetloginfo 159
dbgetlusername 161
dbgetmaxprocs 162
dbgetnatlang 163
dbgetoff 163
dbgetpacket 165
dbgetrow 166
DBGETTIME 168
dbgetuserdata 168
dbhasretstat 169
dbinit 171
DBIORDESC 171
DBIOWDESC 173
DBISAVAIL 174
dbisopt 174
DBLASTROW 175
dbload_xlate 176

dbloadsort 177
dblogin 178
dbloginfree 180
dbmny4add 180
dbmny4cmp 181
dbmny4copy 182
dbmny4divide 183
dbmny4minus 184
dbmny4mul 185
dbmny4sub 186
dbmny4zero 187
dbmnyadd 188
dbmnycmp 189
dbmnycopy 190
dbmnydec 191
dbmnydivide 192
dbmnydown 193
dbmnyinc 195
dbmnyinit 195
dbmnymaxneg 197
dbmnymaxpos 198
dbmnyminus 199
dbmnymul 200
dbmnyndigit 201
dbmnyscale 207
dbmnysub 209
dbmnyzero 210
dbmonthname 210
DBMORECMDS 212
dbmoretext 212
dbmsghandle 213
dbname 217
dbnextrow 218
dbnpcreate 220
dbnpdefine 223
dbnullbind 225
dbnumalts 226
dbnumcols 226
dbnumcompute 228
DBNUMORDERS 228
dbnumrets 229
dbopen 231
dbordercol 235
dbpoll 235

dbprhead 241
dbprrow 241
dbprtype 242
dbqual 244
DBRBUF 247
dbreadpage 248
dbreadtext 249
dbrecftos 251
dbrecvpassthru 252
dbregdrop 254
dbregexec 256
dbreghandle 258
dbreginit 262
dbreglist 264
dbregnowatch 265
dbregparam 266
dbregwatch 270
dbregwatchlist 275
dbresults 276
dbretdata 279
dbretlen 283
dbretname 285
dbretstatus 286
dbrettype 288
DBROWS 290
DBROWTYPE 291
dbrpcinit 291
dbrpcparam 293
dbrpcsend 295
dbrpwclr 296
dbrpwset 297
dbsafestr 298
dbsechandle 300
dbsendpassthru 304
dbservcharset 306
dbsetavail 307
dbsetbusy 307
dbsetdefcharset 310
dbsetdeflang 310
dbsetidle 311
dbsetifile 312
dbsetinterrupt 313
DBSETLAPP 317
DBSETLCHARSET 317

DBSETLENCRYPT 318
DBSETLHOST 319
DBSETLNATLANG 320
dbsetloginfo 321
dbsetlogintime 323
DBSETLPACKET 324
DBSETLPWD 325
DBSETLUSER 326
dbsetmaxprocs 327
dbsetnull 327
dbsetopt 329
dbsetrow 331
dbsettime 333
dbsetuserdata 333
dbsetversion 336
dbspid 337
dbspr1row 338
dbspr1rowlen 340
dbsprhead 341
dbsprline 343
dbsqlexec 344
dbsqlok 346
dbsqlsend 351
dbstrbuild 352
dbstrcmp 354
dbstrcpy 356
dbstrlen 358
dbstrsort 359
dbtabbrowse 360
dbtabcount 361
dbtabname 362
dbtabsource 363
DBTDS 365
dbtextsize 365
dbtsnewlen 366
dbtsnewval 367
dbtsput 368
dbtxptr 369
dbtxtimestamp 371
dbtxtsnewval 372
dbtxtsput 372
dbuse 373
dbvarylen 374
dbversion 375

dbwillconvert 376
dbwritepage 377
dbwritetext 378
dbxlate 383
Errors 386
Options 402
Types 408

CHAPTER 3 Bulk copy 413
Introduction to bulk copy 413
Transferring data into the database 413
Transferring data out of the database to a flat file 415
List of bulk copy routines 416
bcp_batch 417
bcp_bind 418
bcp_colfmt 423
bcp_colfmt_ps 426
bcp_collen 430
bcp_colptr 431
bcp_columns 432
bcp_control 433
bcp_done 435
bcp_exec 436
bcp_getl 437
bcp_init 438
bcp_moretext 440
bcp_options 443
bcp_readfmt 444
bcp_sendrow 445
BCP_SETL 446
bcp_setxlate 446
bcp_writefmt 448

CHAPTER 4 Two-phase commit service 449
Programming distributed transactions 449
The commit service and the application program 450
The probe process 452
Two-phase commit routines 452
Specifying the commit server 454
Two-phase commit example program 454
Program Notes 461
Program note 1 461
Program note 2 461

Program note 3 461
Program note 4 461
Program note 5 462
Program note 6 463
Program note 7 463
Program note 8 464
abort_xact 464
build_xact_string 465
close_commit 466
commit_xact 466
open_commit 467
remove_xact 467
scan_xact 468
start_xact 469
stat_xact 470

APPENDIX A **Cursors 471**
Cursor overview 471
DB-Library cursor capability 471
Differences between DB-Library cursors and browse mode . 472
Differences between DB-Library and Client-Library cursors . 472
Sensitivity to change 473
Static cursor 474
Keyset-driven cursor 474
Dynamic cursor 475
Concurrency control 475
DB-Library cursor functions 476
Holding locks 476
Stored procedures used by DB-Library cursors 477

About This Book

This document, the *Open Client DB-Library/C Reference Manual*, contains reference information for the C version of Open Client™ DB-Library™.

Audience

This book is intended to serve as a reference manual for programmers who are writing DB-Library applications. It is written for application programmers familiar with the C programming language.

How to use this book

When writing a DB-Library application, use the *DB-Library Reference Manual* as a source of reference information.

Chapter 1, “Introducing DB-Library” contains a brief introduction to DB-Library.

Chapter 2, “Routines” contains specific information about each DB-Library routine, such as what parameters the routine takes and what it returns.

Chapter 3, “Bulk copy” contains an introduction to bulk copy and specific information about each bulk copy routine.

Chapter 4, “Two-phase commit service” contains a brief description of two-phase commit service and specific information about each two-phase commit service routine.

Appendix A, “Cursors” introduces DB-Library’s cursor routines.

DB-Library and Open Client version 12.5

Open Client version 12.5 includes DB-Library version 10.0.x. This version of DB-Library includes no new functionality but does include bug fixes made since the previous version.

Sybase® includes DB-Library with Open Client to provide source-code compatibility for existing applications that are coded in DB-Library. Sybase recommends that you develop new applications with Client-Library™ instead of DB-Library. DB-Library does not support the new features added to Client-Library 11.x and later releases.

For information on new features in the latest Client-Library version, see the *Open Client Client-Library/C Reference Manual*.

The *Open Client Client-Library Migration Guide* contains instructions for converting DB-Library code to Client-Library code. Developers who are familiar with DB-Library will find this document a useful introduction to Client-Library.

Related documents

The *Open Client/Server Programmer's Supplement* contains platform-specific material for Open Client/Server™ developers. This document includes information about:

- Compiling and linking an application
- The example programs that are included online with Open Client/Server products
- Routines that have platform-specific behaviors

The *Open Client/Server Configuration Guide* contains platform-specific material about configuring the environment for Open Client/Server applications. This document includes information about:

- The interfaces file
- Environment variables used at application run time by DB-Library routines

The *Open Client/Server Release Bulletin* contains information about DB-Library compatibility with Sybase Adaptive Server™ and Open Server™ releases, as well as important last-minute information for DB-Library developers.

- The *Adaptive Server Enterprise Reference Manual* describes the Transact-SQL® database language, which an application uses to create and manipulate Adaptive Server database objects.
- The *Open Client Client-Library Reference Manual* describes Client-Library. Like DB-Library, Client-Library is a collection of routines for use in writing client applications.
- The *Open Server Server-Library Reference Manual* contains reference information for Server-Library, a collection of routines for use in writing Open Server applications.
- The *Open Client and Open Server Common-Libraries Reference Manual* contains reference information for the functions in CS-Library and Bulk Library. CS-Library is a collection of routines which are called in both Client-Library and Server-Library applications. Bulk Library is a collection of routines for performing bulk-copy operations from Client-Library or Server-Library applications.

Other sources of information

Use the Sybase Technical Library CD and the Technical Library Product Manuals Web site to learn more about your product:

- Technical Library CD contains product manuals and is included with your software. The DynaText browser (downloadable from Product Manuals at http://www.sybase.com/detail/1,3693,1010661,00.html) allows you to access technical information about your product in an easy-to-use format.

  Refer to the *Technical Library Installation Guide* in your documentation package for instructions on installing and starting the Technical Library.

- Technical Library Product Manuals Web site is an HTML version of the Technical Library CD that you can access using a standard Web browser. In addition to product manuals, you will find links to the Technical Documents Web site (formerly known as Tech Info Library), the Solved Cases page, and Sybase/Powersoft newsgroups.

  To access the Technical Library Product Manuals Web site, go to Product Manuals at http://www.sybase.com/support/manuals/.

Sybase documents include a wide range of user guides and reference manuals that describe all aspects of the Sybase relational database management system. Because application development can draw on a number of different parts of the Sybase system, you may encounter most of the Sybase document set at some time or another. A few manuals, however, will prove to be particularly useful:

Conventions

DB-Library routine syntax is shown in a bold, monospace font:

```
DBPROCESS    *dbopen(login, server)

LOGINREC     *login;
char         *server;
```

Program text and computer output are shown in regular monospace font:

```
dbproc = dbopen(login, SERVERNAME)
```

Routine names and Transact-SQL keywords are written in a narrow, bold font:

dbopen, the **select** statement

Online help

If you have access to SQL Server 10.0 or later, you can use **sp_syntax**, a Sybase system procedure, to retrieve the syntax of DB-Library routines.

For information on how to install **sp_syntax**, see the *System Administration Guide Supplement* for your platform. For information on how to run **sp_syntax**, see its reference page in the *Adaptive Server Enterprise Reference Manual*.

If you need help

Each Sybase installation that has purchased a support contract has one or more designated people who are authorized to contact Sybase Technical Support. If you cannot resolve a problem using the information in the Sybase documentation, ask the designated person to contact Sybase Technical Support.

CHAPTER 1

Introducing DB-Library

This chapter gives an overview of DB-Library. It includes the following sections:

Topic	Page
Client/server architecture	1
The Open Client and Open Server products	3
Data structures for communicating with servers	6
DB-Library/C programming	7
DB-Library/C routines	12
Summary of changes for version 10.0	36

Client/server architecture

Client/server architecture divides the work of computing between "clients" and "servers."

Clients make requests of servers and process the results of those requests. For example, a client application might request data from a database server. Another client application might send a request to an environmental control server to lower the temperature in a room.

Servers respond to requests by returning data or other information to clients, or by taking some action. For example, a database server returns tabular data and information about that data to clients, and an electronic mail server directs incoming mail toward its final destination.

Client/server architecture has several advantages over traditional program architectures:

- Application size and complexity can be significantly reduced because common services are handled in a single location, a server. This simplifies client applications, reduces duplicate code, and makes application maintenance easier.

- Client/server architecture facilitates communication between varied applications. Client applications that use dissimilar communications protocols cannot communicate directly, but can communicate through a server that “speaks” both protocols.
- Client/server architecture enables applications to be developed with distinct components, which can be modified or replaced without affecting other parts of the application.

Types of clients

A client is any application that makes requests of a server. Clients include:

- Sybase SQL Toolset™ products such as APT-Edit and Data Workbench
- Stand-alone utilities provided with Adaptive Server, such as **isql** and **bcp**
- Applications written using Open Client libraries
- Applications written using Open Client Embedded SQL™

Types of servers

The Sybase product line includes servers and tools for building servers:

- Sybase Adaptive Server is a database server. Adaptive Servers manage information stored in one or more databases.
- Sybase Open Server provides the tools and interfaces needed to create a custom server, also called an “Open Server application.”

An Open Server application can be any type of server. For example, an Open Server application can perform specialized calculations, provide access to real time data, or interface with services such as electronic mail. An Open Server application is created individually, using the building blocks provided by Open Server Server-Library.

Adaptive Server and Open Server applications are similar in some ways:

- Adaptive Server and Open Server applications are both servers, responding to client requests.
- Clients communicate with both Adaptive Server and Open Server applications through Open Client products.

But they also differ:

- An application programmer must create an Open Server application using Server-Library's building blocks and supplying custom code. Adaptive Server is complete and does not require custom code.
- An Open Server application can be any kind of server, and can be written to understand any language. Adaptive Server is a database server, and understands only Transact-SQL.
- An Open Server can communicate with "foreign" applications and servers that are not based on the TDS protocol, as well as Sybase applications and servers. Adaptive Server can communicate directly only with Sybase applications and servers, although Adaptive Server can communicate with foreign applications and servers by using an Open Server gateway application as an intermediary.

The Open Client and Open Server products

Sybase provides two families of products to enable customers to write client and server application programs. They are Open Client and Open Server.

Open Client

Open Client provides customer applications, third-party products, and other Sybase products with the interfaces needed to communicate with Adaptive Server and Open Server.

Open Client can be thought of as having two components: programming interfaces and network services.

The programming interfaces component of Open Client is made up of libraries designed for use in writing client applications: Client-Library, DB-Library, and CS-Library. (Both Open Client and Open Server include CS-Library, which contains utility routines that are useful to both client and server applications.

Open Client network services include Net-Library, which provides support for specific network protocols, such as TCP/IP or DECnet.

Open Server

Open Server provides the tools and interfaces needed to create custom server applications.

Like Open Client, Open Server has a programming interfaces component and a network services component.

The programming interfaces component of Open Server contains Server-Library and CS-Library. (Both Open Client and Open Server include CS-Library, which contains utility routines that are useful to both client and server applications.

Open Server network services are generally transparent.

Open Client libraries

The libraries that make up Open Client are:

- DB-Library, a collection of routines for use in writing client applications. DB-Library includes a bulk copy library and the two-phase commit special library. DB-Library provides source-code compatibility for older Sybase applications.
- Client-Library, a collection of routines for use in writing client applications. Client-Library is a library designed to accommodate cursors and other advanced features.
- CS-Library, a collection of utility routines that are useful to both client and server applications. All Client-Library applications will include at least one call to CS-Library, because Client-Library routines use a structure which is allocated in CS-Library.

What is in DB-Library/C?

Note DB-Library provides source code compatibility for older Sybase applications. Sybase encourages programmers to implement new applications with Client-Library or Embedded SQL.

DB-Library/C includes C routines and macros that allow an application to interact with Adaptive Server and Open Server applications.

It includes routines that send commands to Adaptive Server and Open Server applications and others that process the results of those commands. Other routines handle error conditions, perform data conversion, and provide a variety of information about the application's interaction with a server.

DB-Library/C also contains several header files that define structures and values used by the routines. Versions of DB-Library have been developed for a number of languages besides C, including COBOL, FORTRAN, Ada, and Pascal.

Comparing the library approach to Embedded SQL

Either an Open Client library application or an Embedded SQL application can be used to send SQL commands to Adaptive Server.

Generally, Embedded SQL is a superset of Transact-SQL. An Embedded SQL application includes Embedded SQL commands intermixed with the application's host language statements. The host language precompiler processes the Embedded SQL commands into calls to Client-Library routines and leaves the existing host-language statements as is. All version 10.0 or later precompilers use a runtime library composed solely of documented Client-Library and CS-Library calls.

In a sense, then, the precompiler transforms an Embedded SQL application into a Client-Library application.

An Open Client library application sends SQL commands through library routines, and does not require a precompiler.

Generally, an Embedded SQL application is easier to write and debug, but a library application can take fuller advantage of the flexibility and power of Open Client routines.

Other Sybase development tools

Sybase provides a complete development environment. In addition to control over server and database access through DB-Library/C, Sybase offers optional tools for creating forms-based, event-driven applications. These tools are APT-Edit and the APT-Library™/C set of routines. APT-Edit is an editor for creating forms and specifying event-driven processing. APT-Library/C is a set of C routines that support forms-based, event-driven processing. Just as DB-Library/C routines provide an application with access to servers and database objects, APT-Library/C routines provide access to a form and its objects. A forms-based application typically contains both APT-Library/C and DB-Library/C routines. For more information, see the *APT Workbench User's Guide*.

Data structures for communicating with servers

A DB-Library/C application communicates with a server through one or more DBPROCESS structures. Through the DBPROCESS, commands are sent to the server and query results are returned to the application. One of the first routines an application typically calls is dbopen, which logs the application into the server and allocates and initializes a DBPROCESS. This DBPROCESS then serves as a connection between the application and the server. Most DB-Library/C routines require a DBPROCESS as the first parameter.

An application can have multiple open DBPROCESSes, connected to one or more servers. For instance, an application that has to perform database updates in the midst of processing the results of a query needs a separate DBPROCESS for each task. As another example, to select data from one server and update a database on another server, an application needs two DBPROCESSes—one for each server. Each DBPROCESS in an application functions independently of any other DBPROCESS.

The DBPROCESS structure points to a command buffer that contains language commands for transmission to the server. It also points to result rows returned from the server—either single rows or buffers of rows if buffering has been specified. In addition, it points to a message buffer that contains error and informational messages returned from the server.

The DBPROCESS also contains a wealth of information on various aspects of server interaction. Many of the DB-Library/C routines deal with extracting information from the DBPROCESS. Applications should access and manipulate components of the DBPROCESS structure only through DB-Library/C routines, and not directly.

One other important structure is the LOGINREC. It contains typical login information, such as the user name and password, which the dbopen routine uses when logging into a server. DB-Library/C routines can specify the information in the LOGINREC.

DB-Library/C programming

An application programmer writes a DB-Library program, using calls to DB-Library routines to set up DB-Library structures, connect to servers, send commands, process results, and clean up. A DB-Library program is compiled and run in the same way as any other C language program.

Programming with DB-Library/C typically involves a few basic steps:

1 Logging into a server.

2 Placing language commands into a buffer and sending them to the server.

3 Processing the results, if any, returned from the server, one command at a time and one result row at a time. The results can be placed in program variables, where they can be manipulated by the application.

4 Handling DB-Library/C errors and server messages.

5 Closing the connection with the server.

The example below shows the basic framework of many DB-Library/C applications. The program opens a connection to a Adaptive Server, sends a Transact-SQL **select** command to the server, and processes the set of rows resulting from the **select**. For brevity's sake, this program does not include the error or message handling routines; those routines are illustrated in the example programs included online with DB-Library.

```
#include <sybfront.h>
 #include <sybdb.h>
 #include <syberror.h>

 /* Forward declarations of the error handler and message
 ** handler.
```

```
*/
interr_handler();
intmsg_handler();

main()
{
     DBPROCESS    *dbproc;     /* The connection with */
                               /* SQL Server */
     LOGINREC     *login;      /* The login information */
     DBCHAR       name[40];
     DBCHAR       city[20];
     RETCODE      return_code;

     /* Initialize DB-Library */
     if (dbinit() == FAIL)
     exit(ERREXIT);

     /*
     ** Install user-supplied error-handling and message-
     ** handling routines. The code for these is omitted
     ** from this example for conciseness.
     */
     dberrhandle(err_handler);
     dbmsghandle(msg_handler);

     /* Get a LOGINREC */
     login = dblogin();
     DBSETLPWD(login, "server_password");
     DBSETLAPP(login, "example");

     /* Get a DBPROCESS structure for communication */
     /* with SQL Server. */
     dbproc = dbopen(login, NULL);

     /*
     ** Retrieve some columns from the "authors" table
     ** in the "pubs2" database.
     */

     /* First, put the command into the command buffer. */
     dbcmd(dbproc, "select au_lname, city from
           pubs2..authors");
   dbcmd(dbproc, "
        where state = 'CA' ");

     /*
```

```
        ** Send the command to SQL Server and start execution
        */
        dbsqlexec(dbproc);

        /* Process the command */
        while ((return_code = dbresults(dbproc)) !=
             NO_MORE_RESULTS)
        {
             if (return_code == SUCCEED)
        {
             /* Bind results to program variables. */
             dbbind(dbproc, 1, STRINGBIND, (DBINT)0, name);
             dbbind(dbproc, 2, STRINGBIND, (DBINT)0, city);

             /* Retrieve and print the result rows. */
             while (dbnextrow(dbproc) != NO_MORE_ROWS)
             {
             printf ("%s: %s\n", name, city);
             }
         }
    }

    /* Close the connection to SQL Server */
    dbexit();
}
```

The example illustrates features common to most DB-Library/C applications:

- *Header files* – Two header files, *sybfront.h* and *sybdb.h*, are required in all source files that contain calls to DB-Library/C routines. *sybfront.h* must appear first in the file. It defines symbolic constants such as function return values, described in the reference pages in Chapter 2, “Routines” and the exit values STDEXIT and ERREXIT. These exit values can be used as the argument for the C standard library function **exit**. Since they are defined appropriately for the operating system running the program, their use provides a system-independent approach to exiting the program. *sybfront.h* also includes type definitions for datatypes that can be used in program variable declarations. These datatypes are described a little later.

sybdb.h contains additional definitions, most of which are meant to be used only by the DB-Library/C routines and should not be directly accessed by the program. Of chief importance in *sybdb.h* is the definition of the DBPROCESS structure. As discussed earlier, the DBPROCESS structure should be manipulated only through DB-Library/C routines; you should not access its components directly. To ensure compatibility with future releases of DB-Library/C, use the contents of *sybdb.h* only as documented in the reference pages in Chapter 2, "Routines"

The third header file in the example, *syberror.h*, contains error severity values and should be included if the program refers to those values.

- dbinit – This routine initializes DB-Library/C. It must be the first DB-Library/C routine in the program. Not all DB-Library/C environments currently require the **dbinit** call. However, to ensure future compatibility and portability, you should include this call at the start of all DB-Library/C programs.
- dberrhandle and dbmsghandle – **dberrhandle** installs a user-supplied error-handling routine, which gets called automatically whenever the application encounters a DB-Library/C error. Similarly, **dbmsghandle** installs a message-handling routine, which gets called in response to informational or error messages returned from the server. The error and message handling routines are user-supplied. Sample handlers have not been supplied with this example, but are included with the online example programs. For more information on the example programs, see the *Open Client/Server Programmer's Supplement*.
- dblogin – This routine allocates a LOGINREC structure, which DB-Library/C will use to log in to the server. The two macros that follow set certain components of the LOGINREC. DBSETLUSER and DBSETLPWD set the user name and password that DB-Library/C will use when logging in. DBSETLAPP sets the name of the application, which will appear in Adaptive Server's *sysprocesses* table. Routines are available for setting other aspects of the LOGINREC. However, in most environments these routines are optional; the LOGINREC contains default values for each of the values they set.
- dbopen – The **dbopen** routine opens a connection between the application and a server. It uses the LOGINREC supplied by **dblogin** to log in to the server. It returns a DBPROCESS structure, which serves as the conduit for information between the application and the server. After this routine has been called, the application is connected with Adaptive Server and can now send Transact-SQL commands to Adaptive Server and process any results.

- dbcmd – This routine fills the command buffer with Transact-SQL commands, which can then be sent to Adaptive Server. Each succeeding call to **dbcmd** simply adds the supplied text to the end of any text already in the buffer. It is the programmer's responsibility to supply necessary blanks between words, such as the blank at the beginning of the text in the second **dbcmd** call in the example. Multiple commands can be included in the buffer. This example only shows how to send and process a single command, but DB-Library/C is designed to allow an application to send multiple commands to a server and process each command's set of results separately.
- dbsqlexec – This routine executes the command buffer; that is, it sends the contents of the buffer to Adaptive Server, which parses and executes them.
- dbresults – This routine gets the results of the current Transact-SQL command ready for processing. In this case, the buffer contains a single command that returns rows, so the program is required to call **dbresults** one time. **dbresults** is called in a loop, however, because it is good programming practice to do so. It is recommended that **dbresults** always be called in a loop, as it is in this example, even when it is not strictly necessary.
- dbbind – **dbbind** binds result columns to program variables. In the example, the first call to **dbbind** binds the first result column to the program variable *city*. In other words, when the program reads a result row by calling **dbnextrow**, the contents of the first result column ("au_lname") will get placed in the program variable *name*. The second **dbbind** call binds the second result column to the variable *city*.

  The bind type of both bindings is STRINGBIND, one of several binding types available for character data. The binding type must correspond to the datatype of the specified program variable. In this example the variable has a DBCHAR datatype, a DB-Library/C-defined datatype that accepts a STRINGBIND result. By means of the binding type parameter, **dbbind** supports a wide variety of type conversions, allowing the datatype of the receiving variable to differ from the datatype of the result column.
- dbnextrow – This routine reads a row and places the results in the program variables specified by the earlier **dbbind** calls. Each successive call to **dbnextrow** reads another result row, until the last row has been read and NO_MORE_ROWS is returned. Processing of the results must take place inside the **dbnextrow** loop, because each call to **dbnextrow** overwrites the previous values in the program variables. This example program merely prints each row's contents.

- dbexit – This routine closes the server connection and deallocates the DBPROCESS. It also cleans up any structures initialized by **dbinit**. It must be the last DB-Library/C routine in the program.

Although DB-Library/C contains a great number of routines, much can be accomplished with just the few routines shown in this example.

DB-Library/C datatypes

DB-Library/C defines datatypes for Adaptive Server data. These datatypes begin with "SYB" (for example, SYBINT4, SYBCHAR, SYBMONEY). Various routines require these datatypes as parameters. DB-Library/C and Server-Library/C also provide type definitions for use in program variable declarations. These types begin with the prefix "DB" (for example, DBINT, DBCHAR, DBMONEY, and so on) for DB-Library/C, and "SRV_" for Server-Library/C (for example, SRV_INT4, SRV_CHAR, SRV_MONEY). By using them, you ensure that your program variables will be compatible.

See Types on page 408 for a list of Adaptive Server datatypes and corresponding DB-Library/C program variable types. See the *Open Server Server-Library/C Reference Manual* for a list of Server-Library datatypes.

The dbconvert_ps routine provides a way to convert data from one server datatype to another. It supports conversion between most datatypes. Since Adaptive Server and Open Server datatypes correspond directly to the DB-Library/C datatypes, you can use **dbconvert_ps** widely within your application. The routines that bind server result columns to program variables—**dbbind** and **dbaltbind**—also provide type conversion.

DB-Library/C routines

The DB-Library/C routines and macros handle a large variety of tasks, which are divided in this section into a number of categories:

- Initialization
- Command processing
- Results processing
- Message and error handling

- Information retrieval
- Browse mode
- text and image handling
- Datatype conversion
- Process control flow
- Remote procedure call processing
- Registered procedure call processing
- Datetime and money
- Cleanup
- Miscellaneous routines

The routines and macros are described in individual reference pages in Chapter 2, “Routines” They all begin with the prefix “db”. The routines are named with lowercase letters; the macros are capitalized.

In addition, DB-Library/C includes two special libraries:

- Bulk Copy, described in Chapter 3, “Bulk copy”
- Two-Phase Commit Service, described in Chapter 4, “Two-phase commit service”

The bulk copy routines begin with the prefix “bcp”. The two-phase commit routines have no standard prefix.

Initialization

These routines set up and define the connection between the application program and a server. They handle such tasks as allocating and defining a LOGINREC structure, opening a connection to a server, and allocating a DBPROCESS structure. Only a few of the routines are absolutely necessary in every DB-Library/C program; in particular, an application requires **dbinit**, **dblogin** and **dbopen**. The lists below specify the initialization routines in the approximate order in which a program is likely to call them.

Initializing DB-Library/C

These are the top level routines that set up DB-Library’s internal environment:

- dbinit – initializes underlying structures used by DB-Library/C
- dbsetversion – specifies a DB-Library version level
- dbsetmaxprocs – sets the maximum number of simultaneously open DBPROCESS structures
- dbgetmaxprocs – indicates the current maximum number of simultaneously open DBPROCESS structures

Setting up the LOGINREC

These routines place data in a LOGINREC. The LOGINREC contains the user information that DB-Library sends to the server when the program calls **dbopen** to open a connection.

- dblogin – allocates a LOGINREC structure for subsequent use by **dbopen**.
- DBSETLUSER – sets the server user name in the LOGINREC.
- DBSETLPWD – sets the server password in the LOGINREC.
- DBSETLAPP – sets the application name in the LOGINREC.
- DBSETLHOST – sets the host name in the LOGINREC.
- DBSETLCHARSET – sets the character set in the LOGINREC.
- DBSETLPACKET – sets the Tabular Data Stream™ (TDS) packet size for an application
- dbgetpacket – returns the current TDS packet size
- dbrpwset – adds a remote password to a LOGINREC structure. The server will use this password when it performs a remote procedure call on another server.
- dbrpwclr – clears all remote passwords from a LOGINREC structure.
- dbloginfree – frees a LOGINREC structure.

Establishing a server connection

The application calls the following routines to set up and open a connection to a remote server:

- dbsetifile – specifies the interfaces file that **dbopen** will use to connect to a server.

- dbsetlogintime – sets the number of seconds DB-Library/C will wait for a server to respond to a request by **dbopen** for a DBPROCESS connection.
- dbopen – sets up communication with the network, logs into a server using the LOGINREC, initializes any options specified in the LOGINREC, and allocates a DBPROCESS. An application can open multiple connections to a server, each connection having its own DBPROCESS. An application can also open multiple connections to multiple servers.
- dbuse – sets the current database. This routine is equivalent to the Transact-SQL **use** command and can be called repeatedly in an application, any time when the connection is open.

Command processing

An application can communicate with a server through language commands. For Adaptive Server, the language is Transact-SQL. For Open Server, the language is whatever the Open Server has been programmed to understand. The application enters the commands into a command buffer, which the DBPROCESS points to. The application can place multiple commands in the command buffer, and the set of commands in the buffer is known as the command batch. The application then sends the command batch to the server, which executes the commands in the order entered in the buffer.

Building the command batch

These routines add commands to the buffer or clear the buffer:

- dbcmd – adds text to the command buffer. It may be called repeatedly to add multiple commands, or parts of commands. The text added with each successive call is concatenated to the previous text.
- dbfcmd – adds text to the command buffer using **sprintf**-type formatting. This routine is the same as **dbcmd**, except that it allows arguments to be substituted into the text.
- dbfreebuf – clears the command buffer. The command buffer is automatically cleared before a batch of commands is entered. To clear it at other times or when the DBNOAUTOFREE option has been set, use **dbfreebuf**.

Accessing the command batch

These routines may be used to examine and copy parts of the command buffer:

- dbgetchar – returns a pointer to a particular character in the command buffer.
- dbstrlen – returns the length of the command buffer.
- dbstrcpy – copies a portion of the command buffer to a program variable. This routine is particularly valuable for debugging, because it can tell you exactly what was sent to the server.

Executing the command batch

Once language commands have been entered in the buffer, they can be sent to a server for execution.

- dbsqlsend – sends the contents of the command buffer to a server for execution. Unlike **dbsqlexec**, this routine does not wait for a response from the server. When **dbsqlsend** has returned SUCCEED, **dbsqlok** must be called to verify the correctness of the command batch.
- dbpoll – when called between **dbsqlsend** (or **dbrpcsend**) and **dbsqlok**, checks if a server response has arrived for a DBPROCESS.
- dbsqlok – waits for results from the server and verifies the correctness of the instructions the server is responding to. This routine is used in conjunction with **dbsqlsend**, **dbrpcsend**, and **dbmoretext**. After a successful **dbsqlok** call, the application must call **dbresults** to process the results.
- dbsqlexec – sends the contents of the command buffer to a server for execution. Once **dbsqlexec** has returned SUCCEED, **dbresults** must be called to process the results. Calling **dbsqlexec** is equivalent to calling **dbsqlsend** followed by **dbsqlok**.

Setting and clearing command options

The application can set a number of Adaptive Server and DB-Library/C command options. Among them are DBPARSEONLY, which causes Adaptive Server to parse but not execute the command batch, and DBBUFFER, which provides buffering of result rows. For a list of all available options and their significance, see Options on page 402.

- dbsetopt – sets an option

- dbclropt – clears an option
- dbisopt – determines whether a particular option is set

Results processing

Once a command batch has been executed in the server, indicated by **dbsqlexec** or **dbsqlok** returning SUCCEED, the application must process any results. Results can include:

- Success or failure indications from the server
- Result rows

Result rows are returned by **select** commands and **execute** commands on stored procedures that contain **select** commands.

There are two types of result rows: regular rows and compute rows. Regular rows are generated from columns in a **select** command's select list; compute rows are generated from columns in a **select** command's **compute** clause. Since these two types of rows contain very different data, the application must process them separately.

The results for each Transact-SQL command in a batch are returned to the application separately. Within each command's set of results, the result rows are processed one at a time.

If a command batch contains only a single Transact-SQL command and that command returns rows (for example, a **select** command), an application must call dbresults to process the results of the command.

If a command batch contains only a single Transact-SQL command and that command does not return rows (for example, a **use database** command or an **insert** command), an application does not have to call **dbresults** to process the results of the command. However, calling **dbresults** in these situations causes no harm. It may result in easier code maintenance if, after every command, you consistently call **dbresults** until it returns NO_MORE_RESULTS.

If the command batch contains more than one Transact-SQL command, an application must call **dbresults** once for every command in the batch, whether or not the command returns rows. For this reason, it is recommended that a DB-Library/C application always call **dbresults** in a loop after sending a command or commands to a server.

The following table lists Transact-SQL commands and the DB-Library/C functions required to process the results that they return:

Table 1-1: DB-Library/C functions required to process Transact-SQL commands

Transact-SQL command	Required DB-Library/C functions
All Transact-SQL commands not listed elsewhere in this table.	**dbresults**. In some cases, for example **dbcc**, the command's normal output is considered by DB-Library/C to consist of errors and messages. The output is thus processed within a DB-Library/C application's error and message handlers instead of in the main program via **dbnextrow** or other DB-Library/C routines.
execute	A DB-Library/C application must call **dbresults** once for every set of results that the stored procedure returns. In addition, if the stored procedure returns rows, the application must call **dbnextrow** or other DB-Library/C result-row routines.
select	**dbresults**. In addition, a DB-Library/C application must call **dbnextrow** or other DB-Library/C result-row routines.

Setting up the results

dbresults sets up the results of the next command in the batch. **dbresults** must be called after **dbsqlexec** or **dbsqlok** has returned SUCCEED, but before calls to **dbbind** or **dbnextrow**.

Getting result data

The simplest way to get result data is to bind result columns to program variables, using **dbbind** and **dbaltbind**. Then, when the application calls **dbnextrow** to read a result row (see "Reading result rows" on page 19), DB-Library/C will automatically place copies of the columns' data into the program variables to which they are bound. The application must call **dbbind** and **dbaltbind** after a **dbresults** call but before the first call to **dbnextrow**.

You can also access a result column's data directly with **dbdata** and **dbadata**, which return pointers to the data. **dbdata** and **dbadata** have the advantage of providing access to the actual data, not a copy of the data. These routines are frequently used in conjunction with **dbdatlen** and **dbadlen**, which return the length of the data and are described in the section "Information retrieval" on page 24. When you are accessing data directly with these routines, you do not perform any preliminary binding of result columns to program variables. Simply call **dbdata** or **dbadata** after a **dbnextrow** call.

The following routines are used to retrieve result columns:

- dbbind – binds a regular row result column to a program variable.
- dbbind_ps – binds a regular row result column to a program variable, with precision and scale support for numeric and decimal variables.
- dbaltbind – binds a compute row result column to a program variable.
- dbaltbind_ps – binds a compute row result column to a program variable, with precision and scale support for numeric and decimal variables.
- dbdata – returns a pointer to the data for a regular row result column.
- dbadata – returns a pointer to the data for a compute row result column.
- dbnullbind – associates an indicator variable with a regular row result column.
- dbanullbind – associates an indicator variable with a compute-row column.
- dbsetnull – defines substitution values to be used when binding null values.
- dbprtype – converts a server type token into a readable string. Tokens are returned by various routines such as **dbcoltype** and **dbaltop**.

Reading result rows

Once **dbresults** has returned SUCCEED and any binding of columns to variables has been specified, the application is ready to process the results. The first step is to make the result rows available to the application. The **dbnextrow** routine accomplishes this. Each call to **dbnextrow** reads the next row returned from the server. The row is read directly from the network.

Once a row has been read in by **dbnextrow**, the application can perform any processing desired on the data in the row. If the result columns have been bound to program variables, the data in the row will have been automatically copied into the variables. Alternatively, the data is accessible through **dbdata** or **dbadata**.

Rows read in by **dbnextrow** may be automatically saved in a row buffer, if desired. The application accomplishes this by setting the DBBUFFER option with the **dbsetopt** routine. Row buffering is useful for applications that need to process result rows in a non-sequential manner. Without row buffering, the application must process each row as it is read in by **dbnextrow**, because the next call to **dbnextrow** will overwrite the row. If the application has enabled row buffering, the rows are added to a row buffer as they are read in by **dbnextrow**. The application can then use the **dbgetrow** routine to skip around in the buffer and return to previously read rows. Since row buffering carries a memory and performance penalty, use it with discretion. Note that row buffering has nothing to do with network buffering and is a completely independent issue.

Routines are also available to print result rows in a default format. Because the format is predetermined, these routines are of limited usefulness and are appropriate primarily for debugging.

Note that DB-Library/C processes results one command at a time. When the application has read all the results for one command, it must call **dbresults** again to set up the results for the next command in the command buffer. To ensure that all results are handled, Sybase strongly recommends that **dbresults** be called in a loop.

The following routines are used to process result rows:

- dbnextrow – reads in the next row. The return value from **dbnextrow** tells the application whether the row is a regular row or a compute row, whether the row buffer is full, and whether the last result row has been read.
- DBCURROW – returns the number of the row currently being read.
- dbprhead – prints default column headings for result rows. This routine is used in conjunction with **dbprrow**.
- dbprrow – prints all the result rows in a default format. When this routine is used, the program does not need to bind results or call **dbnextrow**.

Canceling results

The following routines cancel results:

- dbcancel – cancels results from the current command batch. This routine cancels *all* the commands in the current batch.
- dbcanquery – cancels any rows pending from the most recently executed query.

As an example of the difference between these routines, consider an application that is processing the results of the language batch:

```
select * from pubs.titles
 select * from pubs.authors
```

If the application calls **dbcanquery** while processing the *titles* rows, then the *titles* rows are discarded and the application must continue to call **dbresults** and process the rows from the next statement. If the application calls **dbcancel** while processing the *titles* rows, then DB-Library discards the *titles* rows and the results of all remaining, unprocessed commands in the batch. The application does not need to continue calling **dbresults** after calling **dbcancel**.

Handling stored procedure results

A call to a stored procedure is made through either a remote procedure call, discussed in "Remote procedure call processing" on page 30, or a Transact-SQL **execute** command. The call can generate several types of results. First of all, a stored procedure that contains **select** statements will return result rows in the usual fashion. Each successive call to **dbresults** will access the set of rows from the next **select** statement in the stored procedure. These rows can be processed, as usual, with **dbnextrow**.

Second, stored procedures can contain "return parameters." Return parameters, also called output parameters, provide stored procedures with a "call-by-reference" capability. Any change that a stored procedure makes internally to the value of an output parameter is available to the calling program. The calling program can retrieve output parameter values once it has processed all of the stored procedure's result rows by calling **dbresults** and **dbnextrow**. A number of routines, described below, process return parameter values.

Third, stored procedures can return a status number.

To access a stored procedure's output parameters and return status through the following routines:

- dbnumrets – returns the number of return parameter values generated by a stored procedure. If **dbnumrets** returns less than or equal to zero, no return parameter values are available.
- dbretdata – returns a pointer to a return parameter value.
- dbretlen – returns the length of a return parameter value.
- dbrettype – returns the datatype of a return parameter value.

- dbretname – returns the name of the return parameter associated with a particular value.
- dbretstatus – returns the stored procedure's status number.
- dbhasretstat – indicates whether the current command or remote procedure call generated a stored procedure status number. If **dbhasretstat** returns FALSE, then no stored procedure status number is available.

Setting results timeouts

By default, DB-Library will wait indefinitely for the results of a server command to arrive. Applications can use the routines below to specify a finite timeout period:

- dbsettime – sets the number of seconds that DB-Library/C will wait for a server response
- DBGETTIME – gets the number of seconds that DB-Library/C will wait for a server response

Message and error handling

DB-Library/C applications must handle two types of messages and errors:

- Server messages and errors, which range in severity from informational messages to fatal errors. Server messages and errors are known to DB-Library/C applications as "messages." To list all possible Adaptive Server messages, use the Transact-SQL command:

  ```
  select * from sysmessages
  ```

  Also see the *System Administration Guide*. See the *Open Server Server-Library/C Reference Manual* for a list of Open Server messages.
- DB-Library/C warnings and errors, known to DB-Library/C applications as "errors." For a list of DB-Library/C errors, see Errors on page 386.

Also, success or failure indications are returned by most DB-Library/C routines.

To handle server messages, DB-Library/C errors, and success or failure indications, a DB-Library/C application can:

- Test DB-Library/C routine return codes in the mainline code, handling failures on a case-by-case basis
- Centralize message and error handling by installing a message handler and an error handler, which are then automatically called by DB-Library/C when a message or error occurs

Sybase strongly recommends that all DB-Library/C applications use centralized message and error handling in addition to mainline error testing. Centralized message and error handling has substantial benefits for large or complex applications. For example:

- Centralized message and error handling reduces the need for mainline error-handling logic. This is because DB-Library/C calls an application's message and error handlers automatically whenever a message or error occurs.

  Note, however, that even an application that uses centralized error and message handling will need some mainline error logic, depending on the nature of the application.
- Centralized message and error handling provides a mechanism for gracefully handling unexpected errors. An application using only mainline error-handling logic may not successfully trap errors which have not been anticipated.

To provide a DB-Library/C application with centralized message and error handling, the application programmer must write a message handler and an error handler and install them via **dbmsghandle** and **dberrhandle**.

The DB-Library/C routines for message and error handling are:

- dbmsghandle – installs a user function to handle server informational and error messages.
- dberrhandle – installs a user function to handle DB-Library/C error messages.
- DBDEAD – determines whether a particular DBPROCESS is dead. When a DBPROCESS is dead, the current DB-Library/C routine fails, causing the error handler to be called.

Information retrieval

Information covering several areas, including regular result columns, compute result columns, row buffers, and the command state, can be retrieved from the DBPROCESS structure. As mentioned earlier, regular result columns correspond to columns in the **select** command's select list and compute result columns correspond to columns in the **select** command's optional **compute** clause.

Regular result column information

These routines can be called after **dbsqlexec** returns SUCCEED:

- dbnumcols – determines the number of columns in the current set of results.
- dbcolname – returns the name of a regular result column.
- dbcollen – returns the maximum length for a regular column's data.
- dbcoltype – returns the server datatype for a regular result column.
- dbdatlen – returns the actual length of a regular column's data. This routine is often used in conjunction with **dbdata**. The value returned by **dbdatlen** is different for each regular row read by **dbnextrow**.
- dbvarylen – indicates whether the column's data can vary in length.

Compute result column information

These routines can be called after **dbsqlexec** returns SUCCEED:

- DBROWTYPE – indicates whether the current result row is a regular row or a compute row
- dbnumcompute – returns the number of **compute** clauses in the current set of results.
- dbnumalts – returns the number of columns in a compute row.
- dbbylist – returns the bylist for a compute row.
- dbaltop – returns the type of aggregate operator for a compute column.
- dbalttype – returns the datatype for a compute column.
- dbaltlen – returns the maximum length for a compute column's data.
- dbaltcolid – returns the column id for a compute column.

- dbadlen – returns the actual length of a compute column's data. This routine is often used in conjunction with **dbadata**. The value returned by **dbadlen** is different for each compute row read by **dbnextrow**.

Row buffer information

These macros return information that can be useful when manipulating result rows in buffers:

- DBFIRSTROW – returns the number of the first row in the buffer.
- DBLASTROW – returns the number of the last row in the buffer.
- dbgetrow – reads the specified row in the row buffer. This routine provides the application with access to buffered rows that have been previously read by **dbnextrow**.
- dbclrbuf – drops rows from the row buffer.

Command state information

These routines return information about the current state of the command batch. Several of them return information about the "current" command—that is, the command currently being processed by **dbresults**.

- DBCURCMD – returns the number of the current command in a batch.
- dbgetoff – checks for the existence of specified Transact-SQL constructs in the command buffer. This routine is used in conjunction with the DBOFFSET option.
- DBMORECMDS – indicates whether there are more commands in the batch.
- DBCMDROW – indicates whether the current command is one that can return rows (that is, a **select** or a stored procedure containing a **select**).
- DBROWS – indicates whether the current command actually did return rows.
- DBCOUNT – returns the number of rows affected by a command.
- DBNUMORDERS – returns the number of columns specified in a **select** command's **order by** clause.
- dbordercol – returns the id of a column appearing in a **select** command's **order by** clause.

Browse mode

Browse mode provides a means for browsing through database rows and updating their values a row at a time. From the standpoint of the program, the process involves several steps, because each row must be transferred from the database into program variables before it can be browsed and updated.

Since a row being browsed is not the actual row residing in the database, but is instead a copy residing in program variables, the program must be able to ensure that changes to the variables' values can be reliably used to update the original database row. In particular, in multiuser situations, the program needs to ensure that updates made to the database by one user do not unwittingly overwrite updates recently made by another user. This can be a problem because the application typically selects a number of rows from the database at one time, but the application's users browse and update the database one row at a time. A timestamp column in browsable database tables provides the information necessary to regulate this type of multiuser updating.

Browse mode routines also allow an application to handle ad hoc queries. Several routines return information that an application can use to examine the structure of a complicated ad hoc query to update the underlying database tables.

Conceptually, browse mode involves three steps:

1 Select result rows containing columns derived from one or more database tables.

2 Where appropriate, change values in columns of the result rows (*not* the actual database rows), one row at a time.

3 Update the original database tables, one row at a time, using the new values in the result rows.

These steps are implemented in a program as follows:

1 Execute a **select** command, generating result rows containing result columns. The **select** command must include the **for browse** option.

2 Copy the result column values into program variables, one row at a time.

3 If appropriate, change the values of the variables (possibly in response to user input).

4 If appropriate, execute an **update** command that updates the database row corresponding to the current result row. To handle multiuser updates, the **where** clause of the **update** command must reference the timestamp column. Such a **where** clause can be obtained via the **dbqual** function.

5 Repeat steps 2, 3, and 4 for each result row.

To use browse mode, the following conditions must be true:

- The **select** command must end with the key words **for browse**.
- The table(s) to be updated must be "browsable" (that is, each must have a unique index and a timestamp column). Note that because a browse mode table has unique rows, the keyword **distinct** has no effect in a **select** against a browse-mode table.
- The result columns to be used in the updates must be "updatable"—they must be derived from browsable tables and cannot be the result of SQL expressions, such as "max(colname)". In other words, there must be a valid correspondence between the result column and the database column to be updated. In addition, browse mode usually requires two connections (DBPROCESS pointers)—one for selecting the data and another for performing updates based on the selected data.

For examples of browse-mode programming, see the sample programs included online.

The following constitute the browse-mode routines:

- dbqual – returns a pointer to a **where** clause suitable for use in updating the current row in a browsable table
- dbfreequal – frees the memory allocated by **dbqual**
- dbtsnewval – returns the new value of the timestamp column after a browse-mode update
- dbtsnewlen – returns the length of the new value of the timestamp column after a browse-mode update
- dbtsput – puts the new value of the timestamp column into the given table's current row in the DBPROCESS
- dbcolbrowse – indicates whether the source of a result column is updatable via browse mode
- dbcolsource – returns a pointer to the name of the database column from which the specified result column was derived
- dbtabbrowse – indicates whether a particular table is updatable via browse mode
- dbtabcount – returns the number of tables involved in the current **select** command

- dbtabname – returns the name of a table based on its number
- dbtabsource – returns the name and number of the table from which a particular result column was derived

text and image handling

text and image are Adaptive Server datatypes designed to hold large text or image values. The text datatype will hold up to 2,147,483,647 bytes of printable characters. The image datatype will hold up to 2,147,483,647 bytes of binary data.

Because they can be so large, text and image values are not actually stored in database tables. Instead, a pointer to the text or image value is stored in the table. This pointer is called a “text pointer.”

To ensure that competing applications do not wipe out one another’s modifications to the database, a timestamp is associated with each text or image column. This timestamp is called a “text timestamp.”

A DB-Library/C application that uses **dbwritetext** to insert text or image data into a table must perform the following steps:

1 Use the **insert** command to insert all data into the row except the text or image value.

2 Use the **update** command to update the row, setting the value of the text or image column to NULL. This step is necessary because a text or image column row that contains a null value will have a valid text pointer only if the null value was explicitly entered with the **update** statement.

3 Use the **select** command to select the row. You must specifically select the column that is to contain the text or image value. This step is necessary to provide the application’s DBPROCESS with correct text pointer and text timestamp information. The application should throw away the data returned by this **select**.

4 Call **dbtxtptr** to retrieve the text pointer from the DBPROCESS.

5 Call **dbtxtimestamp** to retrieve the text timestamp from the DBPROCESS.

6 Write the text or image value to Adaptive Server. An application can either:

   - Write the value with a single call to **dbwritetext**, or

- Write the value in chunks, using **dbwritetext** and **dbmoretext**.

7 If the application plans to make another update to this text or image value, it may want to save the new text timestamp that is returned by Adaptive Server at the conclusion of a successful **dbwritetext** operation. The new text timestamp may be accessed via **dbtxtsnewval** and stored for later retrieval via **dbtxtsput**.

Several routines are available to facilitate the process of updating text and image columns in database tables:

- dbreadtext – reads a text or image value from Adaptive Server
- dbwritetext – sends a text or image value to Adaptive Server
- dbmoretext – sends part of a text or image value to Adaptive Server
- dbtxptr – returns the text pointer for a column in the current results row
- dbtxtimestamp – returns the value of the text timestamp for a column in the current results row
- dbtxtsnewval – returns the new value of a text timestamp after a call to **dbwritetext**
- dbtxtsput – puts the new value of a text timestamp into the specified column of the current row in the DBPROCESS

Datatype conversion

DB-Library/C supports conversions between most server datatypes with the **dbconvert** and **dbconvert_ps** routines. For information on server datatypes, see “Types” on page 408.

The **dbbind**, **dbbind_ps**, **dbaltbind**, and **dbaltbind_ps** routines, which bind result columns to program variables, can also be used to perform type conversion. Those routines each contain a parameter that specifies the datatype of the receiving program variable. If the data being returned from the server is of a different datatype, DB-Library/C will usually convert it automatically to the type specified by the parameter.

These routines are used to perform datatype conversion:

- dbconvert_ps – converts data from one server datatype to another, with precision and scale support for numeric and decimal datatypes
- dbconvert – converts data from one server datatype to another

- dbwillconvert – indicates whether a specified datatype conversion is supported

Process control flow

These routines allow the application to schedule its actions around its interaction with a server:

- dbsetbusy – calls a user-supplied function when DB-Library/C is reading or waiting to read results from the server
- dbsetidle – calls a user-supplied function when DB-Library/C is finished reading from the server
- dbsetinterrupt – calls user-supplied functions to handle interrupts while waiting on a read from the server
- DBIORDESC (UNIX and AOS/VS only) – provides access to the UNIX file descriptor used to read data coming from the server, allowing the application to respond to multiple input data streams
- DBIOWDESC (UNIX and AOS/VS only) – provides access to the UNIX file descriptor used to write data to the server, allowing the application to effectively utilize multiple input and output data streams
- DBRBUF (UNIX and AOS/VS only) – determines whether the DB-Library/C network buffer contains any unread bytes

Remote procedure call processing

A remote procedure call is simply a call to a stored procedure residing on a remote server. Either an application or another server makes the call. A remote procedure call made by an application has the same effect as an **execute** command—it executes the stored procedure, generating results accessible through **dbresults**. However, a remote procedure call is often more efficient than an **execute** command. Note that if the procedure being executed resides on a server other than the one to which the application is directly connected, commands executed within the procedure cannot be rolled back.

A server can make a remote procedure call to another server. This occurs when a stored procedure being executed on one server contains an **execute** command for a stored procedure on another server. The **execute** command causes the first server to log in to the second server and perform a remote procedure call on the procedure. This happens without any intervention from the application, although the application can specify the remote password that the first server uses to log in.

The following routines are used to perform remote procedure calls:

- dbrpcinit – initializes a remote procedure call to a stored procedure.
- dbrpcparam – adds a parameter to a remote procedure call.
- dbrpcsend – signals the end of a remote procedure call, causing the server to begin executing the specified procedure.
- dbpoll – when called between **dbsqlsend** (or **dbrpcsend**) and **dbsqlok**, checks if a server response has arrived for a DBPROCESS.
- dbsqlok – waits for results from the server and verifies the correctness of the instructions the server is responding to. This routine is used in conjunction with **dbsqlsend**, **dbrpcsend**, and **dbmoretext**. After a successful **dbsqlok** call, the application must call **dbresults** to process the results.

Registered procedure call processing

A registered procedure is a procedure that is defined and installed in a running Open Server. Registered procedures require Open Server version 2.0 or later. At this time, registered procedures are not supported by Adaptive Server.

For DB-Library/C applications, registered procedures provide a means for inter-application communication and synchronization. This is because DB-Library/C applications connected to an Open Server can “watch” for a registered procedure to execute. When the registered procedure executes, applications watching for it receive a notification that includes the procedure’s name and the arguments it was called with.

Note DB-Library/C applications may create only a special type of registered procedure, known as a “notification procedure.” A notification procedure differs from a normal Open Server registered procedure in that it contains no executable statements.

For example, suppose the following:

- **stockprice** is a real-time DB-Library/C application monitoring stock prices.
- **price_change** is a notification procedure created in Open Server by the **stockprice** application. **price_change** takes as parameters a stock name and a price differential.
- **sellstock**, an application that puts stock up for sale, has requested to be notified when **price_change** executes.

When **stockprice**, the monitoring application, becomes aware that the price of Extravagant Auto Parts stock has risen $1.10, it executes **price_change** with the parameters "Extravagant Auto Parts" and "+1.10".

When **price_change** executes, Open Server sends **sellstock** a notification containing the name of the procedure (**price_change**) and the arguments passed to it ("Extravagant Auto Parts" and "+1.10"). **sellstock** uses the information contained in the notification to decide to put 100 shares of Extravagant Auto Parts stock up for sale.

price_change is the means through which the **stockprice** and **sellstock** applications communicate.

Registered procedures as a means of communication have the following advantages:

- A single call to execute a registered procedure can result in many client applications being notified that the procedure has executed. The application executing the procedure does not need to know how many, or which, clients have requested notifications.
- The registered procedure communication mechanism is server-based. Open Server acts as a central repository for connection addresses. Because of this, client applications can communicate without having to connect directly to each other. Instead, each client simply connects to the server.

A DB-Library/C application can:

- Create a registered procedure in Open Server
- Drop a registered procedure
- List all registered procedures defined in Open Server
- Request to be notified when a particular registered procedure is executed
- Drop a request to be notified when a particular registered procedure is executed

- List all registered procedure notifications
- Execute a registered procedure
- Install a user-supplied handler to be called when an application receives notification that a registered procedure has executed
- Poll Open Server to see if any registered procedure notifications are pending

The registered procedure routines are:

- dbnpcreate – creates a notification procedure
- dbnpdefine – defines a notification procedure
- dbregdrop – drops a registered procedure
- dbreglist – returns a list of all registered procedures currently defined in Open Server
- dbreghandle – installs a handler routine for a registered procedure notification
- dbreginit – initiates execution of a registered procedure
- dbregnowatch – cancels a request to be notified when a registered procedure executes
- dbregparam – defines a parameter for a registered procedure
- dbregexec – executes a registered procedure
- dbregwatch – requests to be notified when a registered procedure executes
- dbregwatchlist – returns a list of registered procedures that a DBPROCESS is watching for
- dbpoll – in an application that uses registered procedure notifications, used to check whether any notifications have arrived

Gateway passthrough routines

These routines can be called in Open Server gateway applications. They enable a DB-Library/C application to send and receive whole Tabular Data Stream (TDS) packets and set TDS packet size.

TDS is an application protocol used for the transfer of requests and request results between clients and servers. These routines are used with the **srvrecvpassthru** and **srvsendpassthru** Open Server Server-Library routines:

- dbrecvpassthru – receives a TDS packet from Open Server
- dbsendpassthru – sends a TDS packet to Open Server

See the *Open Server Server-Library/C Reference Manual* for descriptions of **srvrecvpassthru** and **srvsendpassthru**.

Datetime and money

These routines manipulate datetime and money datatypes. datetime and money datatypes come in long versions, DBDATETIME and DBMONEY, and short (4-byte) versions, DBDATETIME4 and DBMONEY4. All of the DBDATETIME4 routines listed below are also available for DBDATETIME, and all DBMONEY4 routines are available for DBMONEY. For example, **dbmny4add**, listed below, is also available as **dbmnyadd**.

- dbdate4cmp – compares two DATETIME4 values
- dbdate4zero – initializes a DBDATETIME4 value
- dbmny4add – adds two DBMONEY4 values
- dbmny4cmp – compares two DBMONEY4 values
- dbmny4copy – copies a DBMONEY4 value
- dbmny4divide – divides one DBMONEY4 value by another
- dbmny4minus – negates a DBMONEY4 value
- dbmny4mul – multiplies a DBMONEY4 value
- dbmny4sub – subtracts a DBMONEY4 value
- dbmny4zero – initializes a DBMONEY4 value
- dbmnydec – decrements a DBMONEY value
- dbmnydown – divides a DBMONEY value by a positive integer
- dbmnyinc – increments a DBMONEY value
- dbmnyinit – prepares a DBMONEY value for calls to **dbmnyndigit**
- dbmnymaxneg – returns the maximum negative DBMONEY value
- dbmnymaxpos – returns the maximum positive DBMONEY value
- dbmnyndigit – returns the rightmost digit of a DBMONEY value as a DBCHAR

- dbmnyscale – multiplies a DBMONEY value and add a specified amount

Cleanup

These routines sever the connection between the application and a server:

- dbexit – closes and deallocates all DBPROCESS structures. This routine also cleans up any structures initialized by **dbinit**.
- dbclose – closes and deallocates a single DBPROCESS structure.

Secure support

These routines provide security for DB-Library applications running against SQL Server 10.0 or later, or Adaptive Server:

- DBSETLENCRYPT – specifies whether or not password encryption is to be used when logging into SQL Server 10.0 or later
- dbsechandle – installs user functions to handle secure logins
- bcp_options – sets bulk copy options, including BCPLABELED, the security label option

---

Note Calling **DBSETLENCRYPT** causes an error unless you first set the DB-Library version to 10.0. Use **dbsetversion** to set the DB-Library version to 10.0 before calling **DBSETLENCRYPT**.

Miscellaneous routines

These routines may be useful in some applications:

- dbsetavail – marks a DBPROCESS as being available for general use.
- DBISAVAIL – indicates whether a DBPROCESS is available for general use.
- dbname – returns the name of the current database.
- dbchange – indicates whether a command batch has changed the current database.

- dbsetuserdata – uses a DBPROCESS structure to save a pointer to user-allocated data. This routine, along with **dbgetuserdata**, allows the application to associate user data with a particular DBPROCESS. One important use for these routines is to transfer information between a server message handler and the program code that triggered it.
- dbgetuserdata – returns a pointer to user-allocated data from a DBPROCESS structure.
- dbreadpage – reads in a page of binary data from Adaptive Server.
- dbwritepage – writes a page of binary data to Adaptive Server.

Bulk copy special library

The routines in this library allow an application to bulk copy data in and out of Adaptive Server. They provide a facility for high-speed loading of data into Adaptive Server from either files or program variables. They also allow you to copy data out of Adaptive Server into files, using a pre-defined format. The bulk copy routines are similar in function to the **bcp** utility program, described in *Adaptive Server Enterprise Utility Guide*.

For more information on bulk copy, see Chapter 3, “Bulk copy”

Two-phase commit service special library

The routines in this library allow an application to coordinate updates among two or more SQL Servers.

For more information on two-phase commit, see Chapter 4, “Two-phase commit service”

Summary of changes for version 10.0

This section contains information on changes to DB-Library/C between version 4.6 and version 10.0. No changes have been made to DB-Library since version 10.0 other than bug fixes.

Version 10.0 of DB-Library/C includes:

- Support for the 10.0 server numeric and decimal datatypes.

  Note Applications which use the numeric or decimal datatypes must call dbsetversion to specify version 10.0 behavior.

- Routines to support Adaptive Server security features.
- DB-Library cursor routines. For more information on cursors, see Appendix A, “Cursors”

DB-Library routines added to version 10.0

The following table lists DB-Library routines that were added for version 10.0:

Table 1-2: New DB-Library routines

Routine Name	Description
bcp_options	Sets bulk copy options
dbaltbind_ps	Binds a compute column to a program variable, with precision and scale support for numeric and decimal datatypes
dbbind_ps	Binds a regular result column to a program variable, with precision and scale support for numeric and decimal datatypes
dbconvert_ps	Converts data from one datatype to another, with precision and scale support for numeric and decimal datatypes
dbcursor	Inserts, updates, deletes, locks, or refreshes a particular row in the fetch buffer
dbcursorbind	Registers the binding information on the cursor columns
dbcursorclose	Closes the cursor associated with the given handle, releasing all the data belonging to it
dbcursorcolinfo	Returns column information for the specified column number in the open cursor
dbcursorfetch	Fetches a block of rows into the program variables declared by the user in **dbcursorbind**

Routine Name	Description
dbcursorinfo	Returns the number of columns and the number of rows in the keyset if the keyset hit the end of the result set
dbcursoropen	Opens a cursor, specifying the scroll option, concurrency option, and the size of the fetch buffer (the number of rows retrieved with a single fetch)
dbsechandle	Installs user functions to handle secure logins
dbsetdefcharset	Sets the default character set for an application
dbsetdeflang	Sets the default language name for an application
DBSETLENCRYPT	Specifies whether or not password encryption is to be used when logging into SQL Server 10.0 or later
dbsetversion	Specifies a DB-Library version level
dbtextsize	Returns the number of text or image bytes that remain to be read for the current row

DB-Library 12.0 and later

Note DB-Library 12.0 and later contains only bug fixes; that is, it has no new features and does not take advantage of directory services.

DB-Library 12.0 and later is non-reentrant and does not use native threads. You cannot use DB-Library 12.0 and later with an Open Server that uses multithreaded or reentrant (*_r.*) libraries and that runs on any UNIX system.

Recompiling DB-Library 4.x, 10.0.x, and 11.x applications

Applications written for DB-Library versions 4.x, 10.0.x, and 11.x must be recompiled to use DB-Library version 12.0 and later.

Sample programs and sample databases

Examples 3, 9, 10, and 11 require the *pubs2* database to be installed on the server. Example 12 requires the *interpubs* database to be installed on the server. For more details, see the *README* file in the *$SYBASE/$SYBASE_OCS/sample/dblibrary* directory or on NT, *%SYBASE%\%SYBASE_OCS%\sample\dblibrary* directory.

CHAPTER 2

Routines

This chapter contains a reference page for each DB-Library routine:

Routines	Description	Page
db12hour	Determines whether the specified language uses 12-hour or 24-hour time.	49
dbadata	Returns a pointer to the data for a compute column.	51
dbadlen on page 54	Returns the actual length of the data for a compute column.	54
dbaltbind	Binds a compute column to a program variable.	55
dbaltbind_ps	Binds a compute column to a program variable, with precision and scale support for numeric and decimal datatypes.	62
dbaltcolid	Returns the column id for a compute column.	68
dbaltlen	Returns the maximum length of the data for a particular compute column.	69
dbaltop	Returns the type of aggregate operator for a particular compute column.	70
dbalttype	Returns the datatype for a compute column.	71
dbaltutype	Returns the user-defined datatype for a compute column.	72
dbanullbind	Associates an indicator variable with a compute-row column.	73
dbbind	Bind a regular result column to a program variable.	74
dbbind_ps	Binds a regular result column to a program variable, with precision and scale support for numeric and decimal datatypes.	80
dbbufsize	Returns the size of a DBPROCESS row buffer.	84
dbbylist	Returns the bylist for a compute row.	85
dbcancel	Cancels the current command batch.	86
dbcanquery	Cancels any rows pending from the most recently executed query.	88
dbchange	Determines whether a command batch has changed the current database.	89
dbcharsetconv	Indicates whether the server is performing character set translation.	90
dbclose	Closes and deallocate a single DBPROCESS structure.	90
dbclrbuf	Drops rows from the row buffer.	91

Routines	Description	Page
dbclropt	Clears an option set by **dbsetopt**.	92
dbcmd	Adds text to the DBPROCESS command buffer.	94
DBCMDROW	Determines whether the current command can return rows.	95
dbcolbrowse	Determines whether the source of a regular result column is updatable via the DB-Library browse-mode facilities.	96
dbcollen	Returns the maximum length of the data in a regular result column.	97
dbcolname	Returns the name of a regular result column.	98
dbcolsource	Returns a pointer to the name of the database column from which the specified regular result column was derived.	99
dbcoltype	Returns the datatype for a regular result column.	100
dbcoltypeinfo	Returns precision and scale information for a regular result column of type *numeric* or decimal.	102
dbcolutype	Returns the user-defined datatype for a regular result column.	103
dbconvert	Converts data from one datatype to another.	104
dbconvert_ps	Converts data from one datatype to another, with precision and scale support for numeric and decimal datatypes.	109
DBCOUNT	Returns the number of rows affected by a Transact-SQL command.	115
DBCURCMD	Returns the number of the current command.	116
DBCURROW	Returns the number of the row currently being read.	116
dbcursor	Inserts, updates, deletes, locks, or refreshes a particular row in the fetch buffer.	117
dbcursorbind	Registers the binding information on the cursor columns.	119
dbcursorclose	Closes the cursor associated with the given handle and release all the data belonging to it.	121
dbcursorcolinfo	Returns column information for the specified column number in the open cursor.	122
dbcursorfetch	Fetches a block of rows into the program variables declared by the user in **dbcursorbind.**	123
dbcursorinfo	Returns the number of columns and the number of rows in the keyset if the keyset hit the end of the result set.	125
dbcursoropen	Opens a cursor and specify the scroll option, concurrency option, and the size of the fetch buffer (the number of rows retrieved with a single fetch).	125
dbdata	Returns a pointer to the data in a regular result column.	129
dbdate4cmp	Compares two DBDATETIME4 values.	130

Routines	Description	Page
dbdate4zero	Initializes a DBDATETIME4 variable to Jan 1, 1900 12:00AM.	131
dbdatechar	Converts an integer component of a DBDATETIME value into character format.	132
dbdatecmp	Compares two DBDATETIME values.	133
dbdatecrack	Converts a machine-readable DBDATETIME value into user-accessible format.	134
dbdatename	Converts the specified component of a DBDATETIME structure into its corresponding character string.	136
dbdateorder	Returns the date component order for a given language.	139
dbdatepart	Returns the specified part of a DBDATETIME value as a numeric value.	140
dbdatezero	Initializes a DBDATETIME value to Jan 1, 1900 12:00:00:000AM.	141
dbdatlen	Returns the length of the data in a regular result column.	142
dbdayname	Determines the name of a specified weekday in a specified language.	143
DBDEAD	Determines whether a particular DBPROCESS is dead.	144
dberrhandle	Installs a user function to handle DB-Library errors.	145
dbexit	Closes and deallocate all DBPROCESS structures, and clean up any structures initialized by **dbinit**.	149
dbfcmd	Adds text to the DBPROCESS command buffer using C runtime library **sprintf**-type formatting.	150
DBFIRSTROW	Returns the number of the first row in the row buffer.	153
dbfree_xlate	Frees a pair of character set translation tables.	154
dbfreebuf	Clears the command buffer.	155
dbfreequal	Frees the memory allocated by **dbqual**.	156
dbfreesort	Frees a sort order structure allocated by **dbloadsort**.	156
dbgetchar	Returns a pointer to a character in the command buffer.	158
dbgetcharset	Gets the name of the client character set from the DBPROCESS structure.	158
dbgetloginfo	Transfers TDS (Tabular Data Stream) login response information from a DBPROCESS structure to a newly allocated DBLOGINFO structure.	159
dbgetlusername	Returns the user name from a LOGINREC structure.	161
dbgetmaxprocs	Determines the current maximum number of simultaneously open DBPROCESSes.	162
dbgetnatlang	Gets the national language from the DBPROCESS structure.	163

Routines	Description	Page
dbgetoff	Checks for the existence of Transact-SQL constructs in the command buffer.	163
dbgetpacket	Returns the TDS packet size currently in use.	165
dbgetrow	Reads the specified row in the row buffer.	166
DBGETTIME	Returns the number of seconds that DB-Library will wait for a server response to a SQL command.	168
dbgetuserdata	Returns a pointer to user-allocated data from a DBPROCESS structure.	168
dbhasretstat	Determines whether the current command or remote procedure call generated a return status number.	169
dbinit	Initialize DB-Library.	171
DBIORDESC	(UNIX and AOS/VS only) Provides program access to the UNIX or AOS/VS file descriptor used by a DBPROCESS to read data coming from the server.	171
DBIOWDESC	(UNIX and AOS/VS only) Provides program access to the UNIX or AOS/VS file descriptor used by a DBPROCESS to write data to the server.	173
DBISAVAIL	Determines whether a DBPROCESS is available for general use.	174
dbisopt	Checks the status of a server or DB-Library option.	174
DBLASTROW	Returns the number of the last row in the row buffer.	175
dbload_xlate	Loads a pair of character set translation tables.	176
dbloadsort	Loads a server sort order.	177
dblogin	Allocates a login record for use in **dbopen**.	178
dbloginfree	Frees a login record.	180
dbmny4add	Adds two DBMONEY4 values.	180
dbmny4cmp	Compares two DBMONEY4 values.	181
dbmny4copy	Copies a DBMONEY4 value.	182
dbmny4divide	Divides one DBMONEY4 value by another.	183
dbmny4minus	Negates a DBMONEY4 value.	184
dbmny4mul	Multiplies two DBMONEY4 values.	185
dbmny4sub	Subtracts one DBMONEY4 value from another.	186
dbmny4zero	Initializes a DBMONEY4 variable to $0.0000.	187
dbmnyadd	Adds two DBMONEY values.	188
dbmnycmp	Compares two DBMONEY values.	189
dbmnycopy	Copies a DBMONEY value.	190
dbmnydec	Decrements a DBMONEY value by one ten-thousandth of a dollar.	191

Routines	Description	Page
dbmnydivide	Divides one DBMONEY value by another.	192
dbmnydown	Divides a DBMONEY value by a positive integer.	193
dbmnyinc	Increments a DBMONEY value by one ten-thousandth of a dollar.	195
dbmnyinit	Prepares a DBMONEY value for calls to **dbmnyndigit**.	195
dbmnymaxneg	Returns the maximum negative DBMONEY value supported.	197
dbmnymaxpos	Returns the maximum positive DBMONEY value supported.	198
dbmnyminus	Negates a DBMONEY value.	199
dbmnymul	Multiplies two DBMONEY values.	200
dbmnyndigit	Returns the rightmost digit of a DBMONEY value as a DBCHAR.	201
dbmnyscale	Multiplies a DBMONEY value by a positive integer and add a specified amount.	207
dbmnysub	Subtracts one DBMONEY value from another.	209
dbmnyzero	Initializes a DBMONEY value to $0.0000.	210
dbmonthname	Determines the name of a specified month in a specified language.	210
DBMORECMDS	Indicates whether there are more commands to be processed.	212
dbmoretext	Sends part of a text or image value to the server.	212
dbmsghandle	Installs a user function to handle server messages.	213
dbname	Returns the name of the current database.	217
dbnextrow	Reads the next result row into the row buffer and into any program variables that are bound to column data.	218
dbnpcreate	Creates a notification procedure.	220
dbnpdefine	Defines a notification procedure.	223
dbnullbind	Associates an indicator variable with a regular result row column.	225
dbnumalts	Returns the number of columns in a compute row.	226
dbnumcols	Determines the number of regular columns for the current set of results.	226
dbnumcompute	Returns the number of **compute** clauses in the current set of results.	228
DBNUMORDERS	Returns the number of columns specified in a Transact-SQL **select** statement's **order by** clause.	228
dbnumrets	Determines the number of return parameter values generated by a stored procedure.	229

Routines	Description	Page
dbopen	Creates and initialize a DBPROCESS structure.	231
dbordercol	Returns the id of a column appearing in the most recently executed query's **order by** clause.	235
dbpoll	Checks if a server response has arrived for a DBPROCESS.	235
dbprhead	Prints the column headings for rows returned from the server.	241
dbprrow	Prints all the rows returned from the server.	241
dbprtype	Converts a token value to a readable string.	242
dbqual	Returns a pointer to a **where** clause suitable for use in updating the current row in a browsable table.	244
DBRBUF	(UNIX and AOS/VS only) Determines whether the DB-Library network buffer contains any unread bytes.	247
dbreadpage	Reads a page of binary data from the server.	248
dbreadtext	Reads part of a text or image value from the server.	249
dbrecftos	Records all SQL commands sent from the application to the server.	251
dbrecvpassthru	Receive a TDS packet from a server.	252
dbregdrop	Drop a registered procedure.	254
dbregexec	Execute a registered procedure.	256
dbreghandle	Install a handler routine for a registered procedure notification.	258
dbreginit	Initiate execution of a registered procedure.	262
dbreglist	Return a list of registered procedures currently defined in Open Server.	264
dbregnowatch	Cancel a request to be notified when a registered procedure executes.	265
dbregparam	Define or describe a registered procedure parameter.	266
dbregwatch	Request to be notified when a registered procedure executes.	270
dbregwatchlist	Return a list of registered procedures that a DBPROCESS is watching for.	275
dbresults	Set up the results of the next query.	276
dbretdata	Return a pointer to a return parameter value generated by a stored procedure.	279
dbretlen	Determine the length of a return parameter value generated by a stored procedure.	283
dbretname	Determine the name of the stored procedure parameter associated with a particular return parameter value.	285

Routines	Description	Page
dbretstatus	Determine the stored procedure status number returned by the current command or remote procedure call.	286
dbrettype	Determine the datatype of a return parameter value generated by a stored procedure.	288
DBROWS	Indicate whether the current command actually returned rows.	290
DBROWTYPE	Return the type of the current row.	291
dbrpcinit	Initialize a remote procedure call.	291
dbrpcparam	Add a parameter to a remote procedure call.	293
dbrpcsend	Signal the end of a remote procedure call.	295
dbrpwclr	Clear all remote passwords from the LOGINREC structure.	296
dbrpwset	Add a remote password to the LOGINREC structure.	297
dbsafestr	Double the quotes in a character string.	298
dbsechandle	Install user functions to handle secure logins.	300
dbsendpassthru	Send a TDS packet to a server.	304
dbservcharset	Get the name of the server character set.	306
dbsetavail	Mark a DBPROCESS as being available for general use.	307
dbsetbusy	Call a user-supplied function when DB-Library is reading from the server.	307
dbsetdefcharset	Set the default character set for an application.	310
dbsetdeflang	Set the default language name for an application.	310
dbsetidle	Call a user-supplied function when DB-Library is finished reading from the server.	311
dbsetifile	Specify the name and location of the Sybase interfaces file.	312
dbsetinterrupt	Call user-supplied functions to handle interrupts while waiting on a read from the server.	313
DBSETLAPP	Set the application name in the LOGINREC structure.	317
DBSETLCHARSET	Set the character set in the LOGINREC structure.	317
DBSETLENCRYPT	Specify whether or not network password encryption is to be used when logging into SQL Server version 10.0 or later.	318
DBSETLHOST	Set the host name in the LOGINREC structure.	319
DBSETLNATLANG	Set the national language name in the LOGINREC structure.	320
dbsetloginfo	Transfer TDS login information from a DBLOGINFO structure to a LOGINREC structure.	321
dbsetlogintime	Set the number of seconds that DB-Library waits for a server response to a request for a DBPROCESS connection.	323
DBSETLPACKET	Set the TDS packet size in an application's LOGINREC structure.	324

Routines	Description	Page
DBSETLPWD	Set the user server password in the LOGINREC structure.	325
DBSETLUSER	Set the user name in the LOGINREC structure.	326
dbsetmaxprocs	Set the maximum number of simultaneously open DBPROCESS structures.	327
dbsetnull	Define substitution values to be used when binding null values.	327
dbsetopt	Set a server or DB-Library option.	329
dbsetrow	Set a buffered row to “current.”	331
dbsettime	Set the number of seconds that DB-Library will wait for a server response to a SQL command.	333
dbsetuserdata	Use a DBPROCESS structure to save a pointer to user-allocated data.	333
dbsetversion	Specify a DB-Library version level.	336
dbspid	Get the server process id for the specified DBPROCESS.	337
dbspr1row	Place one row of server query results into a buffer.	338
dbspr1rowlen	Determine how large a buffer to allocate to hold the results returned by **dbsprhead, dbsprline,** and **dbspr1row**.	340
dbsprhead	Place the server query results header into a buffer.	341
dbsprline	Get a formatted string that contains underlining for the column names produced by **dbsprhead**.	343
dbsqlexec	Send a command batch to the server.	344
dbsqlok	Wait for results from the server and verify the correctness of the instructions the server is responding to.	346
dbsqlsend	Send a command batch to the server and do not wait for a response.	351
dbstrbuild	Build a printable string from text containing placeholders for variables.	352
dbstrcmp	Compare two character strings using a specified sort order.	354
dbstrcpy	Copy all or a portion of the command buffer.	356
dbstrlen	Return the length, in characters, of the command buffer.	358
dbstrsort	Determine which of two character strings should appear first in a sorted list.	359
dbtabbrowse	Determine whether the specified table is updatable via the DB-Library browse-mode facilities.	360
dbtabcount	Return the number of tables involved in the current **select** query.	361
dbtabname	Return the name of a table based on its number.	362
dbtabsource	Return the name and number of the table from which a particular result column was derived.	363

Routines	Description	Page
DBTDS	Determine which version of TDS (the Tabular Data Stream protocol) is being used.	365
dbtextsize	Returns the number of bytes of text or image data that remain to be read for the current row.	365
dbtsnewlen	Return the length of the new value of the *timestamp* column after a browse-mode update.	366
dbtsnewval	Return the new value of the *timestamp* column after a browse-mode update.	367
dbtsput	Put the new value of the *timestamp* column into the given table's current row in the DBPROCESS.	368
dbtxptr	Return the value of the text pointer for a column in the current row.	369
dbtxtimestamp	Return the value of the text timestamp for a column in the current row.	371
dbtxtsnewval	Return the new value of a text timestamp after a call to **dbwritetext**.	372
dbtxtsput	Put the new value of a text timestamp into the specified column of the current row in the DBPROCESS.	372
dbuse	Use a particular database.	373
dbvarylen	Determine whether the specified regular result column's data can vary in length.	374
dbversion	Determine which version of DB-Library is in use.	375
dbwillconvert	Determine whether a specific datatype conversion is available within DB-Library.	376
dbwritepage	Write a page of binary data to the server.	377
dbwritetext	Send a text or image value to the server.	378
dbxlate	Translate a character string from one character set to another.	383
Errors	The complete collection of DB-Library errors and error severities.	386
Options	The complete list of DB-Library options.	402
Types	Datatypes and symbolic constants for datatypes used by DB-Library.	408

db12hour

Description Determine whether the specified language uses 12-hour or 24-hour time.

Syntax

DBBOOL db12hour(dbproc, language)

```
DBPROCESS    *dbproc;
char         *language;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

language

The name of the language of interest.

Return value

TRUE if *language* uses 12-hour time, FALSE otherwise.

Usage

- **db12hour** returns TRUE if *language* uses 12-hour time, and FALSE if it uses 24-hour time.
- If *language* is NULL, *dbproc*'s current language is signified. If both *language* and *dbproc* are NULL, then DB-Library's default language (for any future calls to **dbopen**) is signified.
- **db12hour** is useful when retrieving and manipulating DBDATETIME values via **dbsqlexec**. When converting DBDATETIME values to character strings, **dbconvert** and **dbbind** always return the month component of the DBDATETIME value in the local language, but use the U.S. English date and time order (month-day-year, 12-hour time). **db12hour**'s return value informs the application that some further manipulation is necessary if 24-hour rather than 12-hour time is desired.
- The following code fragment illustrates the use of **db12hour**:

```
DBBOOL time_format;
 DBCHAR s_date[40];

 /*
 ** Find out whether 12-hour or 24-hour time is
 ** used.
 */
 time_format = db12hour(dbproc, "FRANCAIS");

 /* Put a command into a command buffer */
dbcmd(dbproc, "select start_date from info_table");

 /* Send the command to the SQL Server */
 dbsqlexec(dbproc);

 /* Process the command results */
```

```
dbresults(dbproc);

/*
** Bind column data (start_date) to the program
** variable (s_date)
*/
dbbind(dbproc, 1, NTBSTRINGBIND, 0, s_date);

while (dbnextrow(dbproc) != NO_MORE_ROWS)
{
      /*
      ** If we want 24-hour time, re-format
      ** s_date accordingly.
      */
      if (time_format == TRUE)
           format_24(s_date);

      printf("Next start date: %s\n", s_date);
}
```

See also dbdateorder, dbdayname, dbmonthname, dbsetopt

dbadata

Description Return a pointer to the data for a compute column.

Syntax BYTE *dbadata(dbproc, computeid, column)

```
DBPROCESS      *dbproc;
int             computeid;
int             column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

column

The number of the column of interest. The first column returned is number 1. Note that the order in which compute columns are returned is determined by the order of the corresponding columns in the select list, not by the order in which the compute columns were originally specified. For example, in the following query the result of "sum(price)" is referenced by giving *column* a value of 1, not 2:

```
select price, advance from titles
    compute sum(advance), sum(price)
```

The relative order of compute columns in the select list, rather than their absolute position, determines the value of *column*. For instance, given the following variation of the previous **select**:

```
select title_id, price, advance from titles
   compute sum(advance), sum(price)
```

the *column* for "sum(price)" still has a value of 1 and not 2, because the "title_id" column in the select list is not a compute column and therefore is ignored when determining the compute column's number.

Return value

A BYTE pointer to the data for a particular column in a particular compute. Be sure to cast this pointer into the proper type. A BYTE pointer to NULL is returned if there is no such column or compute or if the data has a null value.

DB-Library allocates and frees the data space that the BYTE pointer points to. Do not overwrite this space.

Usage

- After each call to **dbnextrow**, you can use this routine to return a pointer to the data for a particular column in a compute row. The data is not null-terminated. You can use **dbadlen** to get the length of the data.
- When a column of integer data is summed or averaged, the server always returns a 4-byte integer, regardless of the size of the column. Therefore, be sure that the variable that is to contain the result from such a compute is declared as DBINT.
- Here is a short program fragment which illustrates the use of **dbadata**:

```
DBPROCESS     *dbproc;
 int           rowinfo;
 DBINT         sum;

 /*
 ** First, put the commands into the command
 ** buffer
 */
```

```
dbcmd(dbproc, "select db_name(dbid), dbid, size
     from sysusages");
dbcmd(dbproc, " order by dbid");
dbcmd(dbproc, " compute sum(size) by dbid");

/*
** Send the commands to SQL Server and start
** execution
*/
dbsqlexec(dbproc);

/* Process the command */
dbresults(dbproc);

/* Examine the results of the compute clause */
while((rowinfo = dbnextrow(dbproc)) !=
    NO_MORE_ROWS)
{
    if (rowinfo == REG_ROW)
       printf("regular row returned.\n");
    else
    {
       /*
       ** This row is the result of a compute
       ** clause, and "rowinfo" is the computeid
       ** of this compute clause.
       */

       sum = *(DBINT *)(dbadata(dbproc, rowinfo,
           1));
       printf("sum = %ld\n", sum);
    }
 }
```

- The function **dbaltbind** automatically binds compute data to your program variables. It does a copy of the data, but is often easier to use than **dbadata**. Furthermore, it includes a convenient type conversion capability. By means of this capability, the application can, among other things, easily add a null terminator to a result string or convert money and datetime data to printable strings.

See also dbadlen, dbaltbind, dbaltlen, dbalttype, dbgetrow, dbnextrow, dbnumalts

dbadlen

Description

Return the actual length of the data for a compute column.

Syntax

DBINT dbadlen(dbproc, computeid, column)

```
DBPROCESS    *dbproc;
int          computeid;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

column

The number of the column of interest. The first column is number 1.

Return value

The length, in bytes, of the data for a particular compute column. If there is no such column or **compute** clause, **dbadlen** returns -1. If the data has a null value, **dbadlen** returns 0.

Usage

- This routine returns the actual length of the data for a particular compute column.
- Use the **dbaltlen** routine to determine the maximum possible length for the data. Use **dbadata** to get a pointer to the data.
- Here is a program fragment that illustrates the use of **dbadlen**:

```
DBPROCESS     *dbproc;
 char          biggest_name[MAXNAME+1];
 int           namelen;
 int           rowinfo;

 /* put the command into the command buffer */
 dbcmd(dbproc, "select name from sysobjects");
 dbcmd(dbproc, " order by name");
 dbcmd(dbproc, " compute max(name)");

 /*
```

```
** Send the command to SQL Server and start
** execution.
*/
dbsqlexec(dbproc);

/* process the command */
dbresults(dbproc);

/* examine each row returned by the command */
while ((rowinfo = dbnextrow(dbproc)) !=
     NO_MORE_ROWS)
{
     if (rowinfo == REG_ROW)
        printf("regular row returned.\n");
     else
     {
        /*
        ** This row is the result of a compute
        ** clause, and "rowinfo" is the computeid
        ** of this compute clause.
        */
         namelen = dbadlen(dbproc, rowinfo, 1);
        strncpy(biggest_name,
          (char *)dbadata(dbproc, rowinfo, 1),
          namelen);

        /*
        ** Data pointed to by dbadata() is not
        ** null-terminated.
        */
        biggest_name[namelen] = '\0';

        printf("biggest name = %s\n",
biggest_name);
     }
}
```

See also dbadata, dbaltlen, dbalttype, dbgetrow, dbnextrow, dbnumalts

dbaltbind

Description Bind a compute column to a program variable.

Syntax

```
RETCODE dbaltbind(dbproc, computeid, column, vartype,
        varlen, varaddr)

DBPROCESS       *dbproc;
int              computeid;
int              column;
int              vartype;
DBINT            varlen;
BYTE            *varaddr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1.

column

The column number of the row data that is to be copied to a program variable. The first column is column number 1. Note that the order in which compute columns are returned is determined by the order of the corresponding columns in the select list, not by the order in which the compute columns were originally specified. For example, in the following query the result of "sum(price)" is referenced by giving *column* a value of 1, not 2:

```
select price, advance from titles
     compute sum(advance), sum(price)
```

The relative order of compute columns in the select list, rather than their absolute position, determines the value of *column*. For instance, given the following variation of the previous **select**:

```
select title_id, price, advance from titles
    compute sum(advance), sum(price)
```

the *column* for "sum(price)" still has a value of 1 and not 2, because the "title_id" column in the select list is not a compute column and therefore is ignored when determining the compute column's number.

vartype

This describes the datatype of the binding. It must correspond to the datatype of the program variable that will receive the copy of the data from the DBPROCESS. The table below shows the correspondence between *vartype* values and program variable types.

dbaltbind supports a wide range of type conversions, so the *vartype* can be different from the type returned by the SQL query. For instance, a SYBMONEY result may be bound to a DBFLT8 program variable via FLT8BIND, and the appropriate data conversion will happen automatically. For a list of the data conversions provided by DB-Library, see the reference page for dbwillconvert.

For a list of the type definitions used by DB-Library, see “Types” on page 408.

Here is a list of the legal *vartype* values recognized by **dbaltbind**, along with the server and program variable types that each one refers to:

Table 2-1: Bind types (dbaltbind)

Vartype	Program variable type	Server datatype
CHARBIND	DBCHAR	SYBCHAR
STRINGBIND	DBCHAR	SYBCHAR
NTBSTRINGBIND	DBCHAR	SYBCHAR
VARYCHARBIND	DBVARYCHAR	SYBCHAR
BINARYBIND	DBBINARY	SYBBINARY
VARYBINBIND	DBVARYBIN	SYBBINARY
TINYBIND	DBTINYINT	SYBINT1
SMALLBIND	DBSMALLINT	SYBINT2
INTBIND	DBINT	SYBINT4
FLT8BIND	DBFLT8	SYBFLT8
REALBIND	DBREAL	SYBREAL
NUMERICBIND	DBNUMERIC	SYBNUMERIC
DECIMALBIND	DBDECIMAL	SYBDECIMAL
BITBIND	DBBIT	SYBBIT
DATETIMEBIND	DBDATETIME	SYBDATETIME
SMALLDATETIMEBIND	DBDATETIME4	SYBDATETIME4
MONEYBIND	DBMONEY	SYBMONEY
SMALLMONEYBIND	DBMONEY4	SYBMONEY4
BOUNDARYBIND	DBCHAR	SYBBOUNDARY
SENSITIVITYBIND	DBCHAR	SYBSENSITIVITY

Warning! It is an error to use any of the following values for *vartype* if the library version has not been set (with **dbsetversion**) to DBVERSION_100 or higher: BOUNDARYBIND, DECIMALBIND, NUMERICBIND, or SENSITIVITYBIND.

Note **dbaltbind** does not offer explicit precision and scale support for numeric and decimal datatypes. When handling numeric or decimal data, **dbaltbind** uses a default precision and scale of 18 and 0, respectively, unless the bind is to a numeric or decimal column, in which case **dbaltbind** uses the precision and scale of the source data. Use dbaltbind_ps to explicitly specify precision and scale values—calling **dbaltbind** is equivalent to calling dbaltbind_ps with a NULL *typeinfo* value.

Since SYBTEXT and SYBIMAGE data are never returned through a

compute row, those datatypes are not listed above.

Note that the server type in the table above is listed merely for your information. The *vartype* you specify does not necessarily have to correspond to a particular server type, because, as mentioned earlier, **dbaltbind** will convert server data into the specified *vartype*.

The available representations for character data are shown below. They differ according to whether the data is blank-padded or null-terminated:

Vartype	Program type	Padding	Terminator
CHARBIND	DBCHAR	blanks	none
STRINGBIND	DBCHAR	blanks	\0
NTBSTRINGBIND	DBCHAR	none	\0
VARYCHARBIND	DBVARYCHAR	none	none
BOUNDARYBIND	DBCHAR	none	\0
SENSITIVITYBIND	DBCHAR	none	\0

Note that the "\0" in the table above is the null terminator character.

If overflow occurs when converting integer or float data to a character binding type, the first character of the resulting value will contain an asterisk ("*") to indicate the error.

Binary data may be stored in two different ways:

Vartype	Program Type	Padding
BINARYBIND	DBBINARY	nulls
VARYBINBIND	DBVARBINARY	none

When a column of integer data is summed or averaged, the server always returns a 4-byte integer, regardless of the size of the column. Therefore, be sure that the variable which is to contain the result from such a compute is declared as DBINT and that the *vartype* of the binding is INTBIND.

varlen
The length of the program variable in bytes.

For *vartype* values that represent fixed-length types, such as MONEYBIND or FLT8BIND, this length is ignored.

For character and binary types, *varlen* must describe the total length of the available destination buffer space, including any space that may be required for special terminating bytes, such as a null terminator. If *varlen* is 0, the total number of bytes available will be copied into the program variable. (For char and binary server data, the total number of bytes available is equal to the defined length of the database column, including any blank padding. For varchar and varbinary data, the total number of bytes available is equal to the actual data contained in the column.) Therefore, if you are sure that your program variable is large enough to handle the results, you can just set *varlen* to 0.

varaddr
The address of the program variable to which the data will be copied.

Return value

SUCCEED or FAIL.

dbaltbind returns FAIL if the column number is not valid, if the data conversion specified by *vartype* is not legal, or if *varaddr* is NULL.

Usage

- This routine directs DB-Library to copy compute column data returned by the server into a program variable. (A compute column results from the **compute** clause of a Transact-SQL **select** statement.) When each new row containing compute data is read via **dbnextrow** or **dbgetrow**, the data from the designated *column* in that compute row is copied into the program variable with the address *varaddr*. There must be a separate **dbaltbind** call for each compute column that is to be copied. It is not necessary to bind every compute column to a program variable.
- The server can return two types of rows: regular rows containing data from columns designated by a **select** statement's select list, and compute rows resulting from the **compute** clause. **dbaltbind** binds data from compute rows. Use **dbbind** for binding data from regular rows.
- You must make the calls to **dbaltbind** after a call to **dbresults** and before the first call to **dbnextrow**.
- The typical sequence of calls is:

```
DBCHAR     name[20];
 DBINT      namecount;

 /* read the query into the command buffer */
```

```
dbcmd(dbproc, "select name from emp compute
     count(name)");

/* send the query to SQL Server */
dbsqlexec(dbproc);

/* get ready to process the query results */
dbresults(dbproc);

/* bind the regular row data (name) */
dbbind(dbproc, 1, STRINGBIND, (DBINT) 0, name);

/* bind the compute column data (count of name) */
dbaltbind(dbproc, 1, 1, INTBIND, (DBINT) 0,
     (BYTE *) &namecount);

/* now process each row */
while (dbnextrow(dbproc) != NO_MORE_ROWS)
{
     C-code to print or process row data
}
```

- **dbaltbind** incurs a little overhead because it causes the data to be copied into a program variable. To avoid this copying, you can use the **dbadata** routine to directly access the returned data.
- You can only bind a result column to a single program variable. If you bind a result column to multiple variables, only the last binding takes effect.
- The server can return null column values, and DB-Library provides the following aids for handling null values:
  - A pre-defined set of default values, one for each datatype, that DB-Library automatically substitutes when a bound column contains a null value. The **dbsetnull** function allows you to explicitly set your own null substitution values. See the reference page for the dbsetnull function for a list of the default substitution values.
  - The ability to bind an indicator variable to a column with **dbnullbind** (or **dbanullbind** for compute rows). As rows are fetched, the value of the indicator variable will be set to indicate whether or not the column value was null. See the reference page for the dbnullbind function for indicator values and meanings.

See also

dbadata, dbaltbind_ps, dbanullbind, dbbind, dbbind_ps, dbconvert, dbconvert_ps, dbnullbind, dbsetnull, dbsetversion, dbwillconvert, Types

dbaltbind_ps

Description

Bind a compute column to a program variable, with precision and scale support for numeric and decimal datatypes.

Syntax

```
RETCODE dbaltbind_ps(dbproc, computeid, column,
            vartype, varlen, varaddr,
            typeinfo)

DBPROCESS    *dbproc;
int          computeid;
int          column;
int          vartype;
DBINT        varlen;
BYTE         *varaddr;
DBTYPEINFO   *typeinfo;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1.

column

The column number of the row data that is to be copied to a program variable. The first column is column number 1. Note that the order in which compute columns are returned is determined by the order of the corresponding columns in the select list, not by the order in which the compute columns were originally specified. For example, in the following query the result of "sum(price)" is referenced by giving *column* a value of 1, not 2:

```
select price, advance from titles
     compute sum(advance), sum(price)
```

The relative order of compute columns in the select list, rather than their absolute position, determines the value of *column*. For instance, given the following variation of the previous **select**:

```
select title_id, price, advance from titles
      compute sum(advance), sum(price)
```

the *column* for “sum(price)” still has a value of 1 and not 2, because the “title_id” column in the select list is not a compute column and therefore is ignored when determining the compute column’s number.

vartype

This describes the datatype of the binding. It must correspond to the datatype of the program variable that will receive the copy of the data from the DBPROCESS. The table below shows the correspondence between *vartype* values and program variable types.

dbaltbind_ps supports a wide range of type conversions, so the *vartype* can be different from the type returned by the SQL query. For instance, a SYBMONEY result may be bound to a DBFLT8 program variable via FLT8BIND, and the appropriate data conversion will happen automatically. For a list of the data conversions provided by DB-Library, see the reference page for dbwillconvert.

For a list of the type definitions used by DB-Library, see “Types” on page 408.

Here is a list of the legal *vartype* values recognized by **dbaltbind_ps**, along with the server and program variable types that each one refers to:

Table 2-2: Bind types (dbaltbind_ps)

Vartype	Program variable type	Server datatype
CHARBIND	DBCHAR	SYBCHAR
STRINGBIND	DBCHAR	SYBCHAR
NTBSTRINGBIND	DBCHAR	SYBCHAR
VARYCHARBIND	DBVARYCHAR	SYBCHAR
BINARYBIND	DBBINARY	SYBBINARY
VARYBINBIND	DBVARYBIN	SYBBINARY
TINYBIND	DBTINYINT	SYBINT1
SMALLBIND	DBSMALLINT	SYBINT2
INTBIND	DBINT	SYBINT4
FLT8BIND	DBFLT8	SYBFLT8
REALBIND	DBREAL	SYBREAL
NUMERICBIND	DBNUMERIC	SYBNUMERIC
DECIMALBIND	DBDECIMAL	SYBDECIMAL
BITBIND	DBBIT	SYBBIT
DATETIMEBIND	DBDATETIME	SYBDATETIME
SMALLDATETIMEBIND	DBDATETIME4	SYBDATETIME4
MONEYBIND	DBMONEY	SYBMONEY
SMALLMONEYBIND	DBMONEY4	SYBMONEY4
BOUNDARYBIND	DBCHAR	SYBBOUNDARY
SENSITIVITYBIND	DBCHAR	SYBSENSITIVITY

Warning! It is an error to use any of the following values for *vartype* if the library version has not been set (with **dbsetversion**) to DBVERSION_!00 or higher: BOUNDARYBIND, DECIMALBIND, NUMERICBIND, or SENSITIVITYBIND.

Note **dbaltbind_ps**'s parameters are identical to **dbaltbind**'s, except that **dbaltbind_ps** has the additional parameter *typeinfo*, which contains information about precision and scale for DBNUMERIC or DBDECIMAL variables.

Since SYBTEXT and SYBIMAGE data are never returned through a compute row, those datatypes are not listed above.

Note that the server type in the table above is listed merely for your

information. The *vartype* you specify does not necessarily have to correspond to a particular server type, because, as mentioned earlier, **dbaltbind_ps** will convert server data into the specified *vartype*.

The available representations for character data are shown below. They differ according to whether the data is blank-padded or null-terminated:

Vartype	Program type	Padding	Terminator
CHARBIND	DBCHAR	blanks	none
STRINGBIND	DBCHAR	blanks	\0
NTBSTRINGBIND	DBCHAR	none	\0
VARYCHARBIND	DBVARYCHAR	none	none
BOUNDARYBIND	DBCHAR	none	\0
SENSITIVITYBIND	DBCHAR	none	\0

Note that the "\0" in the table above is the null terminator character.

If overflow occurs when converting integer or float data to a character binding type, the first character of the resulting value will contain an asterisk ("*") to indicate the error.

Binary data may be stored in two different ways:

Vartype	Program type	Padding
BINARYBIND	DBBINARY	nulls
VARYBINBIND	DBVARBINARY	none

When a column of integer data is summed or averaged, the server always returns a 4-byte integer, regardless of the size of the column. Therefore, be sure that the variable which is to contain the result from such a compute is declared as DBINT and that the *vartype* of the binding is INTBIND.

varlen
The length of the program variable in bytes.

For values of *vartype* that represent a fixed-length type, such as MONEYBIND or FLT8BIND, this length is ignored.

For character and binary types, *varlen* must describe the total length of the available destination buffer space, including any space that may be required for special terminating bytes, such as a null terminator. If *varlen* is 0, the total number of bytes available will be copied into the program variable. (For char and binary server data, the total number of bytes available is equal to the defined length of the database column, including any blank padding. For varchar and varbinary data, the total number of bytes available is equal to the actual data contained in the column.) Therefore, if you are sure that your program variable is large enough to handle the results, you can just set *varlen* to 0.

varaddr
The address of the program variable to which the data will be copied.

typeinfo
A pointer to a DBTYPEINFO structure containing information about the precision and scale of decimal or numeric data. An application sets a DBTYPEINFO structure with values for precision and scale before calling **dbaltbind_ps** to bind columns to DBDECIMAL or DBNUMERIC variables.

If *typeinfo* is NULL:

- If the result column is of type numeric or decimal, **dbaltbind_ps** picks up precision and scale values from the result column.
- If the result column is not numeric or decimal, **dbaltbind_ps** uses a default precision of 18 and a default scale of 0.

If *vartype* is not DECIMALBIND or NUMERICBIND, *typeinfo* is ignored.

A DBTYPEINFO structure is defined as follows:

```
typedef struct typeinfo {
      DBINT   precision;
      DBINT   scale;
 } DBTYPEINFO;
```

Legal values for *precision* are from 1 to 77. Legal values for *scale* are from 0 to 77. *scale* must be less than or equal to *precision*.

Return value

SUCCEED or FAIL.

dbaltbind_ps returns FAIL if the column number is not valid, if the data conversion specified by *vartype* is not legal, or if *varaddr* is NULL.

Usage

- **dbaltbind_ps** is the equivalent of **dbaltbind**, except that **dbaltbind_ps** provides precision and scale support for numeric and decimal datatypes, which **dbaltbind** does not. Calling **dbaltbind** is equivalent to calling **dbaltbind_ps** with *typeinfo* as NULL.
- **dbaltbind_ps** directs DB-Library to copy compute column data returned by the server into a program variable. (A compute column results from the **compute** clause of a Transact-SQL **select** statement.) When each new row containing compute data is read via **dbnextrow** or **dbgetrow**, the data from the designated *column* in that compute row is copied into the program variable with the address *varaddr*. There must be a separate **dbaltbind_ps** call for each compute column that is to be copied. It is not necessary to bind every compute column to a program variable.
- The server can return two types of rows: regular rows containing data from columns designated by a **select** statement's select list, and compute rows resulting from the **compute** clause. **dbaltbind_ps** binds data from compute rows. Use **dbbind_ps** for binding data from regular rows.
- You must make the calls to **dbaltbind_ps** after a call to **dbresults** and before the first call to **dbnextrow**.
- **dbaltbind_ps** incurs a little overhead because it causes the data to be copied into a program variable. To avoid this copying, you can use the **dbadata** routine to directly access the returned data.
- You can only bind a result column to a single program variable. If you bind a result column to multiple variables, only the last binding takes effect.
- Since the server can return null values, DB-Library provides a set of default values, one for each datatype, that it will automatically substitute when binding null values. The **dbsetnull** function allows you to explicitly set your own null substitution values. (See the reference page for the dbsetnull function for a list of the default substitution values.)

See also

dbaltbind, dbanullbind, dbbind, dbbind_ps, dbconvert, dbconvert_ps, dbdata, dbnullbind, dbsetnull, dbsetversion, dbwillconvert, Types

dbaltcolid

Description

Return the column id for a compute column.

Syntax

int dbaltcolid(dbproc, computeid, column)

```
DBPROCESS    *dbproc;
int          computeid;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

column

The number of the compute column of interest. The first column in a select list is 1.

Return value

The select list id for the compute column. The first column in a select list is 1. If either the *computeid* or the *column* value is invalid, **dbaltcolid** returns -1.

Usage

- This routine returns the select list id for a compute column. For example, given the SQL statement:

```
select dept, name from employee
      order by dept, name
      compute count(name) by dept
```

  the call **dbaltcolid**(*dbproc*, 1, 1) will return 2, since “name” is the second column in the select list.

See also

dbadata, dbadlen, dbaltlen, dbgetrow, dbnextrow, dbnumalts, dbprtype

dbaltlen

Description

Return the maximum length of the data for a particular compute column.

Syntax

DBINT dbaltlen(dbproc, computeid, column)

```
DBPROCESS      *dbproc;
int            computeid;
int            column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

column

The number of the column of interest. The first column is number 1.

Return value

The maximum length, in bytes, possible for the data in a particular compute column. **dbaltlen** returns -1 if there is no such column or **compute** clause.

Usage

- This routine returns the maximum length for a column in a compute row. In the case of variable length data, this is not necessarily the actual length of the data, but rather the maximum length. For the actual data length, use **dbadlen**.

  For example, given the SQL statement:

  ```
  select dept, name from employee
        order by dept, name
        compute count(name) by dept
  ```

  the call **dbaltlen**(*dbproc*, 1, 1) returns 4 because counts are of SYBINT4 type, which is 4 bytes long.

See also

dbadata, dbadlen, dbalttype, dbgetrow, dbnextrow, dbnumalts

dbaltop

Description

Return the type of aggregate operator for a particular compute column.

Syntax

int dbaltop(dbproc, computeid, column)

```
DBPROCESS    *dbproc;
int          computeid;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

column

The number of the column of interest. The first column is number 1.

Return value

A token value for the type of the compute column's aggregate operator. In case of error, **dbaltop** returns -1.

Usage

- This routine returns the type of aggregate operator for a particular column in a compute row. For example, given the SQL statement:

```
select dept, name from employee
     order by dept, name
     compute count(name) by dept
```

  the call **dbaltop**(*dbproc*, 1, 1) will return the token value for **count** since the first aggregate operator in the first **compute** clause is **count**.

- You can convert the token value to a readable token string with **dbprtype**. See the dbprtype reference page for a list of all token values and their equivalent token strings.

See also

dbadata, dbadlen, dbaltlen, dbnextrow, dbnumalts, dbprtype

dbalttype

Description

Return the datatype for a compute column.

Syntax

int dbalttype(dbproc, computeid, column)

```
DBPROCESS       *dbproc;
int             computeid;
int             column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

column

The number of the column of interest. The first column is number 1.

Return value

A token value for the datatype for a particular compute column.

In a few cases, the token value returned by this routine may not correspond exactly with the column's server datatype:

- SYBVARCHAR is returned as SYBCHAR.
- SYBVARBINARY is returned as SYBBINARY.
- SYBDATETIMN is returned as SYBDATETIME.
- SYBMONEYN is returned as SYBMONEY.
- SYBFLTN is returned as SYBFLT8.
- SYBINTN is returned as SYBINT1, SYBINT2, or SYBINT4, depending on the actual type of the SYBINTN.

dbalttype returns -1 if either the *computeid* or the *column* value is invalid.

Usage

- This routine returns the datatype for a compute column. For a list of server datatypes, see "Types" on page 408.

- **dbalttype** actually returns an integer token value for the datatype (SYBCHAR, SYBFLT8, and so on). To convert the token value into a readable token string, use **dbprtype**. See the dbprtype reference page for a list of all token values and their equivalent token strings.
- For example, given the SQL statement:

```
select dept, name from employee
     order by dept, name
     compute count(name) by dept
```

  the call **dbalttype**(*dbproc*, 1, 1) returns the token value SYBINT4, because counts are of SYBINT4 type. **dbprtype** will convert SYBINT4 into the readable token string "int".

See also

dbadata, dbadlen, dbaltlen, dbnextrow, dbnumalts, dbprtype, Types

dbaltutype

Description

Return the user-defined datatype for a compute column.

Syntax

DBINT dbaltutype(dbproc, computeid, column)

```
DBPROCESS   *dbproc;
int         computeid;
int         column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow**.

column

The number of the column of interest. The first column is number 1.

Return value

The user-defined datatype of the specified compute column on success; a negative integer on error.

Usage

- **dbaltutype** returns the user-defined datatype for a compute column.
- For a description of how to add user-defined datatypes to the server databases or Server-Library programs, see the *Adaptive Server Enterprise Reference Manual* or the *Open Server Server-Library Reference Manual*.
- **dbaltutype** is defined as type DBINT, since both the DB-Library datatype DBINT and user-defined datatypes are 32 bits long.

See also

dbalttype, dbcolutype

dbanullbind

Description

Associate an indicator variable with a compute-row column.

Syntax

```
RETCODE dbanullbind(dbproc, computeid, column,
        indicator)

DBPROCESS   *dbproc;
int         computeid;
int         column;
DBINT       *indicator;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The compute row of interest. A **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1.

column

The number of the column that is to be associated with the indicator variable.

indicator

A pointer to the indicator variable.

Note *indicator* is just the pointer to the indicator variable. It is the variable itself that is set.

Return value

SUCCEED or FAIL.

dbanullbind returns FAIL if either *computeid* or *column* is invalid.

Usage

- **dbanullbind** associates a compute-row column with an indicator variable. The indicator variable indicates whether a particular compute-row column has been converted and copied to a program variable successfully or unsuccessfully, or whether it is null.
- The indicator variable is set when compute rows are processed via **dbnextrow**. The possible values are:
  - *-1* if the column is NULL.
  - *The full length of the column's data, in bytes* if the column was bound to a program variable via **dbaltbind**, the binding did not specify any data conversions, and the bound data was truncated because the program variable was too small to hold the column's data.
  - *0* if the column was bound and copied to a program variable successfully.

Note Detection of character string truncation is implemented only for CHARBIND and VARYCHARBIND.

See also

dbadata, dbadlen, dbaltbind, dbnextrow, dbnullbind

dbbind

Description

Bind a regular result column to a program variable.

Syntax

RETCODE dbbind(dbproc, column, vartype, varlen, varaddr)

```
DBPROCESS    *dbproc;
int          column;
int          vartype;
DBINT        varlen;
BYTE         *varaddr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The column number of the row data that is to be copied to a program variable. The first column is column number 1.

vartype

This describes the datatype of the binding. It must correspond to the datatype of the program variable that will receive the copy of the data from the DBPROCESS. The table below shows the correspondence between *vartype* values and program variable types.

dbbind supports a wide range of type conversions, so the *vartype* can be different from the type returned by the SQL query. For instance, a SYBMONEY result may be bound to a DBFLT8 program variable via FLT8BIND, and the appropriate data conversion will happen automatically. For a list of the data conversions provided by DB-Library, see the reference page for dbwillconvert.

For a list of the type definitions used by DB-Library, see Types on page 408.

Here is a list of the legal *vartype* values recognized by **dbbind**, along with the server and program variable types that each one refers to:

Table 2-3: Bind types (dbbind)

Vartype	Program variable type	Server datatype
CHARBIND	DBCHAR	SYBCHAR or SYBTEXT
STRINGBIND	DBCHAR	SYBCHAR or SYBTEXT
NTBSTRINGBIND	DBCHAR	SYBCHAR or SYBTEXT
VARYCHARBIND	DBVARYCHAR	SYBCHAR or SYBTEXT
BINARYBIND	DBBINARY	SYBBINARY or SYBIMAGE
VARYBINBIND	DBVARYBIN	SYBBINARY or SYBIMAGE
TINYBIND	DBTINYINT	SYBINT1
SMALLBIND	DBSMALLINT	SYBINT2
INTBIND	DBINT	SYBINT4
FLT8BIND	DBFLT8	SYBFLT8
REALBIND	DBREAL	SYBREAL
NUMERICBIND	DBNUMERIC	SYBNUMERIC
DECIMALBIND	DBDECIMAL	SYBDECIMAL
BITBIND	DBBIT	SYBBIT
DATETIMEBIND	DBDATETIME	SYBDATETIME
SMALLDATETIMEBIND	DBDATETIME4	SYBDATETIME4
MONEYBIND	DBMONEY	SYBMONEY
SMALLMONEYBIND	DBMONEY4	SYBMONEY4
BOUNDARYBIND	DBCHAR	SYBBOUNDARY
SENSITIVITYBIND	DBCHAR	SYBSENSITIVITY

Warning! It is an error to use any of the following values for *vartype* if the library version has not been set (with **dbsetversion**) to DBVERSION_!00 or higher: BOUNDARYBIND, DECIMALBIND, NUMERICBIND, or SENSITIVITYBIND.

The server type in the table above is listed merely for your information. The *vartype* you specify does not necessarily have to correspond to a particular server type, because, as mentioned earlier, **dbbind** will convert server data into the specified *vartype*.

Note The server types nchar and nvarchar are converted internally to char and varchar types, which correspond to the DB-Library type constant SYBCHAR.

Note **dbbind** does not offer explicit precision and scale support for numeric and decimal datatypes. When handling numeric or decimal data, **dbbind** uses a default precision and scale of 18 and 0, respectively, unless the bind is to a numeric or decimal column, in which case **dbbind** uses the precision and scale of the source data. Use dbbind_ps to explicitly specify precision and scale values—calling **dbbind** is equivalent to calling dbbind_ps with a NULL *typeinfo* value.

The available representations for character and text data are shown below. They differ according to whether the data is blank-padded or null-terminated. Note that if *varlen* is 0, no padding takes place:

Vartype	Program type	Padding	Terminator
CHARBIND	DBCHAR	blanks	none
STRINGBIND	DBCHAR	blanks	\0
NTBSTRINGBIND	DBCHAR	none	\0
VARYCHARBIND	DBVARYCHAR	none	none
BOUNDARYBIND	DBCHAR	none	\0
SENSITIVITYBIND	DBCHAR	none	\0

Note that the "\0" in the table above is the null terminator character.

If overflow occurs when converting integer or float data to a character/text binding type, the first character of the resulting value will contain an asterisk ("*") to indicate the error.

binary and image data may be stored in two different ways:

Vartype	Program type	Padding
BINARYBIND	DBBINARY	nulls
VARYBINBIND	DBVARBINARY	none

varlen
The length of the program variable in bytes.

For values of *vartype* that represent a fixed-length type, such as MONEYBIND or FLT8BIND, this length is ignored.

For char, text, binary, and image types, *varlen* must describe the total length of the available destination buffer space, including any space that may be required for special terminating bytes, such as a null terminator. If *varlen* is 0, the total number of bytes available will be copied into the program variable. (For char and binary server data, the total number of bytes available is equal to the defined length of the database column, including any blank padding. For varchar, varbinary, text, and image data, the total number of bytes available is equal to the actual data contained in the column.) Therefore, if you are sure that your program variable is large enough to handle the results, you can just set *varlen* to 0.

Note that if *varlen* is 0, no padding takes place.

In some cases, DB-Library issues a message indicating that data conversion resulted in an overflow. This can be caused by a *varlen* specification that is too small for the server data.

varaddr
The address of the program variable to which the data will be copied.

Return value

SUCCEED or FAIL.

dbbind returns FAIL if the column number is not valid, if the data conversion specified by *vartype* is not legal, or if *varaddr* is NULL.

Usage

- Data comes back from the server one row at a time. This routine directs DB-Library to copy the data for a regular column (designated in a **select** statement's select list) into a program variable. When each new row containing regular (*not* compute) data is read via **dbnextrow** or **dbgetrow**, the data from the designated *column* in that row is copied into the program variable with the address *varaddr*. There must be a separate **dbbind** call for each regular column that is to be copied. It is not necessary to bind every column to a program variable.
- The server can return two types of rows: regular rows and compute rows resulting from the **compute** clause of a **select** statement. **dbbind** binds data from regular rows. Use **dbaltbind** for binding data from compute rows.
- You must make the calls to **dbbind** after a call to **dbresults** and before the first call to **dbnextrow**.

- The typical sequence of calls is:

```
DBINT     xvariable;
 DBCHAR    yvariable[10];

 /* read the query into the command buffer */
 dbcmd(dbproc, "select x = 100, y = 'hello'");

 /* send the query to SQL Server */
 dbsqlexec(dbproc);

 /* get ready to process the query results */
 dbresults(dbproc);

 /* bind column data to program variables */
 dbbind(dbproc, 1, INTBIND, (DBINT) 0,
     (BYTE *) &xvariable);
 dbbind(dbproc, 2, STRINGBIND, (DBINT) 0,
     yvariable);

 /* now process each row */
 while (dbnextrow(dbproc) != NO_MORE_ROWS)
 {
     C-code to print or process row data
 }
```

- **dbbind** incurs a little overhead, because it causes the data to be copied into a program variable. To avoid this copying, you can use the **dbdata** routine to directly access the returned data.
- You can only bind a result column to a single program variable. If you bind a result column to multiple variables, only the last binding takes effect.
- Since the server can return null values, DB-Library provides a set of default values, one for each datatype, that it will automatically substitute when binding null values. The **dbsetnull** function allows you to explicitly set your own null substitution values. (See the reference page for the dbsetnull function for a list of the default substitution values.)

See also

dbaltbind, dbaltbind_ps, dbnullbind, dbbind_ps, dbconvert, dbconvert_ps, dbdata, dbnullbind, dbsetnull, dbsetversion, dbwillconvert, Types

dbbind_ps

Description

Bind a regular result column to a program variable, with precision and scale support for numeric and decimal datatypes.

Syntax

```
RETCODE dbbind_ps(dbproc, column, vartype, varlen,
        varaddr, typeinfo)

DBPROCESS    *dbproc;
int          column;
int          vartype;
DBINT        varlen;
BYTE         *varaddr;
DBTYPEINFO   *typeinfo;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The column number of the row data that is to be copied to a program variable. The first column is column number 1.

vartype

This describes the datatype of the binding. It must correspond to the datatype of the program variable that will receive the copy of the data from the DBPROCESS. The table below shows the correspondence between *vartype* values and program variable types.

dbbind_ps supports a wide range of type conversions, so the *vartype* can be different from the type returned by the SQL query. For instance, a SYBMONEY result may be bound to a DBFLT8 program variable via FLT8BIND, and the appropriate data conversion will happen automatically. For a list of the data conversions provided by DB-Library, see the reference page for dbwillconvert.

For a list of the type definitions used by DB-Library, see Types on page 408.

Here is a list of the legal *vartype* values recognized by **dbbind_ps**, along with the server and program variable types that each one refers to:

Table 2-4: Bind types (dbbind_ps)

Vartype	Program variable type	Server type
CHARBIND	DBCHAR	SYBCHAR or SYBTEXT
STRINGBIND	DBCHAR	SYBCHAR or SYBTEXT
NTBSTRINGBIND	DBCHAR	SYBCHAR or SYBTEXT
VARYCHARBIND	DBVARYCHAR	SYBCHAR or SYBTEXT
BINARYBIND	DBBINARY	SYBBINARY or SYBIMAGE
VARYBINBIND	DBVARYBIN	SYBBINARY or SYBIMAGE
TINYBIND	DBTINYINT	SYBINT1
SMALLBIND	DBSMALLINT	SYBINT2
INTBIND	DBINT	SYBINT4
FLT8BIND	DBFLT8	SYBFLT8
REALBIND	DBREAL	SYBREAL
NUMERICBIND	DBNUMERIC	SYBNUMERIC
DECIMALBIND	DBDECIMAL	SYBDECIMAL
BITBIND	DBBIT	SYBBIT
DATETIMEBIND	DBDATETIME	SYBDATETIME
SMALLDATETIMEBIND	DBDATETIME4	SYBDATETIME4
MONEYBIND	DBMONEY	SYBMONEY
SMALLMONEYBIND	DBMONEY4	SYBMONEY4
BOUNDARYBIND	DBCHAR	SYBBOUNDARY
SENSITIVITYBIND	DBCHAR	SYBSENSITIVITY

Warning! It is an error to use any of the following values for *vartype* if the library version has not been set (with **dbsetversion**) to DBVERSION_!00 or higher: BOUNDARYBIND, DECIMALBIND, NUMERICBIND, or SENSITIVITYBIND.

The server type in the table above is listed merely for your information. The *vartype* you specify does not necessarily have to correspond to a particular server type, because, as mentioned earlier, **dbbind_ps** will convert server data into the specified *vartype*.

Note The server types nchar and nvarchar are converted internally to char and varchar types, which correspond to the DB-Library type constant SYBCHAR.

The available representations for character and text data are shown below.

They differ according to whether the data is blank-padded or null-terminated. Note that if *varlen* is 0, no padding takes place:

Vartype	Program type	Padding	Terminator
CHARBIND	DBCHAR	blanks	none
STRINGBIND	DBCHAR	blanks	\0
NTBSTRINGBIND	DBCHAR	none	\0
VARYCHARBIND	DBVARYCHAR	none	none
BOUNDARYBIND	DBCHAR	none	\0
SENSITIVITYBIND	DBCHAR	none	\0

Note that the "\0" in the table above is the null terminator character.

If overflow occurs when converting integer or float data to a character/text binding type, the first character of the resulting value will contain an asterisk ("*") to indicate the error.

binary and image data may be stored in two different ways:

Vartype	Program variable type	Padding
BINARYBIND	DBBINARY	nulls
VARYBINBIND	DBVARBINARY	none

varlen

The length of the program variable in bytes.

For values of *vartype* that represent a fixed-length type, such as MONEYBIND or FLT8BIND, this length is ignored.

For char, text, binary, and image types, *varlen* must describe the total length of the available destination buffer space, including any space that may be required for special terminating bytes, such as a null terminator. If *varlen* is 0, the total number of bytes available will be copied into the program variable. (For char and binary server data, the total number of bytes available is equal to the defined length of the database column, including any blank padding. For varchar, varbinary, text, and image data, the total number of bytes available is equal to the actual data contained in the column.) Therefore, if you are sure that your program variable is large enough to handle the results, you can just set *varlen* to 0.

Note If *varlen* is 0, no padding takes place.

---

varaddr

The address of the program variable to which the data will be copied.

typeinfo

A pointer to a DBTYPEINFO structure containing information about the precision and scale of decimal or numeric data. An application sets a DBTYPEINFO structure with values for precision and scale before calling **dbbind_ps** to bind columns to DBDECIMAL or DBNUMERIC variables.

If *typeinfo* is NULL:

- If the result column is of type numeric or decimal, **dbbind_ps** picks up precision and scale values from the result column.
- If the result column is not numeric or decimal, **dbbind_ps** uses a default precision of 18 and a default scale of 0.

If *vartype* is not DECIMALBIND or NUMERICBIND, *typeinfo* is ignored.

A DBTYPEINFO structure is defined as follows:

```
typedef struct typeinfo {
      DBINTprecision;
      DBINTscale;
 } DBTYPEINFO;
```

Legal values for *precision* are from 1 to 77. Legal values for *scale* are from 0 to 77. *scale* must be less than or equal to *precision*.

Return value

SUCCEED or FAIL.

dbbind_ps returns FAIL if the column number is not valid, if the data conversion specified by *vartype* is not legal, or if *varaddr* is NULL.

Usage

Note **dbbind_ps**'s parameters are identical to **dbbind**'s, except that **dbbind_ps** has the additional parameter *typeinfo*, which contains information about precision and scale for DBNUMERIC or DBDECIMAL variables.

- **dbbind_ps** is the equivalent of **dbbind**, except that **dbbind_ps** provides scale and precision support for numeric and decimal datatypes, which **dbbind** does not. Calling **dbbind** is equivalent to calling **dbbind_ps** with *typeinfo* as NULL.
- Data comes back from the server one row at a time. This routine directs DB-Library to copy the data for a regular column (designated in a **select** statement's select list) into a program variable. When each new row containing regular (*not* compute) data is read via **dbnextrow** or **dbgetrow**, the data from the designated *column* in that row is copied into the program variable with the address *varaddr*. There must be a separate **dbbind** or **dbbind_ps** call for each regular column that is to be copied. It is not necessary to bind every column to a program variable.

- The server can return two types of rows: regular rows and compute rows resulting from the **compute** clause of a **select** statement. **dbbind_ps** binds data from regular rows. Use **dbaltbind_ps** for binding data from compute rows.
- You must make the calls to **dbbind_ps** after a call to **dbresults** and before the first call to **dbnextrow**.
- **dbbind_ps** incurs a little overhead, because it causes the data to be copied into a program variable. To avoid this copying, you can use the **dbdata** routine to directly access the returned data.
- You can only bind a result column to a single program variable. If you bind a result column to multiple variables, only the last binding takes effect.
- Since the server can return null values, DB-Library provides a set of default values, one for each datatype, that it will automatically substitute when binding null values. The **dbsetnull** function allows you to explicitly set your own null substitution values. See the reference page for the dbsetnull function for a list of the default substitution values.

See also

dbaltbind, dbaltbind_ps, dbanullbind, dbbind, dbconvert, dbconvert_ps, dbdata, dbnullbind, dbsetnull, dbsetversion, dbwillconvert, Types

dbbufsize

Description

Return the size of a DBPROCESS row buffer.

Syntax

```
int dbbufsize(dbproc)

 DBPROCESS      *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

Return value

An integer representing the size, in rows, of the DBPROCESS row buffer.

If *dbproc* is NULL or if row buffering is not enabled, **dbbufsize** returns 0.

Usage

- **dbbufsize** returns the size of a DBPROCESS row buffer.

- Row buffering provides a way for an application to keep a specified number of server result rows in program memory. To enable row buffering, call **dbsetopt**(*dbproc*, DBBUFFER, *n*), where *n* is the number of rows to buffer. An application that is buffering result rows can access rows non-sequentially, via **dbgetrow**. See the dbgetrow reference page for a discussion of the benefits and penalties of row buffering.

See also

dbclrbuf, dbgetrow, dbsetopt, Options

dbbylist

Description

Return the bylist for a compute row.

Syntax

BYTE *dbbylist(dbproc, computeid, size)

```
DBPROCESS   *dbproc;
int         computeid;
int         *size;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid

The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

size

A pointer to an integer, which **dbbylist** sets to the number of elements in the bylist.

Return value

A pointer to an array of bytes containing the numbers of the columns that compose the bylist for the specified compute. The array of BYTEs is part of the DBPROCESS, so you must not free it. If the *computeid* is out of range, NULL is returned.

Call **dbcolname** to derive the name of a column from its number.

The size of the array is returned in the *size* parameter. A *size* of 0 indicates that either there is no bylist for this particular compute or the *computeid* is out of range.

Usage

- **dbbylist** returns the bylist for a compute row. (A **select** statement's **compute** clause may contain the keyword **by**, followed by a list of columns. This list, known as the "bylist," divides the results into subgroups, based on changing values in the specified columns. The **compute** clause's row aggregate is applied to each subgroup, generating a compute row for each subgroup.)
- **dbresults** must return SUCCEED before the application calls this routine.
- Assume the following command has been executed:

```
select dept, name, year, sales from employee
     order by dept, name, year
     compute count(name) by dept,name
```

  The call **dbbylist** (*dbproc*, 1, &*size*) sets *size* to 2, because there are two items in the bylist. It returns a pointer to an array of two BYTEs, which contain the values 1 and 2, indicating that the bylist is composed of columns 1 and 2 from the select list.

See also dbadata, dbadlen, dbaltlen, dbalttype, dbcolname, dbgetrow, dbnextrow

dbcancel

Description Cancel the current command batch.

Syntax

```
RETCODE dbcancel(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value SUCCEED or FAIL.

The most common reasons for failure are a dead DBPROCESS or a network error. **dbcancel** will also return FAIL if the server is dead.

Usage

- This routine cancels execution of the current command batch on the server and flushes any pending results. The application can call it after calling **dbsqlexec**, **dbsqlsend**, **dbsqlok**, **dbresults**, or **dbnextrow**. **dbcancel** sends an attention packet to the server which causes the server to cease execution of the command batch. Any pending results are read and discarded.
- **dbcancel** cancels *all* the commands in the current command batch. To cancel only the results from the current command, call **dbcanquery** instead.
- Some applications may need the ability to cancel a long-running query while DB-Library is reading from the network. In this case, the application should use one of these methods:
  - Set a time limit for server reads with **dbsettime** and add a special case to your error handler function to respond to SYBETIME errors. See the reference pages for dberrhandle and dbsettime for details.
  - Use **dbsetinterrupt** to install custom interrupt handling. See the reference page for dbsetinterrupt for details.
- If you have set your own interrupt handler using **dbsetinterrupt**, you cannot call **dbcancel** in your interrupt handler. This would cause output from the server to DB-Library to become out of sync. See the reference page for dbsetinterrupt for an explanation of how to cancel from an interrupt handler.

See also

dbcanquery, dbnextrow, dbresults, dbsetinterrupt, dbsqlexec, dbsqlok, dbsqlsend

dbcanquery

Description

Cancel any rows pending from the most recently executed query.

Syntax

```
RETCODE dbcanquery(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

SUCCEED or FAIL.

The most common reasons for failure are a dead DBPROCESS or a network error.

Usage

- This routine is an efficient way to throw away any unread rows that result from the most recently executed SQL query. Calling **dbcanquery** is equivalent to calling **dbnextrow** until it returns NO_MORE_ROWS, but **dbcanquery** is faster because it allocates no memory and executes no bindings to user data.
- If you have set your own interrupt handler using **dbsetinterrupt**, you cannot call **dbcanquery** in your interrupt handler. This would cause output from the server to DB-Library to become out of sync. If you want to ignore any unread rows from the current query, the interrupt handler should set a flag that you can check before the next call to **dbnextrow**.
- **dbresults** must return SUCCEED before an application can call **dbcanquery**.
- If you want to ignore all of the results from all of the commands in the current command batch, call **dbcancel** instead.

See also

dbcancel, dbnextrow, dbresults, dbsetinterrupt, dbsqlexec

dbchange

Description

Determine whether a command batch has changed the current database.

Syntax

```
char *dbchange(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

A pointer to the null-terminated name of the new database, if any. If the database has not changed, NULL will be returned.

Usage

- **dbchange** informs the program of a change in the current database. It does so by catching any instance of the Transact-SQL **use** command.
- Although a **use** command can appear anywhere in a command batch, the database change does not actually take effect until the end of the batch. **dbchange** is therefore useful only in determining whether the current command batch has changed the database for subsequent command batches.
- The internal DBPROCESS flag that **dbchange** monitors to determine whether the database has changed is cleared when the program executes a new command batch by calling either **dbsqlexec** or **dbsqlsend**. Therefore, the simplest way to keep track of database changes is to call **dbchange** when **dbresults** returns NO_MORE_RESULTS at the end of each command batch.
- Alternatively, you can always get the name of the current database by calling **dbname**.

See also

dbname, dbresults, dbsqlexec, dbsqlsend, dbuse

dbcharsetconv

Description — Indicate whether the server is performing character set translation.

Syntax

```
DBBOOL dbcharsetconv(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

Return value — TRUE if the server is performing translation; FALSE if it is not.

Usage

- **dbcharsetconv** returns TRUE if the server is performing character set translations, and FALSE if it is not.
- Note that if a client and a server are using the same character set, the server is not performing translation. As expected in this case, **dbcharsetconv** returns FALSE.
- To get the name of its own character set, a client can call **dbgetcharset**.
- To get the name of the server's character set, a client can call **dbservcharset**.

See also — dbservcharset, dbgetcharset, DBSETLCHARSET

dbclose

Description — Close and deallocate a single DBPROCESS structure.

Syntax

```
void dbclose(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value — None.

Usage

- **dbclose** is the inverse of **dbopen**. It cleans up any activity associated with one DBPROCESS structure and deallocates the space. It also closes the corresponding network connection.
- To close every open DBPROCESS structure, use **dbexit** instead.
- **dbclose** does not deallocate space associated with a LOGINREC. To deallocate a LOGINREC, an application can call **dbloginfree**.
- Calling **dbclose** with an argument not returned by **dbopen** is sure to cause trouble.

See also

dbexit, dbopen

dbclrbuf

Description

Drop rows from the row buffer.

Syntax

```
void dbclrbuf(dbproc, n)

DBPROCESS*   dbproc;
DBINT        n;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

n

The number of rows you want cleared from the row buffer. If you make *n* equal to or greater than the number of rows in the buffer, all but the newest row will be removed. If *n* is less than 1, the function call is ignored.

Return value

None.

Usage

- DB-Library provides a row-buffering service to application programs. You can turn row buffering on by calling **dbsetopt**(*dbproc*, DBBUFFER, *n*) where *n* is the number of rows you would like DB-Library to buffer. If buffering is on, you can then randomly refer to rows that have been read from the server, using **dbgetrow**. See the dbgetrow reference page for a discussion of the benefits and penalties of row buffering.

- The row buffer can become full for two reasons. Either the server has returned more than the *n* rows you said you wanted buffered, or sufficient space could not be allocated to save the row you wanted. When the row buffer is full, **dbnextrow** returns BUF_FULL and refuses to read in the next row from the server. Once the row buffer is full, subsequent calls to **dbnextrow** will continue to return BUF_FULL until at least one row is freed by calling **dbclrbuf**. **dbclrbuf** always frees the oldest rows in the buffer first.
- Once a result row has been cleared from the buffer, it is no longer available to the program.
- For an example of row buffering, see Example 4 of the online sample programs.

See also dbgetrow, dbnextrow, dbsetopt, Options

dbclropt

Description Clear an option set by **dbsetopt**.

Syntax RETCODE dbclropt(dbproc, option, param)

```
DBPROCESS    *dbproc;
int          option;
char*        param;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server. If *dbproc* is NULL, the option will be cleared for all active DBPROCESS structures.

option

The option that is to be turned off. See Options on page 402 for a list of options.

param

Certain options take parameters. The DBOFFSET option, for example, takes as a parameter the SQL construct for which offsets are to be returned. Options on page 402 lists those options that take parameters. If an option does not take a parameter, *param* must be NULL.

If the option you are clearing takes a parameter, but there can be only one instance of the option, **dbclropt** ignores the *param* argument. For example, **dbclropt** ignores the value of *param* when clearing the DBBUFFER option, because row buffering can have only one setting at a time. On the other hand, the DBOFFSET option can have several settings, each with a different parameter. It may have been set twice—to look for offsets to **select** statements and offsets to **order by** clauses. In that case, **dbclropt** needs the *param* argument to determine whether to clear the **select** offset or the **order by** offset.

If an invalid parameter is specified for one of the server options, this will be discovered the next time a command buffer is sent to the server. The **dbsqlexec** or **dbsqlsend** call will fail, and DB-Library will invoke the user-installed message handler. If an invalid parameter is specified for one of the DB-Library options (DBBUFFER or DBTEXTLIMIT), the **dbclropt** call itself will fail.

Return value

SUCCEED or FAIL.

Usage

- This routine clears the server and DB-Library options that have been set with **dbsetopt**. Although server options may be set and cleared directly through SQL, the application should instead use **dbsetopt** and **dbclropt** to set and clear options. This provides a uniform interface for setting both server and DB-Library options. It also allows the application to use the **dbisopt** function to check the status of an option.
- **dbclropt** does not immediately clear the option. The option is cleared the next time a command buffer is sent to the server (by invoking **dbsqlexec** or **dbsqlsend**).
- For a complete list of options, see Options on page 402.

See also

dbisopt, dbsetopt, Options

dbcmd

Description

Add text to the DBPROCESS command buffer.

Syntax

RETCODE dbcmd(dbproc, cmdstring)

```
DBPROCESS    *dbproc;
char         *cmdstring;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

cmdstring

A null-terminated character string that **dbcmd** copies into the command buffer.

Return value

SUCCEED or FAIL.

Usage

- This routine adds text to the Transact-SQL command buffer in the DBPROCESS structure. It adds to the existing command buffer—it does not delete or overwrite the current contents except after the buffer has been sent to the server (see "Clearing the command buffer" on page 95). A single command buffer may contain multiple commands; in fact, this represents an efficient use of the command buffer.
- **dbfcmd** is a related function. **dbfcmd** interprets the *cmdstring* as a format string that is passed to **sprintf** along with any additional arguments. The application can intermingle calls to **dbcmd** and **dbfcmd**.

Consecutive calls to dbcmd

- The application may call **dbcmd** repeatedly. The command strings in sequential calls are just concatenated together. It is the application's responsibility to ensure that any necessary blanks appear between the end of one string and the beginning of the next.
- Here is a small example of using **dbcmd** to build up a multiline SQL command:

```
DBPROCESS       *dbproc;

 dbcmd(dbproc, "select name from sysobjects");
 dbcmd(dbproc, " where id < 5");
 dbcmd(dbproc, " and type='S'");
```

Note the required spaces at the start of the second and third command strings.

- At any time, the application can access the contents of the command buffer through calls to **dbgetchar**, **dbstrlen**, and **dbstrcpy**.
- Available memory is the only constraint on the size of the DBPROCESS command buffer created by calls to **dbcmd** and **dbfcmd**.

Clearing the command buffer •After a call to **dbsqlexec** or **dbsqlsend**, the first call to either **dbcmd** or **dbfcmd** automatically clears the command buffer before the new text is entered. If this situation is undesirable, set the DBNOAUTOFREE option. When DBNOAUTOFREE is set, the command buffer is cleared only by an explicit call to **dbfreebuf**.

See also dbfcmd, dbfreebuf, dbgetchar, dbstrcpy, dbstrlen, Options

DBCMDROW

Description Determine whether the current command can return rows.

Syntax RETCODE DBCMDROW(dbproc)

```
DBPROCESS     *dbproc;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value SUCCEED or FAIL, to indicate whether the command can return rows.

Usage

- **DBCMDROW** determines whether the command currently being processed by **dbresults** is one that can return rows—that is, a Transact-SQL **select** statement or an **execute** on a stored procedure containing a **select**. The application can call it after **dbresults** returns SUCCEED.
- Even if **DBCMDROW** macro returns SUCCEED, the command does not return any rows if none have qualified. To determine whether any rows are actually being returned, use **DBROWS**.

See also dbnextrow, dbresults, DBROWS, DBROWTYPE

dbcolbrowse

Description

Determine whether the source of a regular result column is updatable via the DB-Library browse-mode facilities.

Syntax

DBBOOL dbcolbrowse(dbproc, colnum)

```
DBPROCESS   *dbproc;
int         colnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

colnum

The number of the result column of interest. Column numbers start at 1.

Return value

TRUE or FALSE.

Usage

- **dbcolbrowse** is one of the DB-Library browse mode routines. See Chapter 1, “Introducing DB-Library” for a detailed discussion of browse mode.
- **dbcolbrowse** provides a way to determine whether the database column that is the source of a regular (that is, non-compute) result column in a select list is updatable via the DB-Library browse-mode facilities. This routine is useful in examining ad hoc queries. If the query has been hard-coded into the program, **dbcolbrowse** obviously is unnecessary.
- To be updatable, a column must be derived from a browsable table (that is, the table must have a unique index and a timestamp column) and cannot be the result of a SQL expression. For example, in the following select list:

  ```
  select title, category=type,
       wholesale=(price * 0.6) ... for browse
  ```

  result columns 1 and 2 (“title” and “category”) are updatable, but column 3 (“wholesale”) is not, because it is the result of an expression.
- The application can call **dbcolbrowse** anytime after **dbresults**.
- To determine the name of the source column given the name of the result column, use **dbcolsource**.
- Example 7 of the online sample programs contains a call to **dbcolbrowse**.

See also

dbcolsource, dbqual, dbtabbrowse, dbtabcount, dbtabname, dbtabsource, dbtsnewlen, dbtsnewval, dbtsput

dbcollen

Description

Return the maximum length of the data in a regular result column.

Syntax

DBINT dbcollen(dbproc, column)

```
DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column of interest. The first column is number 1.

Return value

The maximum length, in bytes, of the data for the particular column. If the column number is not in range, **dbcollen** returns -1.

Usage

- This routine returns the maximum length of the data in a regular (that is, non-compute) result column. In the case of variable length data, this is not necessarily the actual length of the data, but rather the maximum length that the data can be. For the actual data length, use **dbdatlen**.
- The value that **dbcollen** returns is not affected by Transact-SQL string functions such as **rtrim** and **ltrim**. For example, if the column *au_lname* has a maximum length of 20 characters, and the first row instance of *au_lname* is "Goodman " (a value padded with 13 spaces),

  dbcollen returns 20 as the length of *au_lname*, even though the Transact-SQL command **select rtrim(au_lname) from authors** returns a string that is 5 characters long.
- Here is a small program fragment that uses **dbcollen**:

```
DBPROCESS     *dbproc;
 int           colnum;
 DBINT         column_length;

 /* Put the command into the command buffer */
 dbcmd(dbproc, "select name, id, type from
       sysobjects");

 /*
 ** Send the command to SQL Server and begin
 ** execution
 */
```

```
dbsqlexec(dbproc);

/* process the command results */
dbresults(dbproc);

/* examine the column lengths */
for (colnum = 1; colnum < 4; colnum++)
{
    column_length = dbcollen(dbproc, colnum);
    printf("column %d, length is %ld.\n", colnum,
          column_length);
}
```

See also dbcolname, dbcoltype, dbdata, dbdatlen, dbnumcols

dbcolname

Description Return the name of a regular result column.

Syntax char *dbcolname(dbproc, column)

```
DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column
The number of the column of interest. The first column is number 1.

Return value A CHAR pointer to the null-terminated name of the particular column. If the column number is not in range, **dbcolname** returns NULL.

Usage

- This routine returns a pointer to the null-terminated name of a regular (that is, non-compute) result column.
- Here is a small program fragment that uses **dbcolname**:

```
DBPROCESS     *dbproc;

/* Put the command into the command buffer */
dbcmd(dbproc, "select name, id, type from
       sysobjects");
```

```
/*
** Send the command to SQL Server and begin
** execution
*/
dbsqlexec(dbproc);

/* Process the command results */
dbresults(dbproc);

/* Examine the column names */
printf("first column name is %s\n",
                dbcolname(dbproc, 1));
printf("second column name is %s\n",
                dbcolname(dbproc, 2));
printf("third column name is %s\n",
                dbcolname(dbproc, 3));
```

See also dbcollen, dbcoltype, dbdata, dbdatlen, dbnumcols

dbcolsource

Description Return a pointer to the name of the database column from which the specified regular result column was derived.

Syntax char *dbcolsource(dbproc, colnum)

```
DBPROCESS      *dbproc;
int              colnum;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

colnum

The number of the result column of interest. Column numbers start at 1.

Return value A pointer to a null-terminated column name. This pointer will be NULL if the column number is out of range or if the column is the result of a SQL expression, such as **max(colname)**.

Usage

- **dbcolsource** is one of the DB-Library browse mode routines. It is usable only with results from a browse-mode **select** (that is, a **select** containing the key words **for browse**). See Chapter 1, "Introducing DB-Library" for a detailed discussion of browse mode.
- **dbcolsource** provides an application with information it needs to update a database column, based on an ad hoc query. **select** statements may optionally specify header names for regular (that is, non-compute) result columns:

  ```
  select author = au_lname from authors for browse
  ```

  When updating a table, you must use the database column name, not the header name (in this example, "au_lname", not "author"). You can use the **dbcolsource** routine to get the underlying database column name:

  ```
  dbcolsource(dbproc, 1)
  ```

  This call will return a pointer to the string "au_lname".
- **dbcolsource** is useful for ad hoc queries. If the query has been hard-coded into the program, this routine obviously is unnecessary.
- The application can call **dbcolsource** anytime after **dbresults**.
- Example 7 of the online sample programs contains a call to **dbcolsource**.

See also

dbcolbrowse, dbqual, dbtabbrowse, dbtabcount, dbtabname, dbtabsource, dbtsnewlen, dbtsnewval, dbtsput

dbcoltype

Description

Return the datatype for a regular result column.

Syntax

int dbcoltype(dbproc, column)

```
DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column of interest. The first column is number 1.

Return value

A token value for the datatype for a particular column.

In a few cases, the token value returned by this routine may not correspond exactly with the column's server datatype:

- SYBVARCHAR is returned as SYBCHAR.
- SYBVARBINARY is returned as SYBBINARY.
- SYBDATETIMN is returned as SYBDATETIME.
- SYBMONEYN is returned as SYBMONEY.
- SYBFLTN is returned as SYBFLT8.
- SYBINTN is returned as SYBINT1, SYBINT2, or SYBINT4, depending on the actual type of the SYBINTN.

If the column number is not in range, **dbcoltype** returns -1.

Usage

- This routine returns the datatype for a regular (that is, non-compute) result column. For a list of server datatypes, see Types on page 408.
- **dbcoltype** actually returns an integer token value for the datatype (SYBCHAR, SYBFLT8, and so on). To convert the token value into a readable token string, use **dbprtype**. See the dbprtype reference page for a list of all token values and their equivalent token strings.
- You can use **dbvarylen** to determine whether a column's datatype is variable length.
- Here is a program fragment that uses **dbcoltype**:

```
DBPROCESS      *dbproc;
 int            colnum;
 int            coltype;

 /* Put the command into the command buffer */
 dbcmd(dbproc, "select name, id, type from
      sysobjects");

 /* Send the command to SQL Server and begin
 ** execution.
 */
 dbsqlexec(dbproc);

 /* Process the command results */
 dbresults(dbproc);

 /* Examine the column types */
 for (colnum = 1; colnum < 4; colnum++)
```

```
{
coltype = dbcoltype(dbproc, colnum);
printf("column %d, type is %s.\n", colnum,
dbprtype(coltype));
}
```

See also dbcollen, dbcolname, dbdata, dbdatlen, dbnumcols, dbprtype, dbvarylen, Types

dbcoltypeinfo

Description Return precision and scale information for a regular result column of type *numeric* or decimal.

Syntax DBTYPEINFO * dbcoltypeinfo(dbproc, column)

```
DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column of interest. The first column is number 1.

Return value A pointer to a DBTYPEINFO structure that contains precision and scale values for a particular *numeric* or decimal column, or NULL if the specified column number is not in the result set.

A DBTYPEINFO structure is defined as follows:

```
typedef struct typeinfo {
      DBINT    precision;
      DBINT    scale;
 } DBTYPEINFO;
```

If the datatype of the column is not numeric or decimal, the returned structure will contain meaningless values. Check that **dbcoltype** returns SYBNUMERIC or SYBDECIMAL before calling this function.

Usage

- This routine returns a pointer to a DBTYPEINFO structure that provides precision and scale information for a regular (that is, non-compute) result column of datatype numeric or decimal.

- The precision and scale values returned for columns with other datatypes will be meaningless. Check that **dbcoltype** returns SYBNUMERIC or SYBDECIMAL before calling **dbcoltypeinfo**.

See also

dbcollen, dbcolname, dbcoltype, dbdata, dbdatlen, dbnumcols, dbprtype, dbvarylen, Types

dbcolutype

Description

Return the user-defined datatype for a regular result column.

Syntax

DBINT dbcolutype(dbproc, column)

```
DBPROCESS      *dbproc;
int              column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column of interest. The first column is number 1.

Return value

column's user-defined datatype or a negative integer if *column* is not in range.

Usage

- **dbcolutype** returns the user-defined datatype for a regular result column. For a description of how to add user-defined datatypes to Adaptive Server databases, see **sp_addtype** in the *Adaptive Server Enterprise Reference Manual*.
- **dbcolutype** is defined as datatype DBINT to accommodate the size of user-defined datatypes. Both DBINT and user-defined datatypes are 32 bits long.
- The following code fragment illustrates the use of **dbcolutype**:

```
DBPROCESS     *dbproc;
 int           colnum;
 int           numcols;

 /* Put the command into the command buffer */
 dbcmd(dbproc, "select * from mytable");

 /*
```

```
** Send the command to the SQL Server and begin
** execution.
*/
dbsqlexec(dbproc);

/* Process the command results */
dbresults(dbproc);

/* Examine the user-defined column types */
numcols =  dbnumcols(dbproc);
for (colnum = 1; colnum < numcols; colnum++)
{
     printf ("column %d, user-defined type is \
          %ld.\n", colnum, dbcolutype(dbproc,
          colnum));
}
```

See also dbaltutype, dbcoltype

dbconvert

Description Convert data from one datatype to another.

Syntax

```
DBINT dbconvert(dbproc, srctype, src, srclen,
        desttype, dest, destlen)

DBPROCESS    *dbproc;
int          srctype;
BYTE         *src;
DBINT        srclen;
int          desttype;
BYTE         *dest;
DBINT        destlen;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server. In **dbconvert**, the DBPROCESS is used only to supply any custom null values that the program may have specified via **dbsetnull**. If *dbproc* is NULL, **dbconvert** uses the default values for null value data conversions.

srctype

The datatype of the data that is to be converted. This parameter can be any of the server datatypes, as listed below in Table 2-7 on page 112.

src
A pointer to the data which is to be converted. If this pointer is NULL, **dbconvert** will place an appropriate null value in the destination variable. You can use **dbdata** to get the server data.

srclen
The length, in bytes, of the data to be converted. If the *srclen* is 0, the source data is assumed to be null and **dbconvert** will place an appropriate null value in the destination variable. Otherwise, this length is ignored for all datatypes except char, text, binary, and image. For SYBCHAR, SYBBOUNDARY, and SYBSENSITIVITY data, a length of -1 indicates that the string is null-terminated. You can use **dbdatlen** to get the length of the server data.

desttype
The datatype that the source data is to be converted into. This parameter can be any of the server datatypes, as listed below in Table 2-7 on page 112.

dest
A pointer to the destination variable that will receive the converted data. If this pointer is NULL, **dbconvert** will call the user-supplied error handler (if any) and return -1.

destlen
The length, in bytes, of the destination variable. *destlen* is ignored for fixed-length datatypes. For a SYBCHAR, SYBBOUNDARY or SYBSENSITIVITY destination, the value of *destlen* must be the total length of the destination buffer space.

The following table describes special values for *destlen*:

Table 2-5: Special values for destlen (dbconvert)

Value of destlen	Applicable to	Meaning
-1	SYBCHAR, SYBBOUNDARY, SBYSENSITIVITY	There is sufficient space available. The string will be trimmed of trailing blanks and given a terminating null.
-2	SYBCHAR	There is sufficient space available. The string will not be trimmed of trailing blanks, but will be given a terminating null.

Return value

The length of the converted data, in bytes, if the datatype conversion succeeds.

If the conversion fails, **dbconvert** returns either -1 or FAIL, depending on the cause of the failure. **dbconvert** returns -1 to indicate a NULL destination pointer or an illegal datatype. **dbconvert** returns FAIL to indicate other types of failures.

If **dbconvert** fails, it will first call a user-supplied error handler (if any) and set the global DB-Library error value.

This routine may fail for several reasons: the requested conversion was not available; the conversion resulted in truncation, overflow, or loss of precision in the destination variable; or a syntax error occurred in converting a character string to some numeric type.

Usage

- This routine allows the program to convert data from one representation to another. To determine whether a particular conversion is permitted, the program can call **dbwillconvert** before attempting a conversion.
- **dbconvert** can convert data stored in any of the server datatypes (although, of course, not all conversions are legal). See Table 2-7 on page 112 for a list of type constants and corresponding program variable types.
- It is an error to use the following datatypes with **dbconvert** if the library version has not been set (with **dbsetversion**) to DBVERSION_100 or higher: SYBNUMERIC, SYBDECIMAL, SYBBOUNDARY, and SYBSENSITIVITY.
- Table 2-8 on page 113 lists the datatype conversions that **dbconvert** supports. The source datatypes are listed down the leftmost column and the destination datatypes are listed along the top row of the table. (For brevity, the prefix “SYB” has been eliminated from each datatype.) T (TRUE) indicates that the conversion is supported; F (FALSE) indicates that the conversion is not supported.
- A conversion to or from the datatypes SYBBINARY and SYBIMAGE is a straight bit-copy, except when the conversion involves SYBCHAR or SYBTEXT. When converting SYBCHAR or SYBTEXT data to SYBBINARY or SYBIMAGE, DBCONVERT interprets the SYBCHAR or SYBTEXT string as hexadecimal, whether or not the string contains a leading “0x”. When converting SYBBINARY or SYBIMAGE data to SYBCHAR or SYBTEXT, **dbconvert** creates a hexadecimal string without a leading “0x”.
- Note that SYBINT2 and SYBINT4 are signed types. When converting these types to character, conversion error can result if the quantity being converted is unsigned and uses the high bit.

- Converting a SYBMONEY, SYBCHAR, or SYBTEXT value to SYBFLT8 may result in some loss of precision. Converting a SYBFLT8 value to SYBCHAR or SYBTEXT may also result in some loss of precision.
- Converting a SYBFLT8 value to SYBMONEY can result in overflow, because the maximum value for SYBMONEY is $922,337,203,685,477.58.
- If overflow occurs when converting integer or float data to SYBCHAR or SYBTEXT, the first character of the resulting value will contain an asterisk (*) to indicate the error.
- A conversion to SYBBIT has the following effect: If the value being converted is not 0, the SYBBIT value will be set to 1; if the value is 0, the SYBBIT value will be set to 0.
- **dbconvert** does not offer precision and scale support for numeric and decimal datatypes. When converting to SYBNUMERIC or SYBDECIMAL, **dbconvert** uses a default precision and scale of 18 and 0, respectively. To specify a different precision and scale, an application can use **dbconvert_ps**.
- SYBBOUNDARY and SYBSENSITIVITY destinations are always null-terminated.
- In certain cases, it can be useful to convert a datatype to itself. For instance, a conversion of SYBCHAR to SYBCHAR with a *destlen* of -1 serves as a useful way to append a null terminator to a string, as the example below illustrates.
- Here is a short example that illustrates how to convert server data obtained with **dbdata**:

```
DBCHAR      title[81];
 DBCHAR      price[9];

 /* Read the query into the command buffer */
 dbcmd(dbproc, "select title, price, royalty from \
     pubs2..titles");

 /* Send the query to SQL Server */
 dbsqlexec(dbproc);

 /* Get ready to process the query results */
 dbresults(dbproc);

 /* Process each row */
```

```
while (dbnextrow(dbproc) != NO_MORE_ROWS)
{
    /*
    ** The first dbconvert() adds a null
    ** terminator to the string.
    */
    dbconvert(dbproc, SYBCHAR, (dbdata(dbproc,1)),
        (dbdatlen(dbproc,1)), SYBCHAR, title,
        (DBINT)-1);
    /*
    ** The second dbconvert() converts money to
    ** string.
    */
    dbconvert(dbproc, SYBMONEY,
       (dbdata(dbproc,2)), (DBINT)-1, SYBCHAR,
        price, (DBINT)-1);

    if (dbdatlen(dbproc,3) != 0)
        printf ("%s\n $%s  %ld\n", title, price,
        *((DBINT *)dbdata(dbproc,3)));
 }
```

In the **dbconvert** calls it was not necessary to cast the returns from **dbdata**, because **dbdata** returns a BYTE pointer—precisely the datatype **dbconvert** expects in the third parameter.

- If you are binding data to variables with **dbbind** rather than accessing the data directly with **dbdata**, **dbbind** can perform the conversions itself, making **dbconvert** unnecessary.

Note *OS/2 and NetWare programmers:* When **dbconvert** is used to convert a money value to a character type, the resulting string only contains two decimal places. On other platforms, money-to-character conversions result in a string containing four decimal places.

- Example 5 of the online sample programs illustrates several more types of conversions using **dbconvert**.
- See “Types” on page 408 for a list of DB-Library datatypes and the corresponding Adaptive Server datatypes. For more information on Adaptive Server datatypes, see the *Adaptive Server Enterprise Reference Manual*.

See also

dbaltbind, dbaltbind_ps, dbbind, dbbind_ps, dbconvert_ps, dberrhandle, dbsetnull, dbsetversion, dbwillconvert, “Errors,” Types

dbconvert_ps

Description

Convert data from one datatype to another, with precision and scale support for numeric and decimal datatypes.

Syntax

```
DBINT dbconvert_ps(dbproc, srctype, src, srclen,
        desttype, dest, destlen, typeinfo)

DBPROCESS      *dbproc;
int            srctype;
BYTE           *src;
DBINT          srclen;
int            desttype;
BYTE           *dest;
DBINT          destlen;
DBTYPEINFO     *typeinfo;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server. In **dbconvert_ps**, the DBPROCESS is used only to supply any custom null values that the program may have specified via **dbsetnull**. If *dbproc* is NULL, **dbconvert_ps** uses the default values for null value data conversions.

srctype

The datatype of the data which is to be converted. This parameter can be any of the server datatypes, as listed in Table 2-6 on page 110.

src

A pointer to the data that is to be converted. If this pointer is NULL, **dbconvert_ps** will place an appropriate null value in the destination variable. You can use **dbdata** to get the server data.

srclen

The length, in bytes, of the data to be converted. If the *srclen* is 0, the source data is assumed to be NULL and **dbconvert_ps** will place an appropriate null value in the destination variable. Otherwise, this length is ignored for all datatypes except char, text, binary, and image. For SYBCHAR data, a length of -1 indicates that the string is null-terminated. You can use **dbdatlen** to get the length of the server data.

desttype

The datatype that the source data is to be converted into. This parameter can be any of the server datatypes, as listed in Table 2-6 on page 110.

dest

A pointer to the destination variable that will receive the converted data. If this pointer is NULL, **dbconvert_ps** will call the user-supplied error handler (if any) and return -1.

destlen

The length, in bytes, of the destination variable. *destlen* is ignored for fixed-length datatypes. For a SYBCHAR, SYBBOUNDARY, or SYBSENSITIVITY destination, the value of *destlen* must be the total length of the destination buffer space.

The following table describes special values for *destlen*:

Table 2-6: Special values for destlen (dbconvert_ps)

Value of destlen	Applicable to	Meaning
-1	SYBCHAR, SYBBOUNDARY, SBYSENSITIVITY	There is sufficient space available. The string will be trimmed of trailing blanks and given a terminating null.
-2	SYBCHAR	There is sufficient space available. The string will not be trimmed of trailing blanks, but will be given a terminating null.

typeinfo

A pointer to a DBTYPEINFO structure containing information about the precision and scale of decimal or numeric values. An application sets a DBTYPEINFO structure with values for precision and scale before calling **dbconvert_ps** to convert data into DBDECIMAL or DBNUMERIC variables.

If *typeinfo* is NULL:

- If the source value is of type SYBNUMERIC or SYBDECIMAL, **dbconvert_ps** picks up precision and scale values from the source. In effect, the source data is copied to the destination space.
- If the source value is not SYBNUMERIC or SYBDECIMAL, **dbconvert_ps** uses a default precision of 18 and a default scale of 0.

If *srctype* is not SYBDECIMAL or SYBNUMERIC, *typeinfo* is ignored.

A DBTYPEINFO structure is defined as follows:

```
typedef struct typeinfo {
```

```
        DBINT    precision;
        DBINT    scale;
} DBTYPEINFO;
```

Legal values for *precision* are from 1 to 77. Legal values for *scale* are from 0 to 77. *scale* must be less than or equal to *precision*.

Return value

The length of the converted data, in bytes, if the datatype conversion succeeds.

If the conversion fails, **dbconvert_ps** returns either -1 or FAIL, depending on the cause of the failure. **dbconvert_ps** returns -1 to indicate a NULL destination pointer or an illegal datatype. **dbconvert_ps** returns FAIL to indicate other types of failures.

If **dbconvert_ps** fails, it will first call a user-supplied error handler (if any) and set the global DB-Library error value.

This routine may fail for several reasons: the requested conversion was not available; the conversion resulted in truncation, overflow, or loss of precision in the destination variable; or a syntax error occurred in converting a character string to some numeric type.

Usage

- **dbconvert_ps** is the equivalent of **dbconvert**, except that **dbconvert_ps** provides precision and scale support for numeric and decimal datatypes, which **dbconvert** does not. Calling **dbconvert** is equivalent to calling **dbconvert_ps** with *typeinfo* as NULL.
- **dbconvert_ps** allows a program to convert data from one representation to another. To determine whether a particular conversion is permitted, the program can call **dbwillconvert** before attempting a conversion.
- **dbconvert_ps** can convert data stored in any of the server datatypes (but not all conversions are legal—see Table 2-8 on page 113). The following table shows type constants for server datatypes and the corresponding program variable types:

Table 2-7: Type constants and program variable types

Type constant	Program variable type
SYBCHAR	DBCHAR
SYBTEXT	DBCHAR
SYBBINARY	DBBINARY
SYBIMAGE	DBBINARY
SYBINT1	DBTINYINT
SYBINT2	DBSMALLINT
SYBINT4	DBINT
SYBFLT8	DBFLT8
SYBREAL	DBREAL
SYBNUMERIC	DBNUMERIC
SYBDECIMAL	DBDECIMAL
SYBBIT	DBBIT
SYBMONEY	DBMONEY
SYBMONEY4	DBMONEY4
SYBDATETIME	DBDATETIME
SYBDATETIME4	DBDATETIME4
SYBBOUNDARY	DBCHAR
SYBSENSITIVITY	DBCHAR

Warning! It is an error to use the following datatypes with **dbconvert_ps** if the library version has not been set (with **dbsetversion**) to DBVERSION_100 or higher: SYBNUMERIC, SYBDECIMAL, SYBBOUNDARY, and SYBSENSITIVITY.

- Figure 2-8 lists the datatype conversions that **dbconvert_ps** supports. The source datatypes are listed down the leftmost column and the destination datatypes are listed along the top row of the table. (For brevity, the prefix “SYB” has been eliminated from each datatype.) T (TRUE) indicates that the conversion is supported; F (FALSE) indicates that the conversion is not supported.

Table 2-8: Supported datatype conversions

From	To: CHAR	TEXT	BINARY	IMAGE	INT1	INT2	INT4	FLT8	REAL	NUMERIC	DECIMAL	BIT	MONEY	MONEY4	DATETIME	DATETIME4	BOUNDARY	SENSITIVITY
CHAR	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T
TEXT	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T	T
BINARY	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
IMAGE	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
INT1	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
INT2	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
INT4	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
FLT8	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
REAL	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
NUMERIC	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
DECIMAL	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
BIT	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
MONEY	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
MONEY4	T	T	T	T	T	T	T	T	T	T	T	T	T	T	F	F	F	F
DATETIME	T	T	T	T	F	F	F	F	F	F	F	F	F	F	T	T	F	F
DATETIME4	T	T	T	T	F	F	F	F	F	F	F	F	F	F	T	T	F	F
BOUNDARY	T	T	F	F	F	F	F	F	F	F	F	F	F	F	F	F	T	F
SENSITIVITY	T	T	F	F	F	F	F	F	F	F	F	F	F	F	F	F	F	T

- A conversion to or from the datatypes SYBBINARY and SYBIMAGE is a straight bit-copy, except when the conversion involves SYBCHAR or SYBTEXT. When converting SYBCHAR or SYBTEXT data to SYBBINARY or SYBIMAGE, **dbconvert_ps** interprets the SYBCHAR or SYBTEXT string as hexadecimal, whether or not the string contains a leading “0x”. When converting SYBBINARY or SYBIMAGE data to SYBCHAR or SYBTEXT, **dbconvert_ps** creates a hexadecimal string without a leading “0x”.
- Note that SYBINT2 and SYBINT4 are signed types. When converting these types to character, conversion error can result if the quantity being converted is unsigned and uses the high bit.

- Converting a SYBMONEY, SYBCHAR, or SYBTEXT value to SYBFLT8 may result in some loss of precision. Converting a SYBFLT8 value to SYBCHAR or SYBTEXT may also result in some loss of precision.
- Converting a SYBFLT8 value to SYBMONEY can result in overflow, because the maximum value for SYBMONEY is $922,337,203,685,477.58.
- If overflow occurs when converting integer or float data to SYBCHAR or SYBTEXT, the first character of the resulting value will contain an asterisk (*) to indicate the error.
- A conversion to SYBBIT has the following effect: If the value being converted is not 0, the SYBBIT value will be set to 1; if the value is 0, the SYBBIT value will be set to 0.
- SYBBOUNDARY and SYBSENSITIVITY destinations are always null-terminated.
- In certain cases, it can be useful to convert a datatype to itself. For instance, a conversion of SYBCHAR to SYBCHAR with a *destlen* of -1 serves as a useful way to append a null terminator to a string.
- If you are binding data to variables with **dbbind** or **dbbind_ps** rather than accessing the data directly with **dbdata**, **dbbind** can perform the conversions itself, making **dbconvert_ps** unnecessary.

Note *OS/2 and NetWare programmers:* When **dbconvert_ps** is used to convert a money value to a character type, the resulting string only contains two decimal places. On other platforms, money-to-character conversions result in a string containing four decimal places.

- Example 5 of the online sample programs illustrates several more types of conversions using **dbconvert_ps**.
- See Types on page 408 for a list of DB-Library datatypes and the corresponding Adaptive Server datatypes. For more information on Adaptive Server datatypes, see the *Adaptive Server Reference Manual*.

See also

dbaltbind, dbaltbind_ps, dbbind, dbbind_ps, dbconvert, dberrhandle, dbsetnull, dbsetversion, dbwillconvert, “Errors,” Types

DBCOUNT

Description

Return the number of rows affected by a Transact-SQL command.

Syntax

```
DBINT DBCOUNT(dbproc)

 DBPROCESS      *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of rows affected by the command, or -1. DBCOUNT will return -1 if any of the following are true:

- The Transact-SQL command fails for any reason, such as a syntax error.
- The command is one that never affects rows, such as a **print** command.
- The command executes a stored procedure that does not execute any **select** statements.
- The DBNOCOUNT option is on.

Usage

- Once the results of a command have been processed, you can call **DBCOUNT** to find out how many rows were affected by the command. For example, if a **select** command was sent to the server and you have read all the rows by calling **dbnextrow** until it returned NO_MORE_ROWS, you can call this macro to find out how many rows were retrieved.
- If the current command is one that does not return rows, (for example, a **delete**), you can call DBCOUNT immediately after **dbresults**.
- If the command is one that executes a stored procedure, for example an **exec** or remote procedure call, DBCOUNT returns the number of rows returned by the latest **select** statement executed by the stored procedure, or -1 if the stored procedure does not execute any **select** statements. Note that a stored procedure that contains no **select** statements may execute a **select** by calling another stored procedure that does contain a **select**.

See also

dbnextrow, dbresults, Options

DBCURCMD

Description — Return the number of the current command.

Syntax

```
int DBCURCMD(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value — The number of the current command.

Usage

- This macro returns the number of the command whose results are currently being processed.
- The first command in a batch is number 1. The command number is incremented every time **dbresults** returns SUCCEED or FAIL. (Unsuccessful commands are counted.) The command number is reset by each call to **dbsqlexec** or **dbsqlsend**.

See also — DBCMDROW, DBMORECMDS, DBROWS

DBCURROW

Description — Return the number of the row currently being read.

Syntax

```
DBINT DBCURROW(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value — The number of the current row. This routine returns 0 if no rows have been processed yet.

Usage

- This macro returns the number of the row currently being read. Rows are counted from the first row returned from the server, whose number is 1. DBCURROW counts both regular and compute rows.

- The row number is reset to 0 by each new call to **dbresults**.
- The row number grows by one every time dbnextrow returns REG_ROW or a *computeid*.
- When row buffering is used, the row number *does not* represent the position in the row buffer. Rather, it represents the current row's position in the rows returned by the server. For more information on row buffering, see the reference pages for dbgetrow and dbsetrow.

See also

dbclrbuf, DBFIRSTROW, dbgetrow, DBLASTROW, dbnextrow, dbsetopt, Options

dbcursor

Description

Insert, update, delete, lock, or refresh a particular row in the fetch buffer.

Syntax

RETCODE dbcursor(hc, optype, bufno, table, values)

```
DBCURSOR    *hc;
DBINT       optype;
DBINT       bufno;
BYTE        *table;
BYTE        *values
```

Parameters

hc

Cursor handle previously returned by **dbcursoropen**.

optype

Type of operation to perform:

Table 2-9: Values for optype (dbcursor)

Symbolic value	Operation
CRS_UPDATE	Updates data.
CRS_DELETE	Deletes data.
CRS_INSERT	Inserts data.
CRS_REFRESH	Fetches another row in the buffer.
CRS_LOCKCC	Fetches another row and locks it. The row is actually locked only if inside a transaction block. The lock is released when the application commits or ends the transaction.

bufno

Row number in the fetch buffer to which the operation applies. The specified buffer must contain a valid row. If the value of *bufno* is 0, a CRS_REFRESH operation applies to all rows in the buffer. In an **insert** or **update** operation where no values parameter is given, the values are read from the bound variables array in the corresponding *bufno* value. The number of the first row in the buffer is 1.

table

The table to be inserted, updated, deleted, or locked if the cursor declaration contains more than one table. If there is only one table, this parameter is not required.

values

String values to be updated and/or inserted. Use this parameter only with **update** and **insert** to specify the new column values (that is, *Quantity* = *Quantity* + 1). In most cases, you can set this parameter to NULL and the new values for each column are taken from the fetch buffer (the program variable specified by **dbcursorbind**). If the **select** statement includes a computation (that is, **select 5*5...**) and a function (for example, **select getdate()**, **convert()**, and so on), then updating through the buffer array will surely not work.

There are four possible formats for this parameter: two for updating and two for inserting. The chosen format must match the *optype* (**update** or **insert**). Both contain a full and an abbreviated format. The full format is a complete SQL statement (**update** or **insert**) without a **where** clause. The abbreviated format is just the **set** clause (**update**) or just the **values** clause (**insert**). When the full format is used, the value specified for *tablename* overrides the *table* parameter of **dbcursor**. Because a **where** clause is added automatically, do not include one.

Return value

SUCCEED or FAIL.

This function can fail for the following reasons:

- Cursor is opened as read only, no updates allowed.
- Server or connection failure or timeout.
- No permission to update or change the database.
- A trigger in the database caused the **lock** or **update/insert** operation to fail.
- Optimistic concurrency control.

Usage

- If a column used as a unique index column is updated or changed, the corresponding row appears to be missing the next time it is fetched.
- For more information on cursors, see Appendix A, “Cursors”

See also

dbcursorbind, dbcursorclose, dbcursorcolinfo, dbcursorfetch, dbcursorinfo, dbcursoropen

dbcursorbind

Description

Register the binding information on the cursor columns.

Syntax

```
RETCODE dbcursorbind(hc, col, vartype, varlen,
                     poutlen, pvaraddr, typeinfo)

DBCURSOR     *hc;
int          col;
int          vartype;
DBINT        varlen;
DBINT        *poutlen;
BYTE         *pvaraddr;
DBTYPEINFO   *typeinfo;
```

Parameters

hc

Cursor handle created by **dbcursoropen**.

col

Number of the column to be bound to a program variable.

vartype

Binding type, which uses the same datatypes as the *vartype* parameter for **dbbind** and is bound by the same conversion rules. If this value is set to NOBIND for any column, the data is not bound. Instead, a pointer to the data is returned to the address in the corresponding *pvaraddr* entry for every row, and the length of the data is returned to the corresponding *varlen* array entry. This feature lets the application access the cursor data as it does with **dbdata** and **dbdatalen**.

varlen

Maximum length of variable-length datatype, such as CHARBIND, VARYCHARBIND, BINARYBIND, STRINGBIND, NTBSTRINGBIND, and VARYBINBIND. This parameter is ignored for fixed-length datatypes, such as INTBIND, FLT8BIND, MONEYBIND, BITBIND, SMALLBIND, and so on. It is also ignored if the vartype is NOBIND.

poutlen

Pointer to an array of DBINT integers where the actual length of the column's data is returned for each row. If *poutlen* is set to NULL, the lengths are not returned. The array size must be large enough to hold one DBINT variable for every row to be fetched at a time (as indicated by the *nrows* parameter in **dbcursoropen**).

When using **dbcursor** to update or insert with values from bound program variables, you can specify a null value by setting the corresponding *poutlen* to zero before calling **dbcursor**. Nonzero values are ignored except when NOBIND or one of the variable-length datatypes such as VARYCHARBIND or VARYBINBIND has been specified. In that case *poutlen* must contain the actual item length. If STRINGBIND or NTBSTRINGBIND has been specified, any non-zero value for *poutlen* is ignored, and the length of the string is determined by scanning for the null terminator.

pvaraddr

Pointer to the program variable to which the data is copied. If *vartype* is NOBIND, *pvaraddr* is assumed to point to an array of pointers—to the address of the actual data fetched by **dbcursorfetch**. This array's length must equal the value of *nrows* in **dbcursoropen**. If the cursor was opened with *nrows* > 1, *pvaraddr* is assumed to point to an array of *nrows* elements. Calling **dbcursorbind** with *pvaraddr* set to NULL breaks the existing binding.

typeinfo

A pointer to a DBTYPEINFO structure containing information about the precision and scale of decimal or numeric values. If *vartype* is not DECIMALBIND or NUMERICBIND, *typeinfo* is ignored.

To bind to DBNUMERIC or DBDECIMAL variables, an application initializes a DBTYPEINFO structure with values for precision and scale, then calls **dbcursorbind** with *vartype* as DECIMALBIND or NUMERICBIND.

If *typeinfo* is NULL and *vartype* is DECIMALBIND or NUMERICBIND:

- If the result column is of type numeric or decimal, **dbcursorbind** picks up precision and scale values from the result column.
- If the result column is not numeric or decimal, **dbcursorbind** uses a default precision of 18 and a default scale of 0.

A DBTYPEINFO structure is defined as follows:

```
typedef struct typeinfo {
```

```
        DBINTprecision;
        DBINTscale;
} DBTYPEINFO;
```

Legal values for *precision* are from 1 to 77. Legal values for *scale* are from 0 to 77. *scale* must be less than or equal to *precision*.

Return value

SUCCEED or FAIL.

Usage

- If **dbcursorbind** is called more than once for any column, only the last call is effective.
- This function works almost the same as **dbbind** without cursors.
- For more information on cursors, see Appendix A, “Cursors”

See also

dbcursor, dbcursorclose, dbcursorcolinfo, dbcursorfetch, dbcursorinfo, dbcursoropen

dbcursorclose

Description

Close the cursor associated with the given handle and release all the data belonging to it.

Syntax

void dbcursorclose(hc)

```
DBCURSOR    *hc;
```

Parameters

hc

Cursor handle created by **dbcursoropen.**

Return value

None.

Usage

- Closing a DBPROCESS connection with **dbcursorclose** automatically closes all the cursors associated with it. After issuing **dbcursorclose**, the cursor handle should not be used.
- For more information on cursors, see Appendix A, “Cursors”

See also

dbcursor, dbcursorbind, dbcursorcolinfo, dbcursorfetch, dbcursorinfo, dbcursoropen

dbcursorcolinfo

Description

Return column information for the specified column number in the open cursor.

Syntax

```
RETCODE dbcursorcolinfo(hcursor, column, colname,
                        coltype, collen, usertype)

DBCURSOR    *hcursor
DBINT       column;
DBCHAR      *colname;
DBINT       *coltype;
DBINT       *collen;
DBINT       *usertype;
```

Parameters

hcursor
Cursor handle created by **dbcursoropen**.

column
Column number for which information is to be returned.

colname
Location where the name of the column is returned. The user should allocate space large enough to accommodate the column name.

coltype
Location where the column's datatype is returned.

collen
Location where the column's maximum length is returned.

usertype
Location where the column's user-defined datatype is returned.

Return value

SUCCEED or FAIL.

Usage

- Any of the parameters

  colname, *coltype*, *collen*, or *usertype* can be set to NULL, in which case the information for that variable is not returned.

- For more information on cursors, see Appendix A, "Cursors"

See also

dbcursor, dbcursorbind, dbcursorclose, dbcursorfetch, dbcursorinfo, dbcursoropen

dbcursorfetch

Description

Fetch a block of rows into the program variables declared by the user in **dbcursorbind.**

Syntax

RETCODE dbcursorfetch(hc, fetchtype, rownum)

```
DBCURSOR    *hc;
DBINT       fetchtype;
DBINT       rownum;
```

Parameters

hc

Cursor handle created by **dbcursoropen**.

fetchtype

Type of fetch chosen. The scroll option in **dbcursoropen** determines which of these values are legal:

Table 2-10: Values for fetchtype (dbcursorfetch)

Symbolic value	Meaning	Comment
FETCH_FIRST	Fetch the first block of rows.	Although available for all cursor types, this option is especially useful for returning to the beginning of a keyset when you have selected a forward-only scrolling cursor.
FETCH_NEXT	Fetch the next block of rows.	If the result set exceeds the specified keyset size and if FETCH_RANDOM and/or FETCH_RELATIVE have been issued, a FETCH_NEXT can span a keyset boundary. In this case, the fetch that spans a keyset boundary returns a partial buffer, and the next fetch shifts down the keyset and returns the next full set of rows.
FETCH_PREV	Fetch the previous block of rows.	This option is unavailable with forward-only scrolling cursors. If *rownum* falls within the keyset, the range of rows must stay within the keyset because only the rows within the keyset are returned. This option does not change the keyset to the previous *rownum* rows in the result set.
FETCH_RANDOM	Fetch a block of rows, starting from the specified row number within the keyset.	This option is valid only within the keyset. The buffer is only partially filled when the range spans the keyset boundary.

Symbolic value	Meaning	Comment
FETCH_RELATIVE	Fetch a block of rows, relative to the number of rows indicated in the last fetch.	This option jumps *rownum* rows from the first row of the last fetch and starts fetching from there. The rows must remain within the keyset. The buffer is only partially filled when the range spans the keyset boundary.
FETCH_LAST	Fetch the last block of rows.	This value is available only with totally keyset-driven cursors.

rownum
The specified row for the buffer to start filling. Use this parameter only with a *fetchtype* of FETCH_RANDOM or -FETCH_RELATIVE.

Return value

SUCCEED or FAIL.

If the status array contains a status row for every row fetched, SUCCEED is returned. FAIL is returned if at least one of the following is true.

- FETCH_RANDOM and FETCH_RELATIVE require a keyset driven cursor.
- Forward-only scrolling can use only FETCH_FIRST and FETCH_NEXT.
- The server or a connection fails or takes a timeout.
- The client is out of memory.
- The FETCH_LAST option requires a fully keyset-driven cursor.

Usage

- Specify the size of the fetch buffer in dbcursoropen. **dbcursorfetch** fills the array passed as **dbcursoropen**'s *pstatus* parameter with status codes for the fetched rows. See the reference page for dbcursoropen for these codes.
- Program variables must first be registered, using **dbcursorbind**. Then the data can be transferred into the DB-Library buffers. The bound variables must, therefore, be arrays large enough to hold the specified number of rows. The status array contains status code for every row and contains flags for missing rows.
- When the range of rows specified by FETCH_NEXT, FETCH_RANDOM, or FETCH_RELATIVE spans a keyset boundary, only the rows remaining in the keyset are returned. In this case, the buffer is only partially filled, and the FETCH_ENDOFKEYSET flag is set as the status of the last row. The following FETCH_NEXT shifts the keyset down.

- For more information on cursors, see Appendix A, “Cursors”

See also dbcursor, dbcursorbind, dbcursorclose, dbcursorcolinfo, dbcursorinfo, dbcursoropen

dbcursorinfo

Description Return the number of columns and the number of rows in the keyset if the keyset hit the end of the result set.

Syntax RETCODE dbcursorinfo(hcursor, ncols, nrows);

```
DBCURSOR    *hcursor;
DBINT       *ncols
DBINT       *nrows;
```

Parameters

hcursor
Cursor handle created by **dbcursoropen**.

ncols
Location where the number of columns in the cursor is returned.

nrows
Location where the number of rows in the keyset is returned.

Return value SUCCEED or FAIL.

Usage

- For fully keyset-driven cursors, the *nrows* parameter contains the number of rows in the keyset. For mixed or dynamic cursors, *nrows* is always set to -1, unless the keyset is the last one in the result set. In that case, the number of rows in the keyset is returned. This helps the programmer find out when the keyset has hit the end of the result set.
- For more information on cursors, see Appendix A, “Cursors”

See also dbcursor, dbcursorbind, dbcursorclose, dbcursorcolinfo, dbcursorfetch, dbcursoropen

dbcursoropen

Description Open a cursor and specify the scroll option, concurrency option, and the size of the fetch buffer (the number of rows retrieved with a single fetch).

Syntax

```
DBCURSOR *dbcursoropen(dbproc, stmt, scrollopt,
        concuropt, nrows, pstatus)

DBPROCESS   *dbproc;
BYTE        *stmt;
SHORT       scrollopt;
SHORT       concuropt;
USHORT      nrows;
DBINT       *pstatus
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

stmt

The **select** statement that defines a cursor.

scrollopt

Indicator of the desired scrolling technique.

Keyset driven fixes membership in the result set and order at cursor open time.

Dynamic determines membership in the result set and order at fetch time.

Possible values for *scrollopt* are as follows:

Table 2-11: Values for scrollopt (dbcursoropen)

Symbolic value	Meaning
CUR_FORWARD	Forward scrolling only.
CUR_KEYSET	Keyset driven. A copy of the keyset for the result table is kept locally. Number of rows in result table must be less than or equal to 1000.
CUR_DYNAMIC	Fully dynamic.
int *n*	Keyset-driven cursor within (*n***nrows*) blocks, but fully dynamic outside the keyset.

concuropt

Definition of concurrency control. Following are possible values for *concuropt*:

Table 2-12: Values for concuropt (dbcursoropen)

Symbolic value	Meaning	Explanation
CUR_READONLY	Read-only cursor.	The data cannot be modified.

Symbolic value	Meaning	Explanation
CUR_LOCKCC	Intent to update locking.	All data, if inside a transaction block, is locked out as it is fetched through **dbcursorfetch**.
CUR_OPTCC	Optimistic concurrency control, based on timestamp values.	In a given row, modifications to the data succeed only if the row has not been updated since the last fetch. Changes are detected through timestamps or by comparing all non-text, non-image values in a selected table row.
CUR_OPTCCVAL	Optimistic concurrency based on values.	Same as CUR_OPTCC except changes are detected by comparing the values in all selected columns.

nrows

Number of rows in the fetch buffer (the width of the cursor). For mixed cursors the keyset capacity in rows is determined by this number multiplied by the value of the *scrollopt* parameter.

pstatus

Pointer to the array of row status indicators. The status of every row copied into the fetch buffer is returned to this array. The array must be large enough to hold one DBINT integer for every row in the buffer to be fetched. During the **dbcursorfetch** call, as the rows are filled into the bound variable, the corresponding status is filled with status information. **dbcursorfetch** fills in the status by setting bits in the status value. The application can use the bitmask values below to inspect the status value:

Table 2-13: Bitmask values for pstatus (dbcursoropen)

Symbolic value	Meaning
FTC_SUCCEED	The row was successfully copied. If this flag is not set, the row was not fetched.
FTC_MISSING	The row is missing.
FTC_ENDOFKEYSET	The end of the keyset. The remaining rows in the bind arrays are not used.
FTC_ENDOFRESULTS	The end of the result set. The remaining rows are not used.

Return value

If **dbcursoropen** succeeds, a handle to the cursor is returned. The cursor handle is required in calls to subsequent cursor functions.

If **dbcursoropen** fails, NULL is returned. Several errors, such as the following, can cause the cursor open to fail:

- Not enough memory in the system. Reduce the number of rows in the keyset, use dynamic scrolling, or reduce the number of rows to be fetched at a time.
- The CUR_KEYSET option is used for the *scrollopt* parameter, and there are more than 1000 rows in the result set. Use dynamic scrolling if the **select** statement can return more than 1000 rows.
- A unique row identifier could not be found.

Usage

- This function prepares internal DB-Library data structures based on the contents of the **select** statement and the values of *scrollopt*, *concuropt*, and *nrows*. **dbcursoropen** queries the server for information on unique qualifiers (row keys) for the rows in the cursor result set. If the cursor is keyset-driven, **dbcursoropen** queries the server and fetches row keys to build a keyset for the cursor's rows.
- The cursor definition cannot contain stored procedures or multiple Transact-SQL statements.
- For **dbcursor** to succeed, every table in the **select** statement must have a unique index. The Transact-SQL statements **for browse**, **select into**, **compute**, **union**, or **compute by** are not allowed in the cursor statement. Only fully keyset-driven cursors can have **order**, **having**, or **group by** phrases.
- When the **select** statement given as *stmt* refers to temporary tables, the current database must be *tempdb*. This restriction applies even if the temporary table was created in another database.
- Multiple cursors (as many as the system's memory allows) can be opened within the same *dbproc* connection. There should be no commands waiting to be executed or results pending in the DBPROCESS connection when cursor functions are called.
- For more information on cursors, see Appendix A, "Cursors"

See also

dbcursor, dbcursorbind, dbcursorclose, dbcursorcolinfo, dbcursorfetch, dbcursorinfo, dbcursoropen

dbdata

Description

Return a pointer to the data in a regular result column.

Syntax

BYTE *dbdata(dbproc, column)

```
DBPROCESS   *dbproc;
int         column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column of interest. The first column is number 1.

Return value

A BYTE pointer to the data for the particular column of interest. Be sure to cast this pointer into the proper type. A NULL BYTE pointer is returned if there is no such column or if the data has a null value. To make sure that the data is really a null value, you should always check for a return of 0 from **dbdatlen**.

Usage

- This routine returns a pointer to the data in a regular (that is, non-compute) result column. The data is not null-terminated. You can use **dbdatlen** to get the length of the data.
- Here is a small program fragment that uses **dbdata**:

```
DBPROCESS    *dbproc;
 DBINT        row_number = 0;
 DBINT        object_id;

/* Put the command into the command buffer */
 dbcmd(dbproc, "select id from sysobjects");
 /*
 ** Send the command to SQL Server and begin
 ** execution
 */
 dbsqlexec(dbproc);
 /* Process the command results */
 dbresults(dbproc);
 /* Examine the data in each row */
 while (dbnextrow(dbproc) != NO_MORE_ROWS)
 {
      row_number++;
      object_id = *((DBINT *)dbdata(dbproc, 1));
      printf("row %ld, object id is %ld.\n",
           row_number, object_id);
```

```
}
```

- Do not add a null terminator to string data until you have copied it from the DBPROCESS with a routine such as **strncpy**. For example:

```
char    objname[40];
        ...
 strncpy(objname, (char *)dbdata(dbproc,2),
      (int)dbdatlen(dbproc,2));
 objname[dbdatlen(dbproc,2)] = '\0';
```

- The function **dbbind** will automatically bind result data to your program variables. It does a copy of the data, but is often easier to use than **dbdata**. Furthermore, it includes a convenient type-conversion capability. By means of this capability, the application can, among other things, easily add a null terminator to a result string or convert money and datetime data to printable strings.

See also

dbbind, dbcollen, dbcolname, dbcoltype, dbdatlen, dbnumcols

dbdate4cmp

Description

Compare two DBDATETIME4 values.

Syntax

int dbdate4cmp(dbproc, d1, d2)

```
DBPROCESS    *dbproc;
DBDATETIME4  *d1;
DBDATETIME4  *d2;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL.

d1

A pointer to a DBDATETIME4 value.

d2

A pointer to a DBDATETIME4 value.

Return value

If *d1* = *d2*, **dbdate4cmp** returns 0.

If *d1* < *d2*, **dbdate4cmp** returns -1.

If *d1* > *d2*, **dbdate4cmp** returns 1.

Usage

- **dbdate4cmp** compares two DBDATETIME4 values.
- The range of legal DBDATETIME4 values is from January 1, 1900 to June 6, 2079. DBDATETIME4 values have a precision of one minute.

See also

dbdatecmp, dbmnycmp, dbmny4cmp

dbdate4zero

Description

Initialize a DBDATETIME4 variable to Jan 1, 1900 12:00AM.

Syntax

```
RETCODE dbdate4zero(dbproc, dateptr)

DBPROCESS    *dbproc;
DBDATETIME4  *dateptr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL.

dateptr

A pointer to the DBDATETIME4 variable to initialize.

Return value

SUCCEED or FAIL.

dbdate4zero returns FAIL if *dateptr* is NULL.

Usage

- **dbdate4zero** initializes a DBDATETIME4 variable to Jan 1, 1900 12:00AM.
- The range of legal DBDATETIME4 values is from January 1, 1900 to June 6, 2079. DBDATETIME4 values have a precision of one minute.

See also

dbdatezero

dbdatechar

Description

Convert an integer component of a DBDATETIME value into character format.

Syntax

RETCODE dbdatechar(dbproc, charbuf, datepart, value)

```
DBPROCESS   *dbproc;
char        *charbuf;
int         datepart;
int         value;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

charbuf

A pointer to the character buffer that will contain the null-terminated character representation of *value*.

datepart

A symbolic constant describing *value*'s type. The following table lists the date parts, the date part symbols recognized by DB-Library, and the expected values. Note that the names of the months and the days in this table are those for English.

Table 2-14: Date parts and their character representations (dbdatechar)

Date part	Symbol	Character representation of value
year	DBDATE_YY	1753 – 9999
quarter	DBDATE_QQ	1 – 4
month	DBDATE_MM	January – December
day of year	DBDATE_DY	1 – 366
day	DBDATE_DD	1 – 31
week	DBDATE_WK	1 – 54 (for leap years)
weekday	DBDATE_DW	Monday – Sunday
hour	DBDATE_HH	0 – 23
minute	DBDATE_MI	0 – 59
second	DBDATE_SS	0 – 59
millisecond	DBDATE_MS	0 – 999

value

The numeric value to be converted.

Return value
SUCCEED or FAIL.

Usage
- **dbdatechar** converts integer datetime components to character format. For example, **dbdatechar** associates the month component "3" with its associated character string: "March" if English is used, "mars" if French is used, and so on.
- The language of the associated character string is determined by the *dbproc*.
- **dbdatechar** is often useful in conjunction with **dbdatecrack**.

See also
dbconvert, dbdata, dbdatename, dbdatecrack

dbdatecmp

Description
Compare two DBDATETIME values.

Syntax
int dbdatecmp(dbproc, d1, d2)

```
DBPROCESS   *dbproc;
DBDATETIME  *d1;
DBDATETIME  *d2;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL.

d1
A pointer to a DBDATETIME value.

d2
A pointer to a DBDATETIME value.

Return value
If *d1 = d2*, **dbdatecmp** returns 0.

If *d1 < d2*, **dbdatecmp** returns -1.

If *d1 > d2*, **dbdatecmp** returns 1.

Usage
- **dbdatecmp** compares two DBDATETIME values.

- The range of legal DBDATETIME values is from January 1, 1753 to December 31, 9999. DBDATETIME values have a precision of 1/300th of a second (3.33 milliseconds).

See also

dbdate4cmp, dbmnycmp, dbmny4cmp

dbdatecrack

Description

Convert a machine-readable DBDATETIME value into user-accessible format.

Syntax

RETCODE dbdatecrack(dbproc, dateinfo, datetime)

```
DBPROCESS   *dbproc;
DBDATEREC   *dateinfo;
DBDATETIME  *datetime;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

dateinfo

A pointer to a DBDATEREC structure to contain the parts of *datetime*. DBDATEREC is defined as follows:

```
typedef struct dbdaterec
 {
     long    dateyear;    /* 1900 to the future */
     long    datemonth;   /* 0 - 11 */
     long    datedmonth;  /* 1 - 31 */
     long    datedyear;   /* 1 - 366 */
     long    datedweek;   /* 0 - 6  */
     long    datehour;    /* 0 - 23 */
     long    dateminute;  /* 0 - 59 */
     long    datesecond;  /* 0 - 59 */
     long    datemsecond; /* 0 - 997 */
     long    datetzone;   /* 0 - 127 */
```

```
} DBDATEREC;
```

Month and day names depend on the national language of the DBPROCESS. To retrieve these, use dbdatename or dbdayname plus dbmonthname.

Note The *dateinfo->datetzone* field is not set by **dbdatecrack**.

---

datetime
A pointer to the DBDATETIME value of interest.

Return value

SUCCEED or FAIL.

Usage

- **dbdatecrack** converts a DBDATETIME value into its integer components and places those components into a DBDATEREC structure.
- DBDATETIME structures store date and time values in an internal format. For example, a time value is stored as the number of 300th's of a second since midnight, and a date value is stored as the number of days since January 1, 1900. **dbdatecrack** converts the internal value to something more usable by an application program.
- The integer date parts placed in the DBDATEREC structure may be converted to character strings via **dbdatechar**.
- Calling **dbdatecrack** to convert an internal format datetime value is equivalent to calling **dbdatepart** many times.
- The following code fragment illustrates the use of **dbdatecrack**:

```
dbcmd(dbproc, "select name, crdate from \
      master..sysdatabases");
 dbsqlexec(dbproc);
 dbresults(dbproc);
 while (dbnextrow(dbproc) != NO_MORE_ROWS)
 {
      /*
      ** Print the database name and its date info
      */
      dbconvert(dbproc, dbcoltype(dbproc, 2),
           dbdata(dbproc, 2), dbdatlen(dbproc, 2),
           SYBCHAR, datestring, -1);
      printf("%s: %s\n", (char *)
           (dbdata(dbproc, 1)), datestring);
      /*
      ** Break up the creation date into its
      ** constituent parts.
```

```
        */
        dbdatecrack(dbproc, &dateinfo,
             (DBDATETIME *)(dbdata(dbproc, 2)));

      /* Print the parts of the creation date */
        printf("\tYear = &d.\n", dateinfo.dateyear);
        printf("\tMonth = &d.\n",dateinfo.datemonth);
        printf("\tDay of month = &d.\n",
             dateinfo.datedmonth);
        printf("\tDay of year = &d.\n",
             dateinfo.datedyear);
        printf("\tDay of week = &d.\n",
             dateinfo.datedweek);
        printf("\tHour = &d.\n", dateinfo.datehour);
        printf("\tMinute = &d.\n",
             dateinfo.dateminute);
        printf("\tSecond = &d.\n",
             dateinfo.datesecond);
        printf("\tMillisecond = &d.\n",
             dateinfo.datemsecond);
}
```

See also dbconvert, dbdata, dbdatechar, dbdatename, dbdatepart

dbdatename

Description Convert the specified component of a DBDATETIME structure into its corresponding character string.

Syntax int dbdatename(dbproc, charbuf, datepart, datetime)

```
DBPROCESS    *dbproc;
char         *charbuf;
int          datepart;
DBDATETIME   *datetime;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

charbuf
A pointer to a character buffer that will contain the null-terminated character representation of the *datetime* component of interest. If *datetime* is NULL, *charbuf* will contain a zero-length string.

datepart
The date component of interest. The following table lists the date parts, the date part symbols recognized by DB-Library and the expected values. Note that the names of the months and the days in this table are those for English.

Table 2-15: Date parts and their character representations (dbdatename)

Date part	Symbol	Character representation of value
year	DBDATE_YY	1753 – 9999
quarter	DBDATE_QQ	1 – 4
month	DBDATE_MM	January – December
day of year	DBDATE_DY	1 – 366
day	DBDATE_DD	1 – 31
week	DBDATE_WK	1 – 54 (for leap years)
weekday	DBDATE_DW	Monday – Sunday
hour	DBDATE_HH	0 – 23
minute	DBDATE_MI	0 – 59
second	DBDATE_SS	0 – 59
millisecond	DBDATE_MS	0 – 999

datetime
A pointer to the DBDATETIME value of interest.

Return value

The number of bytes placed into **charbuf*.

In case of error, **dbdatename** returns -1.

Usage

- **dbdatename** converts the specified component of a DBDATETIME structure into a character string.
- The names of the months and weekdays are in the language of the specified DBPROCESS. If *dbproc* is NULL, these names will be in DB-Library's default language.
- This function is very similar to the Transact-SQL **datename** function.
- The following code fragment illustrates the use of **dbdatename**:

```
dbcmd(dbproc, "select name, crdate from \
      master..sysdatabases");
 dbsqlexec(dbproc);
```

```
 dbresults(dbproc);

while (dbnextrow(dbproc) != NO_MORE_ROWS)
 {
      /*
      ** Print the database name and its date info
      */
      dbconvert(dbproc, dbcoltype(dbproc, 2),
           dbdata(dbproc, 2), dbdatlen(dbproc, 2),
           SYBCHAR, datestring, -1);
      printf("%s: %s\n", (char *) (dbdata
           (dbproc, 1)), datestring);

      /* Print the parts of the creation date */
      dbdatename(dbproc, datestring, DBDATE_YY,
           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tYear = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_QQ,
           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tQuarter = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_MM,
           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tMonth = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_DW,
           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tDay of week = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_DD,
           (DBDATETIME *) (dbdata(dbproc, 2)));
     printf("\tDay of month = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_DY,

           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tDay of year = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_HH,
           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tHour = %s.\n", datestring);

      dbdatename(dbproc, datestring, DBDATE_MI,
           (DBDATETIME *) (dbdata(dbproc, 2)));
      printf("\tMinute = %s.\n", datestring);
```

```
dbdatename(dbproc, datestring, DBDATE_SS,
     (DBDATETIME *) (dbdata(dbproc, 2)));
printf("\tSecond = %s.\n", datestring);

dbdatename(dbproc, datestring, DBDATE_MS,
     (DBDATETIME *) (dbdata(dbproc, 2)));
printf("\tMillisecond = %s.\n", datestring);

dbdatename(dbproc, datestring, DBDATE_WK,
     (DBDATETIME *) (dbdata(dbproc, 2)));
printf("\tWeek = %s.\n", datestring);
```

See also dbconvert, dbdata, dbdatechar, dbdatecrack

dbdateorder

Description Return the date component order for a given language.

Syntax char *dbdateorder(dbproc, language)

```
DBPROCESS   *dbproc;
char        *language;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

language

The name of the language of interest.

Return value A pointer to a null-terminated, 3-character string containing the characters “m”, “d”, and “y”, representing the month, day, and year date components, respectively. The order of the characters in the **dbdateorder** string corresponds to their order in *language*’s default datetime format.

dbdateorder returns a NULL pointer on failure.

Usage

- **dbdateorder** returns a character string that describes the order in which the month, day, and year date components appear in the specified language. If *language* is NULL, the current language of the specified DBPROCESS is used. If both *language* and *dbproc* are NULL, DB-Library's default language is used.

  Warning! The date order string returned by **dbdateorder** is a pointer to DB-Library's internal data structures. Application programs should neither modify this string, nor free it.

- The following code fragment illustrates the use of **dbdateorder**:

```
/* Retrieve the date order from SQL Server */
 printf("date-order: %s\n",
     (dbdateorder(DBPROCESS *)NULL, (char *)NULL));
```

See also

dbconvert, dbdata, dbdatechar, dbdatecrack

dbdatepart

Description

Return the specified part of a DBDATETIME value as a numeric value.

Syntax

DBINT dbdatepart(dbproc, datepart, datetime)

```
DBPROCESS    *dbproc;
int          datepart;
DBDATETIME   *datetime;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

datepart

The date component of interest. The following table lists the date parts, the date part symbols recognized by DB-Library and the expected values. Note that the names of the months and the days in this table are those for English.

Table 2-16: Date parts and their character representations (dbdatepart)

Date part	Symbol	Character representation of value
year	DBDATE_YY	1753 – 9999
quarter	DBDATE_QQ	1 – 4

Date part	Symbol	Character representation of value
month	DBDATE_MM	January – December
day of year	DBDATE_DY	1 – 366
day	DBDATE_DD	1 – 31
week	DBDATE_WK	1 – 54 (for leap years)
weekday	DBDATE_DW	Monday – Sunday
hour	DBDATE_HH	0 – 23
minute	DBDATE_MI	0 – 59
second	DBDATE_SS	0 – 59
millisecond	DBDATE_MS	0 – 999

datetime
A pointer to the DBDATETIME value of interest.

Return value

The value of the specified date part.

Usage

- **dbdatepart** returns the specified part of a DBDATETIME value as a numeric value.
- **dbdatepart** is similar to the Transact-SQL **datepart** function.

See also

dbconvert, dbdata, dbdatechar, dbdatecrack, dbdatename

dbdatezero

Description

Initialize a DBDATETIME value to Jan 1, 1900 12:00:00:000AM.

Syntax

RETCODE dbdatezero(dbproc, dateptr)

```
DBPROCESS    *dbproc;
DBDATETIME   *dateptr;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL.

dateptr
A pointer to the DBDATETIME variable to initialize.

Return value

SUCCEED or FAIL.

dbdatezero returns FAIL if *dateptr* is NULL.

Usage

- **dbdatezero** initializes a DBDATETIME value to Jan 1, 1900 12:00:00:000AM.
- The range of legal DBDATETIME values is from January 1, 1753 to December 31, 9999. DBDATETIME values have a precision of 1/300th of a second (3.33 milliseconds).

See also

dbdate4zero

dbdatlen

Description

Return the length of the data in a regular result column.

Syntax

```
DBINT dbdatlen(dbproc, column)

DBPROCESS   *dbproc;
int         column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column of interest. The first column is number 1.

Return value

The length, in bytes, of the data that would be returned for the particular column. If the data has a null value, **dbdatlen** returns 0. If the column number is not in range, **dbdatlen** returns -1.

Usage

- This routine returns the length, in bytes, of data that would be returned by a **select** against a regular (that is, non-compute) result column. In most cases, this is the actual length of data for the column. For text and image columns, however, the integer returned by **dbdatlen** can be less than the actual length of data for the column. This is because the server global variable *@@textsize* limits the amount of text or image data returned by a **select**.
- Use the **dbcollen** routine to determine the maximum possible length for the data. Use **dbdata** to get a pointer to the data itself.
- Here is a small program fragment that uses **dbdatlen**:

```
DBPROCESS    *dbproc;
 DBINT        row_number = 0;
 DBINT        data_length;

* Put the command into the command buffer */
 dbcmd(dbproc, "select name from sysobjects");
 /*
 ** Send the command to SQL Server and begin
 ** execution
 */
 dbsqlexec(dbproc);

 /* Process the command results */
 dbresults(dbproc);

 /* Examine the data lengths of each row */
 while (dbnextrow(dbproc) != NO_MORE_ROWS)
 {
      row_number++;
      data_length = dbdatlen(dbproc, 1);
      printf("row %ld, data length is %ld.\n",
           row_number, data_length);
 }
```

See also dbcollen, dbcolname, dbcoltype, dbdata, dbnumcols

dbdayname

Description Determine the name of a specified weekday in a specified language.

Syntax char *dbdayname(dbproc, language, daynum)

```
DBPROCESS   *dbproc;
char        *language;
int         daynum;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

language

The name of the desired language.

daynum
The number of the desired day. Day numbers range from 1 (Monday) to 7 (Sunday).

Return value
The name of the specified day on success; a NULL pointer on error.

Usage
- **dbdayname** returns the name of the specified day in the specified language. If *language* is NULL, *dbproc*'s current language is used. If both *language* and *dbproc* are NULL, then U.S. English is used.
- The following code fragment illustrates the use of **dbdayname**:

```
/*
 ** Retrieve the name of each day of the week in
 ** U.S. English.
 */
 for (daynum = 1; daynum <= 7; daynum++)
      printf("Day %d: %s\n", daynum,
      dbdayname((DBPROCESS *)NULL, (char *)NULL,
      daynum));
```

See also
db12hour, dbdateorder, dbmonthname, DBSETLNATLANG

DBDEAD

Description
Determine whether a particular DBPROCESS is dead.

Syntax
DBBOOL DBDEAD(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value
TRUE or FALSE.

Usage
- This macro indicates whether or not the specified DBPROCESS has been marked dead. It is particularly useful in user-supplied error handlers.

- If a DBPROCESS is dead, then almost every DB-Library routine that receives it as a parameter will immediately fail, calling the user-supplied error handler.

  Note If there is no user-supplied error handler, a dead DBPROCESS will cause the affected DB-Library routines not to fail, but to abort.

- Note that **DBDEAD** does not communicate with the server, but only checks the current status of a DBPROCESS. If a previously called DB-Library routine has not marked a DBPROCESS as dead, **DBDEAD** reports the DBPROCESS as healthy.

See also dberrhandle, Errors

dberrhandle

Description Install a user function to handle DB-Library errors.

Syntax
```
int (*dberrhandle(handler))()

 int        (*handler)();
```

Parameters handler

A pointer to the user function that will be called whenever DB-Library determines that an error has occurred. DB-Library calls this function with six parameters:

Table 2-17: Error handler parameters

Parameter	Meaning
dbproc	The affected DBPROCESS. If there is no DBPROCESS associated with this error, this parameter will be NULL.
severity	The severity of the error (datatype *int*). Error severities are defined in *syberror.h*.
dberr	The identifying number of the error (datatype *int*). Error numbers are defined in *sybdb.h*.
oserr	The operating-system-specific error number that describes the cause of the error (datatype int). If there is no relevant operating system error, the value of this parameter will be DBNOERR.
dberrstr	A printable description of *dberr* (datatype char *).
oserrstr	A printable description of *oserr* (datatype char *).

The error handler must return one of the following four values, directing DB-Library to perform particular actions:

Table 2-18: Error handler returns

Return	Action
INT_EXIT	Print an error message and abort the program. DB-Library will also return an error indication to the operating system. (Note to UNIX programmers: DB-Library will not leave a core file.
INT_CANCEL	Return FAIL from the DB-Library routine that caused the error. Returning INT_CANCEL on timeout errors will kill the *dbproc*.
INT_TIMEOUT	Cancel the operation that caused the error but leave the *dbproc* in working condition. This return value is meaningful only for timeout errors (SYBETIME). In any other case, this value will be considered an error, and will be treated as an INT_EXIT.
INT_CONTINUE	Continue to wait for one additional timeout period. At the end of that period, call the error handler again. This return value is meaningful only for timeout errors (SYBETIME). In any other case, this value will be considered an error, and will be treated as an INT_EXIT.

If the error handler returns any value besides these four, the program will abort.

Error handlers on the Windows platform must be declared with CS_PUBLIC, as shown in the example below. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well.

The following example shows a typical error handler routine:

```
#include <sybfront.h>
 #include <sybdb.h>
 #include <syberror.h>

int CS_PUBLIC err_handler(dbproc, severity, dberr,
oserr, dberrstr, oserrstr)
DBPROCESS    *dbproc;
 int          severity;
 int          dberr;
 int          oserr;
 char         *dberrstr;
 char         *oserrstr;
{
      if ((dbproc == NULL) || (DBDEAD(dbproc)))
           return(INT_EXIT);
      else
      {
           printf("DB-Library error:\n\t%s\n",
                dberrstr);
           if (oserr != DBNOERR)
                printf("Operating-system \
                     error:\n\t%s\n", oserrstr);
           return(INT_CANCEL);
      }
 }
```

Return value

A pointer to the previously installed error handler. This pointer is NULL if no error handler was installed before.

Usage

- **dberrhandle** installs an error-handler function that you supply. When a DB-Library error occurs, DB-Library will call this error handler immediately. You must install an error handler to handle DB-Library errors properly.
- If an application does not call **dberrhandle** to install an error-handler function, DB-Library ignores error messages. The messages are not printed.
- The user-supplied error handler will completely determine the response of DB-Library to any error that occurs. It must tell DB-Library whether to:
  - Abort the program, or
  - Return an error code and mark the DBPROCESS as “dead” (making it unusable), or
  - Cancel the operation that caused the error, or

- Keep trying (in the case of a timeout error).

- If the user does not supply an error handler (or passes a NULL pointer to **dberrhandle**), DB-Library will exhibit its default error-handling behavior: It will abort the program if the error has made the affected DBPROCESS unusable (the user can call **DBDEAD** to determine whether or not a DBPROCESS has become unusable). If the error has not made the DBPROCESS unusable, DB-Library will simply return an error code to its caller.
- You can "de-install" an existing error handler by calling **dberrhandle** with a NULL parameter. You can also, at any time, install a new error handler. The new handler will automatically replace any existing handler.
- If the program refers to error severity values, its source file must include the header file

  syberror.h.
- See "Errors" on page 386 for a list of DB-Library errors.
- Another routine, **dbmsghandle**, installs a message handler that DB-Library calls in response to the server error messages.
- If the application provokes messages from DB-Library and the server simultaneously, DB-Library calls the server message handler before it calls the DB-Library error handler.
- The DB-Library/C error value SYBESMSG is generated in response to a server error message, but not in response to a server informational message. This means that when a server error occurs, both the server message handler and the DB-Library/C error handler are called, but when the server generates an informational message, only the server message handler is called.

  If you have installed a server message handler, you may want to write your DB-Library error handler so as to suppress the printing of any SYBESMSG error, to avoid notifying the user about the same error twice.

  The following table provides information on when DB-Library/C calls an application's message and error handlers:

Table 2-19: Common errors

Error or message	Message handler called?	Error handler called?
SQL syntax error.	Yes.	Yes (SYBESMSG). (Code your handler to ignore the message.

Error or message	Message handler called?	Error handler called?
SQL **print** statement.	Yes.	No.
SQL **raiserror**.	Yes.	No.
Server dies.	No.	Yes (SYBESEOF). (Code your handler to exit the application.
Timeout from the server. **Note** The default timeout period is infinite. The error handler will not receive timeout notifications unless a timeout period is specified with dbsettime.	No.	Yes (SYBETIME). (To wait for another timeout period, code your handler to return -INT_CONTINUE.
Deadlock on query.	Yes. (Code your handler to test for deadlock. See the "dbsetuserdata" on page 333 for an example.	Yes (SYBESMSG). (Code your handler to ignore the message.
Timeout on login.	No.	Yes (SYBEFCON, SYBECONN).
Login fails (**dbopen**).	Yes.	Yes (SYBEPWD). (Code your handler to exit the application.
Use database message.	Yes. (Code your handler to ignore the message.	No .
Incorrect use of DB-Library/C calls, such as not calling **dbresults** when required.	No.	Yes (SYBERPND, ...) Yes (SYBERPND, .
Fatal Server error (severity greater than 16).	Yes. (Code your handler to exit the application.	Yes (SYBESMSG).

See also DBDEAD, dbmsghandle, Errors

dbexit

Description Close and deallocate all DBPROCESS structures, and clean up any structures initialized by **dbinit**.

Syntax void dbexit()

Parameters

None.

Return value

None.

Usage

- **dbexit** calls **dbclose** repeatedly for all allocated DBPROCESS structures. **dbclose** cleans up any activity associated with a single DBPROCESS structure and deallocates the space.
- You can use **dbclose** directly to close just a single DBPROCESS structure.
- **dbexit** also cleans up any structures initialized by **dbinit**, releasing the memory associated with those structures. It must be the last DB-Library call in any application that calls **dbinit**.
- To ensure future compatibility and portability, Sybase strongly recommends that all applications call **dbinit** and **dbexit**, no matter what their environment.

  For environments requiring **dbinit**, the application must not make any other DB-Library call after calling **dbexit**.

See also

dbclose, dbinit, dbopen

dbfcmd

Description

Add text to the DBPROCESS command buffer using C runtime library **sprintf**-type formatting.

Syntax

RETCODE dbfcmd(dbproc, cmdstring, args...)

```
DBPROCESS    *dbproc;
char         *cmdstring;
???          args...;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

cmdstring

A format string of the form used by the **sprintf** routine.

args

...

There is an optional and variable number of arguments to **dbfcmd**. The number and type of arguments required depends on the format specifiers included in the *cmdstring* argument. The arguments are passed directly to the C-library **sprintf** function. Neither **dbfcmd** nor the C compiler can type check these arguments. As with using **sprintf**, the programmer must ensure that each argument type matches the corresponding format specifier.

Return value

SUCCEED or FAIL.

Usage

- This routine adds text to the Transact-SQL command buffer in the DBPROCESS structure. **dbfcmd** works just like the **sprintf** function in the C language standard I/O library, using % conversion specifiers. If you do not need any of the formatting capability of **sprintf**, you can use **dbcmd** instead.
- The following table lists the conversions supported by **dbfcmd**:

Table 2-20: dbfcmd conversions

Conversion	Program variable type
%s	*char**, null-terminated
%d	*int*, decimal representation
%f	*double*
%g	*double*
%e	*double*
%%	None, the "%" character is written into the command buffer

The datatype SYBDATETIME must be converted to a character string and passed using %s. The datatype SYBMONEY may be converted to a character string and passed using %s, or converted to float and passed using %f.

Note Currently, only eight arguments may be handled in each call to **dbfcmd**. To format commands that require more than eight arguments, call **dbfcmd** repeatedly.

- **dbfcmd** manages the space allocation for the command buffer. It adds to the existing command buffer—it does not delete or overwrite the current contents except after the buffer has been sent to the server (see "Clearing the command buffer" on page 152). A single command buffer may contain multiple commands; in fact, this represents an efficient use of the command buffer.

- The application may call **dbfcmd** repeatedly. The command strings in sequential calls are just concatenated together. It is the program's responsibility to ensure that any necessary blanks appear between the end of one string and the beginning of the next.
- Here is a small program fragment that uses **dbfcmd** to build up a multiline SQL command:

```
char         *column_name;
 DBPROCESS    *dbproc;
 int          low_id;
 char         *object_type;
 char         *tablename;
dbfcmd(dbproc, "select %s from %s", column_name,
      tablename);
 dbfcmd(dbproc, " where id > %d", low_id);
 dbfcmd(dbproc, " and type='%s'", object_type);
```

  Note the required spaces at the start of the second and third command strings.

- When passing character or string variables to **dbfcmd**, beware of variables that contain quotes (single or double) or null characters (ASCII 0).
  - Improperly placed quotes in the SQL command can cause SQL syntax errors or, worse yet, unanticipated query results.
  - NULL characters (ASCII 0) should never be inserted into the command buffer. They can confuse DB-Library and the server, causing SQL syntax errors or unanticipated query results.
- Since **dbfcmd** calls **sprintf**, you must remember that % has a special meaning as the beginning of a format command. If you want to include % in the command string, you must precede it with another %.
- Be sure to guard against passing a null pointer as a string parameter to **dbfcmd**. If a null value is a possibility, you should check for it before using the variable in a **dbfcmd** call.
- The application can intermingle calls to **dbcmd** and **dbfcmd**.
- At any time, the application can access the contents of the command buffer through calls to **dbgetchar**, **dbstrlen**, and **dbstrcpy**.
- Available memory is the only constraint on the size of the DBPROCESS command buffer created by calls to **dbcmd** and **dbfcmd**.

Clearing the command buffer

After a call to **dbsqlexec** or **dbsqlsend**, the first call to either **dbcmd** or **dbfcmd** automatically clears the command buffer before the new text is entered. If this situation is undesirable, set the DBNOAUTOFREE option. When DBNOAUTOFREE is set, the command buffer is cleared only by an explicit call to **dbfreebuf**.

Limitations

Currently, only eight *args* may be handled in each call to **dbfcmd**. To format commands that require more than eight *args*, call **dbfcmd** repeatedly. On some platforms, **dbfcmd** may allow more than eight *args* per call. For portable code, do not pass more than 8 arguments.

Because it makes text substitutions, **dbfcmd** uses a working buffer in addition to the DBPROCESS command buffer. **dbfcmd** allocates this working buffer dynamically. The size of the space it allocates is equal to the maximum of a defined constant (1024) or the string length of *cmdstring* *2 . For example, if the length of *cmdstring* is 600 bytes, **dbfcmd** allocates a working buffer 1200 bytes long. If the length of *cmdstring* is 34 bytes, **dbfcmd** allocates a working buffer 1024 bytes long. To work around this limitation:

```
sprintf (buffer, "%s", SQL commmand");
dbcmd (dbproc, buffer)
```

If the *args* are very big in comparison to the size of *cmdstring*, the working buffer may not be large enough to hold the string after substitutions are made. In this situation, break *cmdstring* up and use multiple calls to **dbfcmd**.

Note that the working buffer is not the same as the DBPROCESS command buffer. The working buffer is a temporary buffer used only by **dbfcmd** when making text substitutions. The DBPROCESS command buffer holds the text after substitutions have been made. There is no constraint, other than available memory, on the size of the DBPROCESS command buffer.

See also dbcmd, dbfreebuf, dbgetchar, dbstrcpy, dbstrlen, Options

DBFIRSTROW

Description Return the number of the first row in the row buffer.

Syntax DBINT DBFIRSTROW(dbproc)

DBPROCESS *dbproc;

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of the first row in the row buffer. Rows are counted from the first row returned from the server, whose number is 1. This routine returns 0 if there is an error.

Usage

- This macro returns the number of the first row in the row buffer.
- If you are not buffering rows, DBFIRSTROW, DBCURROW, and DBLASTROW always have the same value. If you have enabled buffering by setting the DBBUFFER option, DBFIRSTROW returns the number of the first row in the row buffer.
- Note that the first row returned from the server (whose value is 1) is not necessarily the first row in the row buffer. The rows in the row buffer are dependent on manipulation by the application program. See the dbclrbuf reference page for details.

See also

dbclrbuf, DBCURROW, dbgetrow, DBLASTROW, dbnextrow, dbsetopt, Options

dbfree_xlate

Description

Free a pair of character set translation tables.

Syntax

```
RETCODE *dbfree_xlate(dbproc, xlt_tosrv, xlt_todisp)

DBPROCESS   *dbproc;
DBXLATE     *xlt_tosrv;
DBXLATE     *xlt_todisp;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

xlt_tosrv

A pointer to a translation table used to translate display-specific character strings to the server character strings. The translation table is allocated via **dbload_xlate**.

xlt_todisp
A pointer to a translation table used to translate server character strings to display-specific character strings. The translation table is allocated via **dbload_xlate**.

Return value
SUCCEED or FAIL.

Usage
- This routine frees a pair of character set translation tables allocated by **dbload_xlate**.
- Character set translation tables translate characters between the server's standard character set and the display device's character set.
- The following code fragment illustrates the use of **dbfree_xlate**

```
char        destbuf[128];
 int         srcbytes_used;
 DBXLATE     *xlt_todisp;  DBXLATE     *xlt_tosrv;
dbload_xlate((DBPROCESS *)NULL, "iso_1",
"trans.xlt", &xlt_tosrv, &xlt_todisp);
printf("Original string: \n\t%s\n\n",
      TEST_STRING);
dbxlate((DBPROCESS *)NULL, TEST_STRING,
     strlen(TEST_STRING), destbuf, -1, xlt_todisp,
      &srcbytes_used);
printf("Translated to display character set: \
      \n\t%s\n\n", destbuf);
dbfree_xlate((DBPROCESS *)NULL, xlt_tosrv,
    xlt_todisp);
```

See also
dbload_xlate, dbxlate

dbfreebuf

Description
Clear the command buffer.

Syntax
void dbfreebuf(dbproc)

DBPROCESS *dbproc;

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value — None.

Usage

- This routine clears a DBPROCESS command buffer by freeing any space allocated to it. It then sets the command buffer to NULL. Commands are added to the command buffer with the **dbcmd** or **dbfcmd** routine.
- After a call to **dbsqlexec** or **dbsqlsend**, the first call to either **dbcmd** or **dbfcmd** automatically calls **dbfreebuf** to clear the command buffer before the new text is entered. If this situation is undesirable, set the DBNOAUTOFREE option. When DBNOAUTOFREE is set, the command buffer is cleared only by an explicit call to **dbfreebuf**.
- At any time, the application can access the contents of the command buffer through calls to **dbgetchar**, **dbstrlen**, and **dbstrcpy**.

See also — dbcmd, dbfcmd, dbgetchar, dbsqlexec, dbsqlsend, dbstrcpy, dbstrlen, Options

dbfreequal

Description — Free the memory allocated by **dbqual**.

Syntax

```
void dbfreequal(qualptr)

 char       *qualptr;
```

Parameters

qualptr
A pointer to the memory allocated by **dbqual**.

Return value — None.

Usage

- **dbfreequal** is one of the DB-Library browse mode routines. See Chapter 1, “Introducing DB-Library” for a detailed discussion of browse mode.
- **dbqual** provides a **where** clause that the application can use to update a single row in a browsable table. In doing so, it dynamically allocates a buffer to contain the **where** clause. When the **where** clause is no longer needed, the application can use **dbfreequal** to deallocate the buffer.

See also — dbqual

dbfreesort

Description — Free a sort order structure allocated by **dbloadsort**.

Syntax

RETCODE dbfreesort(dbproc, sortorder)

```
DBPROCESS       *dbproc;
DBSORTORDER     *sortorder;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

sortorder

A pointer to a DBSORTORDER structure allocated via **dbloadsort**.

Return value

SUCCEED or FAIL.

Usage

- **dbfreesort** frees a sort order structure that was allocated via **dbloadsort**. DB-Library routines such as **dbstrcmp** and **dbstrsort** use sort orders to determine how character data must be sorted.
- When an application program does sorting or comparing, it automatically sorts character data the same way the server does. If no sort order has been loaded, routines such as **dbstrcmp** and **dbstrsort** sort characters by their binary values.

Warning! Application programs must not attempt to use operating-system facilities to free the **sortorder* structure directly, as it may have been allocated using some mechanism other than **malloc** (on operating systems where **malloc** is not supported), and it may consist of multiple parts, some of which must be freed separately.

- The following code fragment illustrates the use of **dbfreesort**:

```
sortorder = dbloadsort(dbproc);

 retval = dbstrcmp(dbproc, "ABC", 3, "abc", 3,
      sortorder);
 printf("ABC dbstrcmp'ed with abc yields %d.\n",
      retval);

 retval = dbstrcmp(dbproc, "abc", 3, "ABC", 3,
      sortorder);
 printf("abc dbstrcmp'ed with ABC yields %d.\n",
      retval);

 dbfreesort(dbproc, sortorder);
```

See also

dbloadsort, dbstrcmp, dbstrsort

dbgetchar

Description

Return a pointer to a character in the command buffer.

Syntax

```
char *dbgetchar(dbproc, n)

DBPROCESS    *dbproc;
int          n;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

n

The position of the desired character in the command buffer. The first character position is 0.

Return value

dbgetchar returns a pointer to the *n*th character in the command buffer. If *n* is not in range, **dbgetchar** returns NULL.

Usage

- You can use **dbgetchar** to retrieve a pointer to a particular character in the command buffer. **dbgetchar** returns a pointer to a character in the command buffer whose position is indicated by *n*. The first character has position 0.
- Internally, the command buffer is a linked list of non-null-terminated text strings. **dbgetchar**, **dbstrcpy**, and **dbstrlen** together provide a way to locate and copy parts of the command buffer.
- Since the command buffer is not just one large text string, but rather a linked list of text strings, you must use **dbgetchar** to index through the buffer. If you just get a pointer using **dbgetchar** and then increment it yourself, it will probably fall off the end of a string and cause a segmentation fault.

See also

dbcmd, dbfcmd, dbfreebuf, dbstrcpy, dbstrlen

dbgetcharset

Description

Get the name of the client character set from the DBPROCESS structure.

Syntax

```
char *dbgetcharset(dbproc)

DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and server.

Return value

A pointer to the null-terminated name of the client character set, or NULL in case of error.

Usage

- **dbgetcharset** returns the name of the client's character set.
- DB-Library/C clients can use a different character set than the server or servers to which they are connected. If a client and server are using different character sets, and the server supports character translation for the client's character set, it will perform all conversions to and from its own character set when communicating with the client.
- An application can inform the server what character set it is using via DBSETLCHARSET.
- To determine if the server is performing character set translations, an application can call **dbcharsetconv**.
- To get the name of the server character set, an application can call **dbservcharset**.

See also

dbcharsetconv, dblogin, dbopen, dbservcharset, DBSETLCHARSET

dbgetloginfo

Description

Transfer TDS (Tabular Data Stream) login response information from a DBPROCESS structure to a newly allocated DBLOGINFO structure.

Syntax

RETCODE dbgetloginfo(dbproc, loginfo)

```
DBPROCESS    *dbproc;
DBLOGINFO    **loginfo;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

loginfo
The address of a DBLOGINFO pointer variable. **dbgetloginfo** sets the DBLOGINFO pointer to the address of a newly-allocated DBLOGINFO structure.

Return value

SUCCEED or FAIL.

Usage

- **dbgetloginfo** transfers TDS login response information from a DBPROCESS structure to a newly allocated DBLOGINFO structure.
- An application needs to call **dbgetloginfo** only if 1) it is an Open Server gateway application, and 2) it is using TDS passthrough.
- TDS is an application protocol used for the transfer of requests and request results between clients and servers.
- When a client connects directly to a server, the two programs negotiate the TDS format they will use to send and receive data. When a gateway application uses TDS passthrough, the application forwards TDS packets between the client and a remote server without examining or processing them. For this reason, the remote server and the client must agree on a TDS format to use.
- **dbgetloginfo** is the second of four calls, two of them Server Library calls, that allow a client and remote server to negotiate a TDS format. The calls, which can be made only in a SRV_CONNECT event handler, are:
  - **srv_getloginfo** - Allocate a DBLOGINFO structure and fill it with TDS information from a client SRV_PROC.
  - **dbsetloginfo** - Transfer the TDS information retrieved in step 1 from the DBLOGINFO structure to a DB-Library/C LOGINREC structure, and then free the DBLOGINFO structure. After the information is transferred, the application can use this LOGINREC structure in the **dbopen** call which establishes its connection with the remote server.
  - **dbgetloginfo** - Transfer the remote server's response to the client's TDS information from a DBPROCESS structure into a newly-allocated DBLOGINFO structure.
  - **srv_setloginfo** - Send the remote server's response, retrieved in the previous step, to the client, and then free the DBLOGINFO structure.
- This is an example of a SRV_CONNECT handler preparing a remote connection for TDS passthrough:

```
RETCODE connect_handler(srvproc)
 SRVPROC       *srvproc;
 {
```

```
 DBLOGINFO    *loginfo;
 LOGINREC     *loginrec;
 DBPROCESS     *dbproc;
/*
 ** Get the TDS login information from the client
 ** SRV_PROC.
 */
 srv_getloginfo(srvproc, &loginfo);
/* Get a LOGINREC structure */
 loginrec = dblogin();
/*
 ** Initialize the LOGINREC with the login info
 ** from the SRV_PROC.
 */
 dbsetloginfo(loginrec, loginfo);
/* Connect to the remote server */
 dbproc = dbopen(loginrec, REMOTE_SERVER_NAME);
/*
 ** Get the TDS login response information from
 ** the remote connection.
 */
 dbgetloginfo(dbproc, &loginfo);
/*
 ** Return the login response information to the
 ** SRV_PROC.
 */
 srv_setloginfo(srvproc, loginfo);
/* Accept the connection and return */
 srv_senddone(srvproc, 0, 0, 0);
 return(SRV_CONTINUE);
 }
```

See also dbrecvpassthru, dbsendpassthru, dbsetloginfo

dbgetlusername

Description Return the user name from a LOGINREC structure.

Syntax int dbgetlusername(login, name_buffer, buffer_len)

```
LOGINREC      *login;
BYTE          *name_buffer;
int           buffer_len;
```

Parameters

login
A pointer to a LOGINREC structure, which can be passed as an argument to **dbopen**. You can get a LOGINREC structure by calling **dblogin**.

name_buffer
A pointer to a buffer. The user name will be copied from the LOGINREC structure to this buffer.

buffer_len
The length, in bytes, of the destination buffer.

Return value

The number of bytes copied into the destination buffer, not including the null-terminator.

If the user name is more than *buffer_len* -1 bytes long, **dbgetlusername** copies *buffer_len* -1 bytes into the destination buffer and returns DBTRUNCATED.

dbgetlusername returns FAIL if *login* is NULL, *name_buffer* is NULL, or *buffer_len* is less than 0.

Usage

- **dbgetlusername** copies the user name from LOGINREC structure into the *name_buffer* buffer.
- To set the user name in a LOGINREC structure, use DBSETLUSER.
- **dbgetlusername** copies a maximum of *buffer_len* -1 bytes, and null-terminates the user name string. Since the longest user name in a LOGINREC structure is DBMAXNAME bytes, an application will never need a destination buffer longer than DBMAXNAME +1 bytes.
- If the user name is in the LOGINREC is longer than *buffer_len* -1 bytes, **dbgetlusername** truncates the name and returns DBTRUNCATED.

See also

dblogin, DBSETLUSER

dbgetmaxprocs

Description

Determine the current maximum number of simultaneously open DBPROCESSes.

Syntax

int dbgetmaxprocs()

Parameters

None.

Return value

An integer representing the current limit on the number of simultaneously open DBPROCESSes.

Usage

- A DB-Library program has a maximum number of simultaneously open DBPROCESSes. By default, this number is 25. The application program may change this limit by calling **dbsetmaxprocs**.

See also

dbopen, dbsetmaxprocs

dbgetnatlang

Description

Get the national language from the DBPROCESS structure.

Syntax

```
char* dbgetnatlang(dbproc)

DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and server.

Return value

A pointer to a character string representing the national language that the client DBPROCESS is using.

Usage

- **dbgetnatlang** returns a pointer to the name of the national language that a client is using.
- DB-Library/C clients may use a different national language than the server or servers to which they are connected. An application can inform the server what national language it wishes to use via DBSETLNATLANG.

See also

dblogin, dbopen, DBSETLNATLANG

dbgetoff

Description

Check for the existence of Transact-SQL constructs in the command buffer.

Syntax

```
int dbgetoff(dbproc, offtype, startfrom)

DBPROCESS    *dbproc;
DBUSMALLINT  offtype;
int                startfrom;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

offtype

The type of offset you want to find. The types, which are defined in the header file *sybdb.h*, are:

OFF_SELECT
OFF_FROM
OFF_ORDER
OFF_COMPUTE
OFF_TABLE
OFF_PROCEDURE
OFF_STATEMENT
OFF_PARAM
OFF_EXEC

See "Options" on page 402 for details.

startfrom

The point in the buffer at which to start looking. The command buffer begins at 0.

Return value

The character offset into the command buffer for the specified offset. If the offset is not found, -1 is returned.

Usage

- If the DBOFFSET option has been set (see Options on page 402), this routine can check for the location of certain Transact-SQL constructs in the command buffer. As a simple example, assume the program does not know the contents of the command buffer but needs to know where the SQL keyword **select** appears:

```
int     select_offset[10];
 int     last_offset;
 int     i;
/* Set the offset option */
 dbsetopt(dbproc, DBOFFSET, "select");

/*
 ** Assume the command buffer contains the
 ** following selects.
 */
 dbcmd(dbproc, "select x = 100 select y = 5");
```

```
/* Send the query to SQL Server */
 dbsqlexec(dbproc);

/* Get all the offsets to the select keyword */
 for (i = 0, last_offset = 0; last_offset != -1;
      i++)
      if ((last_offset = dbgetoff(dbproc,
           OFF_SELECT, last_offset) != -1)
      select_offset[i] = last_offset++;
```

In this example, select_offset[0] = 0 and select_offset[1] = 15.

- **dbgetoff** does not recognize **select** statements in a subquery. Thus, if the command buffer contained:

```
select pub_name
 from publishers
 where pub_id not in
    (select pub_id
     from titles
     where type = "business")
```

the second “select” would not be recognized.

See also dbcmd, dbgetchar, dbsetopt, dbstrcpy, dbstrlen, Options

dbgetpacket

Description Return the TDS packet size currently in use.

Syntax int dbgetpacket(dbproc)

DBPROCESS *dbproc;

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

Return value The TDS packet size currently in use.

Usage

- **dbgetpacket** returns the TDS packet size currently in use.
- TDS (Tabular Data Stream) is an application protocol used for the transfer of requests and request results between clients and servers.

- TDS data is sent in fixed-size chunks, called “packets”. TDS packets have a default size of 512 bytes.
- An application may change the TDS packet size via **DBSETLPACKET**, which sets the packet size field in the LOGINREC structure. When the application logs in to the server or Open Server, the server sets the TDS packet size for the created DBPROCESS connection to be equal to or less than the value of this field. The packet size is set to a value less than the value of the field if the server is experiencing space constraints. Otherwise, the packet size will be equal to the value of the field.
- If an application sends or receives large amounts of text or image data, a packet size larger than the default 512 bytes may improve efficiency, since it results in fewer network reads and writes.

See also DBSETLPACKET

dbgetrow

Description Read the specified row in the row buffer.

Syntax STATUS dbgetrow(dbproc, row)

```
DBPROCESS    *dbproc;
DBINT        row;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

row

The number of the row to read. Rows are counted from the first row returned from the server, whose number is 1. Note that the first row in the row buffer is not necessarily the first row returned from the server.

Return value **dbgetrow** can return four different types of values:

- If the current row is a regular row, REG_ROW is returned.
- If the current row is a compute row, the *computeid* of the row is returned. (See the dbaltbind reference page for information on the *computeid*.

- If the row is not in the row buffer, NO_MORE_ROWS is returned and the current row is left unchanged.
- If the routine was unsuccessful, FAIL is returned.

Usage

- **dbgetrow** sets the current row in the row buffer to a specific row and reads it. This routine works only if the DBBUFFER option is on, enabling row buffering. When **dbgetrow** is called, any binding of row data to program variables (as specified with **dbbind** or **dbaltbind**) takes effect.
- Row buffering provides a way to keep a specified number of server result rows in program memory. Without row buffering, the result row generated by each new **dbnextrow** call overwrites the contents of the previous result row. Row buffering is therefore useful for programs that need to look at result rows in a non-sequential manner. It does, however, carry a memory and performance penalty because each row in the buffer must be allocated and freed individually. Therefore, use it only if you need to. Specifically, the application should only turn the DBBUFFER option on if it calls **dbgetrow** or **dbsetrow**. Note that row buffering has nothing to do with network buffering and is a completely independent issue.
- When row buffering is not enabled, the application processes each row as it is read from the server, by calling **dbnextrow** repeatedly until it returns NO_MORE_ROWS. When row buffering is enabled, the application can use **dbgetrow** to jump to any row that has already been read from the server with **dbnextrow**. Subsequent calls to **dbnextrow** cause the application to read successive rows in the buffer. When **dbnextrow** reaches the last row in the buffer, it reads rows from the server again, if there are any. Once the buffer is full, **dbnextrow** does not read any more rows from the server until some of the rows have been cleared from the buffer with **dbclrbuf**.
- The macros DBFIRSTROW, DBLASTROW, and DBCURROW are useful in conjunction with **dbgetrow** calls. DBFIRSTROW, for instance, gets the number of the first row in the buffer. Thus, the call:

  ```
  dbgetrow(dbproc, DBFIRSTROW(dbproc))
  ```

  sets the current row to the first row in the buffer.
- The routine **dbsetrow** sets a buffered row to “current” but does not read the row.
- For an example of row buffering, see Example 4 of the online sample programs.

See also dbaltbind, dbbind, dbclrbuf, DBCURROW, DBFIRSTROW, DBLASTROW, dbnextrow, dbsetrow, Options

DBGETTIME

Description Return the number of seconds that DB-Library will wait for a server response to a SQL command.

Syntax int DBGETTIME()

Parameters None.

Return value The timeout value—the number of seconds that DB-Library waits for a server response before timing out. A timeout value of 0 represents an infinite timeout period.

Usage
- This routine returns the length of time in seconds that DB-Library will wait for a server response during calls to **dbsqlexec**, **dbsqlok**, **dbresults**, and **dbnextrow**. The default timeout value is 0, which represents an infinite timeout period.
- The program can call **dbsettime** to change the timeout value.

See also dbsettime

dbgetuserdata

Description Return a pointer to user-allocated data from a DBPROCESS structure.

Syntax
```
BYTE *dbgetuserdata(dbproc)

 DBPROCESS     *dbproc;
```

Parameters dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value A generic BYTE pointer to the user's private data space. This pointer must have been previously saved with the **dbsetuserdata** routine.

Usage

- This routine returns, from a DBPROCESS structure, a pointer to user-allocated data. The application must have previously saved this pointer with the **dbsetuserdata** routine.
- **dbgetuserdata** and **dbsetuserdata** allow the application to associate user data with a particular DBPROCESS. This avoids the necessity of using global variables for this purpose. One use for these routines is to handle deadlock, as shown in the example on the dbsetuserdata reference page. That example reruns the transaction when the application's message handler detects deadlock.
- This routine is particularly useful when the application has multiple DBPROCESSes.

See also dbsetuserdata

dbhasretstat

Description Determine whether the current command or remote procedure call generated a return status number.

Syntax

```
DBBOOL dbhasretstat(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value TRUE or FALSE.

Usage

- This routine determines whether the current Transact-SQL command or remote procedure call generated a return status number. Status numbers are returned by all stored procedures running on SQL Server version 4.0 or later. Since status numbers are a feature of stored procedures, only a remote procedure call or an **execute** command can generate a status number.
- The **dbretstatus** routine actually gets the status number. Stored procedures that complete normally return a status number of 0. For a list of return status numbers, see the *Adaptive Server Enterprise Reference Manual*.

- When executing a stored procedure, the server returns the status number immediately after returning all other results. Therefore, the application can call **dbhasretstat** only after processing the stored procedure's results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbhasretstat** or **dbretstatus**, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.
- The order in which the application processes the status number and any return parameter values is unimportant.
- When a stored procedure has been executed as an RPC command using **dbrpcinit**, **dbrpcparam**, and **dbrpcsend**, then the return status can be retrieved after all other results have been processed. For an example of this usage, see Example 8 of the online sample programs.
- When a stored procedure has been executed from a batch of Transact-SQL commands (with **dbsqlexec** or **dbsqlsend**), then other commands might execute after the stored procedure. This situation makes return-status retrieval a little more complicated.
  - If you are sure that the stored procedure command is the only command in the batch, then you can retrieve the return status after the **dbresults** loop, as shown in Example 8 of the online sample programs.
  - If the batch can contain multiple commands, then the return status should be retrieved inside the **dbresults** loop, after all rows have been fetched with **dbnextrow**. The code below shows the program logic to retrieve the return status value in this situation.

```
while ( (result_code = dbresults(dbproc)
          != NO_MORE_RESULTS)
 {
   if (result_code == SUCCEED)
   {
     ... bind rows here ...
     while ((row_code = dbnextrow(dbproc))
            != NO_MORE_ROWS)
     {
       ... process rows here ...
     }
     /* Now check for a return status */
     if (dbhasretstat(dbproc) == TRUE)
     {
       printf("(return status %d)\n",
```

```
                    dbretstatus(dbproc));
            }
            if (dbnumrets(dbproc) > 0)
            {
              ... get output parameters here ...
            }
          } /* if result_code */
          else
          {
            printf("Query failed.\n");
          }
        } /* while dbresults */
```

See also dbnextrow, dbresults, dbretdata, dbretstatus, dbrpcinit, dbrpcparam, dbrpcsend

dbinit

Description Initialize DB-Library.

Syntax RETCODE dbinit()

Parameters

None.

Return value SUCCEED or FAIL.

Usage

- This routine initializes certain private DB-Library structures. For environments that require it, the application must call **dbinit** before calling any other DB-Library routine. Most DB-Library routines will cause the application to exit if they are called before **dbinit**.
- To ensure future compatibility and portability, Sybase strongly recommends that all applications call **dbinit**, no matter what their operating environment.

See also dbexit

DBIORDESC

Description (UNIX and AOS/VS only) Provide program access to the UNIX or AOS/VS file descriptor used by a DBPROCESS to read data coming from the server.

Syntax

```
int DBIORDESC(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

An integer file descriptor used by the specified DBPROCESS to read data coming from the server.

Usage

- This routine provides a way for an application to respond effectively to multiple input streams. Depending on the nature of your application, the time between a request for information from the server (usually made via a call to **dbsqlsend**) and the server's response (read by calling **dbsqlok**, **dbresults**, or **dbnextrow**) may be significant. You may use this time to service other parts of your application. The DBIORDESC routine provides a way to obtain the I/O descriptor that a DBPROCESS uses to read the data stream from the server. This information may then be used with various operating system facilities (such as the UNIX **select** call) to allow the application to respond effectively to multiple input streams.
- **dbpoll** checks if a server response has arrived for any of an application's server connections (represented by DBPROCESS pointers). **dbpoll** is generally simpler to use than DBIORDESC. For this reason, and because DBIORDESC is non-portable, it is generally preferable to use **dbpoll**.
- The file descriptor returned by DBIORDESC may only be used with operating system facilities that *do not* read data from the incoming data stream. If data is read from this stream by any means other than through a DB-Library routine, communications between the front end and the server will become hopelessly scrambled.
- An application can use the DB-Library DBRBUF routine, in addition to the UNIX **select** function or AOS/VS equivalent, to help determine whether any more data from the server is available for reading.
- A companion routine, DBIOWDESC, provides access to the file descriptor used to write data to the server.

See also

dbcmd, DBIOWDESC, dbnextrow, dbpoll, DBRBUF, dbresults, dbsqlok, dbsqlsend

DBIOWDESC

Description

(UNIX and AOS/VS only) Provide program access to the UNIX or AOS/VS file descriptor used by a DBPROCESS to write data to the server.

Syntax

```
int DBIOWDESC(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

An integer file descriptor used by the specified DBPROCESS to write data to the server.

Usage

- This routine provides a way for an application to effectively utilize multiple input and output streams. Depending on the nature of your application, the time interval between the initiation of an attempt to write information to the server (usually made via a call to **dbsqlsend**) and the completion of that attempt may be significant. You may use this time to service other parts of your application. The DBIOWDESC routine provides a way to obtain the I/O descriptor that a DBPROCESS uses to write the data stream to the server. This information may then be used with various operating system facilities (such as the UNIX **select** function) to allow the application to effectively utilize multiple input and output streams.
- The file descriptor returned by this routine may only be used with operating system facilities that *do not* write data to the outgoing data stream. If data is written to this stream by any means other than through a DB-Library routine, communications between the front-end and the server will become hopelessly scrambled.
- A companion routine, DBIORDESC, provides access to the file descriptor used to read data coming from the server. For some applications, another routine, **dbpoll** may be preferable to DBIORDESC.

See also

dbcmd, DBIORDESC, dbnextrow, dbpoll, dbresults, dbsqlok, dbsqlsend

DBISAVAIL

Description
Determine whether a DBPROCESS is available for general use.

Syntax
```
DBBOOL DBISAVAIL(dbproc)

 DBPROCESS     *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value
TRUE if the DBPROCESS is available for general use, otherwise FALSE.

Usage
- This routine indicates whether the specified DBPROCESS is available for general use. When a DBPROCESS is first opened, it is marked as being available, until some use is made of it. Many DB-Library routines will automatically set the DBPROCESS to "not available," but only **dbsetavail** will reset it to "available." This facility is useful when several parts of a program are attempting to share a single DBPROCESS.

See also
dbsetavail

dbisopt

Description
Check the status of a server or DB-Library option.

Syntax
```
DBBOOL dbisopt(dbproc, option, param)

 DBPROCESS    *dbproc;
 int           option;
 char         *param;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server. Unlike in the functions **dbsetopt** and **dbclropt**, *dbproc* cannot be NULL here.

option
The option to be checked. See "Options" on page 402 for the list of options.

param

Certain options take parameters. The DBOFFSET option, for example, takes as a parameter the SQL construct for which offsets are to be returned. Options lists those options that take parameters. If an option does not take a parameter, *param* must be NULL.

If the option you are checking takes a parameter but there can be only one instance of the option, **dbisopt** ignores the *param* argument. For example, **dbisopt** ignores the value of *param* when checking the DBBUFFER option, because row buffering can have only one setting at a time. On the other hand, the DBOFFSET option can have several settings, each with a different parameter. It may have been set twice—to look for offsets to **select** statements and offsets to **order by** clauses. In that case, **dbisopt** needs the *param* argument to determine whether to check the **select** offset or the **order by** offset.

Return value

TRUE or FALSE.

Usage

- This routine checks the status of the server and DB-Library options. Although server options may be set and cleared directly through SQL, the application should instead use **dbsetopt** and **dbclropt** to set and clear options. This provides a uniform interface for setting both server and DB-Library options. It also allows the application to use the **dbisopt** function to check the status of an option.
- For a list of each option and its default status, see Options on page 402.

See also

dbclropt, dbsetopt, Options

DBLASTROW

Description

Return the number of the last row in the row buffer.

Syntax

```
DBINT DBLASTROW(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of the last row in the row buffer. This routine returns 0 if there is an error.

Usage

- This macro returns the number of the last row in the row buffer. Rows are counted from the first row returned from the server, whose number is 1, and *not* from the top of the row buffer.
- If you are not buffering rows, DBFIRSTROW, DBCURROW, and DBLASTROW will always have the same value. If you have enabled buffering by setting the DBBUFFER option, DBLASTROW will return the number of the row that is the last row in the row buffer.

See also

dbclrbuf, DBCURROW, DBFIRSTROW, dbgetrow, dbnextrow, dbsetopt, Options

dbload_xlate

Description

Load a pair of character set translation tables.

Syntax

```
RETCODE dbload_xlate(dbproc, srv_charset, xlate_name,
            xlt_tosrv, xlt_todisp)

DBPROCESS    *dbproc;
char         *srv_charset;
char         *xlt_name;
DBXLATE      **xlt_tosrv;
DBXLATE      **xlt_todisp;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

srv_charset

A pointer to the name of the server's character set. **dbload_xlate** looks for a directory of this name in the *charsets* directory under the main Sybase installation directory. For example, if the server is using the iso_1 character set, **dbload_xlate** looks for *$SYBASE/charsets/iso_1*.

xlt_name

A pointer to the name of the file containing the display-specific character set. **dbload_xlate** looks for this file in the server character set directory.

xlt_tosrv
A pointer to a pointer to a character set translation table used to translate display-specific character strings to the server character strings. The translation table is allocated via **dbload_xlate**.

xlt_todisp
A pointer to a pointer to a character set translation table used to translate server character strings to display-specific character strings. The translation table is allocated via **dbload_xlate**.

Return value

SUCCEED or FAIL.

Usage

- **dbload_xlate** reads a display-specific localization file and allocates two character set translation tables: one for translations from the server's character set to the display-specific character set, and another for translations from the display-specific character set to the server's character set.
- The following code fragment illustrates the use of **dbload_xlate**:

```
char        destbuf[128];
 int         srcbytes_used;
 DBXLATE*    xlt_todisp;
 DBXLATE     *xlt_tosrv;

dbload_xlate((DBPROCESS *)NULL, "iso_1",
      "trans.xlt", &xlt-tosrv, &xlt-todisp);
printf("Original string: \n\t%s\n\n",
       TEST_STRING);
dbxlate((DBPROCESS *)NULL, TEST_STRING,
     strlen(TEST_STRING), destbuf, -1, xlt_todisp,
      &srcbytes_used);
printf("Translated to display character set: \
      \n\t%s\n\n", destbuf);
dbfree_xlate((DBPROCESS *)NULL, xlt_tosrv,
      xlt_todisp);
```

See also

dbfree_xlate, dbxlate

dbloadsort

Description

Load a server sort order.

Syntax

DBSORTORDER *dbloadsort(dbproc)

DBPROCESS *dbproc;

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

A pointer to a DBSORTORDER structure on success, NULL on error.

Usage

- **dbloadsort** provides information about the sort order of the server's character set. This information can be used by **dbstrcmp** or **dbstrsort** to compare two character strings.
- **dbloadsort** allocates a DBSORTORDER structure to contain the server character set sort order information. The structure is freed via **dbfreesort**.
- The following code fragment illustrates the use of **dbloadsort**:

```
sortorder = dbloadsort(dbproc);
retval = dbstrcmp(dbproc, "ABC", 3, "abc", 3,
      sortorder);
 printf("ABC dbstrcmp'ed with abc yields %d.\n",
      retval);
retval = dbstrcmp(dbproc, "abc", 3, "ABC", 3,
      sortorder);
 printf("abc dbstrcmp'ed with ABC yields %d.\n",
      retval);
dbfreesort(dbproc, sortorder);
```

See also

dbfreesort, dbstrcmp, dbstrsort

dblogin

Description

Allocate a login record for use in **dbopen**.

Syntax

LOGINREC *dblogin()

Parameters

None.

Return value

A pointer to a LOGINREC structure. **dblogin** returns NULL if the structure could not be allocated.

Usage

- This routine allocates a LOGINREC structure for use with **dbopen**.

- There are various routines available to supply components of the LOGINREC. The program may supply the host name, user name, user password, and application name—via **DBSETLHOST**, **DBSETLUSER**, **DBSETLPWD**, and **DBSETAPP**, respectively. It is generally only necessary for the program to supply the user password (and even this can be eliminated if the password is a null value). The other variables in the LOGINREC structure will be set to default values.
- Other components of the LOGINREC may also be changed:
  - The national language name can be set in a LOGINREC structure via **DBSETLNATLANG**. Call **DBSETLNATLANG** only if you do not wish to use the server's default national language.
  - The TDS packet size can be set in a LOGINREC via **DBSETLPACKET**. If not explicitly set, the TDS packet size defaults to 512 bytes. TDS is an application protocol used for the exchange of information between clients and servers.
  - The character set can be set in a LOGINREC via **DBSETLCHARSET**. An application needs to call **DBSETLCHARSET** only if it is not using ISO-8859-1 (known to the server as "iso_1").
- When a connection attempt is made between a client and a server, there are two ways in which the connection can fail (assuming that the system is correctly configured):
  - The machine that the server is supposed to be on is running correctly and the network is running correctly.

    In this case, if there is no server listening on the specified port, the machine the server is supposed to be on will signal the client, via a network error, that the connection cannot be formed. Regardless of **dbsetlogintime**, the connection will fail.
  - The machine that the server is on is down.

    In this case, the machine that the server is supposed to be on will not respond. Because "no response" is not considered to be an error, the network will not signal the client that an error has occurred. However, if **dbsetlogintime** has been called to set a timeout period, a timeout error will occur when the client fails to receive a response within the set period.
- Here is a program fragment that uses **dblogin**:

```
DBPROCESS    *dbproc;
 LOGINREC     *loginrec;
```

```
loginrec = dblogin();
 DBSETLPWD(loginrec, "server_password");
 DBSETLAPP(loginrec, "my_program");
 dbproc = dbopen(loginrec, "my_server");
```

- Once the application has made all its **dbopen** calls, the LOGINREC structure is no longer necessary. The program can then call **dbloginfree** to free the LOGINREC structure.

See also

dbloginfree, dbopen, dbrpwclr, dbrpwset, DBSETLAPP, DBSETLCHARSET, DBSETLHOST, DBSETLNATLANG, DBSETLPACKET, DBSETLPWD, DBSETLUSER

dbloginfree

Description

Free a login record.

Syntax

void dbloginfree(loginptr)

```
LOGINREC    *loginptr;
```

Parameters

loginptr
A pointer to a LOGINREC structure.

Return value

None.

Usage

- **dblogin** provides a LOGINREC structure for use with **dbopen**. Once the application has made all its **dbopen** calls, the LOGINREC structure is no longer necessary. **dbloginfree** frees the memory associated with the specified LOGINREC structure.

See also

dblogin, dbopen

dbmny4add

Description

Add two DBMONEY4 values.

Syntax

RETCODE dbmny4add(dbproc, m1, m2, sum)

```
DBPROCESS   *dbproc;
DBMONEY4    *m1;
DBMONEY4    *m2;
DBMONEY4    *sum;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1
A pointer to a DBMONEY4 value.

m2
A pointer to a DBMONEY4 value.

sum
A pointer to a DBMONEY4 variable to hold the result of the addition.

Return value

SUCCEED or FAIL.

dbmny4add returns FAIL in case of overflow, or if *m1*, *m2*, or *sum* is NULL.

Usage

- **dbmny4add** adds the *m1* and *m2* DBMONEY4 values and places the result in **sum*.
- In case of overflow, **dbmny4add** returns FAIL and sets **sum* to $0.00.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmny4sub, dbmny4mul, dbmny4divide, dbmny4minus, dbmny4add, dbmny4sub, dbmnymul, dbmnydivide, dbmnyminus

dbmny4cmp

Description

Compare two DBMONEY4 values.

Syntax

int dbmny4cmp(dbproc, m1, m2)

```
DBPROCESS   *dbproc;
DBMONEY4    *m1;
DBMONEY4    *m2;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1

A pointer to a DBMONEY4 value.

m2

A pointer to a DBMONEY4 value.

Return value

If *m1 = m2*, **dbmny4cmp** returns 0.

If *m1 < m2*, **dbmny4cmp** returns -1.

If *m1 > m2*, **dbmny4cmp** returns 1.

Usage

- **dbmny4cmp** compares two DBMONEY4 values.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmnycmp

dbmny4copy

Description

Copy a DBMONEY4 value.

Syntax

```
RETCODE dbmny4copy(dbproc, src, dest)

DBPROCESS    *dbproc;
DBMONEY4     *src;
DBMONEY4     *dest;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

src

A pointer to the source DBMONEY4 value.

dest

A pointer to the destination DBMONEY4 variable.

Return value

SUCCEED or FAIL.

dbmny4copy returns FAIL if either *src* or *dest* is NULL.

Usage

- **dbmny4copy** copies the *src* DBMONEY4 value to the *dest* DBMONEY4 variable.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmnycopy, dbmnyminus, dbmny4minus

dbmny4divide

Description

Divide one DBMONEY4 value by another.

Syntax

RETCODE dbmny4divide(dbproc, m1, m2, quotient)

```
DBPROCESS    *dbproc;
DBMONEY4     *m1;
DBMONEY4     *m2;
DBMONEY4     *quotient;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1
A pointer to the DBMONEY4 value serving as dividend.

m2
A pointer to the DBMONEY4 value serving as divisor.

quotient
A pointer to a DBMONEY4 variable to hold the result of the division.

Return value

SUCCEED or FAIL.

dbmny4divide returns FAIL in case of overflow or division by zero, or if *m1*, *m2*, or *quotient* is NULL.

Usage

- **dbmny4divide** divides the *m1* DBMONEY4 value by the *m2* DBMONEY4 value and places the result in **quotient*.
- In case of overflow or division by zero, **dbmny4divide** returns FAIL and sets **quotient* to $0.0000.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmny4add, dbmny4sub, dbmny4mul, dbmny4minus, dbmnyadd, dbmnysub, dbmnymul, dbmnydivide, dbmnyminus

dbmny4minus

Description

Negate a DBMONEY4 value.

Syntax

```
RETCODE dbmny4minus(dbproc, src, dest)

 DBPROCESS    *dbproc;
```

```
DBMONEY4    *src;
DBMONEY4    *dest;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

src

A pointer to a DBMONEY4 value.

dest

A pointer to a DBMONEY4 variable to hold the result of the negation.

Return value

SUCCEED or FAIL.

dbmny4minus returns FAIL in case of overflow, or if *src* or *dest* is NULL.

Usage

- **dbmny4minus** negates the *src* DBMONEY4 value and places the result into **dest*.
- In case of overflow, **dbmny4minus** returns FAIL. **dest* is undefined in this case. An attempt to negate the maximum negative DBMONEY4 value will result in overflow.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmnyminus, dbmnycopy, dbmny4copy

dbmny4mul

Description

Multiply two DBMONEY4 values.

Syntax

RETCODE dbmny4mul(dbproc, m1, m2, product)

```
DBPROCESS   *dbproc;
DBMONEY4    *m1;
DBMONEY4    *m2;
DBMONEY4    *product;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1
A pointer to a DBMONEY4 value.

m2
A pointer to a DBMONEY4 value.

product
A pointer to a DBMONEY4 variable to hold the result of the multiplication.

Return value

SUCCEED or FAIL.

dbmny4mul returns FAIL in case of overflow, or if *m1*, *m2*, or *product* is NULL.

Usage

- **dbmny4mul** multiplies the *m1* DBMONEY4 value by the *m2* DBMONEY4 value and places the result in **product*.
- In case of overflow, **dbmny4mul** returns FAIL and sets **product* to $0.0000.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmny4add, dbmny4sub, dbmny4divide, dbmny4minus, dbmnyadd, dbmnysub, dbmnymul, dbmnydivide, dbmnyminus

dbmny4sub

Description

Subtract one DBMONEY4 value from another.

Syntax

RETCODE dbmny4sub(dbproc, m1, m2, difference)

```
DBPROCESS    *dbproc;
DBMONEY4     *m1;
```

```
DBMONEY4    *m2;
DBMONEY4    *difference;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1

A pointer to the DBMONEY4 value to be subtracted from.

m2

A pointer to the DBMONEY4 value to subtract.

difference

A pointer to a DBMONEY4 variable to hold the result of the subtraction.

Return value

SUCCEED or FAIL.

dbmny4sub returns FAIL in case of overflow, or if *m1*, *m2*, or *difference* is NULL.

Usage

- **dbmny4sub** subtracts the *m2* DBMONEY4 value from the *m1* DBMONEY4 value and places the result in **difference*.
- In case of overflow, **dbmny4sub** returns FAIL and sets **difference* to $0.0000.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmny4sub, dbmny4mul, dbmny4divide, dbmny4minus, dbmny4add, dbmnysub, dbmnymul, dbmnydivide, dbmnyminus

dbmny4zero

Description

Initialize a DBMONEY4 variable to $0.0000.

Syntax

RETCODE dbmny4zero(dbproc, mny4ptr)

```
DBPROCESS    *dbproc;
DBMONEY4     *mny4ptr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mny4ptr

A pointer to the DBMONEY4 value to initialize.

Return value

SUCCEED or FAIL.

dbmny4zero returns FAIL if *mny4ptr* is NULL.

Usage

- **dbmny4zero** initializes a DBMONEY4 value to $0.0000.
- The range of legal DBMONEY4 values is from -$214,748.3648 to $214,748.3647. DBMONEY4 values have a precision of one ten-thousandth of a dollar.

See also

dbmnyzero

dbmnyadd

Description

Add two DBMONEY values.

Syntax

RETCODE dbmnyadd(dbproc, m1, m2, sum)

```
DBPROCESS    *dbproc;
DBMONEY      *m1;
DBMONEY      *m2;
DBMONEY     *sum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1

A pointer to a DBMONEY value.

m2

A pointer to a DBMONEY value.

sum

A pointer to a DBMONEY variable to hold the result of the addition.

Return value

SUCCEED or FAIL.

Usage

- **dbmnyadd** adds the *m1* and *m2* DBMONEY values and places the result in **sum*.
- In case of overflow, **dbmnyadd** returns FAIL and sets **sum* to $0.0000.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.
- **dbmnyadd** returns FAIL in case of overflow, or if *m1*, *m2*, or *sum* is NULL.

See also

dbmnysub, dbmnymul, dbmnydivide, dbmnyminus, dbmny4add, dbmny4sub, dbmny4mul, dbmny4divide, dbmny4minus

dbmnycmp

Description

Compare two DBMONEY values.

Syntax

int dbmnycmp(dbproc, m1, m2)

```
DBPROCESS   *dbproc;
DBMONEY     *m1;
DBMONEY     *m2;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1

A pointer to a DBMONEY value.

m2

A pointer to a DBMONEY value.

Return value

If *m1 = m2* **dbmnycmp** returns 0.

If *m1 < m2* **dbmnycmp** returns -1.

If *m1 > m2* **dbmnycmp** returns 1.

Usage

- **dbmnycmp** compares two DBMONEY values.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmny4cmp

dbmnycopy

Description

Copy a DBMONEY value.

Syntax

```
RETCODE dbmnycopy(dbproc, src, dest)

DBPROCESS    *dbproc;
DBMONEY       *src;
DBMONEY      *dest;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

src

A pointer to the source DBMONEY value.

dest

A pointer to the destination DBMONEY variable.

Return value

SUCCEED or FAIL.

dbmnycopy returns FAIL if either *src* or *dest* is NULL.

Usage

- **dbmnycopy** copies the *src* DBMONEY value to the *dest* DBMONEY value.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnycopy, dbmnyminus, dbmny4minus

dbmnydec

Description

Decrement a DBMONEY value by one ten-thousandth of a dollar.

Syntax

RETCODE dbmnydec(dbproc, mnyptr)

```
DBPROCESS    *dbproc;
DBMONEY      *mnyptr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to the DBMONEY value to decrement.

Return value

SUCCEED or FAIL.

dbmnydec returns FAIL in case of overflow or if *mnyptr* is NULL.

Usage

- **dbmnydec** decrements a DBMONEY value by one ten-thousandth of a dollar.
- An attempt to decrement the maximum negative DBMONEY value will result in overflow. In case of overflow, **dbmnydec** returns FAIL. In this case, the contents of **mnyptr* are undefined.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnyinc, dbmnymaxneg

dbmnydivide

Description

Divide one DBMONEY value by another.

Syntax

RETCODE dbmnydivide(dbproc, m1, m2, quotient)

```
DBPROCESS   *dbproc;
DBMONEY     *m1;
DBMONEY     *m2;
DBMONEY     *quotient;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1
A pointer to the DBMONEY value serving as dividend.

m2
A pointer to the DBMONEY value serving as divisor.

quotient
A pointer to a DBMONEY variable to hold the result of the division.

Return value

SUCCEED or FAIL.

dbmnydivide returns FAIL in case of overflow or division by zero, or if *m1*, *m2*, or *quotient* is NULL.

Usage

- **dbmnydivide** divides the *m1* DBMONEY value by the *m2* DBMONEY value and places the result in **quotient*.
- In case of overflow or division by zero, **dbmnydivide** returns FAIL and sets **quotient* to $0.0000.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnyadd, dbmnysub, dbmnymul, dbmnyminus, dbmny4add, dbmny4sub, dbmny4mul, dbmny4divide, dbmny4minus

dbmnydown

Description

Divide a DBMONEY value by a positive integer.

Syntax

RETCODE dbmnydown(dbproc, mnyptr, divisor, remainder)

```
DBPROCESS    *dbproc;
DBMONEY      *mnyptr;
```

```
int             divisor;
int            *remainder;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to the DBMONEY value to divide. **mnyptr* will also contain the result of the division.

divisor

The integer by which **mnyptr* will be divided. *divisor* must be positive, and must be less than or equal to 65535.

remainder

A pointer to an integer variable to hold the remainder from the division, in ten-thousandths of a dollar. If *remainder* is passed as NULL, no remainder is returned.

Return value

SUCCEED or FAIL.

dbmnydown returns FAIL if *mnyptr* is NULL, or if *divisor* is not between 1 and 65535.

Usage

- **dbmnydown** divides a DBMONEY value by a short integer and places the result back in the original DBMONEY variable.
- **dbmnydown** places the remainder of the division into **remainder*. **remainder* is an integer representing the number of ten-thousandths of a dollar left after the division.
- *divisor* must be greater than or equal to one and less than or equal to 65535.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnyscale, dbmnydivide, dbmny4divide

dbmnyinc

Description

Increment a DBMONEY value by one ten-thousandth of a dollar.

Syntax

RETCODE dbmnyinc(dbproc, mnyptr)

```
DBPROCESS    *dbproc;
DBMONEY      *mnyptr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to the DBMONEY value to increment.

Return value

SUCCEED or FAIL.

dbmnyinc returns FAIL in case of overflow or if *mnyptr* is NULL.

Usage

- **dbmnyinc** increments a DBMONEY value by one ten-thousandth of a dollar.
- An attempt to increment the maximum positive DBMONEY value will result in overflow. In case of overflow **dbmnyinc** returns FAIL. **mnyptr* is undefined in this case.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnydec, dbmnymaxpos

dbmnyinit

Description

Prepare a DBMONEY value for calls to **dbmnyndigit**.

Syntax

RETCODE dbmnyinit(dbproc, mnyptr, trim, negative)

```
DBPROCESS    *dbproc;
```

```
DBMONEY     *mnyptr;
int          trim;
DBBOOL      *negative;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to the DBMONEY value to be initialized. **dbmnyinit** changes the value of **mnyptr*.

trim

The number of digits to trim from **mnyptr*. **dbmnyinit** removes digits from **mnyptr* by dividing it by a power of 10. The value of *trim* determines what power of 10 is used. *trim* cannot be less than 0.

negative

A pointer to a DBBOOL variable. If **mnyptr* is negative, **dbmnyinit** makes it positive and sets **negative* to TRUE.

Return value

SUCCEED or FAIL.

dbmnyinit returns FAIL if *mnyptr* is NULL, *negative* is NULL, or *trim* is less than 0.

Usage

- **dbmnyinit** initializes a DBMONEY value for conversion to character. It eliminates unwanted precision and converts negative values to positive.
- **dbmnyinit** eliminates digits from a DBMONEY value by dividing by a power of 10. The integer *trim* determines what power of 10 is used. **dbmnyinit** modifies **mnyptr*, replacing the original value with the trimmed value. If **mnyptr* is negative, **dbmnyinit** makes it positive and sets **negative* to TRUE.
- **dbmnyinit** and **dbmnyndigit** are useful for writing a custom DBMONEY-to-DBCHAR conversion routine. Such a custom routine might be useful if the accuracy provided by **dbconvert**'s DBMONEY-to-DBCHAR conversion (hundredths of a dollar) is not adequate. Also, **dbconvert** does not build a character string containing commas.

- **dbmnyndigit** returns the rightmost digit of a DBMONEY value as a DBCHAR. To get all the digits of a DBMONEY value, call **dbmnyndigit** repeatedly. See the dbmnyndigit reference page for more details.
- **dbmnyinit** is almost always used in conjunction with **dbmnyndigit**. Used alone, **dbmnyinit** can force negative DBMONEY values positive and divide DBMONEY values by a power of 10, but the real purpose of **dbmnyinit** is to prepare a DBMONEY value for calls to **dbmnyndigit**.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.
- The dbmnyndigit reference page contains an example that demonstrates the use of **dbmnyinit**.

See also dbconvert, dbmnyndigit

dbmnymaxneg

Description Return the maximum negative DBMONEY value supported.

Syntax RETCODE dbmnymaxneg(dbproc,dest)

```
DBPROCESS    *dbproc;
DBMONEY      *dest;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

dest

A pointer to a DBMONEY variable.

Return value SUCCEED or FAIL.

dbmnymaxneg returns FAIL if *dest* is NULL.

Usage

- **dbmnymaxneg** fills **dest* with the maximum negative DBMONEY value supported.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnymaxpos

dbmnymaxpos

Description

Return the maximum positive DBMONEY value supported.

Syntax

```
RETCODE dbmnymaxpos(dbproc, dest)

DBPROCESS    *dbproc;
DBMONEY      *dest;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

dest

A pointer to a DBMONEY variable.

Return value

SUCCEED or FAIL.

dbmnymaxpos returns FAIL if *dest* is NULL.

Usage

- **dbmnymaxpos** fills **dest* with the maximum positive DBMONEY value supported.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnymaxneg

dbmnyminus

Description

Negate a DBMONEY value.

Syntax

RETCODE dbmnyminus(dbproc, src, dest)

```
DBPROCESS    *dbproc;
DBMONEY      *src;
DBMONEY      *dest;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

src

A pointer to a DBMONEY value.

dest

A pointer to a DBMONEY variable to hold the result of the negation.

Return value

SUCCEED or FAIL.

dbmnyminus returns FAIL in case of overflow, or if *src* or *dest* is NULL.

Usage

- **dbmnyminus** negates the *src* DBMONEY value and places the result into **dest*.
- In case of overflow, **dbmnyminus** returns FAIL. **dest* is undefined in this case. An attempt to negate the maximum negative DBMONEY value will result in overflow.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmny4minus, dbmnycopy, dbmny4copy

dbmnymul

Description

Multiply two DBMONEY values.

Syntax

```
RETCODE dbmnymul(dbproc, m1, m2, product)

DBPROCESS    *dbproc;
DBMONEY      *m1;
DBMONEY      *m2;
DBMONEY      *product;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1

A pointer to a DBMONEY value.

m2

A pointer to a DBMONEY value.

product

A pointer to a DBMONEY variable to hold the result of the multiplication.

Return value

SUCCEED or FAIL.

dbmnymul returns FAIL in case of overflow, or if *m1*, *m2*, or *product* is NULL.

Usage

- **dbmnymul** multiplies the *m1* DBMONEY value by the *m2* DBMONEY value and places the result in **product*.
- In case of overflow, **dbmnymul** returns FAIL and sets **product* to $0.0000.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnyadd, dbmnysub, dbmnydivide, dbmnyminus, dbmny4add, dbmny4sub, dbmny4mul, dbmny4divide, dbmny4minus

dbmnyndigit

Description

Return the rightmost digit of a DBMONEY value as a DBCHAR.

Syntax

RETCODE dbmnyndigit(dbproc, mnyptr, value, zero)

```
DBPROCESS   *dbproc;
DBMONEY     *mnyptr;
DBCHAR      *value;
DBBOOL      *zero;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to a DBMONEY value. Each call to **dbmnyndigit** divides this value by 10 and places the result back into **mnyptr*.

value

A pointer to a DBCHAR variable to fill with the character representation of the rightmost digit of the DBMONEY value.

zero

A pointer to a DBBOOL variable. Each call to **dbmnyndigit** divides **mnyptr* by 10 and puts the character representation of the remainder of the division in **value*. If the result of the division is $0.0000, **dbmnyndigit** sets **zero* to TRUE. Otherwise, **zero* is set to FALSE. If *zero* is passed as NULL, this information is not returned.

Return value

SUCCEED or FAIL.

dbmnyndigit returns FAIL if *mnyptr* or *value* is NULL.

Usage

- **dbmnyndigit** returns the rightmost digit of a DBMONEY value as a DBCHAR.
- **dbmnyndigit** divides a DBMONEY value by 10. It places the character representation of the remainder of the division in **value*, and replaces **mnyptr* with the result of the division. If the result of the division is $0.0000, **dbmnyndigit** sets **zero* to TRUE.

- To get all the digits of a DBMONEY value, call **dbmnydigit** repeatedly, until **zero* is TRUE.
- **dbmnyinit** and **dbmnyndigit** are useful for writing a custom DBMONEY-to-DBCHAR conversion routine. Such a custom routine might be useful if the accuracy provided by **dbconvert**'s DBMONEY-to-DBCHAR conversion (hundredths of a dollar) is not adequate. Also **dbconvert** does not build a character string containing commas.
- **dbmnyinit** initializes a DBMONEY value for conversion to character. It eliminates unwanted precision and converts negative values to positive. For more information, see the dbmnyinit reference page.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.
- This code fragment demonstrates the use of **dbmnyndigit** and **dbmnyinit**:

```
/*
 ** This example demonstrates dbmnyinit() and
 ** dbmnyndigit(). It is a conversion routine which
 ** converts a DBMONEY value to a character string.
 ** The conversion provided by this routine is unlike
 ** the conversion provided by dbconvert() in that the
 ** resulting character string includes commas. This
 ** conversion provides precision of two digits after
 ** the decimal point.
 **
 ** For simplicity, the example assumes that all
 ** routines succeed and all parameters passed to it
 ** are valid.
 */

 #define PRECISION       2

 RETCODE     new_mnytochar(mnyptr, buf_ptr)
 DBMONEY     *mnyptr;
 char        *buf_ptr;
 {
     DBMONEY    local_mny;
     DBBOOL     negative;
     int        bytes_written;
     DBCHAR     value;
     DBBOOL     zero;
     int        ret;
     char       temp_buf[32];
```

```
/*
** Since dbmnyinit() and dbmnyndigit() modify the
** DBMONEY value passed to it, and since we do
** not want to modify the DBMONEY value passed
** to us by the user we need to make a local copy.
*/
ret = dbmnycopy((DBPROCESS *)NULL, mnyptr,
     &local_mny);
/* The value of 'ret' should be checked */

/*
** Next we need to call dbmnyinit().
**
** dbmnyinit() eliminates any unwanted precision
** from the DBMONEY value.  DBMONEY values are
** stored with accuracy to four digits after the
** decimal point. For this conversion routine we
** only want accuracy to two digits after the
** decimal.
**
** Passing a value of 2 for the second parameter
** eliminates those two digits of precision we do
** not care about.
**
** dbmnyinit() also turns negative DBMONEY values
** into positive DBMONEY values. The value of
** negative is set to TRUE if dbmnyinit() turns a
** negative DBMONEY value into a positive DBMONEY
** value.
**
** NOTE: dbmnyinit() eliminates unwanted by
** precision by dividing DBMONEY values by a
** power of ten. In this conversion routine it
** divides by 100. If we pass dbmnyinit() a
** DBMONEY value of $1534.1277 the resulting
** DBMONEY value is $15.3413.
*/
negative = FALSE;
ret = dbmnyinit((DBPROCESS *)NULL, &local_mny,
    4 - PRECISION, &negative);
/* The value of 'ret' should be checked */

/*
** dbmnyndigit() extracts the rightmost digit out
** of the DBMONEY value, converts it to a
```

```
** character, places the character into the
** variable “value”, and then divides the DBMONEY
** value by 10. dbmnyndigit() sets ’zero’ to TRUE
** if the result of the division is $0.0000.
**
** By calling dbmnyndigit() until ’zero’ is set to
** TRUE we will be returned all the digits (from
** right to left) of the DBMONEY value.
*/
zero = FALSE;
bytes_written = 0;
while( zero == FALSE )
{
    ret = dbmnyndigit((DBPROCESS *)NULL,
    &local_mny, &value, &zero);
    /* The value of ’ret’ should be checked. */

    /*
    ** As we are getting the digits, we want to
    ** place the decimal point and commas in the
    ** proper positions ...
    */
    temp_buf[bytes_written++] = value;

    /*
    ** If zero == TRUE we got all the digits. We
    ** do not want to call
    ** check_comma_and_decimal() since we might
    ** put a comma before the leftmost digit.
    */
    if( zero == FALSE )
    {
         /*
         ** As we are getting the digits, we want
         ** to place the decimal point and commas
         ** in the proper positions ...
         */
         check_comma_and_decimal(temp_buf,
                              &bytes_written);
    }
}

/*
** If we haven’t written PRECISION bytes into the
** buffer yet, pad with zeros, write the decimal
** point to the buffer, and write a zero after
```

```
    ** the decimal point.
    */
    pad_with_zeros(temp_buf, &bytes_written);

    /*
    ** We've written the money value into the buffer
    ** backwards. Now we have to write it the right
    ** way.
    */
    reverse_money(buf_ptr, temp_buf, bytes_written,
                  negative);

    return(SUCCEED);

}

void check_comma_and_decimal(temp_buf,
                             bytes_written)
char *temp_buf;
int  *bytes_written;
{

   static     int     comma = 0;
   static     DBBOOL  after_decimal = FALSE;

   if( after_decimal )
    {
        /*
        ** When comma is 3 it is time to write a
        ** comma. We do not care about commas until
        ** after we've written the decimal point.
        */
        comma++;
    }

    /*
    ** After we've written PRECISION bytes into the
    ** buffer, it's time to write the decimal point.
    */
    if( *bytes_written == PRECISION )
    {
      temp_buf[(*bytes_written)++] = '.';
      after_decimal = TRUE;
    }

    /*
```

```
    ** When (comma == 3) that means we've written three
    ** digits and it's time to put a comma into the
    ** buffer.
    */
    if( comma == 3 )
    {
      temp_buf[(*bytes_written)++] = ',';
      comma = 0;       /* clear comma */
    }

}

void    pad_with_zeros(temp_buf, bytes_written)
char    *temp_buf;
int     *bytes_written;
{

    /* If we haven't written PRECISION bytes into the
    ** buffer yet, pad with zeros, write the decimal
    ** point to the buffer, and write a zero after the
    ** decimal point.
    */
    while( *bytes_written < PRECISION )
    {
        temp_buf[(*bytes_written)++] = '0';
    }

    if( *bytes_written == PRECISION )
    {
        temp_buf[(*bytes_written)++] = '.';
        temp_buf[(*bytes_written)++] = '0';
    }

          }

void reverse_money(char_buf, temp_buf,
               bytes_written, negative)
char     *char_buf;
char     *temp_buf;
int      bytes_written;
DBBOOL   negative;
{

    int       i;

    /*
```

```
        ** We've written the money value into the buffer
        ** backwards. Now we have to write it the right
        ** way. First check to see if we need to write a
        ** negative sign, then write the dollar sign,
        ** finally write the money value.
        */
        i = 0;
        if( negative == TRUE )
        {
            char_buf[i++] = '-';
        }
       char_buf[i++] = '$';

        while( bytes_written-- )
        {
            char_buf[i++] = temp_buf[bytes_written];
        }
        /* Append null-terminator: */
        char_buf[i] = '\0';

    }
```

See also dbconvert, dbmnyinit

dbmnyscale

Description Multiply a DBMONEY value by a positive integer and add a specified amount.

Syntax RETCODE dbmnyscale(dbproc, mnyptr, multiplier, addend)

```
DBPROCESS   *dbproc;
DBMONEY     *mnyptr;
int          multiplier;
int          addend;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to the DBMONEY value to multiply. **mnyptr* will also contain the result of the **dbmnyscale** operation.

multiplier

The integer by which **mnyptr* will be multiplied. *multiplier* must be positive, and must be greater than or equal to 1, and less than or equal to 65535.

addend

An integer representing the number of ten-thousandths of a dollar to add to **mnyptr* after the multiplication.

Return value

SUCCEED or FAIL.

dbmnyscale returns FAIL if *mnyptr* is NULL, if overflow occurs, or if *multiplier* is not between 1 and 65535.

Usage

- **dbmnyscale** multiplies a DBMONEY value by a short integer, adds *addend* ten-thousandths of a dollar, and places the result back in the original DBMONEY variable.
- *multiplier* must be greater than or equal to 1, and less than or equal to 65535.
- In case of overflow, **dbmnyscale** returns FAIL. **mnyptr* is undefined in this case.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnydown, dbmnymul, dbmny4mul

dbmnysub

Description

Subtract one DBMONEY value from another.

Syntax

RETCODE dbmnysub(dbproc, m1, m2, difference)

```
DBPROCESS    *dbproc;
DBMONEY      *m1;
DBMONEY      *m2;
DBMONEY      *difference;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

m1

A pointer to the DBMONEY value to be subtracted from.

m2

A pointer to the DBMONEY value to subtract.

difference

A pointer to a DBMONEY variable to hold the result of the subtraction.

Return value

SUCCEED or FAIL.

dbmnysub returns FAIL in case of overflow, or if *m1*, *m2*, or *difference* is NULL.

Usage

- **dbmnysub** subtracts the *m2* DBMONEY value from the *m1* DBMONEY value and places the result in **difference*.
- In case of overflow, **dbmnysub** returns FAIL and sets *difference* to $0.0000.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also

dbmnyadd, dbmnymul, dbmnydivide, dbmnyminus, dbmny4add, dbmny4sub, dbmny4mul, dbmny4divide, dbmny4minus

dbmnyzero

Description — Initialize a DBMONEY value to $0.0000.

Syntax

```
RETCODE dbmnyzero(dbproc, mnyptr)

DBPROCESS  *dbproc;
DBMONEY    *mnyptr;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

This parameter may be NULL. The DBPROCESS is used as a parameter to an application's error handler. It also contains information on what language to print error messages in. If a DBPROCESS is not supplied, the default national language is used.

mnyptr

A pointer to the DBMONEY value to initialize.

Return value — SUCCEED or FAIL.

dbmnyzero returns FAIL if *mnyptr* is NULL.

Usage

- **dbmnyzero** initializes a DBMONEY value to $0.0000.
- The range of legal DBMONEY values is between +/- $922,337,203,685,477.5808. DBMONEY values have a precision of one ten-thousandth of a dollar.

See also — dbmny4zero

dbmonthname

Description — Determine the name of a specified month in a specified language.

Syntax

```
char *dbmonthname(dbproc, language, monthnum,
        shortform)

DBPROCESS    *dbproc;
char         *language;
int          monthnum;
DBBOOL       shortform;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

language
The name of the desired language.

monthnum
The number of the desired month. Month numbers range from 1 (January) to 12 (December).

shortform
A Boolean value indicating whether the long or short form of the month name is desired. If *shortform* is TRUE, **dbmonthname** returns the short form of the month name; if *shortform* is FALSE, **dbmonthname** returns the full month name. For example, if the month name desired is the U.S. English short form for January, "Jan" is returned.

Short forms of month names are defined in localization files on a per-localization-file basis.

Return value

The name of the specified month on success; a NULL pointer on error.

Usage

- **dbmonthname** returns the name of the specified month in the specified language. If no language is specified (*language* is NULL), *dbproc*'s current language is used. If both *language* and *dbproc* are NULL, DB-Library's default language (if any) is used.
- The following code fragment illustrates the use of **dbmonthname**:

```
for (monthnum = 1; monthnum <= 12; monthnum++)
      printf("Month %d: %s\n", monthnum,
           dbmonthname((DBPROCESS *)NULL,
                (char *)NULL,
           monthnum, TRUE),
           dbmonthname((DBPROCESS *)NULL,
                (char *)NULL,
           monthnum, FALSE));
```

See also

db12hour, dbdateorder, dbdayname, DBSETLNATLANG, dbsetopt

DBMORECMDS

Description — Indicate whether there are more commands to be processed.

Syntax

```
RETCODE DBMORECMDS(dbproc)

DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value — SUCCEED or FAIL, indicating whether there are more results from the command batch.

Usage

- The application can use this macro to determine whether there are more results to process.
- **DBMORECMDS** can be called after **dbnextrow** returns NO_MORE_ROWS. If you know that the current command is returning no rows, you can call **DBMORECMDS** immediately after **dbresults**.
- Applications rarely need this routine, because they can simply call **dbresults** until it returns NO_MORE_RESULTS.

See also — DBCMDROW, dbresults, DBROWS, DBROWTYPE

dbmoretext

Description — Send part of a text or image value to the server.

Syntax

```
RETCODE dbmoretext(dbproc, size, text)

DBPROCESS    *dbproc;
DBINT         size;
BYTE         *text;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

size
The size, in bytes, of this particular part of the text or image value being sent to the server. It is an error to send more text or image bytes to the server than were specified in the call to **dbwritetext**.

text
A pointer to the text or image portion to be written.

Return value

SUCCEED or FAIL.

Usage

- This routine is used in conjunction with **dbwritetext** to send a large SYBTEXT or SYBIMAGE value to the server in the form of a number of smaller chunks. This is particularly useful with operating systems that are unable to allocate extremely long data buffers.
- **dbmoretext** and **dbwritetext** are used in updates only, and serve to replace the Transact-SQL **update** statement.
- **dbsqlok** and **dbresults** must be called before the first call to **dbmoretext** and after the last call to **dbmoretext**.
- See the dbwritetext reference page for more information, including an example that uses **dbmoretext**.
- The DB-Library/C option DBTEXTSIZE affects the value of the server *@@textsize* global variable, which restricts the size of text or image values that the server returns. *@@textsize* has a default value of 32,768 bytes. An application that retrieves text or image values larger than 32,768 bytes will need to call **dbsetopt** to make *@@textsize* larger.

  The DB-Library/C option DBTEXTLIMIT limits the size of text or image values that DB-Library/C will read.

See also

dbtxptr, dbtxtimestamp, dbwritetext

dbmsghandle

Description

Install a user function to handle server messages.

Syntax

```
int (*dbmsghandle(handler))()

int         (*handler)();
```

Parameters

handler

A pointer to the user function that will be called whenever DB-Library receives an error or informational message from the server. DB-Library calls this function with eight parameters:

Table 2-21: Message handler parameters

Parameter	Meaning
dbproc	The affected DBPROCESS.
msgno	The current message's number (datatype DBINT). These numbers are documented in the *sysmessages* table.
msgstate	The current message's error state number (datatype *int*). These numbers provide Sybase Technical Support with information about the context of the error.
severity	The current message's information class or error severity (datatype *int*). These numbers are documented in the Adaptive Server documentation.
msgtext	The null-terminated text of the current message (datatype *char* *).
srvname	The null-terminated name of the server that generated the message (datatype *char* *). A server's name is stored in the *srvname* column of its *sysservers* system table. It is used in server-to-server communication; in particular, it is used when one server logs into another server to perform a remote procedure call. If the server has no name, *srvname* will be of length 0.
procname	The null-terminated name of the stored procedure that generated the message (datatype char *). If the message was not generated by a stored procedure, *procname* will be of length 0.
line	The number of the command batch or stored procedure line that generated the message (datatype *int*). Line numbers start at 1. The line number pertains to the nesting level at which the message was generated. For instance, if a command batch executes stored procedure A, which then calls stored procedure B, and a message is generated at line 3 of B, then the value of *line* is 3. *line* will be 0 if there is no line number associated with the message. Circumstances that could generate messages without line numbers include a login error or a remote procedure call (performed via **dbrpcsend**) to a stored procedure that does not exist.

The message handler must return a value of 0 to DB-Library.

Message handlers on Windows NT or Windows 98 must be declared with CS_PUBLIC, as shown in the following example. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well.

The following example shows a typical message handler routine:

```
#include <sybfront.h>
 #include <sybdb.h>

 int CS_PUBLIC msg_handler(dbproc, msgno, msgstate,
     severity, msgtext, srvname, procname, line)

 DBPROCESS         *dbproc;
 DBINT             msgno;
 int               msgstate;
 int               severity;
 char              *msgtext;
 char              *srvname;
 char              *procname;
 int               line;

 {
      printf ("Msg %ld, Level %d, State %d\n",
          msgno, severity, msgstate);
       if (strlen(srvname) > 0)
           printf ("Server '%s', ", srvname);
       if (strlen(procname) > 0)
           printf ("Procedure '%s', ", procname);
       if (line > 0)
           printf ("Line %d", line);

       printf("\n\t%s\n", msgtext);

       return(0);
 }
```

Return value

A pointer to the previously installed message handler or NULL if no message handler was installed before.

Usage

- **dbmsghandle** installs a message-handler function that you supply. When DB-Library receives a server error or informational message, it will call this message handler immediately. You must install a message handler to handle server messages properly.
- If an application does not call **dbmsghandle** to install a message-handler function, DB-Library ignores server messages. The messages are not printed.
- If the command buffer contains just a single command and that command provokes a server message, DB-Library will call the message handler during **dbsqlexec**.

If the command buffer contains multiple commands (and the first command in the buffer is ok), a runtime error will not cause **dbsqlexec** to fail. Instead, failure will occur with the **dbresults** call that processes the command causing the runtime error.

- You can "de-install" an existing message handler by calling **dbmsghandle** with a NULL parameter. You can also, at any time, install a new message handler. The new handler will automatically replace any existing handler.
- Refer to the *sysmessages* table for a list of server messages. In addition, the Transact-SQL **print** and **raiserror** commands generate server messages that **dbmsghandle** will catch.
- The routines **dbsetuserdata** and **dbgetuserdata** can be particularly useful when you need to transfer information between the message handler and the program code that triggered it. See the dbsetuserdata reference page for an example of how to handle deadlock in this way.
- Another routine, **dberrhandle**, installs an error handler that DB-Library calls in response to DB-Library errors.
- If the application provokes messages from DB-Library and the server simultaneously, DB-Library calls the server message handler before it calls the DB-Library error handler.
- The DB-Library/C error value SYBESMSG is generated in response to a server error message, but not in response to a server informational message. This means that when a server error occurs, both the server message handler and the DB-Library/C error handler are called, but when the server generates an informational message, only the server message handler is called.

  If you have installed a server message handler, you may want to write your DB-Library error handler so as to suppress the printing of any SYBESMSG error, to avoid notifying the user about the same error twice.
- The following table provides information on when DB-Library/C calls an application's message and error handlers:

Table 2-22: When DB-Library calls message and error handlers

Error or message	Message handler called?	Error handler called?
SQL syntax error	Yes	Yes (SYBESMSG). (Code the handler to ignore the message.
SQL **print** statement	Yes	No.
SQL **raiserror**	Yes	No.

Error or message	Message handler called?	Error handler called?
Server dies	No	Yes (SYBESEOF). (Code your handler to exit the application.
Timeout from the server	No	Yes (SYBETIME). (To wait for another timeout period, code your handler to return - INT_CONTINUE.
Deadlock on query	Yes	No. (Code your handler to test for deadlock.
Timeout on login	No	Yes (SYBEFCON).
Login fails (**dbopen**)	Yes	Yes (SYBEPWD). (Code your handler to exit the application.
Use database message	Yes (Code the handler to ignore the message.	No.
Incorrect use of DB-Library/C calls, such as not calling **dbresults** when required	No	Yes (SYBERPND, .
Fatal Server error (severity greater than 16)	Yes	Yes (SYBESMSG).

See also dberrhandle, dbgetuserdata, dbsetuserdata

dbname

Description Return the name of the current database.

Syntax

```
char *dbname(dbproc)

DBPROCESS    *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

A pointer to the null-terminated name of the current database.

Usage

- **dbname** returns the name of the current database.
- If you need to keep track of when the database changes, use **dbchange**.

See also

dbchange, dbuse

dbnextrow

Description

Read the next result row into the row buffer and into any program variables that are bound to column data.

Syntax

```
STATUS dbnextrow(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

dbnextrow returns:

- REG_ROW if a regular row has been read. A regular row is any row that matches the query's **where** clause.
- A *computeid* if a compute row was read. A compute row is a row that is generated by a **compute** clause. The *computeid* matches the number of the compute row that was read; the first compute row is 1, the second is 2, and so forth. A *computeid* cannot match any other of the return types for this function.
- BUF_FULL is returned if buffering is turned on and reading the next row would cause the buffer to be exceeded. In this case, no row will have been read. To read any more rows, at least one row must first be pruned from the top of the row buffer by calling **dbclrbuf.**

- NO_MORE_ROWS if the last row in the result set has been read. If the query did not generate rows (for example, an **update** or **insert**, or a **select** with no match), then the first call to **dbnextrow** will return NO_MORE_ROWS. Also, **dbnextrow** returns this value if the query failed or if there are no pending results.
- FAIL if an abnormal event, such as a network or out-of-memory error, prevented the routine from completing successfully.

Usage

- **dbnextrow** reads the next row of result data, starting with the first row returned from the server. Ordinarily, the next result row is read directly from the server. If the DBBUFFER option is turned on and rows have been read out of order by calling **dbgetrow**, the next row is read instead from a linked list of buffered rows. When **dbnextrow** is called, any binding of row data to program variables (as specified with **dbbind** or **dbaltbind**) takes effect.
- If program variables are bound to columns, then new values will be written into the bound variables before **dbnextrow** returns.
- In regular rows, column values can be retrieved with **dbdata** or bound to program variables with **dbbind.** In compute rows, column values can be retrieved with **dbadata** or bound to program variables with **dbaltbind.**
- **dbresults** must return SUCCEED before an application can call **dbnextrow**. To determine whether a particular command is one that returns rows and needs results processing with **dbnextrow**, call **DBROWS** after **dbresults**.
- After calling **dbresults**, an application can either call **dbcanquery** or **dbcancel** to cancel the current set of results, or call **dbnextrow** in a loop to process the results row-by-row.
- If it chooses to process the results, an application can either:
  - Process all result rows by calling **dbnextrow** in a loop until it returns NO_MORE_ROWS. After NO_MORE_ROWS is returned, the application can call **dbresults** again to set up the next result set (if any) for processing.
  - Process some result rows by calling **dbnextrow**, and then cancel the remaining result rows by calling **dbcancel** (to cancel all results from the command batch or RPC call) or **dbcanquery** (to cancel only the results associated with the last **dbresults** call).

  An application must either cancel or process all result rows.
- The typical sequence of calls is:

```
DBINT     xvariable;
 DBCHAR    yvariable[10];

 /* Read the query into the command buffer */
 dbcmd(dbproc, "select x = 100, y = 'hello'");

 /* Send the query to SQL Server */
 dbsqlexec(dbproc);

 /* Get ready to process the query results */
 dbresults(dbproc);

 /* Bind column data to program variables */
 dbbind(dbproc, 1, INTBIND, (DBINT) 0,
    (BYTE *) &xvariable);
 dbbind(dbproc, 2, STRINGBIND, (DBINT) 0,
     yvariable);

 /* Now process each row */
 while (dbnextrow(dbproc) != NO_MORE_ROWS)
 {
     C-code to print or process row data
 }
```

- The server can return two types of rows: regular rows containing data from columns designated by a **select** statement's select list, and compute rows resulting from the **compute** clause. To facilitate the processing of result rows from the server, **dbnextrow** returns different values according to the type of row. See the "Returns" section in this reference page for details.
- To display server result data on the default output device, you can use **dbprrow** instead of **dbnextrow**.

See also

dbaltbind, dbbind, dbcanquery, dbclrbuf, dbgetrow, dbprrow, dbsetrow, Options

dbnpcreate

Description

Create a notification procedure.

Syntax

```
RETCODE dbnpcreate(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

Return value

SUCCEED or FAIL.

Usage

- **dbnpcreate** creates a notification procedure. A notification procedure is a special type of Open Server registered procedure. A notification procedure differs from a normal Open Server registered procedure in that it contains no executable statements. Notification procedures are the only type of Open Server registered procedure that a DB-Library/C application can create.
- The notification procedure name and its parameters must have been previously defined via **dbnpdefine** and **dbregparam**.
- To create a notification procedure, a DB-Library/C application must:
  - Define the procedure via **dbnpdefine**
  - Describe the procedure's parameters, if any, via **dbregparam**
  - Create the procedure via **dbnpcreate**
- All DB-Library/C routines that apply to registered procedures apply to notification procedures as well. For example, **dbregexec** executes a registered procedure, which may or may not be a notification procedure. Likewise, **dbreglist** lists all registered procedures currently defined in Open Server, some of which may be notification procedures.
- Like other registered procedures, notification procedures are useful for inter-application communication and synchronization, because applications can request to be advised when a notification procedure executes.
- Notification procedures may be created only in Open Server. At this time, Adaptive Server does not support notification procedures.
- A DB-Library/C application requests to be notified of a registered procedure's execution via **dbregwatch**. The application may request to be notified either synchronously or asynchronously.
- This is an example of creating a notification procedure:

```
DBPROCESS   *dbproc;
 DBINT       status;
```

```
/*
** Let's create a notification procedure called
** "message" which has two parameters:
**     msg     varchar(255)
**     user    idint
*/

/*
** Define the name of the notification procedure
** "message"
*/
dbnpdefine (dbproc, "message", DBNULLTERM);

/*
** The notification procedure has two parameters:
**      msg    varchar(255)
**      user   idint
** So, define these parameters.  Note that
** neither of the parameters is defined with a
** default value.
*/
dbregparam (dbproc, "msg", SYBVARCHAR,
     DBNODEFAULT, NULL);
dbregparam (dbproc, "userid", SYBINT4,
     DBNODEFAULT, 4);

/*  Create the notification procedure: */
status = dbnpcreate (dbproc);
if (status == FAIL)
{
    fprintf(stderr, "ERROR:  Failed to create \
     message!\n");
}
else
{
    fprintf(stdout, "Success in creating \
     message!\n");
}
```

See also dbreginit, dbregparam, dbregwatch, dbregnowatch

dbnpdefine

Description

Define a notification procedure.

Syntax

RETCODE dbnpdefine(dbproc, procedure_name, namelen)

```
DBPROCESS    *dbproc;
DBCHAR       *procedure_name;
DBSMALLINT   namelen;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

procedure_name

A pointer to the name of the notification procedure being defined.

namelen

The length of *procedure_name*, in bytes. If *procedure_name* is null-terminated, pass *namelen* as DBNULLTERM.

Return value

SUCCEED or FAIL.

Usage

- **dbnpdefine** defines a notification procedure. Defining a notification procedure is the first step in creating it.
- A notification procedure is a special type of Open Server registered procedure. A notification procedure differs from a normal Open Server registered procedure in that it contains no executable statements. Notification procedures are the only type of Open Server registered procedure that a DB-Library/C application can create.
- To create a notification procedure, a DB-Library/C application must:
  - Define the procedure via **dbnpdefine**
  - Describe the procedure's parameters, if any, via **dbregparam**
  - Create the procedure via **dbnpcreate**
- All DB-Library/C routines that apply to registered procedures apply to notification procedures as well. For example, **dbregexec** executes a registered procedure, which may or may not be a notification procedure. Likewise, **dbreglist** lists all registered procedures currently defined in Open Server, some of which may be notification procedures.
- This is an example of defining a notification procedure:

```
DBPROCESS    *dbproc;
```

```
DBINT        status;

/*
** Let's create a notification procedure called
** "message" which has two parameters:
**     msg     varchar(255)
**     userid  int
*/

/*
** Define the name of the notification procedure
** "message"
*/
dbnpdefine (dbproc, "message", DBNULLTERM);

/* The notification procedure has two parameters:
**     msg    varchar(255)
**     userid int
** So, define these parameters.  Note that
** neither of the parameters is defined with a
** default value.
*/
dbregparam (dbproc, "msg", SYBVARCHAR,
     DBNODEFAULT, NULL);
dbregparam (dbproc, "userid", SYBINT4,
     DBNODEFAULT, 4);

/*  Create the notification procedure: */
status = dbnpcreate (dbproc);
if (status == FAIL)
{
     fprintf(stderr, "ERROR:  Failed to create \
       message!\n");
}
else
{
     fprintf(stdout, "Success in creating \
       message!\n");
}
```

See also dbregparam, dbnpcreate, dbreglist

dbnullbind

Description

Associate an indicator variable with a regular result row column.

Syntax

RETCODE dbnullbind(dbproc, column, indicator)

```
DBPROCESS   *dbproc;
int          column;
DBINT       *indicator;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the column that is to be associated with the indicator variable.

indicator

A pointer to the indicator variable.

Return value

SUCCEED or FAIL.

dbnullbind returns FAIL if *column* is invalid.

Usage

- **dbnullbind** associates a regular result row column with an indicator variable. The indicator variable indicates whether a particular regular result row's column has been converted and copied to a program variable successfully or unsuccessfully, or whether it is null.
- The indicator variable is set when regular result rows are processed via **dbnextrow**. The possible values are:
  - -1 if the column is NULL.
  - The full length of column's data, in bytes, if *column* was bound to a program variable via **dbbind**, the binding did not specify any data conversions, and the bound data was truncated because the program variable was too small to hold *column*'s data.
  - 0 if *column* was bound and copied successfully to a program variable.

Note Detection of character string truncation is implemented only for CHARBIND and VARYCHARBIND.

See also

dbanullbind, dbbind, dbdata, dbdatlen, dbnextrow

dbnumalts

Description — Return the number of columns in a compute row.

Syntax — int dbnumalts(dbproc, computeid)

```
DBPROCESS   *dbproc;
int         computeid;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

computeid
The id that identifies the particular compute row of interest. A SQL **select** statement may have multiple **compute** clauses, each of which returns a separate compute row. The *computeid* corresponding to the first **compute** clause in a **select** is 1. The *computeid* is returned by **dbnextrow** or **dbgetrow.**

Return value — The number of columns for the particular computeid. **dbnumalts** returns -1 if *computeid* is invalid.

Usage

- **dbnumalts** returns the number of columns in a compute row. The application can call this routine after **dbresults** returns SUCCEED. For example, given the SQL statement:

```
select dept, year, sales from employee
     order by dept, year
     compute avg(sales), min(sales),
     max(sales) by dept
```

  the call **dbnumalts**(*dbproc*, 1) returns 3.

See also — dbadata, dbadlen, dbaltlen, dbalttype, dbgetrow, dbnextrow, dbnumcols

dbnumcols

Description — Determine the number of regular columns for the current set of results.

Syntax — int dbnumcols(dbproc)

```
DBPROCESS   *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of columns in the current set of results. If there are no columns, **dbnumcols** returns 0.

Usage

- **dbnumcols** returns the number of regular (that is, non-compute) columns in the current set of results.
- Here is a program fragment that illustrates the use of **dbnumcols**:

```
int        column_count;
 DBPROCESS  *dbproc;

 /* Put the commands into the command buffer */
 dbcmd(dbproc, "select name, id, type from \
    sysobjects");
 dbcmd(dbproc, " select name from sysobjects");

 /*
 ** Send the commands to SQL Server and start
 ** execution
 */
 dbsqlexec(dbproc);

 /* Process each command until there are no more */
 while (dbresults(dbproc) != NO_MORE_RESULTS)
 {
    column_count = dbnumcols(dbproc);
    printf("%d columns in this SQL Server \
      result.\n", column_count);
    while (dbnextrow(dbproc) != NO_MORE_ROWS)
      printf("row received.\n");
 }
```

See also

dbcollen, dbcolname, dbnumalts

dbnumcompute

Description — Return the number of **compute** clauses in the current set of results.

Syntax — int dbnumcompute(dbproc)

```
DBPROCESS   *dbproc;
```

Parameters — dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value — The number of **compute** clauses in the current set of results.

Usage

- This routine returns the number of **compute** clauses in the current set of results. The application can call it after **dbresults** returns SUCCEED. For example, given the SQL statement:

```
select dept, name from employee
     order by dept, name
     compute count(name) by dept
     compute count(name)
```

  the call **dbnumcompute**(*dbproc*) will return 2 since there are two **compute** clauses in the **select** statement.

See also — dbnumalts, dbresults

DBNUMORDERS

Description — Return the number of columns specified in a Transact-SQL **select** statement's **order by** clause.

Syntax — int DBNUMORDERS(dbproc)

```
DBPROCESS   *dbproc;
```

Parameters — dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of **order by** columns. If there is no **order by** clause, this routine returns 0. If there is an error, it returns -1.

Usage

- Once a **select** statement has been executed and **dbresults** has been called to process it, the application can call DBNUMORDERS to find out how many columns were specified in the statement's **order by** clause.

See also

dbordercol

dbnumrets

Description

Determine the number of return parameter values generated by a stored procedure.

Syntax

```
int dbnumrets(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of return parameter values associated with the most recently-executed stored procedure.

Usage

- **dbnumrets** provides the number of return parameter values returned by the most recent **execute** statement or remote procedure call on a stored procedure. If the number returned by **dbnumrets** is less than or equal to 0, then no return parameters are available.
- Transact-SQL stored procedures can return values for specified "return parameters." Changes made to the value of a return parameter inside the stored procedure are then available to the program that called the procedure. This is analogous to the "pass by reference" facility available in some programming languages.

- For a parameter to function as a return parameter, it must be declared as such within the stored procedure. The **execute** statement or remote procedure call that calls the stored procedure must also indicate that the parameter should function as a return parameter. In the case of a remote procedure call, it is the **dbrpcparam** routine that specifies whether a parameter is a return parameter.
- When executing a stored procedure, the server returns any parameter values immediately after returning all other results. Therefore, the application can call **dbnumrets** only after processing the stored procedure's results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbnumrets** or any other routines that process return parameters, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.
- If the stored procedure is invoked with a remote procedure call, the return parameter values are automatically available to the application. If, on the other hand, the stored procedure is invoked with an **execute** statement, the return parameter values are available only if the command batch containing the **execute** statement uses local variables, not constants, for the return parameters. For more details on return parameters from stored procedures, see the *Adaptive Server Enterprise Reference Manual*.
- Other routines are used to retrieve return parameter values:
  - **dbretdata** returns a pointer to a parameter value.
  - **dbretlen** returns the length of a parameter value.
  - **dbretname** returns the name of a parameter value.
  - **dbrettype** returns the datatype of a parameter value.
  - **dbconvert** can be called to convert the value, if necessary.

  For an example of how these routines can be used together with **dbnumrets**, see the reference page for dbretdata.

See also

dbnextrow, dbresults, dbretdata, dbretlen, dbretname, dbrettype, dbrpcinit, dbrpcparam

dbopen

Description

Create and initialize a DBPROCESS structure.

Syntax

```
DBPROCESS *dbopen(login, server)

LOGINREC     *login;
char         *server;
```

Parameters

login

A pointer to a LOGINREC structure. This pointer will be passed as an argument to **dbopen**. You can get one by calling **dblogin**.

Once the application has made all its **dbopen** calls, the LOGINREC structure is no longer necessary. The program can then call **dbloginfree** to free the LOGINREC structure.

server

The server that you want to connect to. *server* is the alias given to the server in the interfaces file. **dbopen** looks up *server* in the interfaces file to get information for connecting to a server.

If *server* is NULL **dbopen** looks up the interfaces entry that corresponds to the value of the DSQUERY environment variable or logical name. If DSQUERY has not been explicitly set, it has a value of "SYBASE". (For information on designating an interfaces file, see the reference page for dbsetifile. For more information on the interfaces file itself, see the *Open Client/Server Configuration Guide*.

Note On non-UNIX platforms, client applications may use a method to find server address information that is different than the UNIX interfaces file. Consult your *Open Client/Server Configuration Guide* for detailed information on how clients connect to servers.

Return value

A DBPROCESS pointer if everything went well. Ordinarily, **dbopen** returns NULL if a DBPROCESS structure could not be created or initialized, or if your login to the server failed. When **dbopen** returns NULL, it generates a DB-Library error number that indicates the error. The application can access this error number through an error handler. However, if there is an unexpected communications failure during the server login process and an error handler has not been installed, the program will be aborted.

Usage

- This routine allocates and initializes a DBPROCESS structure. This structure is the basic data structure that DB-Library uses to communicate with a server. It is the first argument in almost every DB-Library call. Besides allocating the DBPROCESS structure, this routine sets up communication with the network, logs into the server, and initializes any default options.
- Here is a program fragment that uses **dbopen**:

```
DBPROCESS       *dbproc;
 LOGINREC        *loginrec;

 loginrec = dblogin();
 DBSETLPWD(loginrec, "server_password");
 DBSETLAPP(loginrec, "my_program");
 dbproc = dbopen(loginrec, "my_server");
```

- Once the application has logged into a server, it can change databases by calling the **dbuse** routine.

Multiple query entries in an interfaces file

- It is possible to set up an interfaces file so that if **dbopen** fails to establish a connection with a server, it attempts to establish a connection with an alternate server.
- An application can use the **dbopen** call

```
dbopen(loginrec, MARS);
```

to connect to the server MARS. An interfaces file containing an entry for MARS might look like this:

```
#
 MARS
     query tcp hp-ether violet 1025
     master tcp hp-ether violet 1025
     console tcp hp-ether violet 1026
 #
 VENUS
     query tcp hp-ether plum 1050
     master tcp hp-ether plum 1050
     console tcp hp-ether plum 1051
 #
 NEPTUNE
     query tcp hp-ether mauve 1060
     master tcp hp-ether mauve 1060
     console tcp hp-ether mauve 1061
```

- The application is directed to port number 1025 on the machine “violet”. If MARS is not available, the dbopen call fails. If the interfaces file has multiple query entries in it for MARS, however, and the first connection attempt fails, **dbopen** will automatically attempt to connect to the next server listed. Such an interfaces file might look like this:

```
#
 MARS
      query tcp hp-ether violet 1025
      query tcp hp-ether plum 1050
      query tcp hp-ether mauve 1060
      master tcp hp-ether violet 1025
      console tcp hp-ether violet 1026
 #
 VENUS
      query tcp hp-ether plum 1050
      master tcp hp-ether plum 1050
      console tcp hp-ether plum 1051
 #
 NEPTUNE
      query tcp hp-ether mauve 1060
      master tcp hp-ether mauve 1060
      console tcp hp-ether mauve 1061
```

- Note that the second query entry under MARS is identical to the query entry under VENUS, and that the third query entry is identical to the query entry under NEPTUNE. If this interfaces file is used and the application fails to connect with MARS, it will automatically attempt to connect with VENUS. If it fails to connect with VENUS, it will automatically attempt to connect with NEPTUNE. There is no limit on the number of alternate servers that may be listed under a server’s interfaces file entry, but each alternate server must be listed in the same interfaces file. Two numbers may be added after the server’s name in the interfaces file:

```
#
 MARS retries seconds
      query tcp hp-ether violet 1025
      query tcp hp-ether plum 1050
      query tcp hp-ether mauve 1060
      master tcp hp-ether violet 1025
      console tcp hp-ether violet 1026
```

retries represents the number of additional times to loop through the list of query entries if no connection is achieved during the first pass. *seconds* represents the amount of time, in seconds, that **dbopen** will wait at the top of the loop before going through the list again. These numbers are optional. If they are not included, **dbopen** will try to connect to each query entry only once. Looping through the list and pausing between loops is useful in case any of the candidate servers is in the process of booting. Multiple query lines can be particularly useful when alternate servers contain mirrored copies of the primary server's databases.

Errors **dbopen** will return NULL if any of the following errors occur. These errors can be trapped in the application's error handler (installed with **dberrhandle**.

If **dbopen** is called in the entry funtions of a DLL, a deadlock may arise. **dbopen** creates operating system threads and tries to synchronize them using system utilities. This synchronization conflicts with the operating system's serialization process.

Note The use of SIGALARM in a DB-Library application can cause **dbopen** to fail.

SYBEMEM	Unable to allocate sufficient memory.
SYBEDBPS	Maximum number of DBPROCESSes already allocated. Note that an application can set or retrieve the maximum number of DBPROCESS structures with **dbsetmaxprocs** and **dbgetmaxprocs**.
SYBESOCK	Unable to open socket.
SYBEINTF	Server name not found in interfaces file.
SYBEUHST	Unknown host machine name.
SYBECONN	Unable to connect: Adaptive Server is unavailable or does not exist.
SYBEPWD	Login incorrect.
SYBEOPIN	Could not open interfaces file.

See also dbclose, dbexit, dbinit, dblogin, dbloginfree, dbsetifile, dbuse

dbordercol

Description

Return the id of a column appearing in the most recently executed query's **order by** clause.

Syntax

```
int dbordercol(dbproc, order)

DBPROCESS    *dbproc;
int          order;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

order

The id that identifies the particular **order by** column of interest. The first column named within the **order by** clause is number 1.

Return value

The column id (based on the column's position in the select list) for the column in the specified place in the **order by** clause. If the order is invalid, **dbordercol** returns -1.

Usage

- This routine returns the id of the column that appears in a specified location within the **order by** clause of a SQL **select** command.

  For example, given the SQL statement:

  ```
  select dept, name, salary from employee
       order by salary, name
  ```

  the call **dbordercol**(*dbproc*, 1) will return 3 since the first column named in the **order by** clause refers to the third column in the query's select list.

See also

DBNUMORDERS

dbpoll

Description

Check if a server response has arrived for a DBPROCESS.

Syntax

```
RETCODE dbpoll(dbproc, milliseconds, ready_dbproc,
        return_reason)

DBPROCESS    *dbproc;
long         milliseconds;
```

```
DBPROCESS       **ready_dbproc;
int             *return_reason;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and server.

dbproc represents the DBPROCESS connection that **dbpoll** will check.

If *dbproc* is passed as NULL, **dbpoll** will check all open DBPROCESS connections to see if a response has arrived for any of them.

milliseconds

The maximum number of milliseconds that **dbpoll** should wait for a response before returning.

If *milliseconds* is passed as 0, **dbpoll** returns immediately.

If *milliseconds* is passed as -1, **dbpoll** will not return until either a server response arrives or a system interrupt occurs.

ready_dbproc

A pointer to a pointer to a DBPROCESS structure. **dbpoll** sets **ready_dbproc* to point to the DBPROCESS for which the server response has arrived. If no response has arrived, **dbpoll** sets **ready_dbproc* to NULL.

Note *ready_dbproc* is not a DBPROCESS pointer. It is a pointer to a DBPROCESS pointer.

---

return_reason

A pointer to an integer representing the reason **dbpoll** has returned. The integer will be one of the following symbolic values:

DBRESULT	A response to a server command has arrived. The application may call **dbsqlok** (assuming that **dbsqlsend** has been called) to examine the server's response.
DBNOTIFICATION	A registered procedure notification has arrived. If a handler for this registered procedure has been installed via **dbreghandle**, **dbpoll** invokes this handler before it returns. If a handler for the registered procedure has not been installed and there is no default handler installed for this DBPROCESS, DB-Library raises an error when it reads the notification.
DBTIMEOUT	The time indicated by the milliseconds parameter elapsed before any server response arrived.
DBINTERRUPT	An operating-system interrupt occurred before any server response arrived and before the timeout period elapsed.

Note This list may expand in the future, as more kinds of server responses are recognized by DB-Library/C. It is recommended that application programs be coded to handle unexpected values in *return_reason* without error.

Return value

SUCCEED or FAIL.

dbpoll returns FAIL if any of the server connections it checks has died. If **dbpoll** returns FAIL, *ready_dbproc* and *return_reason* are undefined.

Usage

- **dbpoll** checks the TDS (Tabular Data Stream) buffer to see if it contains any server response not yet read by an application.
- *dbproc* represents the DBPROCESS connection that **dbpoll** will check. If *dbproc* is passed as NULL, **dbpoll** examines all open connections and returns as soon as it finds one that has an unread server response.
- If there is an unread response, **dbpoll** sets **ready_dbproc* and *return_reason* to reflect which DBPROCESS connection the response is for and what the response is.
- Note that *ready_dbproc* is not a pointer to a DBPROCESS structure. It is a pointer to the address of a DBPROCESS. **dbpoll** sets **ready_dbproc* to point to the DBPROCESS for which the server response has arrived. If no server response has arrived, **dbpoll** sets **ready_dbproc* to NULL.
- **dbpoll** can be used for two purposes:

- To enable an application to implement non-blocking reads (calls to **dbsqlok**) from the server
- To check if a registered procedure notification has arrived for a DBPROCESS

Using dbpoll for non-blocking reads

- **dbpoll** can be used to check whether bytes are available for **dbsqlok** to read.
- Depending on the nature of an application, the time between the moment when a command is sent to the server (made via **dbsqlsend** or **dbrpcsend**) and the server's response (initially read with **dbsqlok**) may be significant.
- During this time, the server is processing the command and building the result data. An application may use this time to perform other duties. When ready, the application can call **dbpoll** to check if a server response arrived while it was busy elsewhere. For an example of this usage, see the reference page for dbsqlok.

Note On occasion **dbpoll** may report that data is ready for **dbsqlok** to read when only the first bytes of the server response are present. When this occurs, **dbsqlok** waits for the rest of the response or until the timeout period has elapsed, just like **dbsqlexec**. In practice, however, the entire response is usually available at one time.

- **dbpoll** should not be used with **dbresults** or **dbnextrow**. **dbpoll** cannot determine if calls to these routines will block. This is because **dbpoll** works by checking whether or not bytes are available on a DBPROCESS connection, and these two routines do not always read from the network.
  - If all of the results from a command have been read, **dbresults** returns NO_MORE_RESULTS. In this case, **dbresults** does not block even if no bytes are available to be read.
  - If all of the rows for a result set have been read, **dbnextrow** returns NO_MORE_ROWS. In this case, **dbnextrow** does not block even if no bytes are available to be read.

- For non-blocking reads, alternatives to **dbpoll** are **DBRBUF** and **DBIORDESC**. These routines are specific to the UNIX and AOS/VS-specific platform. They are not portable, so their use should be avoided whenever possible. They do, however, provide a way for application programs to integrate handling of DB-Library/C sockets with other sockets being used by an application.
  - **DBRBUF** is a UNIX and AOS/VS-specific routine. It checks an internal DB-Library network buffer to see if a server response has already been read. **dbpoll** checks one or all connections used by an application's DBPROCESSes, to see if a response is ready to be read.
  - **DBIORDESC**, another UNIX and AOS/VS-specific routine, is similar in function to **dbpoll**. **DBIORDESC** provides the socket handle used for network reads by the DBPROCESS. The socket handle can be used with the UNIX **select** function or its AOS/VS equivalent.

Using dbpoll for registered procedure notifications

- An application may have one or more DBPROCESS connections waiting for registered procedure notifications. A DBPROCESS connection will not be aware that a registered procedure notification has arrived unless it reads results from the server. If a connection is not reading results, it can use **dbpoll** to check if a registered procedure notification has arrived. If so, **dbpoll** reads the registered procedure notification stream and calls the handler for that registered procedure.
- Here is a code fragment that uses **dbpoll** to poll for a registered procedure notification:

```
/*
 ** This code fragment illustrates the use of
 ** dbpoll() to processan event notification.
 **
 ** The code fragment will ask the Server to
 ** notify the Client when the event "shutdown"
 ** occurs. When the event notification is
 ** received from the Server, DB-Library will call
 ** the handler installed for that event. This
 ** event handler routine can then access the
 ** event's parameters, and take any appropriate
 ** action.
 */

DBINT   handlerfunc();
 DBINT   ret;

/* First install the handler for this event */
```

```
dbreghandle(dbproc, "shutdown", handlerfunc);

/*
 ** Now make the asynchronous notification
 ** request.
 */
 ret = dbregwatch(dbproc, "shutdown", DBNULLTERM,
     DBNOWAITONE);
 if (ret == FAIL)
 {
     fprintf(stderr, "ERROR: dbregwatch() \
           failed!!\n");
 }
 else if (ret == DBNOPROC)
 {
      fprintf(stderr, "ERROR: procedure shutdown \
           not defined!\n");
 }

/*
 ** Since we are making use of the asynchronous
 ** event notification mechanism, the application
 ** can continue doing other work. All we have to
 ** do is call dbpoll() once in a while, to deal
 ** with the event notification when it arrives.
 */
 while (1)
 {
      /* Have dbpoll() block for one second */
      dbpoll(NULL, 1000, NULL, &ret);

    /*
      ** If we got the event, then get out of this
      ** loop.
      */
      if (ret == DBNOTIFICATION)
      {
           break;
      }
      /* Deal with our other tasks here */
  }
```

See also DBIORDESC, DBRBUF, dbresults, dbreghandle, dbsqlok

dbprhead

Description
Print the column headings for rows returned from the server.

Syntax
```
void dbprhead(dbproc)

 DBPROCESS    *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value
None.

Usage
- This routine displays, on the default output device and in a default format, the column headings for a set of query results. The format is compatible with the format used by **dbprrow**.
- The application can call **dbprhead** once **dbresults** returns SUCCEED.
- You can specify the maximum number of characters to be placed on one line via the DB-Library option DBPRLINELEN.
- This routine is useful for debugging.
- The routines **dbsprhead**, **dbsprline**, and **dbspr1row** provide an alternative to **dbprhead** and **dbprrow**. These routines print the formatted row results into a caller-supplied character buffer.

See also
dbbind, dbnextrow, dbprrow, dbresults, dbspr1row, dbsprhead, dbsprline

dbprrow

Description
Print all the rows returned from the server.

Syntax
```
RETCODE dbprrow(dbproc)

 DBPROCESS    *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

SUCCEED or FAIL.

Usage

- This routine displays, on the default output device and in a default format, the rows for a set of query results. This routine reads and prints all the rows. It saves the trouble of calling routines such as **dbbind** and **dbnextrow**, but it prints only in a single, predetermined format.
- The application can call **dbprrow** once **dbresults** returns SUCCEED.
- When using this routine, you do not need to call **dbnextrow** to loop through the rows.
- You can specify the maximum number of characters to be placed on one line via the DB-Library option DBPRLINELEN.
- **dbprrow** is useful primarily for debugging.
- If row buffering is turned on, **dbprrow** buffers rows in addition to printing them out. If the buffer is full, the oldest rows are removed as necessary.
- The routines **dbsprhead**, **dbsprline**, and **dbspr1row** provide an alternative to **dbprhead** and **dbprrow**. These routines print the formatted row results into a caller-supplied character buffer.

See also

dbbind, dbnextrow, dbprhead, dbresults, dbspr1row, dbsprhead, dbsprline

dbprtype

Description

Convert a token value to a readable string.

Syntax

```
char *dbprtype(token)

 int        token;
```

Parameters

token

The server token value (SYBCHAR, SYBFLT8, and so on) returned by **dbcoltype**, **dbalttype**, **dbrettype,** or **dbaltop**.

Return value

A pointer to a null-terminated string that is the readable translation of the token value. The pointer points to space that is never overwritten, so it is safe to call this routine more than once in the same statement. If the token value is unknown, the routine returns a pointer to an empty string.

Usage

- Certain routines—**dbcoltype**, **dbalttype**, **dbrettype,** and **dbaltop**—return token values representing server datatypes or aggregate operators. **dbprtype** provides a readable string version of a token value.

- For example, **dbprtype** will take a **dbcoltype** token value representing the server binary datatype (SYBBINARY) and return the string “binary”.
- Here is a list of the token strings that **dbprtype** can return and their token value equivalents:

Table 2-23: Token values and their string equivalents

Token string	Token value	Description
char	SYBCHAR	char datatype
text	SYBTEXT	text datatype
binary	SYBBINARY	binary datatype
image	SYBIMAGE	image datatype
tinyint	SYBINT1	1-byte integer datatype
smallint	SYBINT2	2-byte integer datatype
int	SYBINT4	4-byte integer datatype
float	SYBFLT8	8-byte float datatype
real	SYBREAL	4-byte float datatype
numeric	SYBNUMERIC	numeric type
decimal	SYBDECIMAL	decimal type
bit	SYBBIT	bit datatype
money	SYBMONEY	money datatype
smallmoney	SYBMONEY4	4-byte money datatype
datetime	SYBDATETIME	datetime datatype
smalldatetime	SYBDATETIME4	4-byte datetime datatype
boundary	SYBBOUNDARY	boundary type
sensitivity	SYBSENSITIVITY	sensitivity type
sum	SYBAOPSUM	sum aggregate operator
avg	SYBAOPAVG	average aggregate operator
count	SYBAOPCNT	count aggregate operator
min	SYBAOPMIN	minimum aggregate operator
max	SYBAOPMAX	maximum aggregate operator

See also dbaltop, dbalttype, dbcoltype, dbrettype, Types

dbqual

Description

Return a pointer to a **where** clause suitable for use in updating the current row in a browsable table.

Syntax

char *dbqual(dbproc, tabnum, tabname)

```
DBPROCESS    *dbproc;
int           tabnum;
char         *tabname;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

tabnum

The number of the table of interest, as specified in the **select** statement's **from** clause. Table numbers start at 1. If *tabnum* is -1, the *tabname* parameter will be used to identify the table.

tabname

A pointer to the null-terminated name of a table specified in the **select** statement's **from** clause. *tabname* is ignored unless *tabnum* is passed as -1.

Return value

A pointer to a null-terminated **where** clause for the current row in the specified table. This buffer is dynamically allocated, and it is the application's responsibility to free it via **dbfreequal**.

dbqual will return a NULL pointer if the specified table is not browsable. For a table to be "browsable," it must have a unique index and a timestamp column.

dbqual will also return a NULL pointer if the preceding **select** did not include the **for browse** option.

Usage

- **dbqual** is one of the DB-Library browse mode routines. See "Browse mode" on page 26 for a detailed discussion of browse mode.
- **dbqual** provides a **where** clause that the application can use to update a single row in a browsable table. Columns from this row must have previously been retrieved into the application through a browse-mode **select** query (that is, a **select** that ends with the key words **for browse**).

  The **where** clause produced by **dbqual** begins with the keyword **where** and contains references to the row's unique index and timestamp column. The application simply appends the **where** clause to an **update** or **delete** statement; it does not need to examine it or manipulate it in any way.

The timestamp column indicates the time that the particular row was last updated. An update on a browsable table will fail if the timestamp column in the **dbqual**-generated **where** clause is different from the timestamp column in the table. Such a condition, which provokes Adaptive Server error message 532, indicates that another user updated the row between the time this application selected it for browsing and the time it tried to update it. The application itself must provide the logic for handling the update failure. The following program fragment illustrates one approach:

```
/* This code fragment illustrates a technique for
 ** handling the case where a browse-mode update fails
 ** because the row has already been updated
** by another user. In this example, we simply retrieve
 ** the entire row again, allow the user to examine and
 ** modify it, and try the update again.
 **
 ** Note that "q_dbproc" is the DBPROCESS used to query
 ** the database, and "u_dbproc" is the DBPROCESS used
 ** to update the database.
 */

/* First, find out which employee record the user
 ** wants to update.
 */
 employee_id = which_employee();

while (1)
  {
      /* Retrieve that employee record from the database.
       ** We'll assume that "empid" is a unique index,
       ** so this query will return only one row.
       */
      dbfcmd (q_dbproc, "select * from employees where \
          empid = %d for browse", employee_id);
       dbsqlexec(q_dbproc);
       dbresults(q_dbproc);
       dbnextrow(q_dbproc);

     /* Now, let the user examine or edit the employee's
       ** data, first placing the data into program
       ** variables.
       */
      extract_employee_data(q_dbproc, employee_struct);
       examine_and_edit(employee_struct, &edit_flag);

     if (edit_flag == FALSE)
       {
            /* The user didn't edit this record,
```

```
                    ** so we're done.
                    */
                    break;
               }
               else
               {
                   /* The user edited this record, so we'll use
                    ** the edited data to update the
                    ** corresponding row in the database.
                    */
                   qualptr = dbqual(q_dbproc, -1, "employees");
                    dbcmd(u_dbproc, "update employees");
                    dbfcmd (u_dbproc, " set address = '%s', \
                        salary = %d %s",
                        employee_struct->address,
                        employee_struct->salary, qualptr);
                    dbfreequal(qualptr);
                    if ((dbsqlexec(u_dbproc) == FAIL) ||
                        (dbresults(u_dbproc) == FAIL))
                    {
                       /* Our update failed. In a real program,
                        ** it would be necessary to examine the
                       ** messages returned from the SQL Server
                        ** to determine why it failed.  In this
                       ** example, we'll assume that the update
                      ** failed because someone else has already
                        ** updated this row, thereby changing
                        ** the timestamp.
                        **
                      ** To cope with this situation, we'll just
                      ** repeat the loop, retrieving the changed
                        ** row for our user to examine and edit.
                      ** This will give our user the opportunity
                        ** to decide whether to overwrite the
                        ** change made by the other user.
                         */
                         continue;
                    }
                    else
                    {
                      /* The update succeeded, so we're done. */
                        break;
                    }
               }
          }
```

- **dbqual** can only construct **where** clauses for browsable tables. You can use **dbtabbrowse** to determine whether a table is browsable.
- **dbqual** is usually called after **dbnextrow**.
- For a complete example that uses **dbqual** to perform a browse mode update, see the online sample programs.

See also

dbcolbrowse, dbcolsource, dbfreequal, dbtabbrowse, dbtabcount, dbtabname, dbtabsource, dbtsnewlen, dbtsnewval, dbtsput

DBRBUF

Description

(UNIX and AOS/VS only) Determine whether the DB-Library network buffer contains any unread bytes.

Syntax

```
DBBOOL DBRBUF(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

TRUE (bytes remain in buffer) or FALSE (no bytes in buffer).

Note that **DBRBUF** actually returns TRUE both when there are bytes available in the read buffer, and when no more results are available to be processed.

This is because the purpose of **DBRBUF** is to tell an application when it can read and be assured that it will not hang. If **DBRBUF** did not return TRUE in the case of no more results, then applications that loop while **DBRBUF** returns FALSE could loop indefinitely, if all results had already been processed.

Usage

- This routine lets the application know if the DB-Library network buffer contains any bytes yet unread.
- DBRBUF is ordinarily used in conjunction with **dbsqlok** and DBIORDESC.

- **dbpoll**, a DB-Library/C routine which checks if a server response has arrived for any DBPROCESS, may replace **DBRBUF**. Since the UNIX-specific routines **DBRBUF** and **DBIORDESC** are non-portable, their use should be avoided whenever possible. They do, however, provide a way for application programs to integrate handling of DB-Library/C sockets with other sockets being used by an application.
- An application uses these routines to manage multiple input data streams. To manage these streams efficiently, an application that uses **dbsqlok** should check whether any bytes remain either in the network buffer or in the network itself before calling **dbresults**.
- To test whether bytes remain in the network buffer, the application can call **DBRBUF**. To test whether bytes remain in the network itself, the application can either call the UNIX **select** or AOS/VS equivalent and **DBIORDESC**, or call **dbpoll**.

See also

DBIORDESC, dbpoll, dbsqlok, dbresults

dbreadpage

Description

Read a page of binary data from the server.

Syntax

DBINT dbreadpage(dbproc, dbname, pageno, buf)

```
DBPROCESS    *dbproc;
char         *dbname;
DBINT         pageno;
BYTE          buf[];
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

dbname

The name of the database of interest.

pageno

The number of the database page to be read.

buf

A pointer to a buffer to hold the received page data. Adaptive Server pages are currently 2048 bytes long.

Return value

The number of bytes read from the server. If the operation was unsuccessful, **dbreadpage** returns -1.

Usage

- **dbreadpage** reads a page of binary data from the server. This routine is primarily useful for examining and repairing damaged database pages. After calling **dbreadpage**, the DBPROCESS may contain some error or informational messages from the server. These messages may be accessed through a user-supplied message handler.

Warning! Use this routine only if you are absolutely sure you know what you are doing!

Limitations

Alters the contents of the DBPROCESS command buffer.

See also

dbmsghandle, dbwritepage

dbreadtext

Description

Read part of a text or image value from the server.

Syntax

STATUS dbreadtext(dbproc, buf, bufsize)

```
DBPROCESS    *dbproc;
void          *buf;
DBINT        bufsize;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

buf

A pointer to a caller-allocated buffer that will contain the chunk of text or image data.

bufsize

The size of the caller's buffer, in bytes.

Return value

The following table lists the return values for **dbreadtext**:

dbreadtext returns	To indicate
>0	The number of bytes placed into the caller's buffer
0	The end of a row
-1	An error occurred, such as a network or out of memory error
NO_MORE_ROWS	All rows read

Usage

- **dbreadtext** reads a large SYBTEXT or SYBIMAGE value from the server in the form of a number of smaller chunks. This is particularly useful with operating systems that are unable to allocate extremely long data buffers.
- To read successive chunks of the same SYBTEXT or SYBIMAGE value, call **dbreadtext** until it returns 0 (end of row).
- Use **dbreadtext** in place of **dbnextrow** to read SYBTEXT and SYBIMAGE values.
- **dbreadtext** can process the results of Transact-SQL queries if those queries return only one column and that column contains either text or image data. The Transact-SQL **readtext** command returns results of this type.
- The DB-Library/C option DBTEXTSIZE affects the value of the server *@@textsize* global variable, which restricts the size of text or image values that the server returns. *@@textsize* has a default value of 32,768 bytes. An application that retrieves text or image values larger than 32,768 bytes will need to call **dbsetopt** to make *@@textsize* larger.

  The DB-Library/C option DBTEXTLIMIT limits the size of text or image values that DB-Library/C will read. DB-Library/C will throw away any text that exceeds the limit.
- This code fragment demonstrates the use of **dbreadtext**

```
DBPROCESS   *dbproc;
 long        bytes;
 RETCODE     ret;
 char        buf[BUFSIZE + 1];
/*
 ** Install message and error handlers...
 ** Log in to server...
 ** Send a "use database" command...
 */
/* Select a text column: */
 dbfcmd(dbproc, "select textcolumn from bigtable");
 dbsqlexec(dbproc);
```

```
/* Process the results: */
 while( (ret = dbresults(dbproc)) !=
      NO_MORE_RESULTS )
 {
      if( ret == FAIL )
      {
           /* dbresults() failed */
      }
     while( (bytes = dbreadtext(dbproc,
          (void *)buf, BUFSIZE)) != NO_MORE_ROWS )
      {
           if( bytes == -1 )
           {
                /* dbreadtext() failed */
           }
           else if( bytes == 0 )
           {
               /* We've reached the end of a row*/
                printf("End of Row!\n\n");
           }
           else
           {
                /*
                ** 'bytes' bytes have been placed
                ** into our buffer.
                ** Print them:
                */
                buf[bytes] = '\0';
                printf("%s\n", buf);
           }
      }
 }
```

See also dbmoretext, dbnextrow, dbwritetext

dbrecftos

Description Record all SQL commands sent from the application to the server.

Syntax void dbrecftos(filename)

char *filename;

Parameters

filename
A pointer to a null-terminated character string to be used as the basis for naming SQL session files.

Return value

None.

Usage

- **dbrecftos** causes all SQL commands sent from the front-end application program to the server to be recorded in a human-readable file. This SQL session information is useful for debugging purposes.
- DB-Library creates one SQL session file for each call to **dbopen** that occurs after **dbrecftos** is called. Files are named *filename.n*, where *filename* is the name specified in the call to **dbrecftos** and *n* is an integer, starting with 0.

  For example, if *filename* is "foo," the first file created is named *foo.0*, the next *foo.1*, and so forth.

See also

dbopen

dbrecvpassthru

Description

Receive a TDS packet from a server.

Syntax

```
RETCODE dbrecvpassthru(dbproc, recv_bufp)

DBPROCESS     *dbproc;
DBVOIDPTR     *recv_bufp;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

recv_bufp
A pointer to a variable that **dbrecvpassthru** fills with the address of a buffer containing the TDS packet most recently received by this DBPROCESS connection. The application is not responsible for allocating this buffer.

Return value

DB_PASSTHRU_MORE, DB_PASSTHRU_EOM, or FAIL.

Usage

- **dbrecvpassthru** receives a TDS (Tabular Data Stream) packet from a server.

- TDS is an application protocol used for the transfer of requests and request results between clients and servers. Under ordinary circumstances, a DB-Library/C application does not have to deal directly with TDS, because DB-Library/C manages the data stream.
- **dbrecvpassthru** and **dbsendpassthru** are useful in gateway applications. When an application serves as the intermediary between two servers, it can use these routines to pass the TDS stream from one server to the other, eliminating the process of interpreting the information and re-encoding it.
- **dbrecvpassthru** reads a packet of bytes from the server connection identified by *dbproc* and sets **recv_bufp* to point to the buffer containing the bytes.
- A packet has a default size of 512 bytes. An application can change its packet size via **DBSETLPACKET**. For more information on packet sizes, see the dbgetpacket and DBSETLPACKET reference pages.
- **dbrecvpassthru** returns DB_PASSTHRU_EOM if the TDS packet has been marked by the server as EOM (End Of Message). If the TDS packet is not the last in the stream, **dbrecvpassthru** returns DB_PASSTHRU_MORE.
- A DBPROCESS connection which is used for a **dbrecvpassthru** operation cannot be used for any other DB-Library/C function until DB_PASSTHRU_EOM has been received.
- This is a code fragment using **dbrecvpassthru**:

```
/*
 ** The following code fragment illustrates the
 ** use of dbrecvpassthru() in an Open Server
 ** gateway application. It will continually get
 ** packets from a remote server, and pass them
 ** through to the client.
 **
 ** The routine srv_sendpassthru() is the Open
 ** Server counterpart required to complete
 ** this passthru operation.
 */
DBPROCESS   *dbproc;
 SRV_PROC    *srvproc;
 int         ret;
 BYTE        *packet;

while(1)
 {
      /* Get a TDS packet from the remote server */
```

```
  ret = dbrecvpassthru(dbproc, &packet);

if( ret == FAIL )
  {
       fprintf(stderr, "ERROR - dbrecvpassthru\
            failed in handle_results.\n");
       exit();
  }
 /* Now send the packet to the client */
  if( srv_sendpassthru(srvproc, packet,
       (int *)NULL) == FAIL )
  {
  fprintf(stderr, "ERROR - srv_sendpassthru \
            failed in handle_results.\n");
       exit();
  }
 /*
  ** We've sent the packet, so let's see if
  ** there's any more.
  */
  if( ret == DB_PASSTHRU_MORE )
       continue;
  else
       break;
         }
```

See also dbsendpassthru

dbregdrop

Description Drop a registered procedure.

Syntax RETCODE dbregdrop(dbproc, procedure_name, namelen)

```
DBPROCESS     *dbproc;
DBCHAR        *procedure_name;
DBSMALLINT     namelen;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

procedure_name
A pointer to the name of the registered procedure that the DBPROCESS connection wishes to drop.

namelen
The length of *procedure_name*, in bytes. If *procedure_name* is null-terminated, pass *namelen* as DBNULLTERM.

Return value

SUCCEED, DBNOPROC, or FAIL.

Usage

- **dbregdrop** drops a registered procedure from Open Server. Because a notification procedure is simply a special type of registered procedure, a notification procedure may also be dropped via **dbregdrop**.
- A DBPROCESS connection can drop any registered procedure defined in Open Server, including procedures created by other DBPROCESS connections and procedures created by other applications. Any mechanism to protect registered procedures must be embodied in the server application.
- If the procedure referenced by *procedure_name* is not defined in Open Server, **dbregdrop** returns DBNOPROC. An application can use **dbreglist** to obtain a list of registered procedures currently defined in Open Server.
- This is a code fragment that uses **dbregdrop**:

```
/*
** The following code fragment illustrates
 ** dropping a registered procedure.
 */
DBPROCESS   *dbproc;
 RETCODE     ret;
 char        *procname;

procname = "some_event";
ret = dbregdrop(dbproc, procname, DBNULLTERM);
 if (ret == FAIL)
 {
      fprintf(stderr, "ERROR: dbregdrop() \
           failed!!\n");
 }
 else if (ret == DBNOPROC)
 {
     fprintf(stderr, "ERROR: procedure %s was not\
           registered!\n", procname);
 }
```

See also

dbnpcreate, dbreglist

dbregexec

Description

Execute a registered procedure.

Syntax

```
RETCODE dbregexec(dbproc, options)

DBPROCESS      *dbproc;
DBUSMALLINT    options;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

options

A two-byte bitmask, either DBNOTIFYALL or DBNOTIFYNEXT.

If **options** is DBNOTIFYALL, Open Server will notify all DBPROCESSes watching for the execution of this registered procedure.

If **options** is DBNOTIFYNEXT, Open Server will notify only the DBPROCESS that has been watching the longest.

Return value

SUCCEED or FAIL.

Usage

- **dbregexec** completes the process of executing a registered procedure. Because a notification procedure is simply a special type of registered procedure, a notification procedure may also be executed via **dbregexec**.
- The procedure name and its parameters must have been previously defined via **dbreginit** and **dbregparam**.
- To execute a registered procedure, a DB-Library/C application must:
  - Initiate the call via **dbreginit**.
  - Describe the procedure's parameters, if any, via **dbregparam**.
  - Execute the procedure via **dbregexec**.
- An application cannot execute a registered procedure that is not defined in Open Server. **dbreglist** returns a list of registered procedures that are currently defined.
- Registered procedures are useful for inter-application communication and synchronization, because applications can request to be advised when a registered procedure executes.

- Registered procedures may be created only in Open Server. At this time, Adaptive Server does not support registered procedures. An application can use **dbnpcreate**, **dbregparam**, and **dbnpcreate** to create a registered procedure.
- A DB-Library/C application requests to be notified of a registered procedure's execution via **dbregwatch**. The application may request to be notified either synchronously or asynchronously.
- This is an example of executing a registered procedure:

```
DBPROCESS   *dbproc;
 DBINT       newprice = 55;
 DBINT       status;
/*
 ** Initiate execution of the registered procedure
 ** "price_change"
 */
 dbreginit (dbproc, "price_change", DBNULLTERM);

/*
 ** The registered procedure has two parameters:
 **      name        varchar(255)
 **      newprice    int
 ** So pass these parameters to the registered
 ** procedure.
 */
dbregparam (dbproc, "name", SYBVARCHAR, NULL,
      "sybase");
 dbregparam (dbproc, "newprice", SYBINT4, 4,
      &newprice);
/*  Execute the registered procedure: */
 status = dbregexec (dbproc, DBNOTIFYALL);
 if (status == FAIL)
 {
      fprintf(stderr, "ERROR:  Failed to execute \
           price_change!\n");
 }

else if (status == DBNOPROC)
 {
      fprintf(stderr, "ERROR:  Price_change does \
           not exist!\n");
 }
 else
 {
      fprintf(stdout, "Success in executing \
           price_change!\n");
```

```
}
```

See also dbreginit, dbregparam, dbregwatch, dbregnowatch

dbreghandle

Description Install a handler routine for a registered procedure notification.

Syntax

```
RETCODE dbreghandle(dbproc, procedure_name, namelen,
        handler)

DBPROCESS    *dbproc;
DBCHAR       *procedure_name;
DBSMALLINT   namelen;
INTFUNCPTR   handler;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

procedure_name

A pointer to the name of the registered procedure for which the handler is being installed.

If *procedure_name* is passed as NULL, the handler is installed as a default handler. The default handler will be called for all registered procedure notifications read by this DBPROCESS connection for which no other handler has been installed.

namelen

The length of *procedure_name*, in bytes. If *procedure_name* is null-terminated, pass *namelen* as DBNULLTERM.

handler

A pointer to the function to be called by DB-Library/C when the registered procedure notification is read.

If *handler* is passed as NULL, any handler previously installed for the registered procedure is uninstalled.

Return value SUCCEED or FAIL.

Usage

- **dbreghandle** installs a user-supplied handler routine to be called by DB-Library/C when a DBPROCESS connection reads an asynchronous notification that a registered procedure has been executed.

  Because a notification procedure is simply a special type of registered procedure, a handler for a notification procedure may also be installed via **dbreghandle**.

- An application receives an asynchronous notification only if it has previously called **dbregwatch** with *options* passed as DBNOWAITONE or DBNOWAITALL. This call tells Open Server that the application is interested in the execution of the registered procedure, that it will receive the notification asynchronously, and that it will read the notification through a particular DBPROCESS connection.

- If no handler is installed for a notification, DB-Library/C will raise an error when the DBPROCESS connection reads the notification.

- Either *procedure_name* or *handler* may be NULL:

  - If both *procedure_name* and *handler* are supplied, **dbreghandle** installs the handler specified by *handler* for the registered procedure specified by *procedure_name*.

  - If *procedure_name* is NULL and *handler* is NULL, **dbreghandle** uninstalls all handlers for this DBPROCESS connection.

  - If *procedure_name* is NULL but *handler* is supplied, **dbreghandle** installs the handler specified by *handler* as a “default” handler for this DBPROCESS connection. This default handler will be called whenever the DBPROCESS connection reads a registered procedure notification for which no other handler has been installed.

  - If *procedure_name* is supplied but *handler* is NULL, **dbreghandle** uninstalls any handler previously installed for this registered procedure. If a default handler has been installed for this DBPROCESS connection, it remains in effect and will be called if a *procedure_name* notification is read.

- The same handler may be used by several DBPROCESS connections, but it must be installed for each one by a separate call to **dbreghandle**. Because of the possibility of a single notification handler being called when different DBPROCESSes read notifications, all handlers should be written to be re-entrant.

- A single DBPROCESS connection may be watching for several registered procedures to execute. This connection may have different handlers installed to process the various notifications it may read. Each handler must be installed by a separate call to **dbreghandle**.
- A DBPROCESS connection may be idle, sending commands, reading results, or idle with results pending when a registered procedure notification arrives.
    - If the DBPROCESS connection is idle, it is necessary for the application to call **dbpoll** to enable the connection to read the notification. If a handler for the notification has been installed, it will be called before **dbpoll** returns.
    - If the DBPROCESS connection is sending commands, the notification is read and the notification handler called during **dbsqlexec** or **dbsqlok**. After the notification handler returns, flow of control continues normally.
    - If the DBPROCESS connection is reading results, the notification is read and the notification handler called either in **dbresults** or **dbnextrow**. After the notification handler returns, flow of control continues normally.
    - If the DBPROCESS connection is idle with results pending, the notification is not read until all results in the stream up to the notification have been read and processed by the connection.
- Because a notification may be read while a DBPROCESS connection is in any of several different states, the actions that a notification handler may take are restricted. A notification handler may not use an existing DBPROCESS to send a query to the server, process the results of a query, or call **dbcancel** or **dbcanquery**. A notification handler may, however, open a new DBPROCESS and use this new DBPROCESS to send queries and process results within the handler.
- A notification handler can read the arguments passed to the registered procedure upon execution. To do this, the handler can use the DB-Library/C routines **dbnumrets**, **dbrettype**, **dbretlen**, **dbretname**, and **dbretdata**.
- All notification handlers are called by DB-Library/C with the following parameters:
    - *dbproc*, a pointer to the DBPROCESS connection that has been watching for the notification

- *procedure_name*, a pointer to the name of the registered procedure that has been executed
- *reserved1*, a DBUSMALLINT parameter reserved for future use
- *reserved2*, a DBUSMALLINT parameter reserved for future use

- A notification handler must return INT_CONTINUE to indicate normal completion, or INT_EXIT to instruct DB-Library/C to abort the application and return control to the operating system.
- Notification handlers on the Windows platform must be declared with CS_PUBLIC, as shown in the example below. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well.
- This is an example of a notification handler:

```
DBINT CS_PUBLIC my_procedure_handler(dbproc,
            procedure_name, reserved1, reserved2)
/* The client connection */
 DBPROCESS       *dbproc;
 /* A null-terminated string */
 DBCHAR          *procedure_name;
 /* Reserved for future use */
 DBUSMALLINT     reserved1;
 /* Reserved for future use */
 DBUSMALLINT     reserved2;

{
      int     i, type;
      DBINT   len;
      char    *name;
      BYTE    *data;
      int     params;

     /*
      ** Find out how many parameters this
      ** procedure received.
      */
      params = dbnumrets(dbproc);

     i = 0;    /* Initialize counter */

    /* Now process each parameter in turn */
      while(i++ < params)
      {
           /* Get the parameter's datatype */
           type = dbrettype(dbproc, i);

          /* Get the parameter's length */
```

```
                    len = dbretlen(dbproc, i);
                  /* Get the parameter's name */
                    name = dbretname(dbproc, i);
                  /* Get a pointer to the parameter */
                    data = dbretdata(dbproc, i);
                   /* Process the parameter here */
                }
                return(INT_CONTINUE);
        }
```

See also dbregwatch, dbregnowatch, dbregparam, dbregexec

dbreginit

Description Initiate execution of a registered procedure.

Syntax RETCODE dbreginit(dbproc, procedure_name, namelen)

```
DBPROCESS    *dbproc;
DBCHAR       *procedure_name;
DBSMALLINT    namelen;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and server.

procedure_name
A pointer to the name of the registered procedure being executed.

namelen
The length of *procedure_name*, in bytes. If *procedure_name* is null-terminated, pass *namelen* as DBNULLTERM.

Return value SUCCEED or FAIL.

Usage

- **dbreginit** initiates the execution of a registered procedure. Because a notification procedure is simply a special type of registered procedure, execution of a notification procedure may also be initiated via **dbreginit**.
- To execute a registered procedure, a DB-Library/C application must:
  - Initiate the call via **dbreginit**

- Pass the procedure's parameters, if any, via **dbregparam**
- Execute the procedure via **dbregexec**

- This is an example of executing a registered procedure:

```
DBPROCESS   *dbproc;
 DBINT       newprice = 55;
 DBINT       status;

/*
 ** Initiate execution of the registered procedure
 ** "price_change".
 */
 dbreginit (dbproc, "price_change", DBNULLTERM);

/*
 ** The registered procedure has two parameters:
 **      name        varchar(255)
 **      newprice    int
 ** So pass these parameters to the registered
 ** procedure.
 */

dbregparam (dbproc, "name", SYBVARCHAR, NULL,
     "sybase");
 dbregparam (dbproc, "newprice", SYBINT4, 4, 4,
      &newprice);

/*  Execute the registered procedure: */
 status = dbregexec (dbproc, DBNOTIFYALL);
 if (status == FAIL)
 {
      fprintf(stderr, "ERROR:  Failed to execute \
           price_change!\n");
 }
 else if (status == DBNOPROC)
 {
      fprintf(stderr, "ERROR:  Price_change does \
           not exist!\n");
 }
 else
 {
      fprintf(stdout, "Success in executing \
           price_change!\n");
 }
```

See also dbregparam, dbregexec, dbregwatch, dbreglist, dbregwatchlist

dbreglist

Description

Return a list of registered procedures currently defined in Open Server.

Syntax

RETCODE dbreglist(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

Return value

SUCCEED or FAIL.

Usage

- **dbreglist** returns a list of registered procedures currently defined in Open Server. Because a notification procedure is simply a special type of registered procedure, notification procedures will be included in the list of registered procedures.
- The list of registered procedures is returned as rows that an application must explicitly process after calling **dbreglist**. Each row represents the name of a single registered procedure defined in Open Server. A row contains a single column of type SYBVARCHAR.
- The following code fragment illustrates how **dbreglist** might be used in an application:

```
DBPROCESS    *dbproc;
 DBCHAR       *procedurename;
 DBINT        ret;

/* request the list of procedures */
 if( (ret = dbreglist(dbproc)) == FAIL)
 {
      /* Handle failure here */
 }
 dbresults(dbproc);
 while( dbnextrow(dbproc) != NO_MORE_ROWS )
 {
      procedurename = (DBCHAR *)dbdata(dbproc, 1);
      procedurename[dbdatlen(dbproc, 1)] = '\0';

     fprintf(stdout, "The procedure '%s' is \
          defined.\n", procedurename);
 }
 /* All done */
```

See also

dbregwatchlist, dbregwatch

dbregnowatch

Description

Cancel a request to be notified when a registered procedure executes.

Syntax

```
RETCODE dbregnowatch(dbproc, procedure_name,
        namelen)

DBPROCESS    *dbproc;
DBCHAR       *procedure_name;
DBSMALLINT   namelen;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

procedure_name

A pointer to the name of the registered procedure that the DBPROCESS connection is no longer interested in.

namelen

The length of *procedure_name*, in bytes. If *procedure_name* is null-terminated, pass *namelen* as DBNULLTERM.

Return value

SUCCEED, DBNOPROC, or FAIL.

Usage

- **dbregnowatch** cancels a DBPROCESS connection's request to be notified when a registered procedure executes. Because a notification procedure is simply a special type of registered procedure, **dbregnowatch** also cancels a DBPROCESS connection's request to be notified when a notification procedure executes.
- It is meaningful to call **dbregnowatch** only if the DBPROCESS connection has previously requested an asynchronous notification via **dbregwatch**.
- If the procedure referenced by *procedure_name* is not defined in Open Server, **dbregnowatch** returns DBNOPROC. An application can obtain a list of procedures currently registered in Open Server via **dbreglist**.
- An application can obtain a list of registered procedures it is watching for via **dbregwatchlist**.
- This is an example of canceling a request to be notified:

```
DBPROCESS   *dbproc;
 DBINT       ret;

/*
```

```
** Inform the server that we no longer wish to
** be notified when "price_change" executes:
*/
ret = dbregnowatch (dbproc, "price_change",
     DBNULLTERM);
if (ret == DBNOPROC)
{
     /* The registered procedure must not exist */
     fprintf(stderr, "ERROR:  price_change \
          doesn't exist!\n");
}
```

See also dbregwatch, dbregwatchlist, dbreghandle, dbregexec

dbregparam

Description Define or describe a registered procedure parameter.

Syntax

```
RETCODE dbregparam(dbproc,param_name, type, datalen,
          data)

DBPROCESS    *dbproc;
char         *param_name;
int           type;
DBINT        datalen;
BYTE        *data;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

param_name

A pointer to the parameter name.

When creating a registered procedure, *param_name* is required.

When executing a registered procedure, *param_name* may be NULL. In this case, the registered procedure will expect to receive its parameters in the order in which they were originally defined.

type
A symbolic value indicating the datatype of the parameter. Legal data types are: SYBINT1, SYBINT2, SYBINT4, SYBREAL, SYBFLT8, SYBCHAR, SYBBINARY, SYBVARCHAR, SYBDATETIME4, SYBDATETIME, SYBMONEY4, and SYBMONEY.

Note that SYBTEXT and SYBIMAGE are not legal datatypes for parameters.

datalen
The length of the parameter.

When creating a registered procedure:

- *datalen* can be used to indicate that no default value is being supplied for this parameter. To indicate no default, pass *datalen* as DBNODEFAULT.
- *datalen* can be used to indicate that the default value for a parameter is NULL. This is different from having no default. To indicate a NULL default, pass *datalen* as 0.

When executing a registered procedure:

- *datalen* may be 0. In this case, *data* is ignored and NULL is passed to the registered procedure for this parameter.

data
A pointer to the parameter.

When creating a registered procedure, *data* can be used to provide a default value for the parameter. Pass *data* as pointing to the default value. If no default value is desired, pass *datalen* as DBNODEFAULT.

When executing a registered procedure, *data* may be passed as NULL.

Return value

SUCCEED or FAIL.

Usage

- **dbregparam** defines a registered procedure parameter. Because a notification procedure is simply a special type of registered procedure, **dbregparam** also defines a notification procedure parameter.

- **dbregparam** is called to define registered procedure parameters when a registered procedure is created and to describe the parameters when a registered procedure is executed.

  Note DB-Library/C applications can create only a special type of registered procedure, known as a notification procedure. A notification procedure differs from a normal Open Server registered procedure in that it contains no executable statements. For more information on notification procedures, see the dbnpdefine and dbnpcreate reference pages.

- Either **dbnpdefine**, which initiates the process of creating a notification procedure, or **dbreginit**, which initiates the process of executing a registered procedure, must be called before an application calls **dbregparam**.
- When creating a registered procedure:
  - To indicate that no default value is being supplied, pass *datalen* as DBNODEFAULT. *data* is ignored in this case.
  - To supply a default value of NULL, pass *datalen* as 0. *data* is ignored in this case.
  - To supply a default value that is not NULL pass *datalen* as the length of the value (or -1 if it is a fixed-length type), and *data* as pointing to the value.
- When executing a registered procedure:
  - To pass NULL as the value of the parameter, pass *datalen* as 0. In this case, *data* is ignored.
  - To pass a value for this parameter, pass *datalen* as the length of the value (or -1 if it is a fixed-length type), and *data* as pointing to the value.
- To create a notification procedure, a DB-Library/C application must:
  - Define the procedure via **dbnpdefine**
  - Describe the procedure's parameters, if any, via **dbregparam**
  - Create the procedure via **dbnpcreate**
- This is an example of creating a notification procedure:

```
DBPROCESS   *dbproc;
 DBINT       status;
/*
```

```
 ** Let's create a notification procedure called
 ** "message" which has two parameters:
 **      msg     varchar(255)
 **      userid  int
 */
/*

 ** Define the name of the notification procedure
 ** "message"
 */
 dbnpdefine (dbproc, "message", DBNULLTERM);
/* The notification procedure has two parameters:
 **      msg     varchar(255)
 **      userid  int
 ** So, define these parameters.  Note that both
 ** of these parameters are defined with a default
 ** value of NULL. Passing datalen as 0
 ** accomplishes this.
 */

dbregparam (dbproc, "msg", SYBVARCHAR, 0, NULL);
 dbregparam (dbproc, "userid", SYBINT4, 0, NULL);

/*  Create the notification procedure: */
 status = dbnpcreate (dbproc);
 if (status == FAIL)
 {
      fprintf(stderr, "ERROR:  Failed to create \
           message!\n");
 }
 else
 {
      fprintf(stdout, "Success in creating \
           message!\n");
 }
```

- To execute a registered procedure, a DB-Library/C application must:
  - Initiate the call via **dbreginit**
  - Pass the procedure's parameters, if any, via **dbregparam**
  - Execute the procedure via **dbregexec**
- This is an example of executing a registered procedure:

```
DBPROCESS   *dbproc;
 DBINT       newprice = 55;
 DBINT       status;
```

```
/*
 ** Initiate execution of the registered procedure
 ** "price_change".
 */
 dbreginit (dbproc, "price_change", DBNULLTERM);

/*
 ** The registered procedure has two parameters:
 **      name        varchar(255)
 **      newprice    int
 ** So pass these parameters to the registered
 ** procedure.
 */
 dbregparam (dbproc, "name", SYBVARCHAR, 6,
      "sybase");
 dbregparam (dbproc, "newprice", SYBINT4, -1,
      &newprice);

/*  Execute the registered procedure: */
 status = dbregexec (dbproc, DBNOTIFYALL);
 if (status == FAIL)
 {
      fprintf(stderr, "ERROR:  Failed to execute \
           price_change!\n");
 }
 else if (status == DBNOPROC)
 {
      fprintf(stderr, "ERROR:  Price_change does \
           not exist!\n");
 }
 else
 {
      fprintf(stdout, "Success in executing \
           price_change!\n");
 }
```

See also dbreginit, dbregexec, dbnpdefine, dbnpcreate, dbregwatch

dbregwatch

Description Request to be notified when a registered procedure executes.

Syntax RETCODE dbregwatch(dbproc, procedure_name,namelen, options)

```
DBPROCESS    *dbproc;
DBCHAR       *procedure_name;
DBSMALLINT   namelen;
DBUSMALLINT  options;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

procedure_name

A pointer to the name of a registered procedure. The registered procedure must be defined in Open Server.

namelen

The length of *procedure_name*, in bytes. If *procedure_name* is null-terminated, pass *namelen* as DBNULLTERM.

options

A two-byte bitmask: DBWAIT, DBNOWAITONE, or DBNOWAITALL.

If **options** is passed as DBWAIT, **dbregwatch** will not return until the DBPROCESS connection reads a synchronous notification that the registered procedure has executed.

If **options** is passed as DBNOWAITONE, **dbregwatch** returns immediately. The DBPROCESS connection will receive an asynchronous notification when the registered procedure executes. The connection will receive only a single notification, even if the registered procedure executes multiple times.

If **options** is passed as DBNOWAITALL, **dbregwatch** returns immediately. The DBPROCESS connection will receive an asynchronous notification when the registered procedure executes. The connection will continue to receive notifications, one for each execution of the registered procedure, until it informs Open Server that it no longer wishes to receive them.

Return value

SUCCEED, DBNOPROC, or FAIL.

dbregwatch returns FAIL if no handler is installed for the registered procedure.

Usage

- **dbregwatch** informs Open Server that a DBPROCESS connection should be notified when a particular registered procedure executes. Because a notification procedure is simply a special type of registered procedure, **dbregwatch** also informs Open Server that a DBPROCESS connection should be notified when a particular notification procedure executes.

- The connection can request to be notified synchronously or asynchronously:
  - To request synchronous notification, an application passes *options* as DBWAIT in its call to **dbregwatch**. In this case, **dbregwatch** will not return until the DBPROCESS connection reads the notification that the registered procedure has executed.

    Open Server will send only a single notification as the result of a synchronous notification request. If the registered procedure executes a second time, after the synchronous request has been satisfied, the client will not receive a second notification, unless another notification request is made.
  - To request asynchronous notification, an application passes *options* as DBNOWAITONE or DBNOWAITALL in its call to **dbregwatch**. In this case, **dbregwatch** returns immediately. A return code of SUCCEED indicates that Open Server has accepted the request.

    If *options* is DBNOWAITONE, Open Server will send only a single notification, even if the registered procedure executes multiple times.

    If *options* is DBNOWAITALL, Open Server will continue to send a notification every time the registered procedure executes, until it is informed, via **dbregnowatch**, that the client no longer wishes to receive them.
- A DBPROCESS connection may be idle, sending commands, reading results, or idle with results pending when an asynchronous registered procedure notification arrives.
  - If the DBPROCESS connection is idle, it is necessary for the application to call **dbpoll** to enable the connection to read the notification. If a handler for the notification has been installed, it will be called before **dbpoll** returns.
  - If the DBPROCESS connection is sending commands, the notification is read and the notification handler called during **dbsqlexec** or **dbsqlok**. After the notification handler returns, flow of control continues normally.
  - If the DBPROCESS connection is reading results, the notification is read and the notification handler called either in **dbresults** or **dbnextrow**. After the notification handler returns, flow of control continues normally.

  - If the DBPROCESS connection is idle with results pending, the notification is not read until all results in the stream up to the notification have been read and processed by the connection.

- An application must install a handler to process the registered procedure notification before calling **dbregwatch**. If no handler is installed, **dbregwatch** returns FAIL. An application can install a notification handler via **dbreghandle**.

  If the handler is uninstalled after the application calls **dbregwatch** but before the registered procedure notification is received, DB-Library/C raises an error when the notification is received.

- If the procedure referenced by *procedure_name* is not defined in Open Server, **dbregwatch** returns DBNOPROC. An application can obtain a list of procedures currently registered in Open Server via **dbreglist**.

- An application can obtain a list of registered procedures it is watching for via **dbregwatchlist**.

- This is an example of making a synchronous notification request:

```
DBPROCESS   *dbproc;
 DBINT       handlerfunc;
 DBINT       ret;

/*
 ** The registered procedure is defined in Open
 ** Server as:
 **      shutdown   msg_param   varchar(255)
 */

/*
 ** First install the handler for this registered
 ** procedure:
 */
 dbreghandle(dbproc, "shutdown", DBNULLTERM,
      handlerfunc);

/* Make the notification request and wait: */
 ret = dbregwatch(dbproc, "shutdown", DBNULLTERM,
      DBWAIT);

if (ret == FAIL)
 {
      fprintf (stderr, "ERROR:  dbregwatch() \
           failed!\n");
 }
 else if (ret == DBNOPROC)
```

```
{
     fprintf (stderr, "ERROR: procedure shutdown \
          not defined.\n");
}
else
{
     /*
     ** The registered procedure notification has
     ** been returned, and our registered
     ** procedure handler has already been called.
     */
}
```

- This is an example of making an asynchronous notification request:

```
DBPROCESS   *dbproc;
 DBINT       handlerfunc;
 DBINT       ret;

/*
 ** The registered procedure is defined in Open
 ** Server as:
 **     shutdown   msg_param   varchar(255)
 */

/*
 ** First install the handler for this registered
 ** procedure:
 */
 dbreghandle(dbproc, "shutdown", DBNULLTERM,
      handlerfunc);

/* Make the asynchronous notification request: */
 ret = dbregwatch(dbproc, "shutdown", DBNULLTERM,
      DBNOWAITALL);

if (ret == FAIL)
 {
      fprintf (stderr, "ERROR:  dbregwatch() \
           failed!\n");
 }
 else if (ret == DBNOPROC)
 {
     fprintf (stderr, "ERROR: procedure shutdown \
          not defined.\n");
 }

/*
 ** Since we are making use of the asychronous
```

```
 ** registered procedure notification mechanism,
 ** the application can continue doing other work
 ** while waiting for the notification.  All we
 ** have to do is call dbpoll() once in a while to
 ** read the registered procedure notification
 ** when it arrives.
 */

while (1)
 {
      /* Have dbpoll() block for one second */
      dbpoll (NULL, 1000, NULL, &ret);

     /*
      ** If we got the notification, then exit
      ** the loop
      */
      if (ret == DBNOTIFICATION)
           break;
      /*  Handle other program tasks here */
 }
```

See also dbpoll, dbregexec, dbregparam, dbreglist, dbregwatchlist, dbregnowatch

dbregwatchlist

Description Return a list of registered procedures that a DBPROCESS is watching for.

Syntax RETCODE dbregwatchlist(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value SUCCEED or FAIL.

Usage

- **dbregwatchlist** returns a list of registered procedures that a DBPROCESS connection is watching for. Because a notification procedure is simply a special type of registered procedure, the list returned by **dbregwatchlist** will include notification procedures.

- The list of registered procedures is returned as rows that an application must explicitly process after calling **dbregwatchlist**. Each row represents the name of a single registered procedure for which the DBPROCESS has requested notification. A row contains a single column of type SYBVARCHAR.
- The following code fragment illustrates how **dbregwatchlist** might be used in an application:

```
DBPROCESS    *dbproc;
 DBCHAR       *procedurename;
 DBINT        ret;

/* Request the list of procedures */
 if( (ret = dbregwatchlist(dbproc)) == FAIL)
 {
      /* Handle failure here */
 }
 dbresults(dbproc);
 while( dbnextrow(dbproc) != NO_MORE_ROWS )
 {
      procedurename = (DBCHAR *)dbdata(dbproc, 1);
      procedurename[dbdatlen(dbproc, 1)] = '\0';

   fprintf(stdout, "we're waiting for \
         procedure '%s'.\n", procedurename);
 }
 /* All done */
```

See also dbregwatch, dbresults, dbnextrow

dbresults

Description Set up the results of the next query.

Syntax RETCODE dbresults(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value SUCCEED, FAIL or NO_MORE_RESULTS.

dbresults returns NO_MORE_RESULTS if all commands in the buffer have already been processed. The most common reason for failing is a runtime error, such as a database permission violation.

The number of commands in the command buffer determines whether **dbsqlexec** or **dbresults** traps a runtime error. If the buffer contains only a single command, a runtime error will cause **dbsqlexec** to fail. If the command buffer contains multiple commands, a runtime error will not cause **dbsqlexec** to fail. Instead, the **dbresults** call that processes the command causing the runtime error will fail.

The situation is a bit more complicated for runtime errors and stored procedures. A runtime error on an **execute** command may cause **dbresults** to fail, in accordance with the rule given in the previous paragraph. A runtime error on a statement inside a stored procedure will not cause **dbresults** to fail, however. For example, if the stored procedure contains an **insert** statement and the user does not have insert permission on the database table, the **insert** statement will fail, but **dbresults** will still return SUCCEED. To check for runtime errors inside stored procedures, use the **dbretstatus** routine to look at the procedure's return status, and trap relevant server messages inside your message handler.

Usage

- This routine sets up the next command in the command batch for processing. The application program calls it after **dbsqlexec** or **dbsqlok** returns SUCCEED. The first call to **dbresults** will always return either SUCCEED or NO_MORE_RESULTS if the call to **dbsqlexec** or **dbsqlok** has returned SUCCEED. Once **dbresults** returns SUCCEED, the application typically processes any result rows with **dbnextrow**.
- If a command batch contains only a single command, and that command does not return rows, for example a "use database" command, a DB-Library/C application does not have to call **dbresults** to process the results of the command. However, if the command batch contains more than one command, a DB-Library/C application must call **dbresults** once for every command in the batch, whether or not the command returns rows.

  dbresults must also be called at least once for any stored procedure executed in a command batch, whether or not the stored procedure returns rows. If the stored procedure contains more than one Transact-SQL **select**, then **dbresults** must be called once for each **select**.

To ensure that **dbresults** is called the correct number of times, Sybase strongly recommends that **dbresults** always be called in a loop that terminates when **dbresults** returns NO_MORE_RESULTS.

Note All Transact-SQL commands are considered commands by **dbresults**. For a list of Transact-SQL commands, see the *Adaptive ServerEnterprise Reference Manual*.

- To cancel the remaining results from the command batch (and eliminate the need to continue calling **dbresults** until it returns NO_MORE_RESULTS), call **dbcancel**.
- To determine whether a particular command is one that returns rows and needs results processing with **dbnextrow**, call **DBROWS** after the **dbresults** call.
- The typical sequence of calls for using **dbresults** with **dbsqlexec** is:

```
DBINT         xvariable;
 DBCHAR        yvariable[10];
 RETCODE       return_code;

/* Read the query into the command buffer */
 dbcmd(dbproc, "select x = 100, y = 'hello'");

/* Send the query to SQL Server */
 dbsqlexec(dbproc);

/*
 ** Get ready to process the results of the query.
 ** Note that dbresults is called in a loop even
 ** though only a single set of results is expected.
 ** This is simply because it is good programming
 ** practice to always code dbresults calls in loop.
 */

while ((return_code
 =dbresults(dbproc)!=NO_MORE_RESULTS)
 {
     if ((return_code == SUCCEED)
           & & (DBROWS(dbproc) == SUCCEED))
     {
        /* Bind column data to program variables */
         dbbind(dbproc, 1, INTBIND, (DBINT) 0,
                (BYTE *) &xvariable);
         dbbind(dbproc, 2, STRINGBIND, (DBINT) 0,
                 yvariable);
```

```
            /* Now process each row */
              while (dbnextrow(dbproc) != NO_MORE_ROWS)
             {
                C-code to print or process row data
             }
        }
    }
```

Example 1 of the online sample programs shows how to use **dbresults** to process a multiquery command batch.

- To manage multiple input data streams efficiently, an application can confirm that unread bytes are available, either in the DB-Library network buffer or in the network itself. The application can either:
  - (For UNIX or AOS/VS only) call **DBRBUF** to test whether bytes remain in the network buffer, and call **DBIORDESC** and the UNIX **select** or AOS/VS equivalent to test whether bytes remain in the network itself, or
  - (For all systems) call **dbpoll**.
- Another use for **dbresults** is to process the results of a remote procedure call made with **dbrpcsend**. See the dbrpcsend reference page for details.

See also

dbbind, dbcancel, dbnextrow, dbpoll, DBRBUF, dbretstatus, DBROWS, dbrpcsend, dbsqlexec, dbsqlok

dbretdata

Description

Return a pointer to a return parameter value generated by a stored procedure.

Syntax

BYTE *dbretdata(dbproc, retnum)

```
DBPROCESS   *dbproc;
int         retnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

retnum

The number of the return parameter value of interest. The first return value is 1. Values are returned in the same order as the parameters were originally specified in the stored procedure's **create procedure** statement. (Note that this is not necessarily the same order as specified in the remote procedure call.) When specifying *retnum*, non-return parameters do not count. For example, if the second parameter in a stored procedure is the only return parameter, its *retnum* is 1, not 2.

Return value

A pointer to the specified return value. If *retnum* is out of range, **dbretdata** returns NULL. To determine whether the data really has a null value (and *retnum* is not merely out of range), check for a return of 0 from **dbretlen**.

Usage

- **dbretdata** returns a pointer to a return parameter value generated by a stored procedure. It is useful in conjunction with remote procedure calls and **execute** statements on stored procedures.
- Transact-SQL stored procedures can return values for specified "return parameters." Changes made to the value of a return parameter inside the stored procedure are then available to the program that called the procedure. This is analogous to the "pass by reference" facility available in some programming languages.

  For a parameter to function as a return parameter, it must be declared as such within the stored procedure. The **execute** statement or remote procedure call that calls the stored procedure must also indicate that the parameter should function as a return parameter. In the case of a remote procedure call, it is the **dbrpcparam** routine that specifies whether a parameter is a return parameter.
- When executing a stored procedure, the server returns any parameter values immediately after returning all other results. Therefore, the application can call **dbretdata** only after processing the stored procedure's results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbretdata** or any other routines that process return parameters, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.)
- If a stored procedure is invoked with a remote procedure call, the return parameter values are automatically available to the application. If, on the other hand, the stored procedure is invoked with an **execute** statement, the return parameter values are available only if the command batch containing the **execute** statement uses local variables, not constants, for the return parameters.

- The routine **dbnumrets** indicates how many return parameter values are available. If **dbnumrets** returns less than or equal to 0, no return parameter values are available.
- When a stored procedure is invoked with an RPC command (using **dbrpcinit**, **dbrpcparam**, and **dbrpcsend**), then the return parameter values can be retrieved after all other results have been processed. For an example of this usage, see Example 8 of the online sample programs.
- When a stored procedure has been executed from a batch of Transact-SQL commands (with **dbsqlexec** or **dbsqlsend**), then other commands might execute after the stored procedure. This situation makes retrieval of return parameter values a little more complicated.
  - If you are sure that the stored procedure command is the only command in the batch, then you can retrieve the return parameter values after the **dbresults** loop, as shown in Example 8 of the online sample programs.
  - If the batch can contain multiple commands, then the return parameter values should be retrieved inside the **dbresults** loop, after all rows have been fetched with **dbnextrow**. The code below shows where the return parameters should be retrieved in this situation.

```
while ( (result_code = dbresults(dbproc)
         != NO_MORE_RESULTS)
 {
   if (result_code == SUCCEED)
   {
     ... bind rows here ...
     while ((row_code = dbnextrow(dbproc))
            != NO_MORE_ROWS)
     {
       ... process rows here ...
     }
     /* Now check for a return status */
     if (dbhasretstat(dbproc) == TRUE)
     {
       printf("(return status %d)\n",
              dbretstatus(dbproc));
     }
     /* Now check for return parameter values */
     if (dbnumrets(dbproc) > 0)
     {
       ... retrieve output parameters here ...
     }
   } /* if result_code */
```

```
    else
    {
      printf("Query failed.\n");
    }
  } /* while dbresults */
```

- The routines below are used to retrieve return parameter values:
  - **dbnumrets** returns the total number of return parameter values.
  - **dbretlen** returns the length of a parameter value.
  - **dbretname** returns the name of a parameter value.
  - **dbrettype** returns the datatype of a parameter value.
  - **dbconvert** converts the value to another datatype, if necessary.

  The code fragment below shows how these routines are used together:

```
char    dataval[512];
 char    *dataname;
 DBINT   datalen;
 int i,  numrets;
  numrets = dbnumrets(dbproc);

 for (i = 1; i <= numrets; i++)
 {
   dataname = dbretname(dbproc, i);
   datalen = dbretlen(dbproc, i);
   if (datalen == 0)
   {
     /* The parameter's value is NULL */
     strcpy(dataval, "NULL");
   }
   else
   {
     /*
     ** Convert to char. dbconvert appends a null
     ** terminator because we pass the last
     ** parameter, destlen, as -1.
     */
     result = dbconvert(dbproc,
             dbrettype(dbproc, i),
             dbretdata(dbproc, i), datalen,
             SYBCHAR, (BYTE *)dataval, -1);
   } /* else */
   /* Now print out the converted value */
   if (dataname == NULL || *dataname == '\0')
```

```
            printf("\t%s\n", dataval);
        else
            printf("\t%s: %s\n", dataname, dataval);
    }
```

See also dbnextrow, dbnumrets, dbresults, dbretlen, dbretname, dbrettype, dbrpcinit, dbrpcparam

dbretlen

Description Determine the length of a return parameter value generated by a stored procedure.

Syntax DBINT dbretlen(dbproc, retnum)

```
DBPROCESS   *dbproc;
int         retnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

retnum

The number of the return parameter value of interest. The first return value is 1. Values are returned in the same order as the parameters were originally specified in the stored procedure's **create procedure** statement. (Note that this is not necessarily the same order as specified in the remote procedure call.) When specifying *retnum*, non-return parameters do not count. For example, if the second parameter in a stored procedure is the only return parameter, its *retnum* is 1, not 2.

Return value The length of the specified return parameter value. If *retnum* is out of range, **dbretlen** returns -1. If the return value is null, **dbretlen** returns 0.

Usage

- **dbretlen** returns the length of a particular return parameter value generated by a stored procedure. It is useful in conjunction with remote procedure calls and **execute** statements on stored procedures.

- Transact-SQL stored procedures can return values for specified “return parameters.” Changes made to the value of a return parameter inside the stored procedure are then available to the program that called the procedure. This is analogous to the “pass by reference” facility available in some programming languages.

  For a parameter to function as a return parameter, it must be declared as such within the stored procedure. The **execute** statement or remote procedure call that calls the stored procedure must also indicate that the parameter should function as a return parameter. In the case of a remote procedure call, it is the **dbrpcparam** routine that specifies whether a parameter is a return parameter.

- When executing a stored procedure, the server returns any parameter values immediately after returning all other results. Therefore, the application can call **dbretlen** only after processing the stored procedure’s results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbretlen** or any other routines that process return parameters, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.)

- If the stored procedure is invoked with a remote procedure call, the return parameter values are automatically available to the application. If, on the other hand, the stored procedure is invoked with an **execute** statement, the return parameter values are available only if the command batch containing the **execute** statement uses local variables, not constants, for the return parameters.

- Other routines return additional information about return parameter values:

  - **dbnumrets** returns the total number of return parameter values.
  - **dbretdata** returns a pointer to a parameter value.
  - **dbretname** returns the name of a parameter value.
  - **dbrettype** returns the datatype of a parameter value.
  - **dbconvert** converts the value to another datatype, if necessary.

- For an example of this routine, see the dbretdata reference page.

See also

dbnextrow, dbnumrets, dbresults, dbretdata, dbretname, dbrettype, dbrpcinit, dbrpcparam

dbretname

Description

Determine the name of the stored procedure parameter associated with a particular return parameter value.

Syntax

```
char *dbretname(dbproc, retnum)

DBPROCESS    *dbproc;
int          retnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

retnum

The number of the return parameter value of interest. The first return value is 1. Values are returned in the same order as the parameters were originally specified in the stored procedure's **create procedure** statement. (Note that this is not necessarily the same order as specified in the remote procedure call.) When specifying *retnum*, non-return parameters do not count. For example, if the second parameter in a stored procedure is the only return parameter, its *retnum* is 1, not 2.

Return value

A pointer to the null-terminated parameter name for the specified return value. If *retnum* is out of range, **dbretname** returns NULL.

Usage

- **dbretname** returns a pointer to the null-terminated parameter name associated with a return parameter value from a stored procedure. It is useful in conjunction with remote procedure calls and **execute** statements on stored procedures.
- Transact-SQL stored procedures can return values for specified "return parameters." Changes made to the value of a return parameter inside the stored procedure are then available to the program that called the procedure. This is analogous to the "pass by reference" facility available in some programming languages.

  For a parameter to function as a return parameter, it must be declared as such within the stored procedure. The **execute** statement or remote procedure call that calls the stored procedure must also indicate that the parameter should function as a return parameter. In the case of a remote procedure call, it is the **dbrpcparam** routine that specifies whether a parameter is a return parameter.

- When executing a stored procedure, the server returns any parameter values immediately after returning all other results. Therefore, the application can call **dbretname** only after processing the stored procedure's results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbretname** or any other routines that process return parameters, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.)
- If the stored procedure is invoked with a remote procedure call, the return parameter values are automatically available to the application. If, on the other hand, the stored procedure is invoked with an **execute** statement, the return parameter values are available only if the command batch containing the **execute** statement uses local variables, not constants, for the return parameters.
- Other routines return additional information about return parameter values:
  - **dbnumrets** returns the total number of return parameter values.
  - **dbretdata** returns a pointer to a parameter value.
  - **dbretlen** returns the length of a parameter value.
  - **dbrettype** returns the datatype of a parameter value.
  - **dbconvert** converts the value to another datatype, if necessary.
- For an example of this routine, see the dbretdata reference page.

See also

dbnextrow, dbnumrets, dbresults, dbretdata, dbretlen, dbrettype, dbrpcinit, dbrpcparam

dbretstatus

Description

Determine the stored procedure status number returned by the current command or remote procedure call.

Syntax

```
DBINT dbretstatus(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The return status number for the current command.

Usage

- **dbretstatus** fetches a stored procedure's status number. All stored procedures that are run on SQL Server 4.0 or later return a status number. Stored procedures that complete normally return a status number of 0. For a list of return status numbers, see the *Adaptive Server Enterprise Reference Manual*.
- The **dbhasretstat** routine determines whether the current Transact-SQL command or remote procedure call actually generated a return status number. Since status numbers are a feature of stored procedures, only a remote procedure call or a Transact-SQL command that executes a stored procedure can generate a status number.
- When executing a stored procedure, the server returns the status number immediately after returning all other results. Therefore, the application can call **dbretstatus** only after processing the stored procedure's results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbretstatus** or **dbhasretstat**, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.)
- The order in which the application processes the status number and any return parameter values is unimportant.
- When a stored procedure has been executed from a batch of Transact-SQL commands (with **dbsqlexec** or **dbsqlsend**), then other commands might execute after the stored procedure. This situation makes return-status retrieval a little more complicated.
  - If you are sure that the stored procedure command is the only command in the batch, then you can retrieve the return status after the **dbresults** loop, as shown in Example 8 of the online sample programs.
  - If the batch can contain multiple commands, then the return status should be retrieved inside the **dbresults** loop, after all rows have been fetched with **dbnextrow**. For an example of how return statuses are retrieved in this situation, see the dbhasretstat reference page.

- For an example of this routine, see the dbhasretstat reference page.

See also

dbhasretstat, dbnextrow, dbresults, dbretdata, dbrpcinit, dbrpcparam, dbrpcsend

dbrettype

Description

Determine the datatype of a return parameter value generated by a stored procedure.

Syntax

int dbrettype(dbproc, retnum)

```
DBPROCESS   *dbproc;
int         retnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

retnum

The number of the return parameter value of interest. The first return value is 1. Values are returned in the same order as the parameters were originally specified in the stored procedure's **create procedure** statement. (Note that this is not necessarily the same order as specified in the remote procedure call.) When specifying *retnum*, non-return parameters do not count. For example, if the second parameter in a stored procedure is the only return parameter, its *retnum* is 1, not 2.

Return value

A token value for the datatype of the specified return value.

In a few cases, the token value returned by this routine may not correspond exactly with the column's server datatype:

- SYBVARCHAR is returned as SYBCHAR.
- SYBVARBINARY is returned as SYBBINARY.
- SYBDATETIMN is returned as SYBDATETIME.
- SYBMONEYN is returned as SYBMONEY.
- SYBFLTN is returned as SYBFLT8.

- SYBINTN is returned as SYBINT1, SYBINT2, or SYBINT4, depending on the actual type of the SYBINTN.

  If *retnum* is out of range, -1 is returned.

Usage

- **dbrettype** returns the datatype of a return parameter value generated by a stored procedure. It is useful in conjunction with remote procedure calls and **execute** statements on stored procedures.
- Transact-SQL stored procedures can return values for specified "return parameters." Changes made to the value of a return parameter inside the stored procedure are then available to the program that called the procedure. This is analogous to the "pass by reference" facility available in some programming languages.

  For a parameter to function as a return parameter, it must be declared as such within the stored procedure. The **execute** statement or remote procedure call that calls the stored procedure must also indicate that the parameter should function as a return parameter. In the case of a remote procedure call, it is the **dbrpcparam** routine that specifies whether a parameter is a return parameter.
- When executing a stored procedure, the server returns any parameter values immediately after returning all other results. Therefore, the application can call **dbrettype** only after processing the stored procedure's results by calling **dbresults**, as well as **dbnextrow** if appropriate. (Note that a stored procedure can generate several sets of results—one for each **select** it contains. Before the application can call **dbrettype** or any other routines that process return parameters, it must call **dbresults** and **dbnextrow** as many times as necessary to process all the results.)
- If the stored procedure is invoked with a remote procedure call, the return parameter values are automatically available to the application. If, on the other hand, the stored procedure is invoked with an **execute** statement, the return parameter values are available only if the command batch containing the **execute** statement uses local variables, not constants, for the return parameters.
- **dbrettype** actually returns an integer token value for the datatype (SYBCHAR, SYBFLT8, and so on). To convert the token value into a readable token string, use **dbprtype**. See the dbprtype reference page for a list of all token values and their equivalent token strings.
- For a list of server datatypes, see Types on page 408.
- The routines return additional information about return parameter values:
  - **dbnumrets** returns the total number of return parameter values.

- **dbretdata** returns a pointer to a parameter value.
- **dbretlen** returns the length of a parameter value.
- **dbretname** returns the name of a parameter value.
- **dbconvert** converts the value to another datatype, if necessary.

- For an example of this routine, see the dbretdata reference page.

See also: dbnextrow, dbnumrets, dbprtype, dbresults, dbretdata, dbretlen, dbretname, dbrpcinit, dbrpcparam

DBROWS

Description

Indicate whether the current command actually returned rows.

Syntax

```
RETCODE DBROWS(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

SUCCEED or FAIL, indicating whether the current command returned rows.

Usage

- This macro determines whether the command currently being processed by **dbresults** returned any rows. The application can call it after **dbresults** returns SUCCEED.
- The application must not call DBROWS after **dbnextrow**. The macro may return the wrong result at that time.
- The application can use DBROWS to determine whether it needs to call **dbnextrow** to process result rows. If DBROWS returns FAIL, the application can skip the **dbnextrow** calls.
- The **DBCMDROW** macro determines whether the current command is one that can return rows (that is, a Transact-SQL **select** statement or an **execute** on a stored procedure containing a **select**).

See also

DBCMDROW, dbnextrow, dbresults, DBROWTYPE

DBROWTYPE

Description
Return the type of the current row.

Syntax
STATUS DBROWTYPE(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value
Three different types of values can be returned:

- If the current row is a regular row, REG_ROW is returned.
- If the current row is a compute row, the *computeid* of the row is returned. (See the dbaltbind reference page for information on the *computeid*.)
- If no rows have been read, or if the routine failed for any reason, NO_MORE_ROWS is returned.

Usage
- This macro tells you the type (regular or compute) of the current row. Usually you already know this, since **dbnextrow** also returns the row type.

See also
dbnextrow

dbrpcinit

Description
Initialize a remote procedure call.

Syntax
RETCODE dbrpcinit(dbproc, rpcname, options)

```
DBPROCESS    *dbproc;
char         *rpcname;
DBSMALLINT    options;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

rpcname
A pointer to the name of the stored procedure to be invoked.

options
A 2-byte bitmask of RPC options. So far, the only option available is DBRPCRECOMPILE, which causes the stored procedure to be recompiled before it is executed.

Return value

SUCCEED or FAIL.

Usage

- An application can call a stored procedure in two ways: by executing a command buffer containing a Transact-SQL **execute** statement or by making a remote procedure call (RPC).
- Remote procedure calls have a few advantages over **execute** statements:
  - An RPC passes the stored procedure's parameters in their native datatypes, in contrast to the **execute** statement, which passes parameters as ASCII characters. Therefore, the RPC method is faster and usually more compact than the **execute** statement, because it does not require either the application program or the server to convert between native datatypes and their ASCII equivalents.
  - It is simpler and faster to accommodate stored procedure return parameters with an RPC, instead of an **execute** statement. With an RPC, the return parameters are automatically available to the application. (Note, however, that a return parameter must be specified as such when it is originally added to the RPC via the **dbrpcparam** routine.) If, on the other hand, a stored procedure is called with an **execute** statement, the return parameter values are available only if the command batch containing the **execute** statement uses local variables, not constants, as the return parameters. This involves additional parsing each time the command batch is executed.
- To make a remote procedure call, first call **dbrpcinit** to specify the stored procedure that is to be invoked. Then call **dbrpcparam** once for each of the stored procedure's parameters. Finally, call **dbrpcsend** to signify the end of the parameter list. This causes the server to begin executing the specified procedure. You can then call **dbsqlok**, **dbresults**, and **dbnextrow** to process the stored procedure's results. (Note that you will need to call **dbresults** multiple times if the stored procedure contains more than one **select** statement.) After all of the stored procedure's results have been processed, you can call the routines that process return parameters and status numbers, such as **dbretdata** and **dbretstatus**.
- If the procedure being executed resides on a server other than the one to which the application is directly connected, commands executed within the procedure cannot be rolled back.

- For an example of a remote procedure call, see Example 8 in the online sample programs.

See also dbnextrow, dbresults, dbretdata, dbretstatus, dbrpcparam, dbrpcsend, dbsqlok

dbrpcparam

Description Add a parameter to a remote procedure call.

Syntax

```
RETCODE dbrpcparam(dbproc, paramname, status, type,
        maxlen, datalen, value)

DBPROCESS   *dbproc;
char        *paramname;
BYTE        status;
int         type;
DBINT       maxlen;
DBINT       datalen;
BYTE        *value;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

paramname

A pointer to the name of the parameter to be invoked. This name must begin with the "@" character, which prefixes all stored procedure parameter names. As in the Transact-SQL **execute** statement, the name is optional. If it is not used, it should be specified as NULL. In that case, the order of the **dbrpcparam** calls determines the parameter to which each refers.

status

A 1-byte bitmask of RPC-parameter options. So far, the only option available is DBRPCRETURN, which signifies that the application program would like this parameter used as a return parameter.

type

A symbolic constant indicating the datatype of the parameter (for example, SYBINT1, SYBCHAR, and so on). Parameter values should be sent to the server in a datatype that matches the Adaptive Server datatype with which the corresponding stored procedure parameter was defined—see Types on page 408 for a list of type constants and the corresponding Adaptive Server datatypes.

maxlen

For return parameters, this is the maximum desired byte length for the RPC parameter value returned from the stored procedure. *maxlen* is relevant only for values whose datatypes are not fixed in length—that is, char, text, binary, and image values. If this parameter does not apply (that is, if the *type* is a fixed length datatype such as SYBINT2) or if you do not care about restricting the lengths of return parameters, set *maxlen* to -1. *maxlen* should also be set to -1 for parameters not designated as return parameters.

datalen

The length, in bytes, of the RPC parameter to pass to the stored procedure. This length should not count any null terminator.

If *type* is SYBCHAR, SYBVARCHAR, SYBBINARY, SYBVARBINARY, SYBBOUNDARY, or SYBSENSITIVITY, *datalen* must be specified. Passing *datalen* as -1 for any of these datatypes results in the DBPROCESS referenced by *dbproc* being marked as "dead," or unusable.

If *type* is a fixed length datatype, for example, SYBINT2, pass *datalen* as -1.

If the value of the RPC parameter is NULL, pass *datalen* as 0, even if *type* is a fixed-length datatype.

value

A pointer to the RPC parameter itself. If *datalen* is 0, this pointer will be ignored and treated as NULL. Note that DB-Library does not copy **value* into its internal buffer space until the application calls **dbrpcsend**. An application must not write over **value* until after it has called **dbrpcsend**.

The value of *type* indicates the datatype of **value*. See Types on page 408 for more information. For types that have no C equivalent, such as SYBDATETIME, SYBMONEY, SYBNUMERIC, or SYBDECIMAL, use dbconvert_ps to initialize **value*.

Note An application must not write over **value* until after it has called **dbrpcsend** to send the remote procedure call to the server. This is a functional change from previous versions of DB-Library.

Return value

SUCCEED or FAIL.

Usage

- An application can call a stored procedure in two ways: by executing a command buffer containing a Transact-SQL **execute** statement or by making a remote procedure call (RPC). See the reference page for dbrpcinit for a discussion of the differences between these techniques.
- To make a remote procedure call, first call **dbrpcinit** to specify the stored procedure that is to be invoked. Then call **dbrpcparam** once for each of the stored procedure's parameters. Finally, call **dbrpcsend** to signify the end of the parameter list. This causes the server to begin executing the specified procedure. You can then call **dbsqlok**, **dbresults**, and **dbnextrow** to process the stored procedure's results. (Note that you will need to call **dbresults** multiple times if the stored procedure contains more than one **select** statement.) After all of the stored procedure's results have been processed, you can call the routines that process return parameters and status numbers, such as **dbretdata** and **dbretstatus**.
- If *type* is SYBCHAR, SYBVARCHAR, SYBBINARY, SYBVARBINARY, SYBBOUNDARY, and SYBSENSITIVITY, *datalen* must be specified. Passing *datalen* as -1 for any of these datatypes results in the DBPROCESS referenced by *dbproc* being marked as "dead," or unusable.
- If *type* is SYBNUMERIC or SYBDECIMAL, use dbconvert_ps to initialize the DBNUMERIC or DBDECIMAL value in **value* and specify its precision and scale.
- If the procedure being executed resides on a server other than the one to which the application is directly connected, commands executed within the procedure cannot be rolled back.
- For an example of a remote procedure call, see Example 8 in the online sample programs.

See also

dbnextrow, dbresults, dbretdata, dbretstatus, dbrpcinit, dbrpcsend, dbsqlok

dbrpcsend

Description

Signal the end of a remote procedure call.

Syntax

```
RETCODE dbrpcsend(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

Return value

SUCCEED or FAIL.

Usage

- An application can call a stored procedure in two ways: by executing a command buffer containing a Transact-SQL **execute** statement or by making a remote procedure call (RPC). See the reference page for dbrpcinit for a discussion of the differences between these techniques.
- To make a remote procedure call, first call **dbrpcinit** to specify the stored procedure that is to be invoked. Then call **dbrpcparam** once for each of the stored procedure's parameters. Finally, call **dbrpcsend** to signify the end of the parameter list. This causes the server to begin executing the specified procedure. You can then call **dbsqlok**, **dbresults**, and **dbnextrow** to process the stored procedure's results. (Note that you will need to call **dbresults** multiple times if the stored procedure contains more than one **select** statement.) After all of the stored procedure's results have been processed you can call the routines that process return parameters and status numbers, such as **dbretdata** and **dbretstatus**.
- If the procedure being executed resides on a server other than the one to which the application is directly connected, commands executed within the procedure cannot be rolled back.
- For an example of a remote procedure call, see Example 8 in the online sample programs.

See also

dbnextrow, dbresults, dbretdata, dbretstatus, dbrpcinit, dbrpcparam, dbsqlok

dbrpwclr

Description

Clear all remote passwords from the LOGINREC structure.

Syntax

```
void dbrpwclr(loginrec)

 LOGINREC     *loginrec;
```

Parameters

loginrec

A pointer to a LOGINREC structure. This pointer will serve as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

Return value

None.

Usage

- A Transact-SQL command batch or stored procedure running on one server may call a stored procedure located on another server. To accomplish this server-to-server communication, the first server, connected to the application through **dbopen**, actually logs into the second, remote server.

  dbrpwset allows the application to specify the password to be used when the first server attempts to call the stored procedure on the remote server. Multiple passwords may be specified, one for each server that the first server might need to log in to.

- A single LOGINREC can be used repeatedly, in successive **dbopen** calls to different servers. **dbrpwclr** allows the application to remove any remote password information currently in the LOGINREC, so that successive calls to **dbopen** can contain different remote password information (specified with **dbrpwset**).

See also

dblogin, dbopen, dbrpwset, DBSETLAPP, DBSETLHOST, DBSETLPWD, DBSETLUSER

dbrpwset

Description

Add a remote password to the LOGINREC structure.

Syntax

RETCODE dbrpwset(loginrec, srvname, password, pwlen)

```
LOGINREC    *loginrec;
char         *srvname;
char         *password;
int          pwlen;
```

Parameters

loginrec

A pointer to a LOGINREC structure. This pointer will serve as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

srvname

The name of a server. A server's name is stored in the *srvname* column of its *sysservers* system table. When the first server calls a stored procedure located on the server designated by *srvname*, it will use the specified password to log in. If *srvname* is NULL, the specified password will be considered a "universal" password, to be used with any server that does not have a password explicitly specified for it.

password
The password that the first server will use to log in to the specified server.

pwlen
The length of the password in bytes.

Return value

SUCCEED or FAIL.

This routine may fail if the addition of the specified password would overflow the LOGINREC's remote password buffer. (The remote password buffer is 255 bytes long. Each password's entry in the buffer consists of the password itself, the associated server name, and 2 extra bytes.)

Usage

- A Transact-SQL command batch or stored procedure running on one server may call a stored procedure located on another server. To accomplish this server-to-server communication, the first server, connected to the application through **dbopen**, actually logs into the second, remote server and performs a remote procedure call.

  dbrpwset allows the application to specify the password to be used when the first server attempts to call the stored procedure on the remote server. Multiple passwords may be specified, one for each server that the first server might need to log in to.

- If the application has not specified a remote password for a particular server the password defaults to the one set with **DBSETLPWD** (or a null value, if **DBSETLPWD** has not been called). This behavior may be fine if the application's user has the same password on multiple servers.

- **dbrpwclr** clears all remote passwords from the LOGINREC.

See also

dblogin, dbopen, dbrpwclr, DBSETLAPP, DBSETLHOST, DBSETLPWD, DBSETLUSER

dbsafestr

Description

Double the quotes in a character string.

Syntax

```
RETCODE dbsafestr(dbproc, src, srclen, dest, destlen,
        quotetype)

DBPROCESS    *dbproc;
char         *src;
DBINT        srclen;
char         *dest;
```

```
DBINT        destlen;
int          quotetype;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

src
A pointer to the original string.

srclen
The length of *src*, in bytes. If *srclen* is -1, *src* is assumed to be null-terminated.

dest
A pointer to a programmer-supplied buffer to contain the resulting string. *dest* must be large enough for the resulting string plus a null terminator.

destlen
The length of the programmer-supplied buffer to contain the resulting string. If *destlen* is -1, *dest* is assumed to be large enough to hold the resulting string.

quotetype
The type of quotes to double. The following table lists the possible values for *quotetype*:

Table 2-24: Values for quotetype

Value of quotetype	dbsafestr
DBSINGLE	Doubles all single quotes (’) in *src*
DBDOUBLE	Doubles all double quotes (") in *src*
DBBOTH	Doubles all single and double quotes in *src*

Return value

SUCCEED or FAIL.

dbsafestr fails if the resulting string is too large for *dest*, or if an invalid *quotetype* is specified.

Usage

- **dbsafestr** doubles the single and/or double quotes found in a character string. This is useful when specifying literal quotes within a character string.

See also

dbcmd, dbfcmd

dbsechandle

Description

Install user functions to handle secure logins.

Syntax

```
RETCODE *dbsechandle(type, handler)

DBINT       type;
INTFUNCPTR  (*handler)();
```

Parameters

type

An integer variable with one of these symbolic values:

Table 2-25: Values for type (dbsechandle)

Value of type	dbsechandle
DBENCRYPT	Installs a function to handle password encryption
DBLABELS	Installs a function to handle login security labels

handler

A pointer to the user function that DB-Library will call whenever the corresponding type of secure login needs to be handled.

If *handler* is NULL and *type* is DBENCRYPT, DB-Library will use its default encryption handler.

If *handler* is NULL and *type* is DBLABELS, **dbsechandle** uninstalls any current label handler.

Return value

SUCCEED or FAIL.

Usage

- **dbsechandle** installs user functions to handle secure logins.
- An application can use **dbsechandle** to install functions to handle two types of secure logins:
  - Encrypted password secure logins

    In this type of secure login, the server provides the client with a key. The client uses the key to encrypt a password, which it then returns to the server.
  - Security label secure logins

    In this type of secure login, the server asks the client for identifying security labels, which the client then provides.

Encrypted password secure logins

- If *type* is DBENCRYPT, **dbsechandle** installs the function that DB-Library will call when encrypting user passwords.

- DB-Library will perform password encryption only if **DBSETLENCRYPT** has been called prior to calling **dbopen**.
- DB-Library will call its default encryption handler if a user function has not been installed
- Typically, a user function does not need to be installed for password encryption. This is because DB-Library's default encryption handler allows an application to perform password encryption when connecting to an Adaptive Server.
- A user-defined encryption handler should be installed by applications that are gateways. The encryption handler will be responsible for taking the encryption key returned by the remote server, passing it back to the client, reading the encrypted password from the client, and returning the encrypted password to DB-Library so that DB-Library can pass it on to the remote server.
- An encryption handler should be declared as shown in the example below. Encryption handlers on the Windows platform must be declared with CS_PUBLIC. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well. Here is a sample declaration:

```
RETCODE CS_PUBLIC encryption_handler(dbproc, pwd,
    pwdlen, enc_key, keylen, outbuf, buflen, outlen)
 DBPROCESS   *dbproc;
 BYTE        *pwd;
 DBINT       pwdlen;
 BYTE        *enc_key;
 DBINT       keylen;
 BYTE        *outbuf;
 DBINT       buflen;
 DBINT       *outlen;
```

  where:

  - *dbproc* is the DBPROCESS.
  - *pwd* is the user password to be encrypted.
  - *pwdlen* is the length of the user's password.
  - *enc_key* is the key to be used during encryption.
  - *keylen* is the length of the encryption key.
  - *outbuf* is a buffer in which the callback can place the encrypted password. This buffer will be allocated and freed by DB-Library.
  - *buflen* is the length of the output buffer.

- *outlen* is a pointer to a DBINT. The encryption handler should set **outlen* to the length of the encrypted password.

- An encryption handler should return SUCCEED to indicate that the password was encrypted successfully. If the encryption handler returns a value other than SUCCEED, DB-Library will abort the connection attempt.

Security label secure logins

- If type is DBLABELS, **dbsechandle** installs a function that DB-Library will call to get login security labels.
- DB-Library will send login security labels only if **DBSETLABELLED** has been called prior to calling **dbopen**.
- There are two ways for an application to define security labels:
  - The application can call **dbsetsecurity** one time for each label it wants to define. Most applications will use this method.
  - The application can call **dbsechandle** to install a user-supplied function to generate security labels. Typically, only gateway applications will use this method.

  If an application uses both methods, the labels defined via **dbsetsecurity** and the labels generated by the user-supplied function are sent to the server at the same time.
- DB-Library calls an application's label handler during the connection process, in response to a server request for login security labels. Each time it is called, the label handler returns a single label. DB-Library sends these labels, together with any labels previously defined via **dbsetsecurity**, to the server.
- DB-Library does not have a default label handler.
- A user-defined label handler should be installed by applications that are gateways. The label handler will be responsible for reading the client's login security labels and passing them on to DB-Library so that DB-Library can pass them on to the remote server.
- A label handler should be declared as shown in the example below. Label handlers on the Windows platform must be declared with CS_PUBLIC. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well. Here is a sample declaration:

```
RETCODE CS_PUBLIC label_handler(dbproc, namebuf,
    nbuflen, valuebuf, vbuflen, namelen, valuelen)
```

```
DBPROCESS   *dbproc;
DBCHAR      *namebuf;
DBINT       nbuflen;
DBCHAR      *valuebuf;
DBINT       vbuflen;
DBINT       *namelen;
DBINT       *valuelen;
```

where:

- *dbproc* is the DBPROCESS.
- *namebuf* is a buffer in which the handler can place the name of the login security label. This buffer is allocated and freed by DB-Library.
- *nbuflen* is the length of the *namebuf* buffer.
- *valuebuf* is a buffer in which the handler can place the value of the login security label. This buffer is allocated and freed by DB-Library.
- *vbuflen* is the length of the *valuebuf* buffer.
- *namelen* is a pointer to a DBINT. The label handler should set **namelen* to the length of the label name placed in *namebuf*.
- *valuelen* is a pointer to a DBINT. The label handler should set **valuelen* to the length of the label value placed in *valuebuf*.

- The following table lists the return values that are legal for a security label handler. A security label handler must return one of these values:

Table 2-26: Return values for security label handlers

Label handler return value	Indicates
DBMORELABEL	The label handler has set the name and value of a login security label. DB-Library should call the label handler again to get an additional label.
DBENDLABEL	The label handler has set the name and value of a login security label. DB-Library should not call the label handler again.
DBERRLABEL	A label handler error has occurred. DB-Library should abort the connection attempt.

See also

DBSETLENCRYPT, dbopen.

dbsendpassthru

Description

Send a TDS packet to a server.

Syntax

```
RETCODE dbsendpassthru(dbproc, send_bufp)

DBPROCESS    *dbproc;
DBVOIDPTR    send_bufp;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

send_bufp

A pointer to a buffer containing the TDS packet to be sent to the server. A packet has a default size of 512 bytes. This size may be changed via **DBSETLPACKET**.

Return value

DB_PASSTHRU_MORE, DB_PASSTHRU_EOM, or FAIL.

Usage

- **dbsendpassthru** sends a TDS (Tabular Data Stream) packet to a server.
- TDS is an application protocol used for the transfer of requests and request results between clients and servers. Under ordinary circumstances, a DB-Library/C application does not have to deal directly with TDS, because DB-Library/C manages the data stream.
- **dbrecvpassthru** and **dbsendpassthru** are useful in gateway applications. When an application serves as the intermediary between two servers, it can use these routines to pass the TDS stream from one server to the other, eliminating the process of interpreting the information and re-encoding it.
- **dbsendpassthru** sends a packet of bytes from the buffer to which *send_bufp* points. Most commonly, *send_bufp* will be **recv_bufp* as returned by **dbrecvpassthru**. *send_bufp* may also be the address of a user-allocated buffer containing the packet to be sent.
- A packet has a default size of 512 bytes. An application can change its packet size via **DBSETLPACKET**. For more information on packet sizes, see the dbgetpacket and DBSETLPACKET reference pages.
- **dbsendpassthru** returns DB_PASSTHRU_EOM if the TDS packet in the buffer is marked as EOM (End Of Message). If the TDS packet is not the last in the stream, **dbsendpassthru** returns DB_PASSTHRU_MORE.

- A DBPROCESS connection that is used for a **dbsendpassthru** operation cannot be used for any other DB-Library/C function until DB_PASSTHRU_EOM is received.
- This is a code fragment using **dbsendpassthru**:

```
/*
 ** The following code fragment illustrates the
 ** use of dbsendpassthru() in an Open Server
 ** gateway application. It will continually get
 ** packets from a client, and pass them through
 ** to the remote server.
 **
 ** The routine srv_recvpassthru() is the Open
 ** Server counterpart required to complete this
 ** passthru operation.
 */

DBPROCESS    *dbproc;
 SRV_PROC     *srvproc;
 int          ret;
 BYTE         *packet;
while(1)
 {
      ret = srv_recvpassthru(srvproc, &packet,
           (int *)NULL);

     if( ret == SRV_S_PASSTHRU_FAIL )
      {
           fprintf(stderr, "ERROR - \
                srv_recvpassthru failed in \
                lang_execute.\n");
           exit();
      }
     /*
      ** Now send the packet to the remote server
      */
      if( dbsendpassthru(dbproc, packet) == FAIL )
      {
          fprintf(stderr, "ERROR - dbsendpassthru\
                failed in lang_execute.\n");
           exit();
      }
     /*
      ** We've sent the packet, so let's see if
      ** there's any more.
      */
      if( ret == SRV_S_PASSTHRU_MORE )
```

```
                    continue;
               else
                    break;
          }
```

See also dbrecvpassthru

dbservcharset

Description Get the name of the server character set.

Syntax
```
char *dbservcharset(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library/C uses to manage communications and data between the front end and the server.

Return value A pointer to the null-terminated name of the server's character set, or NULL in case of error.

Usage

- **dbservcharset** returns the name of the server's character set.
- DB-Library/C clients can use a different character set than the server or servers to which they are connected. If a client and server are using different character sets, and the server supports character translation for the client's character set, it will perform all conversions to and from its own character set when communicating with the client.
- An application can inform the server what character set it is using via **DBSETLCHARSET**.
- To determine if the server is performing character set translations, an application can call **dbcharsetconv**.
- To get the name of the client character set, an application can call **dbgetcharset**.

See also dbcharsetconv, dbgetcharset, DBSETLCHARSET

dbsetavail

Description
Mark a DBPROCESS as being available for general use.

Syntax
void dbsetavail(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value
None.

Usage
- This routine marks the DBPROCESS as being available for general use. Any subsequent calls to **DBISAVAIL** will return TRUE, until some use is made of the DBPROCESS. Many DB-Library routines automatically set the DBPROCESS to "not available." This is useful when many different parts of a program are attempting to share a single DBPROCESS.

See also
DBISAVAIL

dbsetbusy

Description
Call a user-supplied function when DB-Library is reading from the server.

Syntax
void dbsetbusy(dbproc, busyfunc)

```
DBPROCESS    *dbproc;
int          (*(*busyfunc)())();
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

busyfunc
The user-supplied function that DB-Library will call whenever it accesses the server. DB-Library calls *busyfunc*() with a single parameter—a pointer to the DBPROCESS from the **dbsetbusy** call.

busyfunc() returns a pointer to a function that returns an integer.

Return value

None.

Usage

- This routine associates a user-supplied function with the specified *dbproc*. The user-supplied function will be automatically called whenever DB-Library is reading or waiting to read output from the server. For example, an application may want to print a message whenever the server is accessed. **dbsetbusy** will cause the user-supplied function *busyfunc*() to be called in this case.
- Similarly, **dbsetidle** may also be used to associate a user-supplied function, *idlefunc*(), with a *dbproc*. *idlefunc*() will be automatically called whenever DB-Library has finished reading output from the server.
- The server sends result data to the application in packets of 512 bytes. (The final packet in a set of results may be less than 512 bytes.) DB-Library calls *busyfunc*() at the beginning of each packet and *idlefunc*() at the end of each packet. If the output from the server spans multiple packets, *busyfunc*() and *idlefunc*() will be called multiple times.
- Here is an example of defining and installing *busyfunc*() and *idlefunc*():

Note The application functions *busyfunc*() and *idlefunc*() are callback event handlers and must be declared as CS_PUBLIC for the Windows platform. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well.

```
/*
 ** busyfunc returns a pointer to a function that
 ** returns an integer.
 */
 int      (*busyfunc())();
 void     idlefunc();

int      counterfunc();
 ...

main()
 {
     DBPROCESS      *dbproc;
     ...

   dbproc = dbopen(login, NULL);

    /*
     ** Now that we have a DBPROCESS, install the
     ** busy-function and the idle-function.
     */
     dbsetbusy(dbproc, busyfunc);
```

```
      dbsetidle(dbproc, idlefunc);

    dbcmd(dbproc, "select * from sysdatabases");
      dbcmd(dbproc, " select * from sysobjects");
      dbsqlexec(dbproc);

    /*
      ** DB-Library calls busyfunc() for the first time
     ** during dbsqlexec(). Depending on the size of the
      ** results, it may call busyfunc() again during
      ** processing of the results.
      */

    while (dbresults(dbproc) != NO_MORE_RESULTS)
        dbprrow(dbproc);

    /*
      ** DB-Library calls idlefunc() each time a packet
      ** of results has been received. Depending on the
      ** size of the results, it may call idlefunc()
    ** multiple times during processing of the results.
      */
      ...
}

int CS_PUBLIC (*busyfunc(dbproc))()
  DBPROCESS  dbproc;
  {
       printf("Waiting for data...\n");

     return(counterfunc);
  }

void CS_PUBLIC idlefunc(procptr, dbproc)

/*
  ** idlefunc's first parameter is a pointer to a
  ** routine that returns an integer. This is the same
  ** pointer that busyfunc returns.
  */
  int           (*procptr)();

DBPROCESS      *dbproc;
  {
      int       count;

    printf("Data is ready.\n");
      count = (*procptr)();

    printf ("Counterfunc has been called %d %s.\n",
      count, (count == 1 ? "time" : "times"));
```

```
}
 int counterfunc()
  {
       static     int     counter = 0;
     return(++counter);
  }
```

See also dbsetidle

dbsetdefcharset

Description Set the default character set for an application.

Syntax RETCODE dbsetdefcharset(charset)

```
char       *charset;
```

Parameters charset

The name of the character set to use. *charset* must be a null-terminated character string.

Return value SUCCEED or FAIL.

Usage

- **dbsetdefcharset** sets an application's default character set.
- DB-Library uses a default character set when no DBPROCESS structure is available or when localization information for a DBPROCESS structure's character set cannot be found.
- If an application does not call **dbsetdefcharset**, its default character set is the character set of the first DBPROCESS connection opened, or iso_1 if no DBPROCESS is open.
- If an application plans to call both **dbsetdefcharset** and **dbsetdeflang**, it must call **dbsetdefcharset** first.

See also dbsetdeflang, dbsetdefcharset, dblogin, dbopen

dbsetdeflang

Description Set the default language name for an application.

Syntax

```
RETCODE dbsetdeflang(language)

 char       *language;
```

Parameters

language

The name of the national language to use. *language* must be a null-terminated character string.

Return value

SUCCEED or FAIL.

Usage

- **dbsetdeflang** sets an application's default national language.
- DB-Library uses a default language when no DBPROCESS structure is available or when localization information for a DBPROCESS structure's language cannot be found.
- If an application does not call **dbsetdeflang**, its default language is the language of the first DBPROCESS connection opened, or us_english if no DBPROCESS is open.

See also

DBSETLNATLANG

dbsetidle

Description

Call a user-supplied function when DB-Library is finished reading from the server.

Syntax

```
void dbsetidle(dbproc, idlefunc)

 DBPROCESS     *dbproc;
 void          (*idlefunc)();
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

idlefunc

The user-supplied function that will be called by DB-Library whenever the server has finished sending data to the host. DB-Library calls *idlefunc*() with two parameters—the return value from *busyfunc*() (a pointer to a function that returns an integer) and a pointer to the DBPROCESS from the **dbsetidle** call.

idlefunc() returns void.

Return value

None.

Usage

- This routine associates a user-supplied function with the specified *dbproc*. The user-supplied function will be automatically called when DB-Library is finished reading or waiting to read a packet of output from the server. For example, an application may want to print a message whenever the server has finished sending data to the host. **dbsetidle** will cause the user-supplied function *idlefunc*() to be called in this case.
- Similarly, **dbsetbusy** may also be used to associate a user-supplied function, *busyfunc*(), with a *dbproc*. *busyfunc*() will be automatically called whenever DB-Library is reading or waiting to read a packet of output from the server.
- The server sends result data to the application in packets of 512 bytes. (The final packet in a set of results may be less than 512 bytes.) DB-Library calls *busyfunc*() at the beginning of each packet and *idlefunc*() at the end of each packet. If the output from the server spans multiple packets, *busyfunc*() and *idlefunc*() will be called multiple times.
- See the dbsetbusy reference page for an example of defining and installing *busyfunc*() and *idlefunc*().

See also

dbsetbusy

dbsetifile

Description

Specify the name and location of the Sybase interfaces file.

Syntax

```
void dbsetifile(filename)

 char       *filename;
```

Parameters

filename

The name of the interfaces file that gets searched during every subsequent call to **dbopen**. If this parameter is NULL, DB-Library will revert to the default file name.

Return value

None.

Usage

- This routine lets the application specify the name and location of the interfaces file that will be searched during every subsequent call to **dbopen**. The interfaces file contains the name and network address of every server available on the network.

- If **dbsetifile** has not been called, a call to **dbopen** will initiate the following default behavior: DB-Library will first attempt to use a file named *interfaces* in the directory named by the SYBASE environment variable or logical name. If SYBASE has not been set, DB-Library will attempt to use a file called *interfaces* in the home directory of the user named "sybase".
- For more information on the interfaces file, see the *Open Client/Server Configuration Guide*.

Note On non-UNIX platforms, client applications may use a method to find server address information that is different than the UNIX interfaces file. Consult your *Open Client/Server Configuration Guide* for detailed information on how clients connect to servers.

See also

dbopen

dbsetinterrupt

Description

Call user-supplied functions to handle interrupts while waiting on a read from the server.

Syntax

void dbsetinterrupt(dbproc, chkintr, hndlintr)

```
DBPROCESS    *dbproc;
int           (*chkintr)();
int          (*hndlintr)();
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

chkintr

A pointer to the user function that DB-Library will call to check whether an interrupt is pending. DB-Library calls it periodically while waiting on a read from the server. DB-Library calls *chkintr*() with a single parameter—a pointer to the DBPROCESS from the **dbsetinterrupt** call.

chkintr() must return TRUE or FALSE.

hndlintr

A pointer to the user function that DB-Library will call if an interrupt is returned. DB-Library calls *hndlintr*() with a single parameter—a pointer to the DBPROCESS from the **dbsetinterrupt** call.

The following table lists *hndlintr*()'s legal return values:

Table 2-27: Return values for the application's hndlintr() function

Return value	To indicate
INT_EXIT	Abort the program. (Note to UNIX programmers: DB-Library will not leave a core file.
INT_CANCEL	Abort the current command batch. Results are not flushed from the DBPROCESS connection.
INT_CONTINUE	Continue to wait for the server response.

Return value

None.

Usage

- DB-Library does non-blocking reads from the server. While waiting for a read from the server, it will call the *chkintr*() function to see if an interrupt is pending. If *chkintr*() returns TRUE and a handler has been installed as the *hndlintr*() for **dbsetinterrupt**, *hndlintr*() will be called. **dbsetinterrupt** is provided so that the programmer can substitute alternative interrupt handling for the time that the host program is waiting on reads from the server.
- Depending on the return value from *hndlintr*(), DB-Library performs one the following actions:
  - Sends an attention to the server, causing the server to discontinue processing (INT_CANCEL). For details, see "Canceling from the interrupt handler" on page 314.
  - Continues reading from the server (INT_CONTINUE).
  - Exits the program (INT_EXIT).

Canceling from the interrupt handler

- If *hndlintr*() returns INT_CANCEL, DB-Library sends an attention token to the server. This causes the server to discontinue command processing. The server may send additional results that have already been computed. When control returns to the mainline code, the mainline code should do one of the following:
  - Flush the results using **dbcancel**
  - Process the results normally

- You cannot call **dbcancel** in your interrupt handler, because this will cause output from the server to DB-Library to become out of sync. The steps below describe a correct method to cancel from the interrupt handler.
  - Associate an *int_canceled* flag with the DBPROCESS structure. Use **dbsetuserdata** to install a pointer to the flag in the DBPROCESS, and **dbgetuserdata** to get the address of the flag.
  - Code *hndlintr*() to set the *int_canceled* flag to indicate whether or not it is returning INT_CANCEL.
  - In the mainline code, check the flag before each call to **dbresults** or **dbnextrow**. When the *int_canceled* flag indicates that *hndlintr*() has aborted the server command, the mainline code should call **dbcancel** and clear the flag.

Example

- Here are example *chkintr*() and *hndlintr*() routines:

Note The applications *chkintr*() and *hndlintr*() routines are callback functions and must be declared as CS_PUBLIC for the Windows platform. For portability, callback handlers on other platforms should be declared CS_PUBLIC as well.

```
int CS_PUBLIC  chkintr(dbproc)
 DBPROCESS      *dbproc;
 {
   /*
   ** This routine assumes that the application
   ** sets the global variable
   ** "OS_interrupt_happened" upon catching
   ** an interrupt via some operating system
   ** facility.
   */
   if (OS_interrupt_happened)
   {
     /*
     ** Clear the interrupt flag, for
     ** future use.
     */
     OS_interrupt_happened = FALSE;
     return(TRUE);
   }
   else
```

```
    return(FALSE);
 }


int CS_PUBLIC  hndlintr(dbproc)
 DBPROCESS      *dbproc;
 {
   char    response[10];
   DBBOOL *int_canceled;
   /*
   ** We assume that a DBBOOL flag has been
   ** attached to dbproc with dbsetuserdata.
   */
   int_canceled = (DBBOOL *) dbgetuserdata(dbproc);
   if (int_canceled == (DBBOOL *)NULL)
   {
     printf("Fatal Error: no int_cancel flag \
           in the DBPROCESS\n");
     return(INT_EXIT);
   }
   *int_canceled = FALSE;
   printf("\nAn interrupt has occurred. Do you \
         want to:\n\n");
   printf("\t1) Abort the program\n");
   printf("\t2) Cancel the current query\n");
   printf("\t3) Continue processing the current\
         query's results\n\n");
   printf("Press 1, 2, or 3, followed by the \
         return key: ");
   gets(response);
   switch(response[0])
   {
     case '1':
       return(INT_EXIT);
       break;
     case '2':
       *int_canceled = TRUE;
       return(INT_CANCEL);
       break;
     case '3':
       return(INT_CONTINUE);
       break;
     default:
       printf("Response not understood. \
             Aborting program.\n");
       return(INT_EXIT);
```

```
            break;
        }
    }
```

See also dbcancel, dbgetuserdata, dbsetuserdata, dbsetbusy, dbsetidle

DBSETLAPP

Description Set the application name in the LOGINREC structure.

Syntax RETCODE DBSETLAPP(loginrec, application)

```
LOGINREC    *loginrec;
char        *application;
```

Parameters

loginrec

A pointer to a LOGINREC structure, which will be passed as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

application

The application name that will be sent to the server. It must be a null-terminated character string. The maximum length of the string, not including the null terminator, is 30 characters.

Return value SUCCEED or FAIL.

Usage

- This macro sets the application field in the LOGINREC structure. For it to have any effect, it must be called before **dbopen**.
- It is not necessary to call this routine. By default, the application name will be a null value.
- The server uses the application name in its *sysprocesses* table to help identify your process. If you set the application name, you will see it if you query the *sysprocesses* table in the *master* database.

See also dblogin, dbopen, DBSETLHOST, DBSETLPWD, DBSETLUSER

DBSETLCHARSET

Description Set the character set in the LOGINREC structure.

Syntax

RETCODE DBSETLCHARSET(loginrec, char_set)

```
LOGINREC    *loginrec;
DBCHAR      *char_set;
```

Parameters

loginrec

A pointer to a LOGINREC structure to be passed as an argument to **dbopen**. LOGINREC structures are obtained via **dblogin**.

char_set

The name of the character set the client will use. *char_set* must be a null-terminated string. Default values for *char_set* include "iso_1" for ISO-8859-1 (most platforms), "cp850" for Code Page 850 (IBM RS/6000), and "roman8" for the Roman8 character set (HP platforms).

To indicate that no character set conversion is desired, pass *char_set* as NULL.

Return value

SUCCEED or FAIL.

Usage

- **DBSETLCHARSET** sets the client character set in a LOGINREC structure.
- DB-Library/C clients may use a different character set than the server or servers to which they are connected. **DBSETLCHARSET** is used to inform the server what character set a client is using.
- Because the LOGINREC is passed as a parameter in the **dbopen** call that establishes the client's connection with a server, **DBSETLCHARSET** must be called before **dbopen** to have any effect.
- The server will perform all conversions to and from its own character set when communicating with a client using a different character set.
- If no conversion is desired, call **DBSETLCHARSET** with *char_set* as NULL.

See also

dbgetcharset, dblogin, dbopen

DBSETLENCRYPT

Description

Specify whether or not network password encryption is to be used when logging into SQL Server version 10.0 or later.

Syntax

RETCODE DBSETLENCRYPT(loginrec, enable)

```
LOGINREC    *loginrec;
DBBOOL      enable;
```

Parameters

loginrec

A pointer to a LOGINREC structure, which will be passed as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

enable

A boolean value (TRUE or FALSE) specifying whether or not the server should request an encrypted password at login time.

Return value

SUCCEED or FAIL.

Usage

- **DBSETLENCRYPT** specifies whether or not network password encryption is to be used when logging into SQL Server version 10.0 or later. If an application does not call **DBSETLENCRYPT**, password encryption is not used.
- Network password encryption provides a protected mechanism for authenticating a user's identity.
- If an application specifies that network password encryption is to be used, then when the application attempts to open a connection:
  - No password is sent with the initial connection request. At this time, the client indicates to the server that encryption is desired.
  - The server replies to the connection request with an encryption key.
  - DB-Library uses the key to encrypt the user's password and remote passwords, if any, and sends the encrypted passwords back to the server.
  - The server uses the key to decrypt the encrypted passwords and either accepts or rejects the login attempt.
- If password encryption is not specified, then when an application attempts to open a connection:
  - A password is included with the connection request.
  - The server either accepts or rejects the login attempt.

See also

dbsechandle

DBSETLHOST

Description

Set the host name in the LOGINREC structure.

Syntax

RETCODE DBSETLHOST(loginrec, hostname)

```
LOGINREC    *loginrec;
char        *hostname;
```

Parameters

loginrec

A pointer to a LOGINREC structure, which will be passed as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

hostname

The host name that will be sent to the server. It must be a null-terminated character string. The maximum length of the string, not including the null terminator, is 30 characters.

Return value

SUCCEED or FAIL.

Usage

- This macro sets the host name in the LOGINREC structure. For it to have any effect, it must be called before **dbopen**.
- The host name will show up in the *sysprocesses* table in the *master* database.
- It is not necessary to call this routine. If it is not called, DB-Library will set the default value for the host name. This default value will generally be a version of the host machine's name provided by the operating system.

See also

dblogin, dbopen, DBSETLAPP, DBSETLPWD, DBSETLUSER

DBSETLNATLANG

Description

Set the national language name in the LOGINREC structure.

Syntax

RETCODE DBSETLNATLANG(loginrec, language)

```
LOGINREC    *loginrec;
char        *language;
```

Parameters

loginrec

A pointer to a LOGINREC structure to be passed as an argument to **dbopen**. LOGINREC structures are obtained via **dblogin**.

language

The name of the national language to use. *language* must be a null-terminated character string.

Return value

SUCCEED or FAIL.

Usage

- This macro sets the user language in the LOGINREC structure. If you wish to set a particular user language, call **DBSETLNATLANG** before **dbopen**.
- Call **DBSETLNATLANG** only if you do not wish to use the server's default national language.

See also

dblogin, dbopen, dbsetdeflang

dbsetloginfo

Description

Transfer TDS login information from a DBLOGINFO structure to a LOGINREC structure.

Syntax

```
RETCODE dbsetloginfo(loginrec, loginfo)

LOGINREC      *login;
DBLOGINFO     *loginfo;
```

Parameters

login

A pointer to a LOGINREC structure. This pointer will be passed as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

loginfo

A pointer to a DBLOGINFO structure that contains login parameter information.

Return value

SUCCEED or FAIL.

Usage

- **dbsetloginfo** transfers TDS login information from a DBLOGINFO structure to a LOGINREC structure. After the information is transferred, **dbsetloginfo** frees the DBLOGINFO structure.
- An application needs to call **dbsetloginfo** only if (1) it is an Open Server gateway application and (2) it is using TDS passthrough.
- TDS (Tabular Data Stream) is an application protocol used for the transfer of requests and request results between clients and servers.
- When a client connects directly to a server, the two programs negotiate the TDS format they will use to send and receive data. When a gateway application uses TDS passthrough, the application forwards TDS packets between the client and a remote server without examining or processing them. For this reason, the remote server and the client must agree on a TDS format to use.

- **dbsetloginfo** is the second of four calls, two of them Server Library calls, that allow a client and remote server to negotiate a TDS format. The calls, which can only be made in a SRV_CONNECT event handler, are described here:
  - **srv_getloginfo** allocates a DBLOGINFO structure and fills it with TDS information from a client SRV_PROC.
  - **dbsetloginfo** transfers the TDS information retrieved by **srv_getloginfo** from the DBLOGINFO structure to a DB-Library/C LOGINREC structure, and then frees the DBLOGINFO structure. After the information is transferred, the application can use this LOGINREC structure in the **dbopen** call that establishes its connection with the remote server.
  - **dbgetloginfo** transfers the remote server's response to the client's TDS information from a DBPROCESS structure into a newly-allocated DBLOGINFO structure.
  - **srv_setloginfo** sends the remote server's response, retrieved by **dbgetloginfo**, to the client, and then frees the DBLOGINFO structure.
- This is an example of a SRV_CONNECT handler preparing a remote connection for TDS passthrough:

```
RETCODE      connect_handler(srvproc)
 SRVPROC      *srvproc;
 {
      SYBLOGINFO   *loginfo;
      LOGINREC     *loginrec;
      DBPROCESS    *dbproc;

    /*
      ** Get the TDS login information from the
      ** client SRV_PROC.
      */
      srv_getloginfo(srvproc, &loginfo);

    /* Get a LOGINREC structure */
      loginrec = dblogin();

    /*
      ** Initialize the LOGINREC with the logininfo
      ** from the SRV_PROC.
      */
      dbsetloginfo(loginrec, loginfo);

    /* Connect to the remote server */
```

```
        dbproc = dbopen(loginrec, REMOTE_SERVER_NAME)

    /*
     ** Get the TDS login response informationfrom
     ** the remote connection.
     */
     dbgetloginfo(dbproc, &loginfo);

    /*
     ** Return the login response information to
     ** the SRV_PROC.
     */
     srv_setloginfo(srvproc, loginfo);

    /* Accept the connection and return */
     srv_senddone(srvproc, 0, 0, 0);
     return(SRV_CONTINUE);
}
```

See also dbgetloginfo, dbrecvpassthru, dbsendpassthru

dbsetlogintime

Description Set the number of seconds that DB-Library waits for a server response to a request for a DBPROCESS connection.

Syntax RETCODE dbsetlogintime(seconds)

```
int     seconds;
```

Parameters seconds

The timeout value—the number of seconds that DB-Library waits for a login response before timing out. A timeout value of 0 represents an infinite timeout period.

Return value SUCCEED or FAIL.

Usage

- This routine sets the length of time in seconds that DB-Library will wait for a login response after calling **dbopen**. The default timeout value is 60 seconds.
- When a connection attempt is made between a client and a server, there are two ways in which the connection can fail (assuming that the system is correctly configured):
  - The machine that the server is supposed to be on is running correctly and the network is running correctly.

In this case, if there is no server listening on the specified port, the machine the server is supposed to be on will signal the client, via a network error, that the connection cannot be formed. Regardless of **dbsetlogintime**, the connection will fail.

- The machine that the server is on is down.

  In this case, the machine that the server is supposed to be on will not respond. Because "no response" is not considered to be an error, the network will not signal the client that an error has occurred. However, if **dbsetlogintime** has been called to set a timeout period, a timeout error will occur when the client fails to receive a response within the set period.

See also dberrhandle, dbsettime

DBSETLPACKET

Description Set the TDS packet size in an application's LOGINREC structure.

Syntax RETCODE DBSETLPACKET(login, packet_size)

```
LOGINREC    *login;
short       packet_size;
```

Parameters

login

A pointer to the LOGINREC structure to be passed as an argument to **dbopen** when logging in to the server. An application can obtain a LOGINREC structure via **dblogin**.

packet_size

The packet size being requested, in bytes. The server will set the actual packet size to a value less than or equal to this requested size.

Return value SUCCEED or FAIL.

Usage

- **DBSETLPACKET** sets the packet size field in an application's LOGINREC structure. When the application logs into the server, the server sets the TDS packet size for that DBPROCESS connection to be equal to or less than the value of this field. The packet size is set to a value less than the value of the packet size field if the server is experiencing space constraints. Otherwise, the packet size will be equal to the value of the field.

- If an application sends or receives large amounts of text or image data, a packet size larger than the default 512 bytes may improve efficiency, since it results in fewer network reads and writes.
- To determine the packet size that the server has set, an application can call **dbgetpacket**.
- TDS (Tabular Data Stream) is an application protocol used for the transfer of requests and request results between clients and servers.
- TDS data is sent in fixed-size chunks, called packets. TDS packets have a default size of 512 bytes. The only way an application can change the TDS packet size is via **DBSETLPACKET**. If **DBSETLPACKET** is not called, all DBPROCESS connections in an application will use the default size.
- Different DBPROCESS connections in an application may use different packet sizes. To set different packet sizes for DBPROCESS connections, an application can either:
  - Change the packet size in a single LOGINREC between the **dbopen** calls that create the DBPROCESS connections, or
  - Set different packet sizes in multiple LOGINREC structures, and use these different LOGINREC structures when creating the DBPROCESS connections.
- Because the actual packet size for a DBPROCESS connection is set when the DBPROCESS is created, calls to **DBSETLPACKET** will have no effect on the packet sizes of DBPROCESSes already allocated via **dbopen**.

See also dblogin, dbopen, dbgetpacket

DBSETLPWD

Description Set the user server password in the LOGINREC structure.

Syntax RETCODE DBSETLPWD(loginrec, password)

```
LOGINREC     *loginrec;
char         *password;
```

Parameters loginrec

A pointer to a LOGINREC structure, which will be passed as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

password
The password that will be sent to the server. It must be a null-terminated character string. The maximum length of the string, not including the null terminator, is 30 characters.

Return value
SUCCEED or FAIL.

Usage
- This macro sets the user server password in the LOGINREC structure. For it to have any effect, it must be called before **dbopen**.
- By default, the password field of the LOGINREC has a null value. Therefore, you do not need to call this routine if the password is a null value.
- DB-Library does not automatically blank out the password in *loginrec* after a call to **dbopen**. Therefore, if you want to minimize the risk of having a readable password in your DB-Library program, you should set *password* to something else after you call **dblogin**.

See also
dblogin, dbopen, DBSETLAPP, DBSETLHOST, DBSETLUSER

DBSETLUSER

Description
Set the user name in the LOGINREC structure.

Syntax
```
RETCODE DBSETLUSER(loginrec, username)

LOGINREC     *loginrec;
char         *username;
```

Parameters
loginrec
A pointer to a LOGINREC structure, which will be passed as an argument to **dbopen**. You can allocate a LOGINREC structure by calling **dblogin**.

username
The user name that will be sent to the server. It must be a null-terminated character string. The maximum length of the string, not including the null terminator, is 30 characters. The server will use *username* to determine who is attempting the connection. The server usernames are defined in the *syslogins* table in the master database.

Return value
SUCCEED or FAIL.

Usage
- This macro sets the user name in the LOGINREC structure. For it to have any effect, it must be called before **dbopen**.

- In most environments, this macro is optional. If it is not called, DB-Library will generally set the default value for the user name.

Note On *UNIX*: the user name defaults to the UNIX login name.

On *OS/2* and *NetWare*: The user name defaults to the SYBUSER environment variable, if any. If the application does not call **DBSETLUSER**, the SYBUSER environment variable must be set. Otherwise, **dbopen** will fail when it tries to log in to the server.

On *MPE/XL*: The user name defaults to the value of the system environment variable HPUSER.

See also dblogin, dbopen, DBSETLHOST, DBSETLPWD, DBSETLAPP

dbsetmaxprocs

Description Set the maximum number of simultaneously open DBPROCESS structures.

Syntax
```
RETCODE dbsetmaxprocs(maxprocs)

int     maxprocs;
```

Parameters

maxprocs
The new limit on simultaneously open DBPROCESS structures for this particular program.

Return value SUCCEED or FAIL.

Usage

- A DB-Library program has a maximum number of simultaneously open DBPROCESS structures. By default, this number is 25. The program may change this limit by calling **dbsetmaxprocs**.
- The program may find out what the current limit is by calling **dbgetmaxprocs**.

See also dbgetmaxprocs, dbopen

dbsetnull

Description Define substitution values to be used when binding null values.

Syntax

RETCODE dbsetnull(dbproc, bindtype, bindlen, bindval)

```
DBPROCESS    *dbproc;
int           bindtype;
int           bindlen;
BYTE         *bindval;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

bindtype

A symbolic value specifying the type of variable binding to which the substitute value will apply. (See the reference page for dbbind for more information about the allowable *bindtype* values.

bindlen

The length in bytes of the substitute value you are supplying. DB-Library ignores it in all cases except CHARBIND and BINARYBIND. All the other types are either fixed length or have a special terminator or embedded byte-count that provides the length of the data.

bindval

A generic BYTE pointer to the value you want to use as a null substitution value. **dbsetnull** makes a copy of the value, so you can free this pointer anytime after this call.

Return value

SUCCEED or FAIL.

dbsetnull returns FAIL if you give it an unknown *bindtype*. It will also fail if the specified DBPROCESS is dead.

Usage

- The **dbbind** and **dbaltbind** routines bind result column values to program variables. After the application calls them, calls to **dbnextrow** and **dbgetrow** automatically copy result values into the variables to which they are bound. If the server returns a null value for one of the result columns, DB-Library automatically places a substitute value into the result variable.
- Each DBPROCESS has a list of substitute values for each of the binding types. The following table lists the default substitution values:

Table 2-28: Default null substitution values

Binding type	Null substitution value
TINYBIND	0
SMALLBIND	0

Binding type	Null substitution value
INTBIND	0
CHARBIND	Empty string (padded with blanks)
STRINGBIND	Empty string (padded with blanks, null-terminated)
NTBSTRINGBIND	Empty string (null-terminated)
VARYCHARBIND	Empty string
BINARYBIND	Empty array (padded with zeros)
VARYBINBIND	Empty array
DATETIMEBIND	8 bytes of zeros
SMALLDATETIMEBIND	8 bytes of zeros
MONEYBIND	$0.00
SMALLMONEYBIND	$0.00
FLT8BIND	0.0
REALBIND	0.0
DECIMALBIND	0.0 (with default scale and precision)
NUMERICBIND	0.0 (with default scale and precision)
BOUNDARYBIND	Empty string (null-terminated)
SENSITIVITYBIND	Empty string (null-terminated)

- **dbsetnull** lets you provide your own null substitution values. When you call **dbsetnull** to change a particular null substitution value, the new value will remain in force for the specified DBPROCESS until you change it with another call to **dbsetnull**.
- The **dbconvert** routine also uses the current null substitution values when it needs to set a destination variable to null.
- The **dbnullbind** routine allows you to associate an indicator variable with a bound column. DB-Library will set the indicator value to indicate null data values or conversion errors.

See also dbaltbind, dbbind, dbconvert, dbnullbind, Types

dbsetopt

Description Set a server or DB-Library option.

Syntax **RETCODE dbsetopt(dbproc, option, char_param, int_param)**

```
DBPROCESS    *dbproc;
int           option;
char         *char_param;
int          int_param;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server. If *dbproc* is NULL, the option will be set for all active DBPROCESS structures.

option

The option that is to be turned on. See Options on page 402 for the list of options.

char_param

Certain options take parameters. For example, the DBOFFSET option takes as its parameter the construct for which offsets are to be returned:

```
dbsetopt(dbproc, DBOFFSET, "compute", -1)
```

The DBBUFFER option takes as its parameter the number of rows to be buffered:

```
dbsetopt(dbproc, DBBUFFER, "500", -1)
```

char_param must always be a character string enclosed in quotes, even in the case of a numeric value, as in the DBBUFFER example. If an invalid parameter is specified for one of the server options, this will be discovered the next time a command buffer is sent to the server. The **dbsqlexec** or **dbsqlsend** call will fail, and DB-Library will invoke the user-installed message handler. If an invalid parameter is specified for one of the DB-Library options (DBBUFFER or DBTEXTLIMIT), the **dbsetopt** call itself will fail.

If the option takes no parameters, *char_param* must be NULL.

int_param

Some options require an additional parameter, *int_param*, which is the length of the character string passed as *char_param*. Currently, only DBPRCOLSEP, DBPRLINESEP, and DBPRPAD require this parameter.

If *int_param* is not required, pass it as -1.

Return value

SUCCEED or FAIL.

dbsetopt will fail if *char_param* is invalid for one of the DB-Library options. However, an invalid *char_param* for a server option will not cause **dbsetopt** to fail, because such a parameter does not get validated until the command buffer is sent to the server.

Usage

- This routine sets server and DB-Library options. Although server options may be set and cleared directly through SQL, the application should instead use **dbsetopt** and **dbclropt** to set and clear options. This provides a uniform interface for setting both server and DB-Library options. It also allows the application to use the **dbisopt** function to check the status of an option.
- **dbsetopt** does not immediately set the option. The option is set the next time a command buffer is sent to the server (by invoking **dbsqlexec** or **dbsqlsend**).
- For a list of each option and its default status, see Options on page 402.

See also

dbclropt, dbisopt, Options

dbsetrow

Description

Set a buffered row to "current."

Syntax

STATUS dbsetrow(dbproc, row)

```
DBPROCESS    *dbproc;
DBINT         row;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

row

An integer representing the row number of the row to make current. Row number 1 is the first row returned from the server. This is not necessarily the first row in the row buffer.

Return value

MORE_ROWS, NO_MORE_ROWS, or FAIL.

dbsetrow returns:

- MORE_ROWS if it found *row* in the row buffer, or

- NO_MORE_ROWS if it did not find *row* in the row buffer or if row buffering is not enabled, or
- FAIL if the *dbproc* DBPROCESS is dead or not enabled.

Usage

- **dbsetrow** sets a buffered row to "current." After **dbsetrow** is called, the application's next call to **dbnextrow** will read this row.
- **dbgetrow**, another DB-Library/C routine, also sets a specific row in the row buffer to "current." However, unlike **dbsetrow**, **dbgetrow** reads the row. Any binding of row data to program variables (as specified with **dbbind** and **dbaltbind**) takes effect.
- **dbsetrow** has no effect unless the DB-Library/C option DBBUFFER is on.
- Row buffering provides a way to keep a specified number of server result rows in program memory. Without row buffering, the result row generated by each new **dbnextrow** call overwrites the contents of the previous result row. Row buffering is therefore useful for programs that need to look at result rows in a non-sequential manner. It does, however, carry a memory and performance penalty because each row in the buffer must be allocated and freed individually. Therefore, use it only if you need to. Specifically, the application should only turn the DBBUFFER option on if it calls **dbgetrow** or **dbsetrow**. Note that row buffering has nothing to do with network buffering and is a completely independent issue.
- When row buffering is *not* enabled, the application processes each row as it reads it from the server by calling **dbnextrow** repeatedly until it returns NO_MORE_ROWS. When row buffering *is* enabled, the application can use **dbsetrow** to jump to any row that has already been read from the server with **dbnextrow**. Subsequent calls to **dbnextrow** will cause the application to read successive rows in the buffer, starting with the row specified by the *row* parameter. When **dbnextrow** reaches the last row in the buffer, it reads rows from the server again, if there are any. Once the buffer is full, **dbnextrow** does not read any more rows from the server until some of the rows have been cleared from the buffer with **dbclrbuf**.
- The macro **DBFIRSTROW**, which returns the number of the first row in the row buffer, is useful in conjunction with **dbsetrow**. Thus, the call:

  ```
  dbsetrow(dbproc, DBFIRSTROW(dbproc))
  ```

  sets the current row so that the next call to **dbnextrow** will read the first row in the buffer.

See also

dbclrbuf, DBCURROW, DBFIRSTROW, dbgetrow, DBLASTROW, dbnextrow, Options

dbsettime

Description
Set the number of seconds that DB-Library will wait for a server response to a SQL command.

Syntax
RETCODE dbsettime(seconds)

```
int     seconds;
```

Parameters
seconds
The timeout value—the number of seconds that DB-Library waits for a server response before timing out. A timeout value of 0 represents an infinite timeout period.

Return value
SUCCEED or FAIL.

Usage
- This routine sets the length of time in seconds that DB-Library will wait for a server response during calls to **dbsqlexec**, **dbsqlok**, **dbresults**, and **dbnextrow**. The default timeout value is 0, which represents an infinite timeout period.
- **dbsettime** can be called at any time during the application—before or after a call to **dbopen**. It takes effect immediately upon being called.
- To set a timeout value for calls to **dbopen**, use **dbsetlogintime**.
- Note that, after sending a query to the server, **dbsqlexec** waits until a response is received or until the timeout period has elapsed. To minimize the time spent in DB-Library waiting for a response from the server, an application can instead call **dbsqlsend**, followed by **dbsqlok**.
- The program can call **DBGETTIME** to learn the current timeout value.
- A timeout generates the DB-Library error “SYBETIME.”

See also
dberrhandle, DBGETTIME, dbsetlogintime, dbsqlexec, dbsqlok, dbsqlsend

dbsetuserdata

Description
Use a DBPROCESS structure to save a pointer to user-allocated data.

Syntax
void dbsetuserdata(dbproc, ptr)

```
DBPROCESS    *dbproc;
BYTE         *ptr;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

ptr
A generic BYTE pointer to the user's private data space.

Return value

None.

Usage

- This routine saves, in a DBPROCESS structure, a pointer to user-allocated data. The application can access the data later with the **dbgetuserdata** routine.
- **dbsetuserdata** allows the application to associate user data with a particular DBPROCESS. This avoids the necessity of using global variables for this purpose. One use for this routine is to handle deadlock, as shown in the example below. This routine is particularly useful when the application has multiple DBPROCESS structures.
- The application must allocate the data that *ptr* points to. DB-Library never manipulates this data; it merely saves the pointer to it for later use by the application.
- Here is an example of using this routine to handle deadlock, a situation which occurs occasionally in high-volume applications. See the *System Administration Guide* for more information on deadlock. This program fragment sends updates to the server. It reruns the transaction when its message handler detects deadlock.

```
        ...
/*
** Deadlock detection:
**      In the DBPROCESS structure, we save a pointer to
**      a DBBOOL variable. The message handler sets the
**      variable when deadlock occurs. The result
**      processing logic checks the variable and resends
**      the transaction in case of deadlock.
*/

 /*
 ** Allocate the space for the DBBOOL variable
 ** and save it in the DBPROCESS structure.
 */
          dbsetuserdata(dbproc, malloc(sizeof(DBBOOL)));

          /* Initialize the variable to FALSE */
```

```
                *((DBBOOL *) dbgetuserdata(dbproc)) = FALSE;
                 ...
                /* Run queries and check for deadlock */
deadlock:
      /*
      ** Did we get here via deadlock?
      ** If so, the server has already aborted the
      ** transaction. We'll just start it again.  In a
      ** real application, the deadlock handling may need
      ** to be somewhat more sophisticated.  For
      ** instance, you may want to keep a counter and
      ** retry the transaction just a fixed number
      ** of times.
      */
      if (*((DBBOOL *) dbgetuserdata(dbproc)) == TRUE)
      {
           /* Reset the variable to FALSE */
           *((DBBOOL *) dbgetuserdata(dbproc)) = FALSE;
      }
      /* Start the transaction */
      dbcmd(dbproc, "begin transaction ");
      /* Run the first update command */
      dbcmd(dbproc, "update ......");
      dbsqlexec(dbproc);
      while (dbresults(dbproc) != NO_MORE_RESULTS)
      {
           /* application code */
      }
      /*  Did we deadlock? */
      if (*((DBBOOL *) dbgetuserdata(dbproc)) == TRUE)
           goto deadlock;
      /*  Run the second update command.  */
      dbcmd(dbproc, "update ......");
      dbsqlexec(dbproc);
      while (dbresults(dbproc) != NO_MORE_RESULTS)
      {
           /* application code */
      }
      /*  Did we deadlock? */
      if (*((DBBOOL *) dbgetuserdata(dbproc)) == TRUE)
           goto deadlock;
      /*  No deadlock -- Commit the transaction */
      dbcmd(dbproc, "commit transaction");
      dbsqlexec(dbproc);
      dbresults(dbproc);
 ...
```

```
/*
** SERVERMSGS
** This is the server message handler.  Assume that
** the dbmsghandle() routine installed it earlier in
** the program.
*/
servermsgs(dbproc, msgno, msgstate, severity, msgtext,
   srvname, procname, line)
DBPROCESS    *dbproc;
DBINT        msgno;
int          msgstate;
int          severity;
char         *msgtext;
char         *srvname;
char         *procname;
DBUSMALLINT  line;
{
     /*  Is this a deadlock message? */
     if (msgno == 1205)
     {
          /* Set the deadlock indicator */
          *((DBBOOL *) dbgetuserdata(dbproc)) = TRUE;
          return (0);
     }
     /*  Normal message handling code here */
}
```

See also dbgetuserdata

dbsetversion

Description Specify a DB-Library version level.

Syntax RETCODE dbsetversion(version)

DBINT version;

Parameters version

The version of DB-Library behavior that the application expects. The following table lists the symbolic values that are legal for *version*:

Table 2-29: Values for version (dbsetversion)

Value of version	To Indicate	Features Supported
DBVERSION_46	4.6 behavior	RPCs, registered procedures, remote procedure calls, text and image datatypes. This is the default version of DB-Library.
DBVERSION_100	10.0 behavior	numeric and decimal datatypes.

Return value

SUCCEED or FAIL.

Usage

- **dbversion** sets the version of DB-Library behavior that an application expects. DB-Library will provide the behavior requested, regardless of the actual version of DB-Library in use.
- An application is not required to call **dbsetversion**. However, if **dbsetversion** is not called, DB-Library provides version 4.6-level behavior.
- If an application calls **dbsetversion**, it must do so before calling any other DB-Library routine, with the exception of **dbinit**.
- It is an error to call **dbsetversion** more than once.

See also

dbinit

dbspid

Description

Get the server process id for the specified DBPROCESS.

Syntax

```
int dbspid(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

dbproc's server process id.

Usage

- **dbspid** yields the server process id of the specified DBPROCESS. The process id appears in the server's *sysprocesses* table.
- You can use the server process id to make queries against the *sysprocesses* table.

See also dbopen

dbspr1row

Description Place one row of server query results into a buffer.

Syntax RETCODE dbspr1row(dbproc, buffer, buf_len)

```
DBPROCESS     *dbproc;
char          *buffer;
DBINT         buf_len;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

buffer
A pointer to a character buffer to contain the **dbspr1row** results.

buf_len
The length of *buffer*, including its null terminator.

Return value SUCCEED or FAIL.

Note If an error occurs, the contents of **buffer* are undefined.

Usage

- **dbspr1row** fills a programmer-supplied buffer with a null-terminated character string containing one server query results row.
- **dbspr1row** is useful when displaying data for debugging and writing applications that scroll data displays.
- **dbspr1row** gives programmers greater control over data display than **dbprrow**. **dbprrow** always writes its output to the display device, while **dbspr1row** writes its output to a buffer, which the programmer may then display at whatever time or location is desired.

- To pad results data to its maximum converted length, specify a pad character via the DB-Library option DBPRPAD. The pad character will be appended to each column's data. The maximum converted column length is equal to the longest possible string that could be the column's displayable data, or the length of the column's name, whichever is greater. See Options on page 402 for more details on the DBPRPAD option.
- You can specify the column separator string via the DB-Library option DBPRCOLSEP. The column separator will be added to the end of each converted column's data except the last. The default separator is an ASCII 0x20 (space). See Options on page 402 for more details on the DBPRCOLSEP option.
- You can specify the maximum number of characters to be placed on one line via the DB-Library option DBPRLINELEN.
- You can specify the line separator string via the DB-Library option DBPRLINESEP. The default line separator is a newline (ASCII 0x0a or 0x0d, depending on the host system). See Options on page 402 for more details on the DBPRLINELEN and DBPRLINESEP options.
- The length of the buffer required by **dbspr1row** can be determined by calling **dbspr1rowlen**.
- The format of results rows returned by **dbspr1row** is determined by the SQL query. **dbspr1row** makes no attempt to format the data beyond converting it to printable characters, padding the columns as necessary, and adding the column and line separators.
- To make the best use of **dbspr1row**, application programs should call it once for every successful call to **dbnextrow**.
- The following code fragment illustrates the use of **dbspr1row**:

```
char     mybuffer[2000];

while (dbnextrow(dbproc) != NO_MORE_ROWS)
{
     dbspr1row(dbproc, mybuffer,sizeof(mybuffer));
     fprintf( stdout, "\n%s", mybuffer);
}
```

- The following code fragment shows the use of the DBPRPAD and DBPRCOLSEP options:

```
    char     mybuffer[2000];

     /*
     ** Specify the pad and column separator
```

```
** characters */

/* Pad = 0x2A */
dbsetopt(dbproc, DBPRPAD, "*", DBPADON);
/* Col. sep. = 0x2C20 */
dbsetopt(dbproc, DBPRCOLSEP, ", ", 2);

while (dbnextrow(dbproc) != NO_MORE_ROWS)
{
     dbspr1row(dbproc, mybuffer,
          sizeof(mybuffer) );
     fprintf( stdout, "\n%s", mybuffer);
}

/* Turn padding off */
dbsetopt(dbproc, DBPRPAD, SS, DBPADOFF );
/* Revert to default */
dbsetopt(dbproc, DBPRCOLSEP, RS, -1 );
```

See also dbclropt, dbisopt, dbprhead, dbprrow, dbspr1rowlen, dbsprhead, dbsprline, Options

dbspr1rowlen

Description Determine how large a buffer to allocate to hold the results returned by **dbsprhead, dbsprline,** and **dbspr1row**.

Syntax DBINT dbspr1rowlen(dbproc)

```
DBPROCESS    *dbproc;
```

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value The size of the buffer, in bytes, required by **dbsprhead, dbsprline**, and **dbspr1row** on success; a negative integer on error.

Usage

- **dbspr1rowlen** determines the size of the buffer (in bytes) required by **dbsprhead, dbsprline**, and **dbspr1row**, including the null terminator.

- **dbspr1rowlen** is useful when printing data for debugging and when scrolling data displays.
- To make the best use of **dbspr1rowlen**, application programs should call it once for every successful call to **dbresults**.
- The following code fragment illustrates the use of **dbspr1rowlen**:

```
dbcmd(dbproc, "select * from sysdatabases");
 dbcmd(dbproc, " order by name");
 dbcmd(dbproc, " compute max(crdate) by name");

dbsqlexec(dbproc);
 dbresults(dbproc);
printf("Maximum row length will be %ld \
      characters.\n", dbspr1rowlen(dbproc));
```

See also dbprhead, dbprrow, dbspr1row, dbsprhead, dbsprline, Options

dbsprhead

Description Place the server query results header into a buffer.

Syntax RETCODE dbsprhead(dbproc, buffer, buf_len)

```
DBPROCESS    *dbproc;
char          *buffer;
DBINT         buf_len;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

buffer
A pointer to a character buffer to contain the query results header.

buf_len
The length of *buffer*, including its null terminator.

Return value SUCCEED or FAIL.

Note If an error occurs, the contents of **buffer* are undefined.

Usage

- **dbsprhead** fills a programmer-supplied buffer with a null-terminated character string containing the header for the current set of query results. The header consists of the column names. The sequence of the column names matches that of the output of **dbspr1row**.
- **dbsprhead** is useful when printing data for debugging, and when scrolling data displays.
- To pad each column name to its maximum converted length, specify a pad character via the DB-Library option DBPRPAD. The pad character will be appended to each column's name. The maximum converted column length is equal to the longest possible string that could be the column's displayable data, or the length of the column's name, whichever is greater. See Options on page 402 for more details on the DBPRPAD option.
- You can specify the column separator string via the DB-Library option DBPRCOLSEP. The column separator will be added to the end of each column name except the last. The default separator is an ASCII 0x20 (space). See Options on page 402 for more details on the DBPRCOLSEP option.
- You can specify the maximum number of characters to be placed on one line via the DB-Library option DBPRLINELEN.
- You can specify the line separator string via the DB-Library option DBPRLINESEP. The default line separator is a newline (ASCII 0x0a or 0x0d, depending on the host system). See Options on page 402 for more details on the DBPRLINELEN and DBPRLINESEP options.
- The length of the buffer required by **dbsprhead** can be determined by calling **dbspr1rowlen**.
- To make the best use of **dbsprhead**, application programs should call it once for every successful call to **dbresults**.
- The following code fragment illustrates the use of **dbsprhead**:

```
dbcmd(dbproc, "select * from sysdatabases");
 dbcmd(dbproc, " order by name");
 dbcmd(dbproc, " compute max(crdate) by name");

dbsqlexec(dbproc);
 dbresults(dbproc);

dbsprhead(dbproc, buffer, sizeof(buffer));
 printf("%s\n", buffer);
```

See also

dbprhead, dbprrow, dbsetopt, dbspr1row, dbspr1rowlen, dbsprline, Options

dbsprline

Description

Get a formatted string that contains underlining for the column names produced by **dbsprhead**.

Syntax

RETCODE dbsprline(dbproc, buffer, buf_len, linechar)

```
DBPROCESS    *dbproc;
char         *buffer;
DBINT         buf_len;
DBCHAR       linechar;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

buffer

A pointer to a character buffer to contain the **dbsprline** results.

buf_len

The length of *buffer*, including its null terminator.

linechar

The character with which to "underline" column names produced by **dbsprhead**.

Return value

SUCCEED or FAIL.

Note If an error occurs, the contents of **buffer* are undefined.

Usage

- **dbsprline** is used to "underline" the column names produced by **dbsprhead**. **dbsprline** fills a programmer-supplied buffer with a null-terminated character string containing one group of the character specified by *linechar* for each column in the current set of query results. The format of this line matches the format of the output of **dbsprhead**.
- You can determine the length of the buffer required by **dbsprline** via **dbspr1rowlen**.
- To make the best use of **dbsprhead**, application programs should call it once for every successful call to **dbresults**.
- **dbsprline** is useful when printing data for debugging, and when scrolling data displays.
- The following code fragment illustrates the use of **dbsprline**:

```
dbcmd(dbproc, "select * from sysdatabases");
 dbcmd(dbproc, " order by name");
 dbcmd(dbproc, " compute max(crdate) by name");

dbsqlexec(dbproc);
 dbresults(dbproc);

/*
 ** Display the column headings, underline them
 ** with "*"
 */
 dbsprhead(dbproc, buffer, sizeof(buffer));
 printf("%s\n", buffer);

dbsprline(dbproc, buffer, sizeof(buffer), '*');
 printf("%s\n", buffer);

/* Process returned rows as usual */
```

See also dbprhead, dbprrow, dbspr1row, dbspr1rowlen, dbsprhead, Options

dbsqlexec

Description Send a command batch to the server.

Syntax RETCODE dbsqlexec(dbproc)

DBPROCESS *dbproc;

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value SUCCEED or FAIL.

The most common reason for failing is a SQL syntax error. **dbsqlexec** will also fail if there are semantic errors, such as incorrect column or table names. Failure occurs if any of the commands in the batch contains a semantic or syntax error. **dbsqlexec** also fails if previous results had not been processed, or if the command buffer was empty.

In addition, a runtime error, such as a database protection violation, can cause **dbsqlexec** to fail. A runtime error will cause **dbsqlexec** to fail:

- If the command causing the error is the only command in the command buffer
- If the command causing the error is the first command in a multiple-command buffer

If the command buffer contains multiple commands (and the first command in the buffer is ok), a runtime error will not cause **dbsqlexec** to fail. Instead, failure will occur with the **dbresults** call that processes the command causing the runtime error.

The situation is a bit more complicated for runtime errors and stored procedures. A runtime error on an **execute** command may cause **dbsqlexec** to fail, in accordance with the rule given in the previous paragraphs. A runtime error on a statement *inside* a stored procedure will not cause **dbsqlexec** to fail, however. For example, if the stored procedure contains an **insert** statement and the user does not have insert permission on the database table, the **insert** statement will fail, but **dbsqlexec** will still return SUCCEED. To check for runtime errors inside stored procedures, use the **dbretstatus** routine to look at the procedure's return status, and trap relevant server messages inside your message handler.

Usage

- This routine sends SQL commands, stored in the command buffer of the DBPROCESS, to the server. Commands may be added to the DBPROCESS structure by calling **dbcmd** or **dbfcmd**.
- Once **dbsqlexec** returns SUCCEED, the application must call **dbresults** to process the results.
- The typical sequence of calls is:

```
DBINT       xvariable;
 DBCHAR      yvariable[10];

/* Read the query into the command buffer */
 dbcmd(dbproc, "select x = 100, y = 'hello'");

/* Send the query to SQL Server */
 dbsqlexec(dbproc);

/* Get ready to process the query results */
 dbresults(dbproc);

/* Bind column data to program variables */
 dbbind(dbproc, 1, INTBIND, (DBINT) 0,
      (BYTE *) &xvariable);
 dbbind(dbproc, 2, STRINGBIND, (DBINT) 0,
      yvariable);

/* Now process each row */
```

```
while (dbnextrow(dbproc) != NO_MORE_ROWS)
{
     C-code to print or process row data
}
```

- **dbsqlexec** is equivalent to **dbsqlsend** followed by **dbsqlok**. However, after sending a query to the server, **dbsqlexec** waits until a response is received or until the timeout period has elapsed. By substituting **dbsqlsend** and **dbsqlok** for **dbsqlexec**, you can sometimes provide a way for the application to respond more effectively to multiple input and output streams. See the reference pages for dbsqlsend and dbsqlok for more information.
- Multiple commands may exist in the command buffer when an application calls **dbsqlexec**. These commands are sent to the server as a unit and are considered to be a single command batch.

See also

dbcmd, dbfcmd, dbnextrow, dbresults, dbretstatus, dbsettime, dbsqlok, dbsqlsend

dbsqlok

Description

Wait for results from the server and verify the correctness of the instructions the server is responding to.

Syntax

RETCODE dbsqlok(dbproc)

DBPROCESS *dbproc;

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

SUCCEED or FAIL.

The most common reason for failing is a SQL syntax error. **dbsqlok** will also fail if there are semantic errors, such as incorrect column or table names. Failure occurs if any of the commands in the batch contains a semantic or syntax error.

In addition, a runtime error, such as a database protection violation, will cause **dbsqlok** to fail *if* the command buffer contains only a single command. If the command buffer contains multiple commands, a runtime error will *not* cause **dbsqlok** to fail. Instead, failure will occur with the **dbresults** call that processes the command causing the runtime error.

The situation is a bit more complicated for runtime errors and stored procedures. A runtime error on an **execute** command may cause **dbsqlok** to fail, in accordance with the rule given in the previous paragraph. A runtime error on a statement *inside* a stored procedure will not cause **dbsqlok** to fail, however. For example, if the stored procedure contains an **insert** statement and the user does not have insert permission on the database table, the **insert** statement will fail, but **dbsqlok** will still return SUCCEED. To check for runtime errors inside stored procedures, use the **dbretstatus** routine to look at the procedure's return status and trap relevant server messages inside your message handler.

Usage

- **dbsqlok** reports the success or failure of a server command and initiates results processing for successful commands.
- A successful **dbsqlok** call must always be followed by a call to **dbresults** to process the results.
- **dbsqlok** is useful in the following situations:
    - After a **dbsqlsend** call

      dbsqlok must be called after a batch of Transact-SQL commands is sent to the server with **dbsqlsend**.
    - After a **dbrpcsend** call

      dbsqlok must be called after an RPC command is sent with **dbrpcinit**, **dbrpcparam**, and **dbrpcsend**.
    - After calls to **dbwritetext** or **dbmoretext**

      dbsqlok must be called after a text update command is sent to the server by a call to **dbwritetext** or **dbmoretext**.

Using dbsqlok with dbsqlsend

- **dbsqlok** initiates results processing after a call to **dbsqlsend**.
- **dbsqlok** and **dbsqlsend** provide an alternative to **dbsqlexec**. **dbsqlexec** sends a command batch and waits for initial results from the server. The application is blocked from doing anything else until results arrive. When **dbsqlsend** and **dbsqlok** are used with **dbpoll**, the application has a non-blocking alternative. The typical control sequence is as follows:

- A call to **dbsqlsend** sends the command to the server.
- The program calls **dbpoll** in a loop to check for the arrival of server results. Non-related work can be performed during each loop iteration. The loop terminates when **dbpoll** indicates results have arrived.
- A call to **dbsqlok** reports success or failure and initiates results processing if successful.

Note On occasion, **dbpoll** may report that data is ready for **dbsqlok** to read when only the first bytes of the server response are present. When this occurs, **dbsqlok** waits for the rest of the response or until the timeout period has elapsed, just like **dbsqlexec**. In practice, however, the entire response is usually available at one time.

- The example below illustrates the use of **dbsqlok** and **dbpoll**. The example calls an application function, **busy_wait,** to execute a **dbpoll** loop. Here is the mainline code that calls **busy_wait**:

```
/*
 ** This is a query that will take some time.
 */
 dbcmd(dbproc, "waitfor delay '00:00:05' select its = 'over'");

 /*
 ** Send the query with dbsqlsend. dbsqlsend does not
 ** wait for a server response.
 */
 retcode = dbsqlsend(dbproc);
 if (retcode != SUCCEED)
 {
   fprintf(stdout, "dbsqlsend failed. Exiting.\n");
   dbexit();
   exit(ERREXIT);
 }

 /*
 ** If we call dbsqlok() now, it might block. But, we can use
 ** a dbpoll() loop to get some other work done while
 ** we are waiting for the results.
 */
 busy_wait(dbproc);

 /*
 ** Now there should be some results waiting to be read, so
```

```
  ** call dbsqlok().
  */
  retcode = dbsqlok(dbproc);
  if (retcode != SUCCEED)
  {
    fprintf(stdout, "Query failed.\n");
  }
  else
  {
    ... dbresults() loop goes here ...
  }
```

busy_wait executes a **dbpoll** loop. During each iteration of the loop, a call to **dbpoll** determines whether results have arrived. If results have arrived, **busy_wait** returns. Otherwise, the function **wait_work** is called. **wait_work** performs a piece of non-related work, then returns. The functions **wait_work_init** and **wait_work_cleanup** perform initialization and cleanup for **wait_work**. Here is the code for these functions:

```
void busy_wait(dbproc)
 DBPROCESS *dbproc;
 {
   RETCODE retcode;
   DBPROCESS *ready_dbproc;
   int poll_ret_reason;

   wait_work_init();
   while(1)
   {
     retcode = dbpoll(dbproc, 0, &ready_dbproc, &poll_ret_reason);
     if (retcode != SUCCEED)
     {
       fprintf(stdout, "dbpoll() failed! Exiting.\n");
       dbexit();
       exit(ERREXIT);
     }
     if (poll_ret_reason == DBRESULT)
     {
       /*
       ** Query results have arrived. Now we break out of
       ** the loop and return. Our caller can then call dbsqlok().
       */
       break; /* while */
     }
     else
     {
       /*
```

```
      ** Here's where we can do some non-related work while we
      ** are waiting.
      */
      wait_work();
    }
  } /* while */
  wait_work_cleanup();
} /* busy_wait */

/* These globals are used by the wait functions. */
static int wait_pos;
static char wait_char;
void wait_work()
{

  /*
  ** "work", as defined here, consists of drawing a 'w' or 'W' to
  ** the terminal. We output one character each time we are called.
  ** When we reach the 65th character position, we switch from
  ** 'w' to 'W' (or vice-versa) and start over.
  */
  fputc(wait_char, stdout);
  ++wait_pos;
  if (wait_pos >= 65)
  {
     /*
     ** Go back to the beginning of the line, then switch from
     ** 'W' to 'w' or vice versa.
     */
     fputc('\r', stdout);
     wait_pos = 0;
     wait_char = (wait_char == 'w' ? 'W' : 'w');
  }
}
void wait_work_init()
{
  wait_pos = 0;
  wait_char = 'w';
}
void wait_work_cleanup()
{
  fputc('\n', stdout);
}
```

Using dbsqlok with dbrpcsend

- **dbsqlok** initiates results processing after an RPC command. RPC commands are constructed and sent with **dbrpcinit**, **dbrpcparam**, and **dbrpcsend**. After **dbrpcsend**, the program must call **dbsqlok**.
- **dbpoll** can be called in a loop to poll for a server response between **dbrpcsend** and **dbsqlok**.
- See the reference pages for dbrpcinit, dbrpcparam, and dbrpcsend for more information. Example 8 of the online sample programs demonstrates an RPC command.

Using dbsqlok with dbwritetext and dbmoretext

- **dbsqlok** initiates results processing after a text update command. For text updates, chunks of text can be sent to the server with **dbwritetext** and **dbmoretext**. After both of these calls, **dbsqlok** must be called.
- See the reference pages for dbwritetext and dbmoretext for more information. **dbwritetext** has an example.

See also

dbcmd, dbfcmd, DBIORDESC, DBIOWDESC, dbmoretext, dbnextrow, dbpoll, DBRBUF, dbresults, dbretstatus, dbrpcsend, dbsettime, dbsqlexec, dbsqlsend, dbwritetext

dbsqlsend

Description

Send a command batch to the server and do not wait for a response.

Syntax

```
RETCODE dbsqlsend(dbproc)

 DBPROCESS      *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

SUCCEED or FAIL.

dbsqlsend may fail if previous results had not been processed, or if the command buffer was empty.

Usage

- This routine sends SQL commands, stored in the command buffer, to the server. The application can add commands to the command buffer by calling **dbcmd** or **dbfcmd**.

- Once **dbsqlsend** returns SUCCEED, the application must call **dbsqlok** to verify the accuracy of the command batch. The application can then call **dbresults** to process the results.
- **dbsqlexec** is equivalent to **dbsqlsend** followed by **dbsqlok**.
- The use of **dbsqlsend** with **dbsqlok** is of particular value in UNIX and AOS/VS applications. After sending a query to the server, **dbsqlexec** waits until a response is received or until the timeout period has elapsed. By substituting **dbsqlsend**, **dbpoll** and **dbsqlok** for **dbsqlexec**, you can sometimes provide a way for an application to respond more effectively to multiple input and output streams. See the dbsqlok reference page for more information and an example.

See also

dbcmd, dbfcmd, DBIORDESC, DBIOWDESC, dbnextrow, dbpoll, dbresults, dbsettime, dbsqlexec, dbsqlok

dbstrbuild

Description

Build a printable string from text containing placeholders for variables.

Syntax

```
int dbstrbuild(dbproc, charbuf, bufsize,
        text [, formats [, arg] ... ])

DBPROCESS    *dbproc;
char         *charbuf;
int          bufsize;
char         *text;
char         *formats;
???          args???;
```

Parameters

dbproc

A pointer to the DBPROCESS that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server. **dbstrbuild** uses it only as a parameter to the programmer-installed error handler (if one exists) when an error occurs.

charbuf

A pointer to the destination buffer that will contain the message built by **dbstrbuild**.

bufsize

The size of the destination buffer, in bytes. This size must include a single byte for the results string's null terminator.

text
A pointer to a null-terminated character string that contains message text and placeholders for variables. Placeholders consist of a percent sign, an integer, and an exclamation point. The integer indicates which argument to substitute for a particular placeholder. Arguments and format strings are numbered from left to right. Argument 1 is substituted for placeholder "%1!", and so on.

formats
A pointer to a null-terminated string containing one **sprintf**-style format specifier for each place holder in the *text* string.

args
The values that will be converted according to the contents of the *formats* string. There must be one argument for each format in the *formats* string. The first value will correspond to the "%1!" parameter, the second the "%2!", and so forth. The results are undefined if there are insufficient arguments for the format. If the format is exhausted while arguments remain, the excess arguments are simply ignored.

Return value

The length of the resulting message string, not including the null terminator, on success; a negative integer on failure.

Usage

- Parameters in error messages may occur in different orders in different languages. **dbstrbuild** allows construction of error messages in a manner similar to the C standard-library **sprintf** routine. Use of **dbstrbuild** ensures easy translation of error messages from one language to another.
- **dbstrbuild** builds a printable string from an error text that contains placeholders for variables, a format string containing information about the types and appearances of those variables, and a variable number of arguments that provide actual values for those variables.
- Placeholders for variables consist of a percent sign, an integer, and an exclamation point. The integer indicates which argument to substitute for a particular placeholder. Arguments and format strings are numbered from left to right. Argument 1 is substituted for placeholder "%1!", and so on.

  For example, consider an error message that complains about a misused keyword in a stored procedure. The message requires three arguments: the misused keyword, the line in which the keyword occurs, and the name of the stored procedure in which the misuse occurs. In the English localization file, the message text might appear as:

```
The keyword '%1!' is misused in line %2! of stored
 procedure '%3!' .
```

In the Yoda-speak localization file, the same message might appear as:

```
In line '%2!' of stored procedure '%3!', the keyword
'%1!' misused is.
```

The **dbstrbuild** line for either of the above messages would be:

```
dbstrbuild(dbproc, charbuf, BUFSIZE, <get the
 message somehow>, "%s %d %s", keyword,
 linenum, sp_name)
```

keyword is substituted for placeholder "%1!", *linenum* is substituted for placeholder "%2!", and *sp_name* is substituted for placeholder "%3!".

- The following code fragment illustrates the use of **dbstrbuild** to build messages. For simplicity, the text of the message is hard-coded. In practice, **dbstrbuild** message texts come from a localization file.

```
char    charbuf[BUFSIZE];
 int     linenum = 15;
 char    *filename = "myfile";
 char    *dirname = "mydir";

dbstrbuild (dbproc, charbuf, BUFSIZE,
 "Unable to read line %1! of file %2! in \
 directory %3!.", "%d %s %s", linenum,
 filename, dirname);
printf(charbuf);
```

- **dbstrbuild** format specifiers may be separated by any other characters, or they may be adjacent to each other. This allows pre-existing English-language message strings to be used as **dbstrbuild** format parameters. The first format specifier describes the "%1!" parameter, the second the "%2!" parameter, and so forth.

See also dbconvert, dbdatename, dbdatepart

dbstrcmp

Description Compare two character strings using a specified sort order.

Syntax

```
int dbstrcmp(dbproc, str1, len1, str2, len2,
        sortorder)

DBPROCESS      *dbproc;
char           *str1;
int            len1;
```

```
char            *str2;
int             len2;
DBSORTORDER     *sortorder;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

str1
A pointer to the first character string to compare. *str1* may be NULL.

len1
The length, in bytes, of *str1*. If *len1* is -1, *str1* is assumed to be null-terminated.

str2
A pointer to the second character string to compare. *str2* may be NULL.

len2
The length, in bytes, of *str2*. If *len2* is -1, *str2* is assumed to be null-terminated.

sortorder
A pointer to a DBSORTORDER structure allocated via **dbloadsort**. If *sortorder* is NULL, **dbstrcmp** compares *str1* and *str2* using their binary values, just as **strcmp** does.

Return value

- 1 if *str1* is lexicographically greater than *str2*.
- 0 if *str1* is lexicographically equal to *str2*.
- -1 if *str1* is lexicographically less than *str2*.

Usage

- **dbstrcmp** compares *str1* and *str2* and returns an integer greater than, equal to, or less than 0, according to whether *str1* is lexicographically greater than, equal to, or less than *str2*.
- **dbstrcmp** uses a sort order that was retrieved from the server via **dbloadsort**. This allows DB-Library application programs to compare strings using the same sort order as the server.

- Note that some languages contain strings that are lexicographically equal according to some specified sort order, but contain different characters. Even though they are "equal," there is a standard order that should be used when placing them into an ordered list. When given two strings like this to compare, **dbstrcmp** returns 0 (indicating the two strings are equal), but **dbstrsort** returns some non-zero value indicating that one of these strings should appear before the other in a sorted list.

  Below is an example of this behavior. The two English-language character strings are used with a case-insensitive sort order that specifies that uppercase letters should appear before lowercase:

```
/* This call returns 0: */
 dbstrcmp(dbproc, "ABC", 3, "abc", 3, mysort);

 /* This call returns a negative value: */
 dbstrsort(dbproc, "ABC", 3, "abc", 3, mysort);
```

See also dbfreesort, dbloadsort, dbstrsort

dbstrcpy

Description Copy all or a portion of the command buffer.

Syntax RETCODE dbstrcpy(dbproc, start, numbytes, dest)

```
DBPROCESS    *dbproc;
int           start;
int           numbytes;
char         *dest;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

start

Character position in the command buffer to start copying from. The first character has position 0. If *start* is greater than the length of the command buffer, **dbstrcpy** inserts a null terminator at *dest[0]*.

numbytes
The number of characters to copy. If *numbytes* is -1, **dbstrcpy** will copy the entire command buffer, whether or not *dest* points to adequate space. It is legal to copy 0 bytes, in which case **dbstrcpy** inserts a null terminator at *dest[0]*. If there are not *numbytes* available to copy, **dbstrcpy** copies the number of bytes available and returns SUCCEED.

dest
A pointer to the destination buffer to copy the source string into. Before calling **dbstrcpy**, the caller must verify that the destination buffer is large enough to hold the copied characters. The function **dbstrlen** returns the size of the entire command buffer.

Return value

SUCCEED or FAIL.

dbstrcpy returns FAIL if *start* is negative.

Usage

- **dbstrcpy** copies a portion of the command buffer to a string buffer supplied by the application. The copy is null-terminated.
- Internally, the command buffer is a linked list of non-null-terminated text strings. **dbgetchar**, **dbstrcpy**, and **dbstrlen** together provide a way to locate and copy parts of the command buffer.
- **dbstrcpy** assumes that the destination is large enough to receive the source string. If not, a segmentation fault is likely.
- When *numbytes* is passed as -1, **dbstrcpy** copies the entire command buffer. Do not pass *numbytes* as -1 unless you are certain that *dest* points to adequate space for this string. The function **dbstrlen** returns the length of the current command string.
- The following fragment shows how to print the entire command buffer to a file:

```
FILE      *outfile;
 DBPROCESS *dbproc;
 char      *prbuf; /* buffer for collecting the command buffer
                   ** contents as a null-terminated string
                   */
 RETCODE return_code;

/*
 ** Allocate sufficient space.  dbstrlen() returns the number of
 ** characters currently in the command buffer.  We need one
 ** more byte because dbstrcpy will append a null terminator.
 ** NOTE that memory allocation and disposal may be done
 ** differently on your platform.
 **/
```

```
prbuf = (char *) malloc(dbstrlen(dbproc) + 1);
if (prbuf == NULL)
{
  fprintf(stderr, "Out of memory.");
  dbexit();
  exit(ERREXIT); /* ERREXIT is defined in the DB-lib headers */
}
/* Copy the command buffer into the allocated space: */
return_code = dbstrcpy(dbproc, 0, -1, prbuf);
assert(return_code == SUCCEED);

/* Print the contents: */
fprintf(outfile, "%s", prbuf);

/* Free the buffer: */
free(prbuf);
```

See also dbcmd, dbfcmd, dbfreebuf, dbgetchar, dbstrlen

dbstrlen

Description Return the length, in characters, of the command buffer.

Syntax int dbstrlen(dbproc)

DBPROCESS *dbproc;

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value The length, in characters, of the command buffer.

Usage

- **dbstrlen** returns the length, in characters, of the SQL command text in the command buffer.
- Internally, the command buffer is a linked list of non-null-terminated text strings. **dbgetchar**, **dbstrcpy**, and **dbstrlen** together provide a way to locate and copy parts of the command buffer.
- Before you copy the command buffer with **dbstrcpy**, use **dbstrlen** to make sure that the destination buffer is large enough.

- The count returned by **dbstrlen** does not include space for a null terminator.

See also

dbcmd, dbfcmd, dbfreebuf, dbgetchar, dbstrcpy

dbstrsort

Description

Determine which of two character strings should appear first in a sorted list.

Syntax

```
int dbstrsort(dbproc, str1, len1, str2, len2,
        sortorder)

 DBPROCESS      *dbproc;
 char           *str1;
 int             len1;
 char           *str2;
 int             len2;
 DBSORTORDER    *sortorder;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

str1

A pointer to the first character string to compare. *str1* may be NULL.

len1

The length, in bytes, of *str1*. If *len1* is -1, *str1* is assumed to be null-terminated.

str2

A pointer to the second character string to compare. *str2* may be NULL.

len2

The length, in bytes, of *str2*. If *len2* is -1, *str2* is assumed to be null-terminated.

sortorder

A pointer to a DBSORTORDER structure allocated via **dbloadsort**. If *sortorder* is NULL, **dbstrsort** compares *str1* and *str2* using their binary values, just as **strcmp** does.

Return value

- 1 if *str1* should appear after *str2*.

- 0 if *str1* is identical to *str2*.
- -1 if *str1* should appear before *str2*.

Usage

- **dbstrsort** compares *str1* and *str2* and returns an integer greater than, equal to, or less than 0, according to whether *str1* should appear after, at the same place (the strings are identical), or before *str2* in a sorted list.
- **dbstrsort** uses a sort order that was retrieved from the server via **dbloadsort**. This allows DB-Library application programs to compare strings using the same sort order as the server.
- Note that some languages contain strings that are lexicographically equal according to some specified sort order, but contain different characters. Even though they are "equal," there is a standard order that should be used when placing them into an ordered list. When given two strings like this to compare, **dbstrcmp** returns 0 (indicating the two strings are equal), but **dbstrsort** returns some non-zero value indicating that one of these strings should appear before the other in a sorted list.

  Below is an example of this behavior. The two English-language character strings are used with a case-insensitive sort order that specifies that uppercase characters should appear before lowercase:

```
/* This call returns 0: */
 dbstrcmp(dbproc, "ABC", 3, "abc", 3, mysort);

 /* This call returns a negative value: */
 dbstrsort(dbproc, "ABC", 3, "abc", 3, mysort);
```

- **dbstrsort** can only be used to examine two character strings that have already been identified as equal via **dbstrcmp**. If **dbstrcmp** has not identified these strings as being equal to each other, **dbstrsort**'s behavior is undefined.

See also

dbfreesort, dbloadsort, dbstrcmp

dbtabbrowse

Description

Determine whether the specified table is updatable via the DB-Library browse-mode facilities.

Syntax

DBBOOL dbtabbrowse(dbproc, tabnum)

```
DBPROCESS    *dbproc;
int          tabnum;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

tabnum
The number of the table of interest, as specified in the **select** statement's **from** clause. Table numbers start at 1.

Return value

TRUE or FALSE.

Usage

- **dbtabbrowse** is one of the DB-Library browse-mode routines. See "Browse mode" on page 26 for a detailed discussion of browse mode.
- **dbtabbrowse** provides a way to identify browsable tables. It is useful when examining ad hoc queries prior to performing browse mode updates based on them. If the query has been hard coded into the program, this routine is obviously unnecessary.
- For a table to be considered "browsable," it must have a unique index and a timestamp column.
- The application can call **dbtabbrowse** anytime after **dbresults**.
- Example 7 of the online sample programs contains a call to **dbtabbrowse**.

See also

dbcolbrowse, dbcolsource, dbqual, dbtabcount, dbtabname, dbtabsource, dbtsnewlen, dbtsnewval, dbtsput

dbtabcount

Description

Return the number of tables involved in the current **select** query.

Syntax

```
int dbtabcount(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The number of tables, including server work tables, involved in the current set of row results.

dbtabcount will return -1 in case of error.

Usage

- **dbtabcount** is one of the DB-Library browse-mode routines. It is usable only with results from a browse-mode **select** (that is, a **select** containing the key words **for browse**). See “Browse mode” on page 26 for a detailed discussion of browse mode.
- A **select** query can generate a set of result rows whose columns are derived from several database tables. To perform browse-mode updates of columns in a query’s select list, the application must know how many tables were involved in the query, because each table requires a separate **update** statement. **dbtabcount** can provide this information for ad hoc queries. If the query has been hard-coded into the program, this routine is obviously unnecessary.
- The count returned by this routine includes any server “work tables” used in processing the query. The server sometimes creates temporary, internal work tables to process a query. It deletes these work tables by the time it finishes processing the statement. Work tables are not updatable and are not available to the application. Therefore, before using a table number, the application must make sure that it does not belong to a work table. **dbtabname** can be used to determine whether a particular table number refers to a work table.
- The application can call **dbtabcount** anytime after **dbresults**.
- Example 7 of the online sample programs contains a call to **dbtabcount**.

See also

dbcolbrowse, dbcolsource, dbqual, dbtabbrowse, dbtabname, dbtabsource, dbtsnewlen, dbtsnewval, dbtsput

dbtabname

Description

Return the name of a table based on its number.

Syntax

```
char *dbtabname(dbproc, tabnum)

DBPROCESS    *dbproc;
int          tabnum;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

tabnum
The number of the table of interest. Table numbers start with 1. Use **dbtabcount** to find out the total number of tables involved in a particular query.

Return value

A pointer to the null-terminated name of the specified table. This pointer will be NULL if the table number is out of range or if the specified table is a server work table. See the dbtabcount reference page for a description of work tables.

Usage

- **dbtabname** is one of the DB-Library browse-mode routines. It is usable only with results from a browse-mode **select** (that is, a **select** containing the key words **for browse**). See "Browse mode" on page 26 for a detailed discussion of browse mode.
- A **select** query can generate a set of result rows whose columns are derived from several database tables. **dbtabname** provides a way for an application to determine the name of each table involved in an ad hoc query. If the query has been hard-coded into the program, this routine obviously is unnecessary.
- The application can call **dbtabname** anytime after **dbresults**.
- Example 7 of the online sample programs contains a call to **dbtabname**.

See also

dbcolbrowse, dbcolsource, dbqual, dbtabbrowse, dbtabcount, dbtabsource, dbtsnewlen, dbtsnewval, dbtsput

dbtabsource

Description

Return the name and number of the table from which a particular result column was derived.

Syntax

char *dbtabsource(dbproc, colnum, tabnum)

```
DBPROCESS     *dbproc;
int            colnum;
int           *tabnum;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

colnum
The number of the result column of interest. Column numbers start at 1.

tabnum
A pointer to an integer, which will be filled in with the table's number. Many DB-Library routines that deal with browse mode accept either a table name or a table number. If **dbtabsource** returns NULL (see the "Returns" section below), **tabnum* will be set to -1.

Return value

A pointer to the name of the table from which this result column was derived. A NULL return value can mean a few different things:

- The DBPROCESS is dead or not enabled. This is an error that will cause an application's error handler to be invoked.
- The column number is out of range.
- The column is the result of an expression, such as **max(colname)**.

Usage

- **dbtabsource** is one of the DB-Library browse-mode routines. It is usable only with results from a browse-mode **select** (that is, a **select** containing the key words **for browse**). See "Browse mode" on page 26 for a detailed discussion of browse mode.
- **dbtabsource** allows an application to determine which tables provided the columns in the current set of result rows. This information is valuable when using **dbqual** to construct **where** clauses for **update** and **delete** statements based on ad hoc queries. If the query has been hard-coded into the program, this routine obviously is unnecessary.
- The application can call **dbtabsource** anytime after **dbresults**.
- Example 7 of the online sample programs contains a call to **dbtabsource**.

See also

dbcolbrowse, dbcolsource, dbqual, dbtabbrowse, dbtabcount, dbtabname, dbtsnewlen, dbtsnewval

DBTDS

Description

Determine which version of TDS (the Tabular Data Stream protocol) is being used.

Syntax

```
int DBTDS(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The version of TDS used by *dbproc* to communicate with the server. Currently, the possible versions are:

- DBTDS_2_0
- DBTDS_3_4
- DBTDS_4_0
- DBTDS_4_2
- DBTDS_4_6
- DBTDS_4_9_5
- DBTDS_5_0

DBTDS returns a negative integer on error.

Usage

- DBTDS returns the version of TDS (Tabular Data Stream protocol) being used by *dbproc* to communicate with the server.

See also

dbversion

dbtextsize

Description

Returns the number of bytes of text or image data that remain to be read for the current row.

Syntax

```
DBINT dbtextsize(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The following table lists the return values for **dbtextsize**

dbtextsize returns	To indicate
>= 0	The number of bytes that remain to be read. Zero indicates NO_MORE_ROWS.
-1	An error has occurred.
-2	**dbtextsize** has been called for RPC data.

Usage

- **dbtextsize** assumes that there is only one column and that this column is of datatype text or image.
- **dbtextsize** is useful when an application does not know how large a text or image value is.
- **dbtextsize** does not work with RPC text data.

See also

dbreadtext

dbtsnewlen

Description

Return the length of the new value of the *timestamp* column after a browse-mode update.

Syntax

```
int dbtsnewlen(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

The length (in bytes) of the updated row's new timestamp value. If no timestamp was returned to the application (possibly because the update was unsuccessful, or because the **update** statement did not contain the **tsequal** built-in function), **dbtsnewlen** will return -1.

Usage

- **dbtsnewlen** is one of the DB-Library browse-mode routines. See "Browse mode" on page 26 for a detailed discussion of browse mode.
- **dbtsnewlen** provides information about the *timestamp* column. The **where** clause returned by **dbqual** contains a call to the **tsequal** built-in function. When such a **where** clause is used in an **update** statement, the **tsequal** function places a new value in the updated row's *timestamp* column and returns the new timestamp value to the application (if the update is successful). The **dbtsnewlen** function allows the application to save the length of the new timestamp value, possibly for use with **dbtsput**.

See also

dbcolbrowse, dbcolsource, dbqual, dbtabbrowse, dbtabcount, dbtabname, dbtabsource, dbtsnewval, dbtsput

dbtsnewval

Description

Return the new value of the *timestamp* column after a browse-mode update.

Syntax

```
DBBINARY *dbtsnewval(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value

A pointer to the updated row's new timestamp value. If no timestamp was returned to the application (possibly because the update was unsuccessful, or because the **update** statement did not contain the **tsequal** built-in function), the pointer will be NULL.

Usage

- **dbtsnewval** is one of the DB-Library browse-mode routines. See "Browse mode" on page 26 for a detailed discussion of browse mode.
- **dbtsnewval** provides information about the *timestamp* column. The **where** clause returned by **dbqual** contains a call to the **tsequal** built-in function. When such a **where** clause is used in an **update** statement, the **tsequal** function places a new value in the updated row's *timestamp* column and returns the new timestamp value to the application (if the update is successful). This routine allows the application to save the new timestamp value, possibly for use with **dbtsput**.

See also

dbtabbrowse, dbtabsource, dbqual, dbtabbrowse, dbtabcount, dbtabname, dbtabsource, dbtsnewlen, dbtsput

dbtsput

Description

Put the new value of the *timestamp* column into the given table's current row in the DBPROCESS.

Syntax

```
RETCODE dbtsput(dbproc, newts, newtslen, tabnum,
        tabname)

DBPROCESS    *dbproc;
DBBINARY      *newts;
int            newtslen;
int            tabnum;
char          *tabname;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

This must be the DBPROCESS used to perform the original **select** query.

newts

A pointer to the new timestamp value. It is returned by **dbtsnewval**.

newtslen

The length of the new timestamp value. It is returned by **dbtsnewlen**.

tabnum

The number of the updated table. Table numbers start at 1. *tabnum* must refer to a browsable table. Use **dbtabbrowse** to determine whether a table is browsable.

If this value is -1, the *tabname* parameter will be used to identify the table.

tabname

A pointer to a null-terminated table name. *tabname* must refer to a browsable table. If this pointer is NULL, the *tabnum* parameter will be used to identify the table.

Return value

SUCCEED or FAIL.

The following situations will cause this routine to return FAIL:

- The application tries to update the timestamp of a non-existent row.
- The application tries to update the timestamp using NULL as the timestamp value (*newts*).
- The specified table is non-browsable.

Usage

- **dbtsput** is one of the DB-Library browse-mode routines. See "Browse mode" on page 26 for a detailed discussion of browse mode.
- **dbtsput** manipulates the timestamp column. The **where** clause returned by **dbqual** contains a call to the **tsequal** built-in function. When such a **where** clause is used in an **update** statement, the **tsequal** function places a new value in the updated row's timestamp column and returns the new timestamp value to the application (if the update is successful). If the same row is updated a second time, the **update** statement's **where** clause must use the latest timestamp value.

  This routine updates the timestamp in the DBPROCESS for the row currently being browsed. Then, if the application needs to update the row a second time, it can call **dbqual** to formulate a new **where** clause that uses the new timestamp.

See also

dbcolbrowse, dbcolsource, dbqual, dbtabbrowse, dbtabcount, dbtabname, dbtabsource, dbtsnewlen, dbtsnewval

dbtxptr

Description

Return the value of the text pointer for a column in the current row.

Syntax

```
DBBINARY *dbtxptr(dbproc, column)

DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the select list column of interest. Column numbers start at 1.

Return value

A DBBINARY pointer to the text pointer for the column of interest. This pointer may be NULL.

Usage

- Every database column row of type SYBTEXT or SYBIMAGE has an associated text pointer, which uniquely identifies the text or image value. This text pointer is used by the **dbwritetext** function to update text and image values.
- It is important that all the rows of the specified text or image column have valid text pointers. A text or image column row will have a valid text pointer if it contains data. However, if the text or image column row contains a null value, its text pointer will be valid only if the null value was explicitly entered with the **update** statement.

  Assume a table *textnull* with columns *key* and *x*, where *x* is a text column that permits nulls. The following statement assigns valid text pointers to the text column's rows:

  ```
  update textnull
   set x = null
  ```

  On the other hand, the **insert** of a null value into a text column does not provide a valid text pointer. This is true for an **insert** of an explicit null or an **insert** of an implicit null, such as the following:

  ```
  insert textnull (key)
   values (2)
  ```

  When dealing with a null text or image value, be sure to use **update** to get a valid text pointer.
- An application must **select** a row containing a text or image value before calling **dbtxptr** to return the associated text pointer. The **select** causes a copy of the text pointer to be placed in the application's DBPROCESS. The application can then retrieve this text pointer from the DBPROCESS via **dbtxptr**.

  If no **select** is performed prior to the call to **dbtxptr**, the call will result in a DB-Library error message.
- For an example that uses **dbtxptr**, see the dbwritetext reference page.

See also

dbtxtimestamp, dbwritetext

dbtxtimestamp

Description

Return the value of the text timestamp for a column in the current row.

Syntax

```
DBBINARY *dbtxtimestamp(dbproc, column)

DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column

The number of the select list column of interest. Column numbers start at 1.

Return value

A DBBINARY pointer to the text timestamp for the column of interest. This pointer may be NULL.

Usage

- Every database column of type SYBTEXT or SYBIMAGE has an associated text timestamp, which marks the time of the column's last modification. The text timestamp is useful in conjunction with the **dbwritetext** function, to ensure that two competing application users do not inadvertently wipe out each other's modifications to the same value in the database. It is returned to the DBPROCESS when a Transact-SQL **select** is performed on a SYBTEXT or SYBIMAGE column.
- The length of a non-NULL text timestamp is always DBTXTSLEN (currently defined as 8 bytes).
- An application must **select** a row containing a text or image value before calling **dbtxtimestamp** to return the associated text timestamp. The **select** causes a copy of the text timestamp to be placed in the application's DBPROCESS. The application can then retrieve this text timestamp from the DBPROCESS via **dbtxtimestamp**.

  If no **select** is performed prior to the call to **dbtxtimestamp**, the call will result in a DB-Library error message.
- For an example that uses **dbtxtimestamp**, see the dbwritetext reference page.

See also

dbtxptr, dbwritetext

dbtxtsnewval

Description
Return the new value of a text timestamp after a call to **dbwritetext**.

Syntax
```
DBBINARY *dbtxtsnewval(dbproc)

 DBPROCESS     *dbproc;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

Return value
A pointer to the new text timestamp value for the SYBTEXT or SYBIMAGE value modified by a **dbwritetext** operation. This pointer may be NULL.

Usage

- Every database column of type SYBTEXT or SYBIMAGE has an associated text timestamp, which is updated whenever the column's value is changed. The text timestamp is useful in conjunction with the **dbwritetext** function to ensure that two competing application users do not inadvertently wipe out each other's modifications to the same value in the database. It is returned to the DBPROCESS when a Transact-SQL **select** is performed on a SYBTEXT or SYBIMAGE column and may be examined by calling **dbtxtimestamp**.
- After each successful **dbwritetext** operation (which may include a number of calls to **dbmoretext**), the server will send the updated text timestamp value back to DB-Library. **dbtxtsnewval** provides a way for the application to get this new timestamp value.
- The application can use **dbtxtsnewval** in two ways. First, the return from **dbtxtsnewval** can be used as the *timestamp* parameter of a **dbwritetext** call. Second, **dbtxtsnewval** and **dbtxtsput** can be used together to put the new timestamp value into the DBPROCESS row buffer, for future access via **dbtxtimestamp**. This is particularly useful when the application is buffering result rows and does not need the new timestamp immediately.

See also
dbmoretext, dbtxtimestamp, dbtxtsput, dbwritetext

dbtxtsput

Description
Put the new value of a text timestamp into the specified column of the current row in the DBPROCESS.

Syntax

RETCODE dbtxtsput(dbproc, newtxts, colnum)

```
DBPROCESS    *dbproc;
DBBINARY     *newtxts;
int          colnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

newtxts

A pointer to the new text timestamp value. It is returned by **dbtxtsnewval**.

colnum

The number of the select list column of interest. Column numbers start at 1.

Return value

SUCCEED or FAIL.

Usage

- Every database column of type SYBTEXT or SYBIMAGE has an associated text timestamp, which is updated whenever the column's value is changed. The text timestamp is useful in conjunction with the **dbwritetext** function, to ensure that two competing application users do not inadvertently wipe out each other's modifications to the same value in the database. It is returned to the DBPROCESS when a Transact-SQL **select** is performed on a SYBTEXT or SYBIMAGE column and may be examined by calling **dbtxtimestamp**.
- After each successful **dbwritetext** operation (which may include a number of calls to **dbmoretext**), the server will send the updated text timestamp value back to DB-Library. **dbtxtsnewval** allows the application to get this new timestamp value. The application can then use **dbtxtsput** to put the new timestamp value into the DBPROCESS row buffer, for future access via **dbtxtimestamp**. This is particularly useful when the application is buffering result rows and does not need the new timestamp immediately.

See also

dbmoretext, dbtxtimestamp, dbtxtsnewval, dbwritetext

dbuse

Description

Use a particular database.

Syntax

RETCODE dbuse(dbproc, dbname)

```
DBPROCESS    *dbproc;
char         *dbname;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

dbname

The name of the database to use.

Return value

SUCCEED or FAIL.

Usage

- This routine issues a Transact-SQL **use** command for the specified database for a particular DBPROCESS. It sets up the command and calls **dbsqlexec** and **dbresults**.
- If the **use** command fails because the requested database has not yet completed a recovery process, **dbuse** will continue to send **use** commands at one second intervals until it either succeeds or encounters some other error.
- The routine uses the *dbproc* provided by the caller. It also uses the command buffer of that dbproc. **dbuse** overwrites any existing commands in the buffer and clears the buffer when it is finished.

See also

dbchange, dbname

dbvarylen

Description

Determine whether the specified regular result column's data can vary in length.

Syntax

```
DBBOOL dbvarylen(dbproc, column)

DBPROCESS    *dbproc;
int          column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

column
The number of the regular result column of interest. The first column is number 1.

Return value
TRUE or FALSE, indicating whether or not the column's data can vary in length. **dbvarylen** also returns FALSE if the column number is out of range.

Usage
- This routine indicates whether a particular regular (that is, non-compute) result column's data can vary in length. It will return TRUE if the result column is derived from a database column of type varchar, varbinary, text, image, boundary, or sensitivity. It will also return TRUE if the source database column is defined as NULL, meaning that it may contain a null value.
- This routine is useful with programs that handle ad hoc queries, if the program needs to be alerted to the possibility of null or variable length data.
- You can use dbcoltype to determine a column's datatype. See Types on page 408 for a list of datatypes.

See also
dbcollen, dbcolname, dbcoltype, dbdata, dbdatlen, dbnumcols, dbprtype

dbversion

Description
Determine which version of DB-Library is in use.

Syntax
char *dbversion()

Parameters
None.

Return value
A pointer to a character string containing the version of DB-Library in use.

Usage
- **dbversion** returns a pointer to a character string that contains the version number for the DB-Library that is currently in use.

See also
DBTDS

dbwillconvert

Description

Determine whether a specific datatype conversion is available within DB-Library.

Syntax

DBBOOL dbwillconvert(srctype, desttype)

```
int     srctype;
int     desttype;
```

Parameters

srctype

The datatype of the data that is to be converted. This parameter can be any of the server datatypes, as listed in Table 2-30.

desttype

The datatype that the source data is to be converted into. This parameter can be any of the server datatypes, as listed in Table 2-30.

Return value

TRUE if the datatype conversion is supported; FALSE if the conversion is not supported.

Usage

- This routine allows the program to determine whether **dbconvert** is capable of performing a specific datatype conversion. When **dbconvert** is asked to perform a conversion that it doesn't support, it calls a user-supplied error handler (if any), sets a global error number, and returns FAIL.
- **dbconvert** can convert data stored in any of the server datatypes (although, of course, not all conversions are legal):

Table 2-30: Server and DB-Library datatypes

Server type	Program variable type
SYBCHAR	DBCHAR
SYBTEXT	DBCHAR
SYBBINARY	DBBINARY
SYBIMAGE	DBBINARY
SYBINT1	DBTINYINT
SYBINT2	DBSMALLINT
SYBINT4	DBINT
SYBFLT8	DBFLT8
SYBREAL	DBREAL
SYBNUMERIC	DBNUMERIC
SYBDECIMAL	DBDECIMAL
SYBBIT	DBBIT
SYBMONEY	DBMONEY

Server type	Program variable type
SYBMONEY4	DBMONEY4
SYBDATETIME	DBDATETIME
SYBDATETIME4	DBDATETIME4
SYBBOUNDARY	DBCHAR
SYBSENSITIVITY	DBCHAR

- Table 2-8 on page 113 lists the datatype conversions that **dbconvert** and **dbconvert_ps** support. The source datatypes are listed down the leftmost column and the destination datatypes are listed along the top row of the table. (For brevity, the prefix "SYB" has been eliminated from each datatype.) If **dbwillconvert** returns TRUE (T), the conversion is supported; if it returns FALSE (F), the conversion is not supported.
- See the reference pages for dbconvert or dbconvert_ps for more information on datatype conversions.

See also

dbaltbind, dbbind, dbconvert, dbconvert_ps, Types

dbwritepage

Description

Write a page of binary data to the server.

Syntax

RETCODE dbwritepage(dbproc, dbname, pageno, size, buf)

```
DBPROCESS    *dbproc;
char         *dbname;
DBINT        pageno;
DBINT        size;
BYTE         buf[];
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

dbname

The name of the database of interest.

pageno

The number of the database page to be written.

size
The number of bytes to be written to the server. Currently, Adaptive Server database pages are 2048 bytes long.

buf
A pointer to a buffer that holds the data to be written.

Return value

SUCCEED or FAIL.

Usage

- **dbwritepage** writes a page of binary data to the server. This routine is useful primarily for examining and repairing damaged database pages. After calling **dbwritepage**, the DBPROCESS may contain some error or informational messages from the server. These messages may be accessed through a user-supplied message handler.

  Warning! Use this routine only if you are absolutely sure you know what you are doing!

See also

dbmsghandle, dbreadpage

dbwritetext

Description

Send a text or image value to the server.

Syntax

```
RETCODE dbwritetext(dbproc, objname, textptr,
            textptrlen, timestamp, log,
            size, text)

DBPROCESS     *dbproc;
char          *objname;
DBBINARY      *textptr;
DBTINYINT      textptrlen;
DBBINARY      *timestamp;
DBBOOL         log;
DBINT          size;
BYTE          *text;
```

Parameters

dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and server.

objname
The database table and column name of interest. The table and the column should be separated by a period (.

textptr
A pointer to the text pointer of the text or image value to be modified. This can be obtained by calling **dbtxptr**. The text pointer must be a valid one, as described on the dbtxptr reference page.

textptrlen
This parameter is included for future compatibility. For now, its value must be the defined constant DBTXPLEN.

timestamp
A pointer to the text timestamp of the text or image value to be modified. This can be obtained via **dbtxtimestamp** or **dbtxtsnewval**. This value changes whenever the text or image value itself is changed. This parameter is optional and may be passed as NULL.

log
A boolean value specifying whether this **dbwritetext** operation should be recorded in the transaction log.

size
The total size, in bytes, of the text or image value to be written. Since **dbwritetext** uses this parameter as its only guide to determining how many bytes to send, *size* must not exceed the actual size of the value.

text
The address of a buffer containing the text or image value to be written. If this pointer is NULL, the application must subsequently call dbmoretext one or more times, until all *size* bytes of data have been sent to the server.

Return value

SUCCEED or FAIL.

A common cause for failure is an invalid *timestamp* parameter. This will occur if, between the time the application retrieves the text column and the time the application calls **dbwritetext** to update it, a second application intervenes with its own update.

Usage

- **dbwritetext** updates SYBTEXT and SYBIMAGE values. It allows the application to send long values to the server without having to copy them into a Transact-SQL **update** statement. In addition, **dbwritetext** gives applications access to the text timestamp mechanism, which can be used to ensure that two competing application users do not inadvertently wipe out each other's modifications to the same value in the database.
- The *timestamp* parameter is optional.

If the *timestamp* parameter is supplied, **dbwritetext** will succeed only if the value of the *timestamp* parameter matches the text column's timestamp in the database. If a match occurs, **dbwritetext** updates the text column and at the same time updates the column's timestamp with the current time. This has the effect of governing updates by competing applications—an application's **dbwritetext** call will fail if a second application updated the text column between the time the first application retrieved the column and the time it made its **dbwritetext** call.

If the *timestamp* parameter is not supplied, **dbwritetext** updates the text column regardless of the value of the column's timestamp.

- The value to use as the *timestamp* parameter is placed in an application's DBPROCESS when the application performs a **select** on a text or image value. It can be retrieved from the DBPROCESS via **dbtxtimestamp**.

  In addition, after each successful **dbwritetext** operation, which may include a number of calls to **dbmoretext**, Adaptive Server sends a new text timestamp value back to DB-Library. **dbtxtsnewval** provides a way for an application to retrieve this new value.

- **dbwritetext** is similar in function to the Transact-SQL **writetext** command. It is usually more efficient to call **dbwritetext** than to send a **writetext** command through the command buffer. In addition, **dbwritetext** can handle columns up to 2GB in length, while **writetext** data is limited to approximately 120K. For more information on **writetext**, see the *Adaptive Server Enterprise Reference Manual*.

- **dbwritetext** can be invoked with or without logging, according to the value of the *log* parameter.

  While logging aids media recovery, logging text data quickly increases the size of the transaction log. If you are logging **dbwritetext** operations, make sure that the transaction log resides on a separate database device. See the *System Administration Guide*, the **create database** reference page in the *Adaptive Server Enterprise Reference Manual*, and the **sp_logdevice** reference page in the *Adaptive Server Enterprise Reference Manual* for details.

  To use **dbwritetext** with logging turned off, the database option **select into/bulkcopy** must be set to "true". The following SQL command will do this:

  ```
  sp_dboption 'mydb', 'select into/bulkcopy', 'true'
  ```

  See the *Adaptive Server Enterprise Reference Manual* for further details on **sp_dboption**.

- The application can send a text or image value to the server all at once or a chunk at a time. **dbwritetext** by itself handles sending an entire text or image value. The use of **dbwritetext** with **dbmoretext** allows the application to send a large text or image value to the server in the form of a number of smaller chunks. This is particularly useful with operating systems unable to allocate extremely long data buffers.
- Sending an entire text or image value requires a non-NULL *text* parameter. Then, **dbwritetext** will execute the data transfer from start to finish, including any necessary calls to **dbsqlok** and **dbresults**. Here is a code fragment that illustrates this use of **dbwritetext**:

```
LOGINREC     *login;
 DBPROCESS    *q_dbproc;
 DBPROCESS    *u_dbproc;
 DBCHAR       abstract_var[512];

 /* Initialize DB-Library. */
 if (dbinit() == FAIL)
      exit(ERREXIT);
 /*
 ** Open separate DBPROCESSes for querying and updating.
 ** This is not strictly necessary in this example,
 ** which retrieves only one row.  However, this
 ** approach becomes essential when performing updates
 ** on multiple rows of retrieved data.
 */
 login = dblogin();
 q_dbproc = dbopen(login, NULL);
 u_dbproc = dbopen(login, NULL);

 /* The database column "abstract" is a text column.
  ** Retrieve the value of one of its rows.
 */
 dbcmd(q_dbproc, "select abstract from articles where \
      article_id = 10");
 dbsqlexec(q_dbproc);
 dbresults(q_dbproc);
 dbbind(q_dbproc, 1, STRINGBIND, (DBINT) 0,
      abstract_var);

 /*
 ** For simplicity, we'll assume that just one row is
 ** returned.
 */
 dbnextrow(q_dbproc);
```

```
 /* Here we can change the value of "abstract_var" */
 /* For instance ... */
 strcpy(abstract_var, "A brand new value.");

 /* Update the text column */
 dbwritetext (u_dbproc, "articles.abstract",
      dbtxptr(q_dbproc, 1), DBTXPLEN,
      dbtxtimestamp(q_dbproc, 1), TRUE,
      (DBINT)strlen(abstract_var), abstract_var);
 /* We're all done */
 dbexit();
```

- To send chunks of text or image, rather than the whole value at once, set the *text* parameter to NULL. Then, **dbwritetext** will return control to the application immediately after notifying the server that a text transfer is about to begin. The actual text will be sent to the server with **dbmoretext**, which can be called multiple times, once for each chunk. Here is a code fragment that illustrates the use of **dbwritetext** with **dbmoretext**:

```
LOGINREC         *login;
 DBPROCESS    *q_dbproc;
 DBPROCESS    *u_dbproc;
 DBCHAR       part1[512];
 static DBCHAR   part2[512] = " This adds another \
      sentence to the text.";

if (dbinit() == FAIL)
      exit(ERREXIT);

login = dblogin();
 q_dbproc = dbopen(login, NULL);
 u_dbproc = dbopen(login, NULL);

dbcmd(q_dbproc, "select abstract from articles where \
      article_id = 10");
 dbsqlexec(q_dbproc);
 dbresults(q_dbproc);
 dbbind(q_dbproc, 1, STRINGBIND, (DBINT) 0, part1);

/*
 ** For simplicity, we'll assume that just one row is
 ** returned.
 */
 dbnextrow(q_dbproc);

/*
 ** Here we can change the value of part of the text
 ** column. In this example, we will merely add a
```

```
 ** sentence to the end of the existing text.
 */

/* Update the text column */
 dbwritetext (u_dbproc, "articles.abstract",
      dbtxptr(q_dbproc, 1), DBTXPLEN,
      dbtxtimestamp(q_dbproc, 1), TRUE,
      (DBINT)(strlen(part1) + strlen(part2)), NULL);

dbsqlok(u_dbproc);
 dbresults(u_dbproc);

/* Send the update value in chunks */
 dbmoretext(u_dbproc, (DBINT)strlen(part1), part1);
 dbmoretext(u_dbproc, (DBINT)strlen(part2), part2);

dbsqlok(u_dbproc);
 dbresults(u_dbproc);
dbexit();
```

Note the required calls to **dbsqlok** and **dbresults** between the call to **dbwritetext** and the first call to **dbmoretext**, and after the final call to **dbmoretext**.

- When **dbwritetext** is used with **dbmoretext**, it locks the specified database text column. The lock is not released until the final **dbmoretext** has sent its data. This ensures that a second application does not read or update the text column in the midst of the first application's update.
- You cannot use **dbwritetext** on text or image columns in views.
- The DB-Library/C option DBTEXTSIZE affects the value of the server *@@textsize* global variable, which restricts the size of text or image values that Adaptive Server returns. *@@textsize* has a default value of 32,768 bytes. An application that retrieves text or image values larger than 32,768 bytes will need to call **dbsetopt** to make *@@textsize* larger.
- The DB-Library/C option DBTEXTLIMIT limits the size of text or image values that DB-Library/C will read.

See also dbmoretext, dbtxptr, dbtxtimestamp, dbwritetext, dbtxtsput

dbxlate

Description Translate a character string from one character set to another.

Syntax

```
int dbxlate(dbproc, src, srclen, dest, destlen, xlt,
        srcbytes_used, srcend, status)

DBPROCESS    dbproc;
char         *src;
int          srclen;
char         *dest;
int          destlen;
DBXLATE      *xlt;
int          *srcbytes_used;
DBBOOL       srcend;
int          *status;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/server process. It contains all the information that DB-Library uses to manage communications and data between the front end and the server.

src

A pointer to the string to be translated.

srclen

The length, in bytes, of *src*. If *srclen* is -1, *src* is assumed to be null-terminated.

dest

A pointer to the buffer to contain the translated string, including a null terminator.

destlen

The size, in bytes, of the buffer to contain the translated string. If *destlen* is -1, *dest* is assumed to be large enough to hold the translated string and its null terminator.

xlt

A pointer to a translation structure used to translate character strings from one character set to another. The translation structure is allocated via **dbload_xlate**.

srcbytes_used

The number of bytes actually translated. If the fully translated string would overflow *dest*, **dbxlate** translates only as much of *src* as will fit. If *destlen* is -1, *srcbytes_used* is *srclen*.

srcend

A boolean value indicating whether or not more data is arriving. If *srcend* is TRUE, no more data is arriving. If *srcend* is FALSE, *src* is part of a larger string of data to be translated, and it is not the end of the string.

status

A pointer to a code indicating the status of the translated character string. The following table lists the possible values for *status*:

Table 2-31: Values for status

Value of status	To indicate
DBXLATE_XOF	The translated string overflowed *dest*.
DBXLATE_XOK	The translation succeeded.
DBXLATE_XPAT	The last bytes of *src* are the beginning of a pattern for which there is a translation. These bytes were not translated.

Return value

The number of bytes actually placed in *dest* on success; a negative integer on error.

Usage

- **dbxlate** translates a character string from one character set to another. It is useful when the server character set differs from the display device's character set.
- The following code fragment illustrates the use of **dbxlate**:

```
char        destbuf[128];
 int         srcbytes_used;
 DBXLATE     *xlt_todisp;
 DBXLATE     *xlt_tosrv;

dbload_xlate((DBPROCESS *)NULL, "iso_1",
      "trans.xlt", &xlt_tosrv, &xlt_todisp);
 printf("Original string: \n\t%s\n\n",
      TEST_STRING);
 dbxlate((DBPROCESS *)NULL, TEST_STRING,
      strlen(TEST_STRING), destbuf, -1, xlt_todisp,
      &srcbytes_used);
 printf("Translated to display character set: \
      \n\t%s\n\n", destbuf);
 dbfree_xlate((DBPROCESS *)NULL, xlt_tosrv,
      xlt_todisp);
```

See also

dbload_xlate, dbfree_xlate

Errors

Description

The complete collection of DB-Library errors and error severities.

Syntax

```
#include <sybfront.h>

#include <sybdb.h>

#include <syberror.h>
```

Usage

- This is the complete list of possible DB-Library errors and error severities.
- The error values are listed alphabetically in Table 2-32 on page 386. The second column of this table gives the error severity for each error as a symbolic value. The third column contains the text associated with the error.
- Table 2-33 on page 402 provides a list of all possible error severities, with their numerical equivalents and an explanation of the type of error.
- When an error or informational event occurs, these numbers are passed to the application's current error handler (if any). An application calls **dberrhandle** to install an error handler.
- Error values are defined in the header file *sybdb.h*. Error severity values are defined in the header file *syberror.h*. Your program needs to include *syberror.h* only if it refers to the symbolic error severities.

Errors The following table lists all the DB-Library errors.

Table 2-32: Errors

Error name	Error severity	Error text
SYBEAAMT	EXPROGRAM	User attempted a **dbaltbind** with mismatched column and variable types.
SYBEABMT	EXPROGRAM	User attempted a **dbbind** with mismatched column and variable types.
SYBEABNC	EXPROGRAM	Attempt to bind to a non-existent column.
SYBEABNP	EXPROGRAM	Attempt to bind using NULL pointers.
SYBEABNV	EXPROGRAM	Attempt to bind to a NULL program variable.
SYBEACNV	EXCONVERSION	Attempt to do data-conversion with NULL destination variable.

Error name	Error severity	Error text
SYBEADST	EXCONSISTENCY	International Release: Error in attempting to determine the size of a pair of translation tables.
SYBEAICF	EXCONSISTENCY	International Release: Error in attempting to install custom format.
SYBEALTT	EXCONSISTENCY	International Release: Error in attempting to load a pair of translation tables.
SYBEAOLF	EXRESOURCE	International Release: Error in attempting to open a localization file.
SYBEAPCT	EXCONSISTENCY	International Release: Error in attempting to perform a character set translation.
SYBEAPUT	EXPROGRAM	Attempt to print unknown token.
SYBEARDI	EXRESOURCE	International Release: Error in attempting to read datetime information from a localization file.
SYBEARDL	EXRESOURCE	International Release: Error in attempting to read the *dblib.loc* localization file.
SYBEASEC	EXPROGRAM	Attempt to send an empty command buffer to the server.
SYBEASNL	EXPROGRAM	Attempt to set fields in a null LOGINREC.
SYBEASTL	EXPROGRAM	Synchronous I/O attempted at AST level.
SYBEASUL	EXPROGRAM	Attempt to set unknown LOGINREC field.
SYBEAUTN	EXPROGRAM	Attempt to update the timestamp of a table that has no timestamp column.
SYBEBADPK	EXINFO	Packet size of %1 not supported-size of %2 used instead!
SYBEBBCI	EXINFO	Batch successfully bulk copied to the server.
SYBEBBL	EXPROGRAM	Bad *bindlen* parameter passed to **dbsetnull**.

Error name	Error severity	Error text
SYBEBCBC	EXPROGRAM	**bcp_columns** must be called before **bcp_colfmt** and **bcp_colfmt_ps**.
SYBEBCBNPR	EXPROGRAM	**bcp_bind**: if *varaddr* is NULL, *prefixlen* must be 0 and no terminator should be specified.
SYBEBCBNTYP	EXPROGRAM	**bcp_bind**: if *varaddr* is NULL and *varlen* greater than 0, the table column type must be SYBTEXT or SYBIMAGE and the program variable type must be SYBTEXT, SYBCHAR, SYBIMAGE or SYBBINARY.
SYBEBCBPREF	EXPROGRAM	Illegal prefix length. Legal values are 0, 1, 2 or 4.
SYBEBCFO	EXUSER	**bcp** host files must contain at least one column.
SYBEBCHLEN	EXPROGRAM	*host_collen* should be greater than or equal to -1.
SYBEBCIS	EXCONSISTENCY	Attempt to bulk copy an illegally-sized column value to the server.
SYBEBCIT	EXPROGRAM	It is illegal to use BCP terminators with program variables other than SYBCHAR, SYBBINARY, SYBTEXT, or SYBIMAGE.
SYBEBCITBLEN	EXPROGRAM	**bcp_init**: *tblname* parameter is too long.
SYBEBCITBNM	EXPROGRAM	**bcp_init**: *tblname* parameter cannot be NULL.
SYBEBCMTXT	EXPROGRAM	**bcp_moretext** may be used only when there is at least one text or image column in the server table.
SYBEBCNL	EXNONFATAL	Negative length-prefix found in BCP data-file.
SYBEBCNN	EXUSER	Attempt to bulk copy a NULL value into a Server column which does not accept null values.
SYBEBCNT	EXUSER	Attempt to use Bulk Copy with a non-existent Server table.
SYBEBCOR	EXCONSISTENCY	Attempt to bulk copy an oversized row to the server.

Error name	Error severity	Error text
SYBEBCPB	EXPROGRAM	**bcp_bind, bcp_moretext** and **bcp_sendrow** may not be used after **bcp_init** has been passed a non-NULL input file name.
SYBEBCPCTYP	EXPROGRAM	**bcp_colfmt**: If *table_colnum* is 0, *host_type* cannot be 0.
SYBEBCPI	EXPROGRAM	**bcp_init** must be called before any other **bcp** routines.
SYBEBCPN	EXPROGRAM	**bcp_bind, bcp_collen, bcp_colptr, bcp_moretext** and **bcp_sendrow** may be used only after **bcp_init** has been called with the copy direction set to DB_IN.
SYBEBCPREC	EXNONFATAL	Column %1!: Illegal precision value encountered.
SYBEBCPREF	EXPROGRAM	Illegal prefix length. Legal values are -1, 0, 1, 2 or 4.
SYBEBCRE	EXNONFATAL	I/O error while reading **bcp** datafile.
SYBEBCRO	EXINFO	The BCP hostfile '%1!' contains only %2! rows. It was impossible to read the requested %3! rows.
SYBEBCSA	EXUSER	The BCP hostfile '%1!' contains only %2! rows. Skipping all of these rows is not allowed.
SYBEBCSET	EXCONSISTENCY	Unknown character-set encountered.
SYBEBCSI	EXPROGRAM	Host-file columns may be skipped only when copying into the Server.
SYBEBCSNDROW	EXPROGRAM	**bcp_sendrow** may not be called unless all text data for the previous row has been sent using **bcp_moretext**.
SYBEBCSNTYP	EXPROGRAM	column number %1!: if *varaddr* is NULL and *varlen* greater than 0, the table column type must be SYBTEXT or SYBIMAGE and the program variable type must be SYBTEXT, SYBCHAR, SYBIMAGE or SYBBINARY.
SYBEBCUC	EXRESOURCE	**bcp**: Unable to close host datafile.
SYBEBCUO	EXRESOURCE	**bcp**: Unable to open host datafile.

Error name	Error severity	Error text
SYBEBCVH	EXPROGRAM	**bcp_exec** may be called only after **bcp_init** has been passed a valid host file.
SYBEBCVLEN	EXPROGRAM	*varlen* should be greater than or equal to -1.
SYBEBCWE	EXNONFATAL	I/O error while writing **bcp** datafile.
SYBEBDIO	EXPROGRAM	Bad bulk copy direction. Must be either IN or OUT.
SYBEBEOF	EXNONFATAL	Unexpected EOF encountered in **bcp** datafile.
SYBEBIHC	EXPROGRAM	Incorrect host-column number found in **bcp** format file.
SYBEBIVI	EXPROGRAM	**bcp_columns, bcp_colfmt** and **bcp_colfmt_ps** may be used only after **bcp_init** has been passed a valid input file.
SYBEBNCR	EXPROGRAM	Attempt to bind user variable to a non-existent compute row.
SYBEBNUM	EXPROGRAM	Bad *numbytes* parameter passed to **dbstrcpy**.
SYBEBPKS	EXPROGRAM	In **DBSETLPACKET**, the packet size parameter must be between 0 and 999999.
SYBEBPREC	EXPROGRAM	Illegal precision specified.
SYBEBPROBADDEF	EXCONSISTENCY	**bcp** protocol error: illegal default column id received.
SYBEBPROCOL	EXCONSISTENCY	**bcp** protocol error: returned column count differs from the actual number of columns received.
SYBEBPRODEF	EXCONSISTENCY	**bcp** protocol error: expected default information and got none.
SYBEBPRODEFID	EXCONSISTENCY	**bcp** protocol error: default column id and actual column id are not same
SYBEBPRODEFTYP	EXCONSISTENCY	**bcp** protocol error: default value datatype differs from column datatype.

Error name	Error severity	Error text
SYBEBPROEXTDEF	EXCONSISTENCY	**bcp** protocol error: more than one row of default information received.
SYBEBPROEXTRES	EXCONSISTENCY	**bcp** protocol error: unexpected set of results received.
SYBEBPRONODEF	EXCONSISTENCY	**bcp** protocol error: default value received for column that does not have default.
SYBEBPRONUMDEF	EXCONSISTENCY	**bcp** protocol error: expected number of defaults differs from the actual number of defaults received.
SYBEBRFF	EXRESOURCE	I/O error while reading **bcp** format file.
SYBEBSCALE	EXPROGRAM	Illegal scale specified.
SYBEBTMT	EXPROGRAM	Attempt to send too much text data via the **bcp_moretext** call.
SYBEBTOK	EXCOMM	Bad token from the server: Datastream processing out of sync.
SYBEBTYP	EXPROGRAM	Unknown bind type passed to DB-Library function.
SYBEBTYPSRV	EXPROGRAM	Datatype is not supported by the server.
SYBEBUCE	EXRESOURCE	**bcp**: Unable to close error file.
SYBEBUCF	EXPROGRAM	**bcp**: Unable to close format file.
SYBEBUDF	EXPROGRAM	**bcp**: Unrecognized datatype found in format file.
SYBEBUFF	EXPROGRAM	**bcp**: Unable to create format file.
SYBEBUFL	EXCONSISTENCY	DB-Library internal error-send buffer length corrupted.
SYBEBUOE	EXRESOURCE	**bcp**: Unable to open error file.
SYBEBUOF	EXPROGRAM	**bcp**: Unable to open format file.
SYBEBWEF	EXNONFATAL	I/O error while writing **bcp** error file.
SYBEBWFF	EXRESOURCE	I/O error while writing **bcp** format file.
SYBECAP	EXCOMM	DB-Library capabilities not accepted by the Server.
SYBECAPTYP	EXCOMM	Unexpected capability type in CAPABILITY datastream.

Error name	Error severity	Error text
SYBECDNS	EXCONSISTENCY	Datastream indicates that a compute column is derived from a non-existent select list member.
SYBECDOMAIN	EXCONVERSION	Source field value is not within the domain of legal values.
SYBECINTERNAL	EXCONVERSION	Internal Conversion error.
SYBECLOS	EXCOMM	Error in closing network connection.
SYBECLPR	EXCONVERSION	Data conversion resulted in loss of precision.
SYBECNOR	EXPROGRAM	Column number out of range.
SYBECNOV	EXCONVERSION	Attempt to set variable to NULL resulted in overflow.
SYBECOFL	EXCONVERSION	Data conversion resulted in overflow.
SYBECONN	EXCOMM	Unable to connect: Adaptive Server is unavailable or does not exist.
SYBECRNC	EXPROGRAM	The current row is not a result of compute clause %1!, so it is illegal to attempt to extract that data from this row.
SYBECRSAGR	EXPROGRAM	Aggregate functions are not allowed in a cursor statement.
SYBECRSBROL	EXPROGRAM	Backward scrolling cannot be used in a forward scrolling cursor.
SYBECRSBSKEY	EXPROGRAM	Keyset cannot be scrolled backward in mixed cursors with a previous fetch type.
SYBECRSBUFR	EXPROGRAM	Row buffering should not be turned on when using cursor APIs.
SYBECRSDIS	EXPROGRAM	Cursor statement contains one of the disallowed phrases **compute**, **union**, **for browse**, or **select into**.
SYBECRSFLAST	EXPROGRAM	Fetch type LAST requires fully keyset driven cursors.
SYBECRSFRAND	EXPROGRAM	Fetch types RANDOM and RELATIVE can only be used within the keyset of keyset driven cursors.

Error name	Error severity	Error text
SYBECRSFROWN	EXPROGRAM	Row number to be fetched is outside valid range.
SYBECRSFTYPE	EXRESOURCE	Unknown fetch type.
SYBECRSINV	EXPROGRAM	Invalid cursor statement.
SYBECRSINVALID	EXRESOURCE	The cursor handle is invalid.
SYBECRSMROWS	EXRESOURCE	Multiple rows are returned, only one is expected while retrieving **dbname**.
SYBECRSNOBIND	EXPROGRAM	Cursor bind must be called prior to **dbcursor** invocation.
SYBECRSNOCOUNT	EXPROGRAM	The DBNOCOUNT option should not be turned on when doing updates or deletes with **dbcursor**.
SYBECRSNOFREE	EXPROGRAM	The DBNOAUTOFREE option should not be turned on when using cursor APIs.
SYBECRSNOIND	EXPROGRAM	One of the tables involved in the cursor statement does not have a unique index.
SYBECRSNOKEYS	EXRESOURCE	The entire keyset must be defined for KEYSET type cursors.
SYBECRSNOLEN	EXRESOURCE	No unique index found.
SYBECRSNOPTCC	EXRESOURCE	No OPTCC was found.
SYBECRSNORDER	EXRESOURCE	The order of clauses must be **from**, **where**, and **order by**.
SYBECRSNORES	EXPROGRAM	Cursor statement generated no results.
SYBECRSNROWS	EXRESOURCE	No rows returned, at least one is expected.
SYBECRSNOTABLE	EXRESOURCE	Table name is NULL.
SYBECRSNOUPD	EXPROGRAM	**Update** or **delete** operation did not affect any rows.
SYBECRSNOWHERE	EXPROGRAM	A **WHERE** clause is not allowed in a cursor **update** or **insert**.
SYBECRSNUNIQUE	EXRESOURCE	No unique keys associated with this view.
SYBECRSORD	EXPROGRAM	Only fully keyset driven cursors can have **order by**, **group by**, or **having** phrases.

Error name	Error severity	Error text
SYBECRSRO	EXPROGRAM	Data locking or modifications cannot be made in a READONLY cursor.
SYBECRSSET	EXPROGRAM	A SET clause is required for a cursor **update** or **insert**.
SYBECRSTAB	EXPROGRAM	Table name must be determined in operations involving data locking or modifications.
SYBECRSVAR	EXRESOURCE	There is no valid address associated with this bind.
SYBECRSVIEW	EXPROGRAM	A view cannot be joined with another table or a view in a cursor statement.
SYBECRSVIIND	EXPROGRAM	The view used in the cursor statement does not include all the unique index columns of the underlying tables.
SYBECRSUPDNB	EXPROGRAM	**Update** or **insert** operations cannot use bind variables when binding type is NOBIND.
SYBECRSUPDTAB	EXPROGRAM	**Update** or **insert** operations using bind variables require single table cursors.
SYBECSYN	EXCONVERSION	Attempt to convert data stopped by syntax error in source field.
SYBECUFL	EXCONVERSION	Data conversion resulted in underflow.
SYBEDBPS	EXRESOURCE	Maximum number of DBPROCESSes already allocated.
SYBEDDNE	EXINFO	DBPROCESS is dead or not enabled.
SYBEDIVZ	EXUSER	Attempt to divide by $0.00 in function %1!.
SYBEDNTI	EXPROGRAM	Attempt to use **dbtxtsput** to put a new text timestamp into a column whose datatype is neither SYBTEXT nor SYBIMAGE.
SYBEDPOR	EXPROGRAM	Out-of-range *datepart* constant.
SYBEDVOR	EXPROGRAM	Day values must be between 1 and 7.

Error name	Error severity	Error text
SYBEECAN	EXINFO	Attempted to cancel unrequested event notification.
SYBEEINI	EXINFO	Must call **dbreginit** before **dbregexec**.
SYBEETD	EXPROGRAM	Failure to send the expected amount of TEXT or IMAGE data via **dbmoretext**.
SYBEEUNR	EXCOMM	Unsolicited event notification received.
SYBEEVOP	EXINFO	Called **dbregwatch** with a bad options parameter.
SYBEEVST	EXINFO	Must initiate a transaction before calling **dbregparam**.
SYBEFCON	EXCOMM	Adaptive Server connection failed.
SYBEFRES	EXFATAL	Challenge-Response function failed.
SYBEFSHD	EXRESOURCE	Error in attempting to find the Sybase home directory.
SYBEFUNC	EXPROGRAM	Functionality not supported at the specified version level.
SYBEICN	EXPROGRAM	Invalid *computeid* or compute column number.
SYBEIDCL	EXCONSISTENCY	Illegal datetime column length returned by Adaptive Server. Legal datetime lengths are 4 and 8 bytes.
SYBEIDECCL	EXCONSISTENCY	Invalid decimal column length returned by the server.
SYBEIFCL	EXCONSISTENCY	Illegal floating-point column length returned by Adaptive Server. Legal floating-point lengths are 4 and 8 bytes.
SYBEIFNB	EXPROGRAM	Illegal field number passed to **bcp_control**.
SYBEIICL	EXCONSISTENCY	Illegal integer column length returned by Adaptive Server. Legal integer lengths are 1, 2, and 4 bytes.
SYBEIMCL	EXCONSISTENCY	Illegal money column length returned by Adaptive Server. Legal money lengths are 4 and 8 bytes.

Error name	Error severity	Error text
SYBEINLN	EXUSER	Interface file: unexpected end-of-line.
SYBEINTF	EXUSER	Server name not found in interface file.
SYBEINUMCL	EXCONSISTENCY	Invalid numeric column length returned by the server.
SYBEIPV	EXINFO	%1! is an illegal value for the %2! parameter of %3!.
SYBEISOI	EXCONSISTENCY	International Release: Invalid sort-order information found.
SYBEISRVPREC	EXCONSISTENCY	Illegal precision value returned by the server.
SYBEISRVSCL	EXCONSISTENCY	Illegal scale value returned by the server.
SYBEITIM	EXPROGRAM	Illegal timeout value specified.
SYBEIVERS	EXPROGRAM	Illegal version level specified.
SYBEKBCI	EXINFO	1000 rows sent to the server.
SYBEKBCO	EXINFO	1000 rows successfully bulk copied to host file.
SYBEMEM	EXRESOURCE	Unable to allocate sufficient memory.
SYBEMOV	EXUSER	Money arithmetic resulted in overflow in function %1!.
SYBEMPLL	EXUSER	Attempt to set maximum number of DBPROCESSes lower than 1.
SYBEMVOR	EXPROGRAM	Month values must be between 1 and 12.
SYBENBUF	EXINFO	Called **dbsendpassthru** with a NULL *buf* parameter.
SYBENBVP	EXPROGRAM	Cannot pass **dbsetnull** a NULL *bindval* pointer.
SYBENDC	EXPROGRAM	Cannot have negative component in date in numeric form.
SYBENDTP	EXPROGRAM	Called **dbdatecrack** with NULL datetime parameter.
SYBENEG	EXCOMM	Negotiated login attempt failed.
SYBENHAN	EXINFO	Called **dbrecvpassthru** with a NULL handle parameter.
SYBENMOB	EXPROGRAM	No such member of **order by** clause.

Error name	Error severity	Error text
SYBENOEV	EXINFO	DBPOLL can not be called when registered procedure notifications have been disabled.
SYBENPRM	EXPROGRAM	NULL parameter not allowed for this **dboption**.
SYBENSIP	EXPROGRAM	Negative starting index passed to **dbstrcpy**.
SYBENTLL	EXUSER	Name too long for LOGINREC field.
SYBENTTN	EXPROGRAM	Attempt to use **dbtxtsput** to put a new text timestamp into a non-existent data row.
SYBENULL	EXINFO	NULL DBPROCESS pointer passed to DB-Library.
SYBENULP	EXPROGRAM	Called %s with a NULL %s parameter.
SYBENXID	EXNONFATAL	The Server did not grant us a distributed-transaction ID.
SYBEONCE	EXPROGRAM	Function can be called only once.
SYBEOOB	EXCOMM	Error in sending out-of-band data to the server.
SYBEOPIN	EXNONFATAL	Could not open interface file.
SYBEOPNA	EXNONFATAL	Option is not available with current server.
SYBEOREN	EXINFO	International Release: Warning: an out-of-range error-number was encountered in *dblib.loc*. The maximum permissible error-number is defined as DBERRCOUNT in *sybdb.h*.
SYBEORPF	EXUSER	Attempt to set remote password would overflow the login record's remote password field.
SYBEPOLL	EXINFO	There is already an active **dbpoll**.
SYBEPRTF	EXINFO	**dbtracestring** may only be called from a *printfunc*.
SYBEPWD	EXUSER	Login incorrect.
SYBERDCN	EXCONVERSION	Requested data conversion does not exist.
SYBERDNR	EXPROGRAM	Attempt to retrieve data from a non-existent row.

Error name	Error severity	Error text
SYBEREAD	EXCOMM	Read from the server failed.
SYBERESP	EXPROGRAM	Response function address passed to **dbresponse** must be non-NULL.
SYBERPCS	EXINFO	Must call **dbrpcinit** before **dbrpcparam** or **dbrpcsend**.
SYBERPIL	EXPROGRAM	It is illegal to pass -1 to **dbrpcparam** for the *datalen* of parameters which are of type SYBCHAR, SYBVARCHAR, SYBBINARY, or SYBVARBINARY.
SYBERPNA	EXNONFATAL	The RPC facility is available only when using a server whose version number is 4.0 or later.
SYBERPND	EXPROGRAM	Attempt to initiate a new Adaptive Server operation with results pending.
SYBERPNULL	EXPROGRAM	*value* parameter for **dbrpcparam** can be NULL, only if the *datalen* parameter is 0.
SYBERPTXTIM	EXPROGRAM	RPC parameters cannot be of type text or image.
SYBERPUL	EXPROGRAM	When passing a SYBINTN, SYBDATETIMN, SYBMONEYN, or SYBFLTN parameter via **dbrpcparam**, it is necessary to specify the parameter's maximum or actual length so that DB-Library can recognize it as a SYINT1, SYBINT2, SYBINT4, SYBMONEY, SYBMONEY4, and so on.
SYBERTCC	EXPROGRAM	**dbreadtext** may not be used to receive the results of a query that contains a COMPUTE clause.
SYBERTSC	EXPROGRAM	**dbreadtext** may be used only to receive the results of a query that contains a single result column.
SYBERXID	EXNONFATAL	The Server did not recognize our distributed-transaction ID.

Error name	Error severity	Error text
SYBESECURE	EXPROGRAM	Secure SQL Server function not supported in this version.
SYBESEFA	EXPROGRAM	DBSETNOTIFS cannot be called if connections are present.
SYBESEOF	EXCOMM	Unexpected EOF from the server.
SYBESFOV	EXPROGRAM	International Release: **dbsafestr** overflowed its destination buffer.
SYBESMSG	EXSERVER	General Adaptive Server error: Check messages from the server.
SYBESOCK	EXCOMM	Unable to open socket.
SYBESPID	EXPROGRAM	Called **dbspid** with a NULL **dbproc**.
SYBESYNC	EXCOMM	Read attempted while out of synchronization with Adaptive Server.
SYBETEXS	EXINFO	Called **dbmoretext** with a bad size parameter.
SYBETIME	EXTIME	Adaptive Server connection timed out.
SYBETMCF	EXPROGRAM	Attempt to install too many custom formats via **dbfmtinstall**.
SYBETMTD	EXPROGRAM	Attempt to send too much TEXT data via the **dbmoretext** call.
SYBETPAR	EXPROGRAM	No SYBTEXT or SYBIMAGE parameters were defined.
SYBETPTN	EXUSER	Syntax error: only two periods are permitted in table names.
SYBETRAC	EXINFO	Attempted to turn off a trace flag that was not on.
SYBETRAN	EXINFO	DBPROCESS is being used for another transaction.
SYBETRAS	EXINFO	DB-Library internal error - trace structure not found.
SYBETRSN	EXINFO	Bad *numbytes* parameter passed to **dbtracestring**.
SYBETSIT	EXINFO	Attempt to call **dbtsput** with an invalid timestamp.

Error name	Error severity	Error text
SYBETTS	EXUSER	The table which bulk copy is attempting to copy to a host file is shorter than the number of rows which bulk copy was instructed to skip.
SYBETYPE	EXINFO	Invalid argument type given to Hyper/DB-Library.
SYBEUCPT	EXUSER	Unrecognized custom-format parameter-type encountered in **dbstrbuild**.
SYBEUCRR	EXCONSISTENCY	Internal software error: Unknown connection result reported by **dbpasswd**.
SYBEUDTY	EXCONSISTENCY	Unknown datatype encountered.
SYBEUFDS	EXUSER	Unrecognized format encountered in **dbstrbuild**.
SYBEUFDT	EXCONSISTENCY	Unknown fixed-length datatype encountered.
SYBEUHST	EXUSER	Unknown host machine name.
SYBEUMSG	EXCOMM	Unknown message-id in MSG datastream.
SYBEUNAM	EXFATAL	Unable to get current user name from operating system.
SYBEUNOP	EXNONFATAL	Unknown option passed to **dbsetopt**.
SYBEUNT	EXUSER	Unknown network type found in interface file.
SYBEURCI	EXRESOURCE	International Release: Unable to read copyright information from the DB-Library localization file.
SYBEUREI	EXRESOURCE	International Release: Unable to read error information from the DB-Library localization file.
SYBEUREM	EXRESOURCE	International Release: Unable to read error mnemonic from the DB-Library localization file.
SYBEURES	EXRESOURCE	International Release: Unable to read error string from the DB-Library localization file.

Error name	Error severity	Error text
SYBEURMI	EXRESOURCE	International Release: Unable to read money-format information from the DB-Library localization file.
SYBEUSCT	EXCOMM	Unable to set communications timer.
SYBEUTDS	EXCOMM	Unrecognized TDS version received from the server.
SYBEUVBF	EXPROGRAM	Attempt to read an unknown version of **bcp** format file.
SYBEUVDT	EXCONSISTENCY	Unknown variable-length datatype encountered.
SYBEVDPT	EXUSER	For bulk copy, all variable-length data must have either a length-prefix or a terminator specified.
SYBEWAID	EXCONSISTENCY	DB-Library internal error: ALTFMT following ALTNAME has wrong id.
SYBEWRIT	EXCOMM	Write to the server failed.
SYBEXOCI	EXNONFATAL	International Release: A character-set translation overflowed its destination buffer while using **bcp** to copy data from a host-file to the server.
SYBEXTDN	EXPROGRAM	Warning: the *xlt_todisp* parameter to **dbfree_xlate** was NULL. The space associated with the **xlt_tosrv** parameter has been freed.
SYBEXTN	EXPROGRAM	The *xlt_tosrv* and *xlt_todisp* parameters to **dbfree_xlate** were NULL.
SYBEXTSN	EXPROGRAM	Warning: the *xlt_tosrv* parameter to **dbfree_xlate** was NULL. The space associated with the *xlt_todisp* parameter has been freed.
SYBEZTXT	EXINFO	Attempt to send zero length TEXT or IMAGE to dataserver via **dbwritetext**.
UNUSED	EXINFO	This error number is unused.

Error severities

The following table lists the meanings for each symbolic error severity value.

Table 2-33: Error severities

Error severity	Numerical equivalent	Explanation
EXINFO	1	Informational, non-error.
EXUSER	2	User error.
EXNONFATAL	3	Non-fatal error.
EXCONVERSION	4	Error in DB-Library data conversion.
EXSERVER	5	The Server has returned an error flag.
EXTIME	6	We have exceeded our timeout period while waiting for a response from the Server—the DBPROCESS is still alive.
EXPROGRAM	7	Coding error in user program.
EXRESOURCE	8	Running out of resources—the DBPROCESS may be dead.
EXCOMM	9	Failure in communication with Server—the DBPROCESS is dead.
EXFATAL	10	Fatal error—the DBPROCESS is dead.
EXCONSISTENCY	11	Internal software error—notify Sybase Technical Support.

See also

DBDEAD, dberrhandle

Options

Description

The complete list of DB-Library options.

Syntax

```
#include <sybfront.h>

 #include <sybdb.h>
```

Usage

- **dbsetopt** and **dbclropt** use the following constants, defined in *sybdb.h*, for setting and clearing options. All options are off by default. These options are available:
- DBARITHABORT —If this option is set, the server will abort a query when an arithmetic exception occurs during its execution.

- DBARITHIGNORE —If this option is set, the server will substitute null values for selected or updated values when an arithmetic exception occurs during query execution. The Adaptive Server will not return a warning message. If neither DBARITHABORT nor DBARITHIGNORE is set, Adaptive Server will substitute null values and print a warning message after the query has been executed.
- DBAUTH —This option sets system administration authorization levels. Possible levels are: “sa”, “sso”, “oper”, and “dbcc_edit”. For information on these levels, see the *Adaptive Server Enterprise Reference Manual*.
- DBBUFFER —This option allows the application to buffer result rows, so that it can access them non-sequentially via the **dbgetrow** function. This option is handled locally by DB-Library and is not a server option. When the option is set, you supply a parameter that is the number of rows you want buffered. If you use 0 as the number of rows to buffer, the buffer will be set to a default size (currently 1000 rows).

  When an application calls **dbclropt** to clear the DBBUFFER option, DB-Library frees the memory associated with the row buffer.
- DBCHAINXACTS — This option is used to select chained or unchained transaction behavior.

  Chained behavior means that each SQL statement that modifies or retrieves data implicitly begins a multi-statement transaction. Any **delete**, **insert**, **open**, **fetch**, **select**, or **update** statement implicitly begins a transaction. An explicit **commit** or **rollback** statement is required to end the transaction. Chained mode provides compatibility with ANSI SQL.

  Unchained behavior means that each SQL statement that modifies or retrieves data is implicitly a distinct transaction. Explicit **begin transaction** and **commit** or **rollback** statements are required to define a multi-statement transaction.

  This option is off (indicating unchained behavior) by default. Applications that operate in chained mode should turn on the option right after a connection has been opened, since this option affects the behavior of all queries.
- DBDATEFIRST — Sets the first weekday to a number from 1 to 7. The us_english default is 1 (Sunday).
- DBDATEFORMAT — Sets the order of the date parts month/day/year for entering datetime or smalldatetime data. Valid arguments are “mdy”, “dmy”, “ymd”, “ydm”, “myd”, or “dym”. The us_english default is “mdy”.

Row buffering provides a way to keep a specified number of server result rows in program memory. Without row buffering, the result row generated by each new **dbnextrow** call overwrites the contents of the previous result row. Row buffering is therefore useful for programs that need to look at result rows in a non-sequential manner. It does, however, carry a memory and performance penalty because each row in the buffer must be allocated and freed individually. Therefore, use it only if you need to. Specifically, the application should only turn the DBBUFFER option on if it calls **dbgetrow** or **dbsetrow**. Note that row buffering has nothing to do with network buffering and is a completely independent issue. (See the dbgetrow**,** dbnextrow, and dbclrbuf reference pages for more information about row buffering.

- DBFIPSFLAG —Setting this option causes the server to flag non-standard SQL commands. This option is off by default.
- DBISOLATION —This option is used to specify the transaction isolation level. Possible levels are 1 and 3. The default level is 1. Setting the level to 3 causes all pages of tables specified in a **select** query inside a transaction to be locked for the duration of the transaction.
- DBNATLANG —This is a DB-Library Internationalization option. Associate the specified DBPROCESS (or all open DBPROCESSes, if a DBPROCESS is not specified) with a national language. If the national language is not set for a particular DBPROCESS, U.S. English is used by default.

  In programs that allow application users to make ad hoc queries, the user may override DBNATLANG with the Transact-SQL **set language** command.

  Note All DBPROCESSes opened using a particular LOGINREC will also use that LOGINREC's associated national language. Use the **DBSETLNATLANG** macro to associate a national language with a LOGINREC.

- DBNOAUTOFREE —This option causes the command buffer to be cleared only by an explicit call to **dbfreebuf**. When DBNOAUTOFREE is not set, after a call to **dbsqlexec** or **dbsqlsend** the first call to either **dbcmd** or **dbfcmd** automatically clears the command buffer before the new text is entered.

- DBNOCOUNT —This option causes the server to stop sending back information about the number of rows affected by each SQL statement. The application can otherwise obtain this information by calling DBCOUNT.
- DBNOEXEC —If this option is set, the server will process the query through the compile step but the query will not be executed. This can be used in conjunction with DBSHOWPLAN.
- DBOFFSET —This option indicates that the server should return offsets for certain constructs in the query. DBOFFSET takes a parameter that specifies the particular construct. The valid parameters for this option are “select”, “from”, “table”, “order”, “compute”, “statement”, “procedure”, “execute”, or “param”. (Note that “param” refers to parameters of stored procedures.) Calls to routines such as **dbsetopt** can specify these option parameters in either lowercase or uppercase. Offsets are returned only if the batch contains no syntax errors.
- DBPARSEONLY —If this option is set, the server only checks the syntax of the query and returns error messages to the host. Offsets are returned if the DBOFFSET option is set and there are no errors.
- DBPRCOLSEP — Specify the column separator character(s). Query results rows formatted via **dbprhead, dbprrow, dbsprhead, dbsprline,** and **dbspr1row** will have columns separated by the specified string. The default separator is an ASCII 0x20 (space). The third parameter, a string, is not necessarily null-terminated. The length of the string used is given as the fourth parameter in the call to **dbsetopt**. To revert to using the default separator, specify a length of -1. In this case, the third parameter is ignored.
- DBPRLINELEN —Specify the maximum number of characters to be placed on one line. This value is used by **dbprhead, dbprrow, dbsprhead, dbsprline**, and **dbspr1row**. The default line length is 80 characters.
- DBPRLINESEP —Specify the row separator character to be used by **dbprhead, dbprrow, dbsprhead, dbsprline**, and **dbspr1row**. The default separator is a newline (ASCII 0x0D or 0x0A, depending on the host system). The third parameter, a string, is not necessarily null-terminated. The length of the string is given as the fourth parameter in the call to **dbsetopt**. To revert to the default terminator, specify a length of -1; in this case, the third parameter is ignored.

- DBPRPADSpecify the pad character used when printing results via **dbprhead, dbprrow, dbsprhead, dbsprline,** and **dbspr1row**. To activate padding, specify DBPADON as the fourth parameter in the **dbsetopt** call. The pad character may be specified as the third parameter in the **dbsetopt** call. If the character is not specified, the ASCII character 0x20 (space) is used. To turn off padding, call **dbsetopt** with DBPADOFF as the fourth parameter; the third parameter is ignored when turning padding off.
- DBROWCOUNT —If this option is set to a value greater than 0, the server limits the number of regular rows returned for **select** statements and the number of table rows affected by **update** or **delete** statements. This option does not limit the number of compute rows returned by a **select** statement.

  DBROWCOUNT works somewhat differently from most options. It is always set on, never off. Setting DBROWCOUNT to 0 sets it back to the default—that is, to return all the rows generated by a **select** statement. Therefore, the way to turn DBROWCOUNT *off* is to set it *on* with a count of 0.
- DBSHOWPLAN —If this option is set, the server will generate a description of the processing plan after compilation and continue executing the query.
- DBSTAT —This option determines when performance statistics (CPU time, elapsed time, I/O, and so on) will be returned to the host after each query. DBSTAT takes one of two parameters: “io”, for statistics about Adaptive Server internal I/O; and “time”, for information about Adaptive Server’s parsing, compilation, and execution times. These statistics are received by DB-Library in the form of informational messages, and application programs can access them through the user-supplied message handler.
- DBSTORPROCID —If this option is set, the server will send the stored procedure id to the host before sending rows generated by the stored procedure.

- DBTEXTLIMIT —This option causes DB-Library to limit the size of returned text or image values. When setting this option, you supply a parameter that is the length, in bytes, of the longest text or image value that your program can handle. DB-Library will read but ignore any part of a text or image value that goes over this limit. DB-Library's default behavior is to read and return all the data sent by the server. To restore this default behavior, set DBTEXTLIMIT to a value less than 1. In the case of huge text values, it may take some time for the entire text value to be returned over the network. To keep the server from sending this extra text in the first place, use the DBTEXTSIZE option instead.
- DBTEXTSIZE —This option changes the value of the server global variable *@@textsize*, which limits the size of text or image values that the server returns. When setting this option, you supply a parameter that is the length, in bytes, of the longest text or image value that the server should return. *@@textsize* has a default value of 32,768 bytes. Note that, in programs that allow application users to make ad hoc queries, the user may override this option with the Transact-SQL **set textsize** command. To set a text limit that the user cannot override, use the DBTEXTLIMIT option instead.
- DBBUFFER, DBNOAUTOFREE, and DBTEXTLIMIT are DB-Library options. That is, they affect DB-Library but are not sent to the server. The other options are Adaptive Server options—they are sent to the server. Adaptive Server options can also be set via Transact-SQL commands.
- As mentioned in the preceding descriptions, certain options take parameters:

Table 2-34: Parameter values for options

Option	Possible parameter values
DBTEXTSIZE	"0" to "2,147,483,647"
DBOFFSET	"select", "from", "table", "order", "compute", "statement", "procedure", "execute", or "param"
DBSTAT	"io" or "time"
DBROWCOUNT	"0" to "2,147,483,647"
DBBUFFER	"0" to either "32,767" or "2,147,483,647", depending on whether your *int* datatype is 2 or 4 bytes long
DBTEXTLIMIT	"0" to "2,147,483,647"

dbsetopt requires that an option parameter be specified when setting any option on the preceding list. **dbclropt** and **dbisopt** require that the parameter be specified only for DBOFFSET and DBSTAT. This is because DBOFFSET and DBSTAT are the only options that can have multiple settings at a time, and thus they require further definition before being cleared or checked.

Note that parameters specified in calls to **dbsetopt**, **dbclropt**, and **dbisopt** are always passed as character strings and must be quoted, even if they are numeric values. See the dbsetopt reference page for more information.

See also

dbclropt, dbisopt, dbsetopt

Types

Description

Datatypes and symbolic constants for datatypes used by DB-Library.

Syntax

```
#include <sybfront.h>

#include <sybdb.h>
```

Usage

- Here is a list of symbolic constants for server datatypes. **dbconvert** and **dbwillconvert** use these constants. In addition, the routines **dbcoltype**, **dbalttype**, and **dbrettype** will return one of these types.

Table 2-35: Symbolic constants for server datatypes

Symbolic constant	Represents
SYBDATETIME	datetime type.
SYBDATETIME4	4-byte datetime type.
SYBMONEY4	4-byte money type.
SYBMONEY	money type.
SYBFLT8	8-byte float type.
SYBDECIMAL	decimal type.
SYBNUMERIC	numeric type.
SYBREAL	4-byte float type.
SYBINT4	4-byte integer.
SYBINT2	2-byte integer.
SYBINT1	1-byte integer.
SYBIMAGE	image type.
SYBTEXT	text type.
SYBCHAR	char type.

Symbolic constant	Represents
SYBBIT	bit type.
SYBBINARY	binary type.
SYBBOUNDARY	Security sensitivity_boundary type. **Note** Use DBCHAR as the type for program variables.
SYBSENSITIVITY	Security sensitivity type. **Note** Use DBCHAR as the type for program variables.

For more information on server datatypes, see the *Transact-SQL Reference Manual*.

- Here is a list of C datatypes used by DB-Library functions. These types are useful for defining program variables, particularly variables used with **dbbind**, **dbaltbind**, **dbconvert**, and **dbdatecrack**.

```
/* char, text, boundary, and sensitivity types */
 typedef char             DBCHAR;

/* binary and image type */
 typedef unsigned char    DBBINARY;

/* 1-byte integer */
 typedef unsigned char   DBTINYINT;

 /* 2-byte integer */
 typedef short           DBSMALLINT;

 / unsigned 2-byte integer */
 typedef unsigned short  DBUSMALLINT;

 /* 4-byte integer */
 typedef long            DBINT;

 /* 4-byte float type */
 typedef float           DBREAL;

 typedef struct          dbnumeric
 {
      char               precision;
      char               scale;
      unsigned char      val[MAXNUMLEN];
 } DBNUMERIC;
```

```
typedef DBNUMERIC       DBDECIMAL;

/* 8-byte float type */
typedef double          DBFLT8;

/* bit type */
typedef unsigned char   DBBIT;

/* SUCCEED or FAIL */
typedef int             RETCODE;

/* datetime type */
typedef struct          datetime
{
     /* number of days since 1/1/1900 */
     long            dtdays;
     /* 300ths of a second since midnight */
     unsigned long   dttime;
} DBDATETIME;

/* 4-byte datetime type */
typedef struct            datetime4
{
     /* number of days since 1/1/1900 */
     unsigned short  numdays;
     /* number of minutes since midnight */
     unsigned short  nummins;
} DBDATETIME4;

    typedef struct            dbdaterec
    {
    /* 1900 to the future */
    long    dateyear;
    /* 0 - 11 */
    long    datemonth;
    /* 1 - 31 */
    long    datedmonth;
    /* 1 - 366 */
    long    datedyear;
    /* 0 - 6 (day names depend on language */
    long    datedweek;
    /* 0 - 23 */
    long    datehour;
    /* 0 - 59 */
    long    dateminute;
    /* 0 - 59 */
    long    datesecond;
```

```
    /* 0 - 997 */
    long    datemsecond;
    /* 0 - 127 -- NOTE: Currently unused.*/
    long    datetzone;
  } DBDATEREC;

/* money type */
typedef struct          money
{
    long            mnyhigh;
    unsigned long   mnylow;
} DBMONEY;

/* 4-byte money type */
typedef signed long     DBMONEY4;

/* Pascal-type string */
typedef struct          dbvarychar
{
    /* character count */
    DBSMALLINT      len;
    /* non-terminated string */
    DBCHAR          str[DBMAXCHAR];
} DBVARYCHAR;

/* Pascal-type binary array */
typedef struct          dbvarybin
{
    /* byte count */
    DBSMALLINT       len;
    /* non-terminated array */
    BYTE             array[DBMAXCHAR];
} DBVARYBIN;

/* Used by DB-Library for indicator variables */
typedef     DBSMALLINT     DBINDICATOR;
```

Note The SYBBOUNDARY and SYBSENSITIVITY symbolic constants correspond to the program variable type DBCHAR.

See also

dbaltbind, dbalttype, dbbind, dbcoltype, dbconvert, dbprtype, dbrettype, dbwillconvert, Options

CHAPTER 3

Bulk copy

This chapter describes the DB-Library bulk copy routines.

Topic	Page
Introduction to bulk copy	413
List of bulk copy routines	416

Introduction to bulk copy

Bulk copy is a tool for high-speed transfer of data between a database table and program variables or a host file. It provides an alternative to SQL **insert** and **select** commands.

The DB-Library/C bulk copy special library is a collection of routines that provide bulk copy functionality to a DB-Library/C application. A DB-Library/C application may find bulk copy useful if it needs to exchange data with a non-database application, load data into a new database, or move data from one database to another.

Transferring data into the database

Data can be copied into a database from program variables or from a flat file on the client's host machine.

When copying data into a database table, the chief advantage of bulk copy over the alternative SQL **insert** command is speed. Also, SQL **insert** requires that the data be in character string format, while bulk copy can transfer native datatypes.

When copying data into a non-indexed table, the “high speed” version of bulk copy is used. High-speed means that no data logging is performed during the transfer. If the system fails before the transfer is complete, no new data will remain in the database. Because high-speed transfer affects the recoverability of the database, it is only enabled if the Adaptive Server option “select into/bulkcopy” has been turned on. If the option is not enabled, and a user tries to copy data into a table that has no indexes, Adaptive Server generates an error message.

After the bulk copy is complete, the System Administrator should dump the database to ensure its future recoverability.

When you copy data into an indexed table, a slower version of **bcp** is automatically used, and row inserts are logged.

To copy data into a database, a DB-Library/C application must perform the following introductory steps:

- Call **dblogin** to acquire a LOGINREC structure for later use with **dbopen**.
- Call **BCP_SETL** to set up the LOGINREC for bulk copy operations into the database.
- Call **dbopen** to establish a connection with Adaptive Server.

Call **bcp_init** to initialize the bulk copy operation and inform Adaptive Server whether the copy will be performed from program variables or a host file. To copy data into the database, the **bcp_init** *direction* parameter must be passed as DB_IN.

At this point, an application copying data from program variables will need to perform different steps than an application copying data from a host file.

To copy data from program variables, a DB-Library/C application must perform the following steps in addition to the introductory ones listed previously:

- Call **bcp_bind** once for each program variable that is to be bound to a database column.
- Transfer a batch of data in a loop:
  - Assign program variables the data values to transfer.
  - Call **bcp_sendrow** to send the row of data.

- After a batch of rows has been sent, call **bcp_batch** to save the rows in Adaptive Server.
- After all the data has been sent, call **bcp_done** to end the bulk copy operation.

To copy data from a host file, a DB-Library/C application needs to perform the following steps in addition to the introductory ones listed previously:

- Call **bcp_control** to set the batch size and change control parameter default settings.
- Call **bcp_columns** to set the total number of columns found in the host file.
- Call **bcp_colfmt** once for each column in the host file. If the host file matches the database table exactly, an application does not have to call **bcp_colfmt**.
- Call **bcp_exec** to start the copy in.

Transferring data out of the database to a flat file

Data can be copied out from a database only into an operating system (host) file. Bulk copy does not allow the transfer of data from a database into program variables.

When transferring data out to a host file from a database table, the chief advantage of bulk copy over SQL **select** is that it allows very specific output file formats to be specified. Bulk copy is not significantly faster than SQL **select**.

To copy data out from a database, a DB-Library/C application must perform the following steps:

1 Call **dblogin** to acquire a LOGINREC structure for later use with **dbopen**.

2 Call **dbopen** to establish a connection with Adaptive Server.

Call **bcp_init** to initialize the bulk copy operation. To copy data out from the database, *direction* must be passed as DB_OUT.

---

3 Call **bcp_control** to set the batch size and change control parameter default settings.

4 Call **bcp_columns** to set the total number of columns found in the host file.

5 Call **bcp_colfmt** once for each column in the host file. If the host file matches the database table exactly, an application does not have to call **bcp_colfmt**.

6 Call **bcp_exec** to start the copy out.

List of bulk copy routines

Routine	Description
bcp_batch	Save any preceding rows in Adaptive Server.
bcp_bind	Bind data from a program variable to a Adaptive Server table.
bcp_colfmt	Specify the format of a host file for bulk copy purposes.
bcp_colfmt_ps	Specify the format of a host file for bulk copy purposes, with precision and scale support for numeric and decimal columns.
bcp_collen	Set the program variable data length for the current bulk copy **into the database**.
bcp_colptr	Set the program variable data address for the current bulk copy **into the database**.
bcp_columns	Set the total number of columns found in the host file.
bcp_control	Change various control parameter default settings.
bcp_done	End a bulk copy from program variables into Adaptive Server.
bcp_exec	Execute a bulk copy of data between a database table and a host file.
bcp_getl	Determine if the LOGINREC has been set for bulk copy operations.
bcp_init	Initialize bulk copy.
bcp_moretext	Send part of a text or image value to Adaptive Server.
bcp_options	Set bulk copy options.

Routine	Description
bcp_readfmt	Read a datafile format definition from a host file.
bcp_sendrow	Send a row of data from program variables to Adaptive Server.
BCP_SETL	Set the LOGINREC for bulk copy operations into the database.
bcp_setxlate	Specify the character set translations to use when retrieving data from or inserting data into a Adaptive Server.
bcp_writefmt	Write a datafile format definition to a host file.

bcp_batch

Description

Save any preceding rows in Adaptive Server.

Syntax

```
DBINT bcp_batch(dbproc)

 DBPROCESS    *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

Return value

The number of rows saved since the last call to **bcp_batch**, or -1 in case of error.

Usage

- When an application uses **bcp_bind** and **bcp_sendrow** to bulk-copy rows from program variables to Adaptive Server tables, the rows are permanently saved in Adaptive Server only when the program calls **bcp_batch** or **bcp_done**.
- You may call bcp_batch once every *n* rows or when there is a lull between periods of incoming data (as in a telemetry application). Of course, you may choose some other criteria, or may decide not to call bcp_batch at all. If bcp_batch is not called, the rows are permanently saved in Adaptive Server when bcp_done is called.

- By default, Adaptive Server copies all the rows specified in one batch. Adaptive Server considers each batch to be a separate **bcp** operation. Each batch is copied in a single **insert** transaction, and if any row in the batch is rejected, the entire **insert** is rolled back. **bcp** then continues to the next batch. You can use **bcp_batch** to break large input files into smaller units of recoverability. For example, if 300,000 rows are bulk copied and **bcp_batch** is called every 100,000 rows, if there is a fatal error after row 200,000, the first two batches—200,000 rows—will have been successfully copied into Adaptive Server.
- **bcp_batch** actually sends two commands to the server. The first command tells the server to permanently save the rows. The second tells the server to begin a new transaction. It is possible that the command to save the rows completes successfully but the command to start a new transaction does not. In this case, **bcp_batch**'s error return of -1 does not indicate that the rows have not been successfully saved. To find out whether this has happened, an application can refer to the messages generated by Adaptive Server or DB-Library/C.

See also

bcp_bind, bcp_done, bcp_sendrow

bcp_bind

Description

Bind data from a program variable to an Adaptive Server table.

Syntax

```
RETCODE bcp_bind (dbproc, varaddr, prefixlen, varlen,
        terminator, termlen, type,
        table_column)

DBPROCESS    *dbproc;
BYTE         *varaddr;
int          prefixlen;
DBINT        varlen;
BYTE         *terminator;
int          termlen;
int          type;
int          table_column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

varaddr

The address of the program variable from which the data will be copied. If type is SYBTEXT or SYBIMAGE, *varaddr* can be NULL. A NULL *varaddr* indicates that text and image values will be sent to Adaptive Server in chunks by **bcp_moretext**, rather than all at once by **bcp_sendrow**.

prefixlen

The length, in bytes, of any length prefix this column may have. For example, strings in some non-C programming languages are made up of a one-byte length prefix, followed by the string data itself. If the data does not have a length prefix, set *prefixlen* to 0.

varlen

The length of the data in the program variable, *not* including the length of any length prefix and/or terminator. Setting *varlen* to 0 signifies that the data is null. Setting *varlen* to -1 indicates that the system should ignore this parameter.

For fixed-length datatypes, such as integer, the datatype itself indicates to the system the length of the data. Therefore, for fixed-length datatypes, *varlen* must always be -1, except when the data is null, in which case *varlen* must be 0.

For char, text, binary, and image datatypes, *varlen* can be -1, 0, or some positive value. If *varlen* is -1, the system will use either a length prefix or a terminator sequence to determine the length. (If both are supplied, the system will use the one that results in the shortest amount of data being copied.) If *varlen* is -1 and neither a prefix length nor a terminator sequence is specified, the system will return an error message. If *varlen* is 0, the system assumes the data is null. If *varlen* is some positive value, the system uses *varlen* as the data length. However, if, in addition to a positive *varlen*, a prefix length and/or terminator sequence is provided, the system determines the data length by using the method that results in the shortest amount of data being copied.

terminator
A pointer to the byte pattern, if any, that marks the end of this program variable. For example, C strings usually have a 1-byte terminator whose value is 0. If there is no terminator for the variable, set *terminator* to NULL.

If you want to designate the C null terminator as the program variable terminator, the simplest way is to use an empty string ("") as *terminator* and set *termlen* to 1, since the null terminator constitutes a single byte. For instance, the second **bcp_bind** call in the "Example" section below uses two tabs as the program variable terminator. It could be rewritten to use a C null terminator instead, as follows:

```
bcp_bind (dbproc, co_name, 0, -1, "", 1, 0, 2)
```

termlen
The length of this program variable's terminator, if any. If there is no terminator for the variable, set *termlen* to 0.

type
The datatype of your program variable, expressed as an Adaptive Server datatype. The data in the program variable will be automatically converted to the type of the database column. If this parameter is 0, no conversion will be performed. See the dbconvert reference page for a list of supported conversions. That reference page also contains a list of Adaptive Server datatypes.

table_column
The column in the database table to which the data will be copied. Column numbers start at 1.

Return value

SUCCEED or FAIL.

Examples

- The following program fragment illustrates **bcp_bind**:

```
LOGINREC            *login;
 DBPROCESS           *dbproc;
 char                co_name[MAXNAME];
 DBINT               co_id;
 DBINT               rows_sent;
 DBBOOL              more_data;
 char                *terminator = "\t\t";

 /* Initialize DB-Library. */
 if (dbinit() == FAIL)
      exit(ERREXIT);

 /* Install error-handler and message-handler. */
 dberrhandle(err_handler);
```

```
dbmsghandle(msg_handler);

/* Open a DBPROCESS. */
login = dblogin();
BCP_SETL(login, TRUE);
dbproc = dbopen(login, NULL);

/* Initialize bcp. */
if (bcp_init(dbproc, "comdb..accounts_info",
     NULL, NULL, DB_IN) == FAIL)
     exit(ERREXIT);

/* Bind program variables to table columns. */
if (bcp_bind(dbproc, &co_id, 0, -1,
     (BYTE *)NULL, 0, 0, 1) == FAIL)
{
     fprintf(stderr, "bcp_bind, column 1, failed.\n");
     exit(ERREXIT);
}

if (bcp_bind
     (dbproc, co_name, 0, -1, (BYTE *)terminator,
     strlen(terminator), 0, 2)
          == FAIL)
{
     fprintf(stderr, "bcp_bind, column 2, failed.\n");
     exit(ERREXIT);
}

while (TRUE)
{
     /* Process/retrieve program data. */
     more_data = getdata(&co_id, co_name);

     if (more_data == FALSE)
           break;

      /* Send the data. */
      if (bcp_sendrow(dbproc) == FAIL)
         exit(ERREXIT);
}

/* Terminate the bulk copy operation. */
if ((rows_sent = bcp_done(dbproc)) == -1)
     printf("Bulk-copy unsuccessful.\n");
else
```

```
        printf("%ld rows copied.\n", rows_sent);
```

Usage

- There may be times when you want to copy data directly from a program variable into a table in Adaptive Server, without having to first place the data in a host file or use the SQL **insert** command. The **bcp_bind** function is a fast and efficient way to do this.
- You must call **bcp_init** before calling this or any other bulk copy functions.
- There must be a separate **bcp_bind** call for every column in the Adaptive Server table into which you want to copy. After the necessary **bcp_bind** calls have been made, you then call **bcp_sendrow** to send a row of data from your program variables to Adaptive Server. The table to be copied into is set by calling **bcp_init**.
- You can override the program variable data length (*varlen*) for a particular column on the current copy in by calling bcp_collen.
- Whenever you want Adaptive Server to checkpoint the rows already received, call **bcp_batch**. For example, you may want to call **bcp_batch** once for every 1000 rows inserted, or at any other interval.
- When there are no more rows to be inserted, call **bcp_done**. Failure to do so will result in an error.
- When using bcp_bind, the host file name parameter used in the call to bcp_init, *hfile*, must be set to NULL, and the direction parameter, *direction*, must be set to DB_IN.
- Prefix lengths should not be used with fixed-length datatypes, such as integer or float. For fixed-length datatypes, since bulk copy can figure out the length of the data from the datatype, pass *prefixlen* as 0 and *varlen* as -1, except when the data is NULL, in which case *varlen* must be 0.
- Control parameter settings, specified with **bcp_control**, have no effect on **bcp_bind** row transfers.
- It is an error to call **bcp_columns** when using **bcp_bind**.

See also

bcp_batch, bcp_colfmt, bcp_collen, bcp_colptr, bcp_columns, bcp_control, bcp_done, bcp_exec, bcp_init, bcp_moretext, bcp_sendrow

bcp_colfmt

Description

Specify the format of a host file for bulk copy purposes.

Syntax

```
RETCODE bcp_colfmt (dbproc, host_colnum, host_type,
        host_prefixlen, host_collen,
        host_term, host_termlen,
        table_colnum)

DBPROCESS    *dbproc;
int           host_colnum;
int           host_type;
int           host_prefixlen;
DBINT         host_collen;
BYTE          *host_term;
int           host_termlen;
int           table_colnum;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

host_colnum

The column in the host file whose format is being specified. The first column is number 1.

host_type

The datatype of this column in the host file, expressed as an Adaptive Server datatype. If it is different from the datatype of the corresponding column in the database table (*table_colnum*), the conversion will be performed automatically. See the dbconvert reference page for a table of allowable data conversions. That reference page also contains a list of Adaptive Server datatypes.

If you want to specify the same datatype as in the corresponding column of the database table (*table_colnum*), this parameter should be set to 0.

Note **bcp_colfmt** does not offer precision and scale support for numeric and decimal types. When setting the format of a numeric or decimal host column, **bcp_colfmt** uses a default precision and scale of 18 and 0, respectively. To specify a different precision and scale, an application can call **bcp_colfmt_ps**.

---

host_prefixlen
The length of the length prefix for this column in the host file. Legal prefix lengths are 1, 2, and 4 bytes. To avoid using a length prefix, this parameter should be set to 0. To let **bcp** decide whether to use a length prefix, this parameter should be set to -1. In such a case, **bcp** will use a length prefix (of whatever length is necessary) if the database column length is variable.

If more than one means of specifying a host file column length is used (such as a length prefix and a maximum column length, or a length prefix and a terminator sequence), **bcp** will look at all of them and use the one that results in the smallest amount of data being copied.

One valuable use for length prefixes is to simplify the specifying of null data values in a host file. For instance, assume you have a 1-byte length prefix for a 4-byte integer column. Ordinarily, the length prefix will contain a value of 4, to indicate that a 4-byte value follows. However, if the value of the column is NULL, the length prefix can be set to 0 to indicate that 0 bytes follow for the column.

host_collen
The maximum length of this column's data in the host file, *not* including the length of any length prefix and/or terminator. Setting *host_collen* to 0 signifies that the data is NULL. Setting *host_collen* to -1 indicates that the system should ignore this parameter (that is, there is no default maximum length).

For fixed-length datatypes, such as integer, the length of the data is constant, except for the special case of null values. Therefore, for fixed-length datatypes, *host_collen* must always be -1, except when the data is null, in which case *host_collen* must be 0.

For char, text, binary, and image datatypes, *host_collen* can be -1, 0, or some positive value. If *host_collen* is -1, the system will use either a length prefix or a terminator sequence to determine the length of the data. (If both are supplied, the system will use the one that results in the shortest amount of data being copied.) If *host_collen* is -1 and neither a prefix length nor a terminator sequence is specified, the system will return an error message. If *host_collen* is 0, the system assumes the data is NULL. If *host_collen* is some positive value, the system uses *host_collen* as the maximum data length. However, if, in addition to a positive *host_collen*, a prefix length and/or terminator sequence is provided, the system determines the data length by using the method that results in the shortest amount of data being copied.

host_term
The terminator sequence to be used for this column. This parameter is mainly useful for char, text, binary, and image datatypes, because all other datatypes are of fixed length. To avoid using a terminator, set this parameter to NULL. To set the terminator to the NULL character, set *host_term* to "\0". To make the tab character the terminator, set *host_term* to "\t". To make the newline character the terminator, set *host_term* to "\n".

If more than one means of specifying a host file column length is used (such as a terminator and a length prefix, or a terminator and a maximum column length), **bcp** will look at all of them and use the one that results in the smallest amount of data being copied.

host_termlen
The length, in bytes, of the terminator sequence to be used for this column. To avoid using a terminator, set this value to -1.

table_colnum
The corresponding column in the database table. If this value is 0, this column will *not* be copied. The first column is column 1.

Return value

SUCCEED or FAIL.

Usage

- **bcp_colfmt** allows you to specify the host file format for bulk copies. For bulk copy purposes, a format contains the following parts:
  - A mapping from host file columns to database columns
  - The datatype of each host file column
  - The length of the optional length prefix of each column
  - The maximum length of the host file column's data
  - The optional terminating byte sequence for each column
  - The length of this optional terminating byte sequence
- Each call to bcp_colfmt specifies the format for one host file column. For example, if you have a table with five columns and want to change the default settings for three of those columns, you should first call bcp_columns(*dbproc*, 5), and then call bcp_colfmt five times, with three of those calls setting your custom format. The remaining two calls should have their *host_type* set to 0, and their *host_prefixlen*, *host_collen*, and *host_termlen* parameters set to -1. The result of this would be to copy all five columns—three with your customized format and two with the default format.
- **bcp_columns** *must* be called before any calls to **bcp_colfmt**.

- You must call **bcp_colfmt** for every column in the host file, regardless of whether some of those columns use the default format or are skipped.
- To skip a column, set the *table_column* parameter to 0.

See also

bcp_batch, bcp_bind, bcp_colfmt_ps, bcp_collen, bcp_colptr, bcp_columns, bcp_control, bcp_done, bcp_exec, bcp_init, bcp_sendrow

bcp_colfmt_ps

Description

Specify the format of a host file for bulk copy purposes, with precision and scale support for numeric and decimal columns.

Syntax

```
RETCODE bcp_colfmt_ps (dbproc, host_colnum, host_type,
                host_prefixlen, host_collen,
                host_term, host_termlen,
                table_colnum, typeinfo)

DBPROCESS    *dbproc;
int           host_colnum;
int           host_type;
int           host_prefixlen;
DBINT         host_collen;
BYTE          *host_term;
int           host_termlen;
int           table_colnum;
DBTYPEINFO   *typeinfo;
```

Note **bcp_colfmt_ps**'s parameters are identical to **bcp_colfmt**'s, except that **bcp_colfmt_ps** has the additional parameter *typeinfo*, which contains information about precision and scale for numeric or decimal columns.

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

host_colnum

The column in the host file whose format is being specified. The first column is number 1.

host_type

The datatype of this column in the host file, expressed as an Adaptive Server datatype. If it is different from the datatype of the corresponding column in the database table (*table_colnum*), the conversion will be performed automatically. See the dbconvert reference page for a table of allowable data conversions. That reference page also contains a list of Adaptive Server datatypes.

If you want to specify the same datatype as in the corresponding column of the database table (*table_colnum*), this parameter should be set to 0.

host_prefixlen

The length of the length prefix for this column in the host file. Legal prefix lengths are 1, 2, and 4 bytes. To avoid using a length prefix, this parameter should be set to 0. To let **bcp** decide whether to use a length prefix, this parameter should be set to -1. In such a case, **bcp** will use a length prefix (of whatever length is necessary) if the database column length is variable.

If more than one means of specifying a host file column length is used (such as a length prefix and a maximum column length, or a length prefix and a terminator sequence), **bcp** will look at all of them and use the one that results in the shortest amount of data being copied.

One valuable use for length prefixes is to simplify the specifying of null data values in a host file. For instance, assume you have a 1-byte length prefix for a 4-byte integer column. Ordinarily, the length prefix will contain a value of 4, to indicate that a 4-byte value follows. However, if the value of the column is null, the length prefix can be set to 0, to indicate that 0 bytes follow for the column.

host_collen
: The maximum length of this column's data in the host file, not including the length of any length prefix and/or terminator. Setting *host_collen* to 0 signifies that the data is NULL. Setting *host_collen* to -1 indicates that the system should ignore this parameter (that is, there is no default maximum length).

  For fixed-length datatypes, such as integer, the length of the data is constant, except for the special case of null values. Therefore, for fixed-length datatypes, *host_collen* must always be -1, except when the data is NULL, in which case *host_collen* must be 0.

  For char, text, binary, and image datatypes, *host_collen* can be -1, 0, or some positive value. If *host_collen* is -1, the system will use either a length prefix or a terminator sequence to determine the length of the data. (If both are supplied, the system will use the one that results in the smallest amount of data being copied.) If *host_collen* is -1 and neither a prefix length nor a terminator sequence is specified, the system will return an error message. If *host_collen* is 0, the system assumes the data is NULL. If *host_collen* is some positive value, the system uses *host_collen* as the maximum data length. However, if, in addition to a positive *host_collen*, a prefix length and/or terminator sequence is provided, the system determines the data length by using the method that results in the smallest amount of data being copied.

host_term
: The terminator sequence to be used for this column. This parameter is mainly useful for char, text, binary, and image datatypes, because all other types are of fixed length. To avoid using a terminator, set this parameter to NULL. To set the terminator to the null character, set *host_term* to "\0". To make the tab character the terminator, set *host_term* to "\t". To make the newline character the terminator, set *host_term* to "\n".

  If more than one means of specifying a host file column length is used (such as a terminator and a length prefix, or a terminator and a maximum column length), **bcp** will look at all of them and use the one that results in the smallest amount of data being copied.

host_termlen
: The length, in bytes, of the terminator sequence to be used for this column. To avoid using a terminator, set this value to -1.

table_colnum
: The corresponding column in the database table. If this value is 0, this column will not be copied. The first column is column 1.

typeinfo

A pointer to a DBTYPEINFO structure containing information about the precision and scale of decimal or numeric host file columns. An application sets a DBTYPEINFO structure with values for precision and scale before calling **bcp_colfmt_ps** to specify the host file format of decimal or numeric columns.

If *typeinfo* is NULL, **bcp_colfmt_ps** is the equivalent of **bcp_colfmt**. That is:

- If the server column is of type numeric or decimal, **bcp_colfmt_ps** picks up precision and scale values from the column.
- If the server column is not numeric or decimal, **bcp_colfmt_ps** uses a default precision of 18 and a default scale of 0.

If *host_type* is not 0, SYBDECIMAL or SYBNUMERIC, *typeinfo* is ignored.

If *host_type* is 0 and the corresponding server column is not numeric or decimal, *typeinfo* is ignored.

A DBTYPEINFO structure is defined as follows:

```
typedef struct typeinfo {
     DBINT   precision;
     DBINT   scale;
 } DBTYPEINFO;
```

Legal values for *precision* are from 1 to 77. Legal values for *scale* are from 0 to 77. *scale* must be less than or equal to *precision*.

Return value

SUCCEED or FAIL.

Usage

- bcp_colfmt_ps is the equivalent of bcp_colfmt, except that bcp_colfmt_ps provides precision and scale support for numeric and decimal datatypes, which bcp_colfmt does not. Calling bcp_colfmt is equivalent to calling bcp_colfmt_ps with *typeinfo* as NULL.
- **bcp_colfmt_ps** allows you to specify the host file format for bulk copies. For bulk copy purposes, a format contains the following parts:
  - A mapping from host file columns to database columns
  - The datatype of each host file column
  - The length of the optional length prefix of each column
  - The maximum length of the host file column's data
  - The optional terminating byte sequence for each column

    - The length of this optional terminating byte sequence
- Each call to bcp_colfmt_ps specifies the format for one host file column. For example, if you have a table with five columns, and want to change the default settings for three of those columns, you should first call bcp_columns(*dbproc*, 5), and then call bcp_colfmt_ps five times, with three of those calls setting your custom format. The remaining two calls should have their *host_type* set to 0, and their *host_prefixlen*, *host_collen*, and *host_termlen* parameters set to -1. The result of this would be to copy all five columns—three with your customized format and two with the default format.
- **bcp_columns** *must* be called before any calls to **bcp_colfmt_ps**.
- You must call **bcp_colfmt_ps** for every column in the host file, regardless of whether some of those columns use the default format or are skipped.
- To skip a column, set the *table_column* parameter to 0.

See also

bcp_batch, bcp_bind, bcp_colfmt, bcp_collen, bcp_colptr, bcp_columns, bcp_control, bcp_done, bcp_exec, bcp_init, bcp_sendrow

bcp_collen

Description

Set the program variable data length for the current bulk copy into the database.

Syntax

RETCODE bcp_collen(dbproc, varlen, table_column)

```
DBPROCESS     *dbproc;
DBINT          varlen;
int           table_column;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

varlen

The length of the program variable, which does *not* include the length of the length prefix or terminator. Setting *varlen* to 0 signifies that the data is NULL. Setting it to -1 signifies that the data is variable-length and that the length will be determined by the length prefix or terminator. If both a length prefix and a terminator exist, **bcp** will use the one that results in the smallest amount of data being copied.

table_column
The column in the Adaptive Server table to which the data will be copied. Column numbers start at 1.

Return value
SUCCEED or FAIL.

Usage
- The **bcp_collen** function allows you to change the program variable data length for a particular column while running a copy **in** via calls to **bcp_bind.**
- Initially, the program variable data length is determined when **bcp_bind** is called. If the program variable data length changes between calls to **bcp_sendrow,** and no length prefix or terminator is being used, you may call **bcp_collen** to reset the length. The next call to **bcp_sendrow** will use the length you just set.
- There must be a separate **bcp_collen** call for every column in the table whose data length you want to modify.

See also
bcp_bind, bcp_colptr, bcp_sendrow

bcp_colptr

Description
Set the program variable data address for the current bulk copy into the database.

Syntax
RETCODE bcp_colptr(dbproc, colptr, table_column)

```
DBPROCESS    *dbproc;
BYTE          *colptr;
int           table_column;
```

Parameters
dbproc
A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

colptr
The address of the program variable.

table_column
The column in the Adaptive Server table to which the data will be copied. Column numbers start at 1.

Return value
SUCCEED or FAIL.

Usage

- The **bcp_colptr** function allows you to change the program variable data address for a particular column while running a copy **in** via calls to **bcp_bind.**
- Initially, the program variable data address is determined when **bcp_bind** is called. If the program variable data address changes between calls to **bcp_sendrow,** you may call **bcp_colptr** to reset the address of the data. The next call to **bcp_sendrow** will use the data at the address you just set.
- There must be a separate **bcp_colptr** call for every column in the table whose data address you want to modify.

See also

bcp_bind, bcp_collen, bcp_sendrow

bcp_columns

Description

Set the total number of columns found in the host file.

Syntax

RETCODE bcp_columns(dbproc, host_colcount)

```
DBPROCESS    *dbproc;
int          host_colcount;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

host_colcount

The total number of columns in the host file. Even if you are preparing to bulk copy data from the host file to an Adaptive Server table and do not intend to copy all columns in the host file, you must still set *host_colcount* to the total number of host file columns.

Return value

SUCCEED or FAIL.

Usage

- This function sets the total number of columns found in a host file for use with bulk copy. This routine may be called only after **bcp_init** has been called with a valid file name.
- You should call this routine only if you intend to use a host file format that differs from the default. The default host file format is described on the bcp_init reference page.

- After calling bcp_columns, you must call bcp_colfmt *host_colcount* times, because you are defining a completely custom file format.

See also

bcp_colfmt, bcp_init

bcp_control

Description

Change various control parameter default settings.

Syntax

RETCODE bcp_control(dbproc, field, value)

```
DBPROCESS    *dbproc;
int           field;
DBINT        value;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

field

A control-parameter identifier consisting of one of the following symbolic values:

Field	Description
BCPMAXERRS	The number of errors allowed before giving up. The default is 10.
BCPFIRST	The first row to copy. The default is 1. A value of less than 1 resets this field to its default value of 1.
BCPLAST	The last row to copy. The default is to copy all rows. A value of less than 1 resets this field to its default value.
BCPBATCH	The number of rows per batch. The default is 0, which means that the entire bulk copy will be done in one batch. This field is only meaningful when copying from a host file into Adaptive Server.

value

The value to change the corresponding control parameter to.

Return value

SUCCEED or FAIL.

Usage

- This function sets various control parameters for bulk copy operations, including the number of errors allowed before aborting a bulk copy, the numbers of the first and last rows to copy, and the batch size.

- These control parameters are only meaningful when the application copies between a host file and an Adaptive Server table. Control parameter settings have no effect on **bcp_bind** row transfers.
- By default, Adaptive Server copies all the rows specified in one batch. Adaptive Server considers each batch to be a separate **bcp** operation. Each batch is copied in a single **insert** transaction, and if any row in the batch is rejected, the entire **insert** is rolled back. **bcp** then continues to the next batch. You can use **bcp_batch** to break large input files into smaller units of recoverability. For example, if 300,000 rows are bulk copied in with a batch size of 100,000 rows, and there is a fatal error after row 200,000, the first two batches—200,000 rows—will have been successfully copied into Adaptive Server. If batching had not been used, no rows would have been copied into Adaptive Server.
- The following program fragment illustrates **bcp_control**:

```
LOGINREC        *login;
 DBPROCESS       *dbproc;
 DBINT           rowsread;

 /* Initialize DB-Library. */
 if (dbinit() == FAIL)
      exit(ERREXIT);

 /* Install error-handler and message-handler. */
 dberrhandle(err_handler);
 dbmsghandle(msg_handler);

 /* Open a DBPROCESS. */
 login = dblogin();
 BCP_SETL(login, TRUE);
 dbproc = dbopen(login, NULL);

 /* Initialize bcp. */
 if (bcp_init(dbproc, "comdb..address", "address.add",
      "addr.error", DB_IN) == FAIL)
      exit(ERREXIT);

 /* Set the number of rows per batch. */
 if (bcp_control(dbproc, BCPBATCH, 1000) == FAIL)
 {
      printf("bcp_control failed to set batching behavior.\n");
      exit(ERREXIT);
 }
```

```
/* Set host column count. */
if (bcp_columns(dbproc, 1) == FAIL)
{
     printf("bcp_columns failed.\n");
     exit(ERREXIT);
}

/* Set the host-file format. */
if (bcp_colfmt(dbproc, 1, 0, 0, -1, (BYTE *)("\n"), 1, 1) == FAIL)
{
     printf("bcp_colformat failed.\n");
     exit(ERREXIT);
}

/* Now, execute the bulk copy. */
if (bcp_exec(dbproc, &rowsread) == FAIL)
{
     printf("Incomplete bulk copy.  Only %ld row%c copied.\n",
             rowsread, (rowsread == 1) ? ' ': 's');

     exit(ERREXIT);
}
```

See also bcp_batch, bcp_bind, bcp_colfmt, bcp_collen, bcp_colptr, bcp_columns, bcp_done, bcp_exec, bcp_init

bcp_done

Description End a bulk copy from program variables into Adaptive Server.

Syntax DBINT bcp_done(dbproc)

DBPROCESS*dbproc;

Parameters dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

Return value The number of rows permanently saved since the last call to bcp_batch, or -1 in case of error.

Usage

- **bcp_done** ends a bulk copy performed with **bcp_bind** and **bcp_sendrow**. It should be called after the last call to **bcp_sendrow** or **bcp_moretext**. Failure to call **bcp_done** after you have completed copying in all your data will result in unpredictable errors.

See also

bcp_batch, bcp_bind, bcp_moretext, bcp_sendrow

bcp_exec

Description

Execute a bulk copy of data between a database table and a host file.

Syntax

RETCODE bcp_exec(dbproc, rows_copied)

```
DBPROCESS    *dbproc;
DBINT        *rows_copied;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

rows_copied

A pointer to a DBINT. **bcp_exec** will fill this DBINT with the number of rows successfully copied. If set to NULL, this parameter will not be filled in by **bcp_exec**.

Return value

SUCCEED or FAIL.

bcp_exec returns SUCCEED if all rows are copied. If a partial or complete failure occurs, bcp_exec returns FAIL. Check the *rows_copied* parameter for the number of rows successfully copied.

Usage

- This routine copies data from a host file to a database table or vice-versa, depending on the value of the *direction* parameter in bcp_init.
- Before calling this function you must call **bcp_init** with a valid host file name. Failure to do so will result in an error.
- The following program fragment illustrates **bcp_exec**:

```
LOGINREC        *login;
 DBPROCESS       *dbproc;
 DBINT           rowsread;

 /* Initialize DB-Library. */
```

```
if (dbinit() == FAIL)
      exit(ERREXIT);

/* Install error-handler and message-handler. */
dberrhandle(err_handler);
dbmsghandle(msg_handler);

/* Open a DBPROCESS. */
login = dblogin();
BCP_SETL(login, TRUE);
dbproc = dbopen(login, NULL);

/* Initialize bcp. */
if (bcp_init(dbproc, "pubs2..authors", "authors.save",
     (BYTE *)NULL, DB_OUT) == FAIL)
     exit(ERREXIT);

/* Now, execute the bulk copy. */
if (bcp_exec(dbproc, &rowsread) == FAIL)
     printf("Incomplete bulk copy.  Only %ld row%c copied.\n",
             rowsread, (rowsread == 1) ? ' ': 's');
```

See also bcp_batch, bcp_bind, bcp_colfmt, bcp_collen, bcp_colptr, bcp_columns, bcp_control, bcp_done, bcp_init, bcp_sendrow

bcp_getl

Description Determine if the LOGINREC has been set for bulk copy operations.

Syntax

```
DBBOOL bcp_getl(loginrec)

LOGINREC      *loginrec;
```

Parameters

loginrec
A pointer to a LOGINREC structure that will be passed as an argument to **dbopen**. You can get a LOGINREC structure by calling **dblogin**.

Return value TRUE or FALSE.

Usage

- bcp_getl returns TRUE if **loginrec* is enabled for bulk copy operations, and FALSE if it is not.

- A DBPROCESS connection cannot be used for bulk copy **in** operations unless the LOGINREC used to open the connection has been set to enable bulk copy. The macro **BCP_SETL** sets a LOGINREC to enable bulk copy. By default, DBPROCESSes are not enabled for bulk copy operations.
- Applications that allow users to make ad hoc queries may want to avoid calling BCP_SETL (or call it with a value of FALSE for the *enable* parameter) to prevent users from initiating a bulk copy sequence via SQL commands. Once a bulk copy sequence has begun, it cannot be stopped by an ordinary SQL command.
- If LOGINREC is NULL, **bcp_getl** returns FALSE.

See also

bcp_init, BCP_SETL, dblogin, dbopen

bcp_init

Description

Initialize bulk copy.

Syntax

```
RETCODE bcp_init(dbproc, tblname, hfile, errfile,
        direction)

DBPROCESS    *dbproc;
char         *tblname;
char         *hfile;
char         *errfile;
int          direction;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

tblname

The name of the database table to be copied in or out. This name may also include the database name or the owner name. For example, *pubs2.gracie.titles*, *pubs2..titles*, *gracie.titles*, and *titles* are all legal table names.

hfile

The name of the host file to be copied in or out. If no host file is involved (the situation when data is being copied directly from variables), *hfile* should be NULL.

errfile

The name of the error file to be used. This error file will be filled with progress messages, error messages, and copies of any rows that, for any reason, could not be copied from a host file to an Adaptive Server table.

If *errfile* is NULL, no error file is used.

If *hfile* is NULL, *errfile* is ignored. This is because an error file is not necessary when bulk-copying from program variables.

direction

The direction of the copy. It must be one of two values—DB_IN or DB_OUT. DB_IN indicates a copy from the host into the database table, while DB_OUT indicates a copy from the database table into the host file.

It is illegal to request a bulk copy from the database table (DB_OUT) without supplying a host file name.

Return value

SUCCEED or FAIL.

Usage

- **bcp_init** performs the necessary initialization for a bulk copy of data between the front-end and an Adaptive Server. It sets the default host file data formats and examines the structure of the database table.
- **bcp_init** must be called before any other bulk copy functions. Failure to do so will result in an error.
- If a host file is being used (see the description of *hfile* in the "Parameters" section above), the default data formats are as follows:
  - The order, type, length and number of the columns in the host file are assumed to be identical to the order, type and number of the columns in the database table.
  - If a given database column's data is fixed-length, then the host file's data column will also be fixed-length. If a given database column's data is variable-length or may contain null values, the host file's data column will be prefixed by a 4-byte length value for SYBTEXT and SYBIMAGE data types, and a 1-byte length value for all other types.
  - There are no terminators of any kind between host file columns.

  Any of these defaults can be overridden by calling **bcp_columns** and **bcp_colfmt.**
- Using the bulk copy routines to copy data to a database table requires the following:
  - The DBPROCESS structure must be usable for bulk copy purposes. This is accomplished by calling **BCP_SETL**:

```
login = dblogin();
 BCP_SETL(login, TRUE);
```

- If the table has no indexes, the database option **select into/bulkcopy** must be set to "true". The following SQL command will do this:

  ```
  sp_dboption 'mydb', 'select into/bulkcopy',
  'true'
  ```

  See the *SQLServer Reference Manual* for further details on **sp_dboption**.

- If no host file is being used, it is necessary to call **bcp_bind** to specify the format and location in memory of each column's data value.
- If no host file is being used, *errfile* is ignored. An error file is not necessary when bulk-copying from program variables because bcp_sendrow returns FAIL if an error occurs. In this case, the application can examine the bulk copy program variables to determine which row values caused the error.

See also

bcp_batch, bcp_bind, bcp_colfmt, bcp_collen, bcp_colptr, bcp_columns, bcp_control, bcp_done, bcp_exec, bcp_sendrow

bcp_moretext

Description

Send part of a text or image value to Adaptive Server.

Syntax

RETCODE bcp_moretext(dbproc, size, text)

```
DBPROCESS    *dbproc;
DBINT         size;
BYTE         *text;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

size

The size of this particular part of the text or image value being sent to Adaptive Server. It is an error to send more text or image bytes to Adaptive Server than were specified in the call to **bcp_bind** or **bcp_collen**.

text

A pointer to the text or image portion to be written.

Return value

SUCCEED or FAIL.

Usage

- This routine is used in conjunction with **bcp_bind** and **bcp_sendrow** to send a large SYBTEXT or SYBIMAGE value to Adaptive Server in the form of a number of smaller chunks. This is particularly useful with operating systems unable to allocate extremely long data buffers.
- If bcp_bind is called with a *type* parameter of SYBTEXT or SYBIMAGE and a non-NULL *varaddr* parameter, bcp_sendrow will send the entire text or image data value, just as it does for all other datatypes. If, however, bcp_bind has a NULL *varaddr* parameter, bcp_sendrow will return control to the application immediately after all non-text or image columns are sent to Adaptive Server. The application can then call bcp_moretext repeatedly to send the text and image columns to Adaptive Server, a chunk at a time.
- Here is an example that illustrates how to use **bcp_moretext** with **bcp_bind** and **bcp_sendrow**:

```
LOGINREC        *login;
 DBPROCESS       *dbproc;

 DBINT           id = 5;
 char            *part1 = "This text value isn't very long,";
 char            *part2 = " but it's broken up into three parts";
 char            *part3 = " anyhow.";

 /* Initialize DB-Library. */
 if (dbinit() == FAIL)
       exit(ERREXIT);

 /* Install error handler and message handler. */
 dberrhandle(err_handler);
 dbmsghandle(msg_handler);

 /* Open a DBPROCESS */
 login = dblogin();
 BCP_SETL(login, TRUE);
 dbproc = dbopen(login, NULL);

 /* Initialize bcp. */
 if (bcp_init(dbproc, "comdb..articles", (BYTE *)NULL,
      (BYTE *)NULL, DB_IN) == FAIL)
      exit(ERREXIT);

 /* Bind program variables to table columns. */
 if (bcp_bind(dbproc, (BYTE *)&id, 0, (DBINT)-1,
      (BYTE *)NULL, 0, SYBINT4, 1) == FAIL)
```

```
{
     fprintf(stderr, "bcp_bind, column 1, failed.\n");
     exit(ERREXIT);
}

if (bcp_bind
      (dbproc, (BYTE *)NULL, 0,
      (DBINT) (strlen(part1) + strlen(part2) + strlen(part3)),
      (BYTE *)NULL, 0, SYBTEXT, 2)
      == FAIL)
{
      fprintf(stderr, "bcp_bind, column 2, failed.\n");
      exit(ERREXIT);
}

/*
** Now send this row, with the text value broken into
** three chunks.
*/
if (bcp_sendrow(dbproc) == FAIL)
     exit(ERREXIT);
if (bcp_moretext(dbproc, (DBINT)strlen(part1), part1) == FAIL)
     exit(ERREXIT);
if (bcp_moretext(dbproc, (DBINT)strlen(part2), part2) == FAIL)
     exit(ERREXIT);
if (bcp_moretext(dbproc, (DBINT)strlen(part3), part3) == FAIL)
     exit(ERREXIT);

/* We're all done. */
bcp_done(dbproc);
dbclose(dbproc);
```

- If you use **bcp_moretext** to send one text or image column in the row, you must also use it to send all other text and image columns in the row.
- If the row contains more than one text or image column, **bcp_moretext** will first send its data to the lowest-numbered (that is, leftmost) text or image column, followed by the next lowest-numbered column, and so on.
- An application will normally call **bcp_sendrow** and **bcp_moretext** within loops, to send a number of rows of data. Here is an outline of how to do this for a table containing two text columns:

```
while (there are still rows to send)
 {
     bcp_sendrow(...);
```

```
            for (all the data in the first text column)
                bcp_moretext(...);

            for (all the data in the second text column)
                bcp_moretext(...);
        }
```

See also bcp_bind, bcp_sendrow, dbmoretext, dbwritetext

bcp_options

Description Set bulk copy options.

Syntax RETCODE bcp_options (dbproc, option, value, valuelen)

```
DBPROCESS    *dbproc;
BYTE          *value;
int           valuelen;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

value

A generic BYTE pointer to the value of the specified option. As the following table describes, what value should point to depends on *option*:

Table 3-1: Values for value (bcp_options)

If option Is	*value should be
BCPLABELED	A DBBOOL value. Set **value* to TRUE to enable a bulk copy with sensitivity labels. Set **value* to FALSE for a normal bulk copy operation.

valuelen

The length of the data to which *value* points. If value points to a fixed-length item (for example a DBBOOL, DBINT, and so on), pass *valuelen* as -1.

Return value SUCCEED or FAIL.

Usage

- **bcp_options** sets bulk copy options.
- Currently the only bulk copy option available is BCPLABELED.

See also bcp_init, bcp_control

bcp_readfmt

Description

Read a datafile format definition from a host file.

Syntax

RETCODE bcp_readfmt(dbproc, filename)

```
DBPROCESS    *dbproc;
char         *filename;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

filename

The full directory specification of the file containing the format definitions.

Return value

SUCCEED or FAIL.

Usage

- **bcp_readfmt** reads a datafile format definition from a host file, then makes the appropriate calls to **bcp_columns** and **bcp_colfmt**. This automates the bulk copy of multiple files that share a common data format.
- **bcp**, the bulk copy utility, copies a database table to or from a host file in a user-specified format. User-specified formats may be saved via **bcp** in datafile format definition files, which can later be used to automate the bulk copy of files that share a common format. See the *Open Client/Server Programmer's Supplement* for more information on the **bcp** utility and datafile format definition files.
- Application programs can call **bcp_writefmt** to create files with datafile format definitions.
- The following code fragment illustrates the use of **bcp_readfmt**:

```
bcp_init(dbproc, "mytable", "bcpdata", "bcperrs", DB_IN);
 bcp_readfmt(dbproc, "my_fmtfile");
 bcp_exec(dbproc, &rows_copied);
```

See also

bcp_colfmt, bcp_columns, bcp_writefmt

bcp_sendrow

Description

Send a row of data from program variables to Adaptive Server.

Syntax

```
RETCODE bcp_sendrow(dbproc)

 DBPROCESS     *dbproc;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

Return value

SUCCEED or FAIL.

Usage

- **bcp_sendrow** builds a row from program variables and sends it to Adaptive Server.
- Before calling **bcp_sendrow**, you must make calls to **bcp_bind** to specify the program variables to be used.
- If bcp_bind is called with a *type* parameter of SYBTEXT or SYBIMAGE and a non-null *varaddr* parameter, bcp_sendrow will send the entire text or image data value, just as it does for all other datatypes. If, however, bcp_bind has a null *varaddr* parameter, bcp_sendrow will return control to the application immediately after all non-text or image columns are sent to Adaptive Server. The application can then call bcp_moretext repeatedly to send the text and image columns to Adaptive Server, a chunk at a time. For an example, see the bcp_moretext reference page.
- After the last call to **bcp_sendrow**, you must call **bcp_done** to ensure proper internal cleanup.
- When **bcp_sendrow** is used to bulk copy rows from program variables into Adaptive Server tables, rows are permanently saved in Adaptive Server only when the user calls **bcp_batch** or **bcp_done**.

  The user may choose to call bcp_batch once every *n* rows, or when there is a lull between periods of incoming data (as in a telemetry application). Of course, the user may choose some other criteria or may decide not to call bcp_batch at all. If bcp_batch is never called, the rows are permanently saved in Adaptive Server when bcp_done is called.

See also

bcp_batch, bcp_bind, bcp_colfmt, bcp_collen, bcp_colptr, bcp_columns, bcp_control, bcp_done, bcp_exec, bcp_init, bcp_moretext

BCP_SETL

Description
Set the LOGINREC for bulk copy operations into the database.

Syntax
RETCODE BCP_SETL(loginrec, enable)

```
LOGINREC    *loginrec;
DBBOOL      enable;
```

Parameters

loginrec
This is a pointer to a LOGINREC structure, which will be passed as an argument to **dbopen**. You can get a LOGINREC structure by calling **dblogin**.

enable
This is a Boolean value (TRUE or FALSE) that specifies whether or not to enable bulk copy operations for the resulting DBPROCESS. By default, DBPROCESSes are *not* enabled for bulk copy operations.

Return value
SUCCEED or FAIL.

Usage

- This macro sets a field in the LOGINREC structure that tells Adaptive Server that the DBPROCESS connection may be used for bulk copy operations. To have any effect, it must be called before **dbopen**, the routine that actually allocates the DBPROCESS structure.
- Applications that allow users to make ad hoc queries may want to avoid calling BCP_SETL (or call it with a value of FALSE for the *enable* parameter) to prevent users from initiating a bulk copy sequence via SQL commands. Once a bulk copy sequence has begun, it cannot be stopped via an ordinary SQL command.
- **BCP_SETL** applies to “copy in” operations only.

See also
bcp_init, bcp_getl, dblogin, dbopen, DBSETLAPP, DBSETLHOST, DBSETLPWD, DBSETLUSER

bcp_setxlate

Description
Specify the character set translations to use when retrieving data from or inserting data into an Adaptive Server.

Syntax
RETCODE bcp_setxlate(dbproc, xlt_tosrv, xlt_todisp)

```
DBPROCESS    *dbproc;
```

```
DBXLATE     *xlt_tosrv;
DBXLATE     *xlt_todisp;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

xlt_tosrv

A pointer to a translation structure. The translation structure is allocated via **dbload_xlate**. *xlt_tosrv* indicates the character set translation to use when moving data from the application program to the Adaptive Server (the copy *in*, or DB_IN, direction).

xlt_todisp

A pointer to a translation structure. The translation structure is allocated via **dbload_xlate**. *xlt_todisp* indicates the character set translation to use when moving data from Adaptive Server to the application program (the copy *out*, or DB_OUT, direction).

Return value

SUCCEED or FAIL.

Usage

- **bcp_setxlate** specifies the character set translations to use when transferring character data between the Adaptive Server and a front-end application program via **bcp**.
- The specified character set translations need not be the same as those being used to display or input data on the user's terminal. The translations may be used to read or write a data file in a completely different character set that is not intended for immediate display.
- The following code fragment illustrates the use of **bcp_setxlate**:

```
bcp_init(dbproc, "mytable", "bcpdata", "bcperrs", DB_OUT);
 bcp_setxlate(dbproc, xlt_tosrv, xlt_todisp);
 bcp_columns(dbproc, 3);
 bcp_colfmt(dbproc, 1, SYBCHAR, 0, -1, "\t", 1, 1);
 bcp_colfmt(dbproc, 2, SYBCHAR, 0, -1, "\t", 1, 2);
 bcp_colfmt(dbproc, 3, SYBCHAR, 0, -1, "\n", 1, 3);
 bcp_exec(dbproc);
```

See also

dbfree_xlate, dbload_xlate, dbxlate

bcp_writefmt

Description — Write a datafile format definition to a host file.

Syntax — RETCODE bcp_writefmt(dbproc, filename)

```
DBPROCESS    *dbproc;
char         *filename;
```

Parameters

dbproc

A pointer to the DBPROCESS structure that provides the connection for a particular front-end/Adaptive Server process. It contains all the information that DB-Library uses to manage communications and data between the front end and Adaptive Server.

filename

The full directory specification of the file that contains the format definitions.

Return value — SUCCEED or FAIL.

Usage

- **bcp_writefmt** writes a datafile format definition to a host file. The format reflects previous calls to **bcp_columns** and **bcp_colfmt**.
- **bcp**, the bulk copy utility, copies a database table to or from a host file in a user-specified format. User-specified formats may be saved via **bcp** in "datafile format definition files," which can later be used to automate the bulk copy of files that share a common format. See the *Open Client/Server Programmer's Supplement* for more information on the **bcp** utility and datafile format definition files.
- Format definition files are read via **bcp_readfmt**.
- The following code fragment illustrates the use of **bcp_writefmt**:

```
bcp_init(dbproc, "mytable", "bcpdata", "bcperrs", DB_OUT);

 bcp_columns(dbproc, 3);
 bcp_colfmt(dbproc, 1, SYBCHAR, 0, -1, "\t", 1, 1);
 bcp_colfmt(dbproc, 2, SYBCHAR, 0, -1, "\t", 1, 2);
 bcp_colfmt(dbproc, 3, SYBCHAR, 0, -1, "\n", 1, 3);

 bcp_writefmt(dbproc, "my_fmtfile");
 bcp_exec(dbproc, &rows_copied);
```

See also — bcp_colfmt, bcp_columns, bcp_readfmt

CHAPTER 4

Two-phase commit service

Adaptive Server provides a two-phase commit service that allows a client application to coordinate transactions that are distributed on two or more SQL Servers.

This chapter describes the two-phase commit process and the DB-Library routines that are involved.

Topic	Page
Programming distributed transactions	449
The commit service and the application program	450
The probe process	452
Two-phase commit routines	452
Specifying the commit server	454
Two-phase commit example program	454
Program Notes	461

Programming distributed transactions

The two-phase commit service allows an application to coordinate updates among two or more SQL Servers. This initial implementation of distributed transactions treats separate transactions (which may be on separate SQL Servers) as if they were a single transaction. The commit service uses one Adaptive Server, the "commit server," as a central record-keeper that helps the application determine whether to commit, or whether to roll back transactions in case of failure. Thus, the two-phase commit guarantees that either all or none of the databases on the participating servers are updated.

A distributed transaction is performed by submitting Transact-SQL statements to the SQL Servers via DB-Library routines. An application program opens a session with each server, issues the update commands, and then prepares to commit the transaction. Through DB-Library, the application issues the following to each participating server:

- A **begin transaction** with identifying information on the application, the transaction, and the commit server
- The Transact-SQL update statements
- A **prepare transaction** statement that indicates that the updates have been performed and that the server is prepared to commit

After the updates have been performed on all the servers participating in the distributed transaction, the two-phase commit begins. In the first phase, all servers agree that they are ready to commit. In the second phase, the application informs the commit service that the transaction is complete (that is, the commit will take place), and a **commit transaction** is then issued to all of the servers, causing them to commit.

If an error occurs between phase one and phase two, all servers coordinate with the commit service to determine whether the transaction should be committed or aborted.

Note If certain types of errors occur during a two-phase transaction, Adaptive Server may need to mark a two-phase process as "infected." Marking the process as infected rather than killing it aids in later error recovery. To ensure that Adaptive Server is able to mark processes as infected, boot Adaptive Server with the flag -T3620 passed on the command line.

The commit service and the application program

The role of the commit service is to be a single place of record that helps the application decide whether the transaction should be committed or aborted.

If the SQL Servers are all prepared to commit, the application notifies the commit service to mark the transaction as committed. Once this happens, the transaction is committed despite any failures that might subsequently happen.

If any Adaptive Server or the application program fails before the **prepare transaction** statement, the Adaptive Server will **rollback** the transaction.

If any Adaptive Server or the application program fails after the **prepare** but before the **commit**, the Adaptive Server will communicate with the server functioning as the commit service and ask it whether to **rollback** or **commit**.

If the Adaptive Server cannot communicate with the server functioning as the commit service, it will mark the user task process as infected in Adaptive Server. At this point, the System Administrator can either kill the infected process immediately, or wait until communication to the commit service is restored to kill the infected process.

- If the System Administrator kills the infected process immediately, two-phase commit protocol is violated and the integrity of the two-phase transaction is not guaranteed. Servers participating in the transaction may be in inconsistent states.
- If the System Administrator kills the infected process after communication with the commit service has been restored, the Adaptive Server will communicate with the commit service to determine whether or not to commit the transaction locally. The integrity of the two-phase transaction is guaranteed.

To decide whether or not to kill the infected process immediately, the System Administrator must consider the estimated downtime of the commit service, the number and importance of locks held by the infected process, and the complexity of the transaction in progress.

The role of the application program is to deliver the Transact-SQL statements to the SQL Servers in the proper order, using the proper DB-Library routines. The role of the commit service is to provide a single place where the commit/rollback status is maintained. The SQL Servers communicate with the commit service only if a failure happens during the two-phase commit.

The commit service needs its own DBPROCESS, separate from the DBPROCESSes used for the distributed transaction, to perform its record-keeping. Note, however, that the server handling the commit service can also be one of the servers participating in the transaction, as long as the commit service has its own DBPROCESS. In fact, all the servers involved in the transaction can be one and the same.

The probe process

If any server must recover the transaction, it initiates a process, **probe**, that determines the last known status of the transaction. After it returns the status of that transaction to the commit service, the **probe** process dies. The **probe** process makes use of **stat_xact**, the same status-checking routine that the commit service uses to check the progress of a distributed transaction.

Note SQL Server version 10.0 is not compatible with pre-10.0 **probe** versions.

Two-phase commit routines

The following routines make up the two-phase commit service:

Routine	Description	Page
abort_xact	Tells the commit service to abort the transaction.	464
build_xact_string	Builds a name string for use by each participating Adaptive Server for its **begin transaction** and **prepare transaction** statements. This string encodes the application's transaction name, the commit service name, and the *commid*.	465
close_commit	Closes the connection with the commit service.	466
commit_xact	Tells the commit service to commit the transaction.	466
open_commit	Opens a connection with the commit service. The routine is given the login ID of the user initiating the session and the name of the commit service. It returns a pointer to a DBPROCESS structure used in subsequent commit service calls.	467
remove_xact	Decrements the count of servers still participating in the transaction.	467
start_xact	Records the start of a distributed transaction and stores initial information about the transaction (DBPROCESS id, application name, transaction name, and number of sites participating) in a lookup table on the commit server. It returns the *commid* identifying number for the transaction.	469

Two additional routines are used for ongoing status reports:

Routine	Description	Page
scan_xact	Returns the status of a single transaction or all distributed transactions.	468
stat_xact	Returns the completion status of a distributed transaction.	470

During the course of a session, the diagnostic routines **scan_xact** and **stat_xact** are used to check that the commit service carried out the request.

The **scan_xact** routine uses the commit service lookup table, **spt_committab**, which holds the following values:

- Transaction ID
- Time the transaction started
- Last time the row was updated
- Number of servers initially involved in the transaction
- Number of servers that have not yet completed
- Status: “a” (abort), “c” (commit), “b” (begin)
- Application name
- Transaction name

The two-phase commit routines call internal stored procedures (for example, **sp_start_xact**) that are created in each server’s master database. The **installmaster** script creates the commit service lookup table and stored procedures in each server’s master database, for use whenever that server becomes a commit server.

Specifying the commit server

The commit server must have an entry in the interfaces file on each machine participating in the distributed transaction. On the machine on which the commit server is actually running, the commit server entry must specify the usual ports described in the *Open Client/Server Configuration Guide*, including a query port. For example:

```
SERVICE
 master tcp sun-ether rose 2001
 query tcp sun-ether rose 2001
```

On any additional machines containing other servers participating in the distributed transaction, the commit server entry need specify only the query port:

```
SERVICE
 query tcp sun-ether rose 2001

 SITEA
 master tcp sun-ether gaia 2011
 query tcp sun-ether gaia 2011
```

The name of the commit server (in these examples, “SERVICE”) is used as a parameter in calls to the **open_commit** and **build_xact_string** routines. The commit server name must be the same on all machines participating in the transaction. The name cannot contain a period (.) or a colon (:).

Two-phase commit example program

An example program illustrating the two-phase commit service is included with DB-Library’s online sample programs. This same example is duplicated below, but with comments added to document how recovery occurs for the different types of failure that may occur at various points in the transaction.

```
/*
 **    twophase.c
 **
 **    Demo of Two-Phase Commit Service
 **
```

```
**    This example uses the two-phase commit service
**    to perform a simultaneous update on two servers.
**    In this example, one of the servers participating
**    in the distributed transaction also functions as
**    the commit service.
**
**    In this particular example, the same update is
**    performed on both servers. You can, however, use
**    the commit server to perform completely different
**    updates on each server.
**
*/

#include <stdio.h>
#include <sybfront.h>
#include <sybdb.h>
#include "sybdbex.h"

int err_handler();
int msg_handler();

char     cmdbuf[256];
char     xact_string[128];

main()
{

DBPROCESS     *dbproc_server1;
DBPROCESS     *dbproc_server2;
DBPROCESS     *dbproc_commit;
LOGINREC      *login;
int           commid;

RETCODE       ret_server1;
RETCODE       ret_server2;

/* Initialize DB-Library. */
if (dbinit() == FAIL)
exit(ERREXIT);


     dberrhandle(err_handler);
     dbmsghandle(msg_handler);

     printf("Demo of Two Phase Commit\n");
```

```
/* Open connections with the servers and the
** commit service. */
login = dblogin();
DBSETLPWD(login, "server_password");
DBSETLAPP(login, "twophase");

dbproc_server1 = dbopen (login, "SERVICE");
dbproc_server2 = dbopen (login, "PRACTICE");
dbproc_commit = open_commit (login, "SERVICE");

if (dbproc_server1 == NULL ||
    dbproc_server2 == NULL ||
    dbproc_commit  == NULL)
{
     printf (" Connections failed!\n");
     exit (ERREXIT);
}

/* Use the "pubs2" database. */
sprintf(cmdbuf, "use pubs2");
dbcmd(dbproc_server1, cmdbuf);
dbsqlexec(dbproc_server1);
dbcmd(dbproc_server2, cmdbuf);
dbsqlexec(dbproc_server2);

/*
** Start the distributed transaction on the
** commit service.
*/
commid = start_xact(dbproc_commit, "demo", "test", 2);
```

Note The application is now in the *begin* phase of the two-phase commit transaction.

```
/* Build the transaction name. */
 build_xact_string ("test", "SERVICE", commid, xact_string);

 /* Build the first command buffer. */
 sprintf(cmdbuf, "begin transaction %s", xact_string);

 /* Begin the transactions on the different servers. */
 dbcmd(dbproc_server1, cmdbuf);
 dbsqlexec(dbproc_server1);
 dbcmd(dbproc_server2, cmdbuf);
 dbsqlexec(dbproc_server2);
```

```
/* Do various updates. */
sprintf(cmdbuf, " update titles set price = $1.50 where");
strcat(cmdbuf, " title_id = 'BU1032'");
dbcmd(dbproc_server1, cmdbuf);
ret_server1 = dbsqlexec(dbproc_server1);
dbcmd(dbproc_server2, cmdbuf);
ret_server2 =dbsqlexec(dbproc_server2);
```

Note See "Program note 1" on page 461.

```
if (ret_server1 == FAIL || ret_server2 == FAIL)
 {
      /* Some part of the transaction failed. */
      printf(" Transaction aborted -- dbsqlexec failed\n");
      abortall(dbproc_server1, dbproc_server2,
               dbproc_commit, commid);
 }

 /* Find out if all servers can commit the transaction. */
 sprintf(cmdbuf, "prepare transaction");
 dbcmd(dbproc_server1, cmdbuf);
 dbcmd(dbproc_server2, cmdbuf);
 ret_server1 = dbsqlexec(dbproc_server1);
```

Note See "Program note 2" on page 461.

```
ret_server2 = dbsqlexec(dbproc_server2);
```

Note See "Program note 3" on page 461.

```
if (ret_server1 == FAIL || ret_server2 == FAIL)
 {
      /* One or both of the servers failed to prepare. */
      printf(" Transaction aborted -- prepare failed\n");
      abortall(dbproc_server1, dbproc_server2,
               dbproc_commit, commid);
 }
```

Note See "Program note 4" on page 461.

```
    /* Commit the transaction. */
     if (commit_xact(dbproc_commit, commid) == FAIL)
     {
          /* The commit server failed to record the commit. */
          printf( " Transaction aborted -- commit_xact failed\n");
          abortall(dbproc_server1, dbproc_server2,
                   dbproc_commit, commid);
          exit(ERREXIT);
     }
```

Note See “Program note 5” on page 462.

```
    /* The transaction has successfully committed.
     ** Inform the servers.
     */
     sprintf(cmdbuf, "commit transaction");
     dbcmd(dbproc_server1, cmdbuf);
     if (dbsqlexec(dbproc_server1) != FAIL)
          remove_xact(dbproc_commit, commid, 1);
```

Note See “Program note 6” on page 463.

```
    dbcmd(dbproc_server2, cmdbuf);
     if (dbsqlexec(dbproc_server2) != FAIL)
     remove_xact(dbproc_commit, commid, 1);
```

Note See “Program note 7” on page 463.

```
    /* Close the connection to the commit server. */
     close_commit(dbproc_commit);
```

Note See “Program note 8” on page 464.

```
    printf( "We made it!\n");
     dbexit();
     exit(STDEXIT);
}

/* Function to abort the distributed transaction. */

abortall( dbproc_server1, dbproc_server2, dbproc_commit, commid )
DBPROCESS      *dbproc_server1;
```

```
DBPROCESS      *dbproc_server2;
DBPROCESS      *dbproc_commit;
int            commid;
{
     /* Some part of the transaction failed. */

     /* Inform the commit server of the failure. */
     abort_xact(dbproc_commit, commid);

     /* Roll back the transactions on the different servers. */
     sprintf(cmdbuf, "rollback transaction");
     dbcmd(dbproc_server1, cmdbuf);
     if (dbsqlexec(dbproc_server1) != FAIL)
          remove_xact(dbproc_commit, commid, 1);
     dbcmd(dbproc_server2, cmdbuf);
     if (dbsqlexec(dbproc_server2) != FAIL)
          remove_xact(dbproc_commit, commid, 1);

     dbexit();
     exit(ERREXIT);
}

/* Message and error handling functions. */

int msg_handler(dbproc, msgno, msgstate, severity, msgtext,
                servername, procname, line)

DBPROCESS       *dbproc;
DBINT           msgno;
int             msgstate;
int             severity;
char            *msgtext;
char            *servername;
char            *procname;
DBUSMALLINT     line;

{
     /*   Msg 5701 is just a use database message, so skip it. */
     if (msgno == 5701)
          return (0);

     /* Print any severity 0 message as is, without extra stuff. */
     if (severity == 0)
     {
          (void) fprintf (ERR_CH, "%s\n",msgtext);
          return (0);
```

```
     }

     (void) fprintf (ERR_CH, "Msg %ld, Level %d, State %d\n",
             msgno, severity, msgstate);

     if (strlen(servername) > 0)
          (void) fprintf (ERR_CH, "Server '%s', ", servername);
     if (strlen(procname) > 0)
          (void) fprintf (ERR_CH, "Procedure '%s', ", procname);
     if (line > 0)
          (void) fprintf (ERR_CH, "Line %d", line);

     (void) fprintf (ERR_CH, "\n\t%s\n", msgtext);

     if (severity >= 16)
     {
          (void) fprintf (ERR_CH, "Program Terminated! Fatal\
                          SQL Server error.\n");
          exit(ERREXIT);
     }

     return (0);
}

int err_handler(dbproc, severity, dberr, oserr, dberrstr, oserrstr)
DBPROCESS    *dbproc;
int          severity;
int          dberr;
int          oserr;
char         *dberrstr;
char         *oserrstr;
{
     if ((dbproc == NULL) || (DBDEAD(dbproc)))
          return (INT_EXIT);
     else
     {
          (void) fprintf (ERR_CH, "DB-Library error: \
                          \n\t %s\n", dberrstr);
          if (oserr != DBNOERR)
               (void) fprintf (ERR_CH, "Operating system error:\
                               \n\t%s\n", oserrstr);
     }

     return (INT_CANCEL);
}
```

Program Notes

This section contains the notes referenced in the sample code.

Program note 1

If any type of failure occurs at this point, it is the application's responsibility to roll back the transactions using **abort_xact**.

Program note 2

The application has entered the *prepare* stage of the two-phase commit transaction. As far as the commit server is aware, however, the application is still in the *begin* phase.

Program note 3

If any type of failure occurs at this point, it is the application's responsibility to roll back the transactions using **abort_xact**.

Program note 4

At this point, the following failures are possible:

- The application's link to the commit server, or the commit server itself, may go down.

  In this case, the following call to **commit_xact** will fail, and the application must roll back the transactions via **abort_xact**.

- The application's link to a participating server may go down.

  In this case, the following call to **commit_xact** will succeed, but the application's **commit transaction** command to the participating server will not. However, the server will be aware that its connection with the application has died. It will communicate with the commit server, using **probe**, to determine whether to commit the transaction locally.

- A participating server may go down.

In this case, the following call to **commit_xact** will succeed, but the application's **commit transaction** to the participating server will not. When the participating server comes back up, it will use **probe** to determine whether to commit the transaction locally.

- Both the application's link to the commit server and the application's link to the participating server may go down.

  In this case, the following call to **commit_xact** will fail. The application must roll back the transactions with **abort_xact**, but will not be able to communicate with the participating server. The participating server will use **probe** to communicate with the commit server. It will learn that the transaction has not been committed in the commit service, and will roll back the transaction locally.

- Both the application's link to the participating server and the participating server's link to the commit server may go down.

  In this case, the following call to **commit_xact** will succeed, but the application will not be able to communicate this to the participating server. When its connection to the application dies, the participating server will attempt to communicate with the commit server via **probe** to determine whether or not to commit the transaction locally. Because its link to the commit server is down, however, it will not be able to.

  Because it cannot resolve the transaction, the participating server marks the user task process as infected.

  If the System Administrator kills the infected process while the commit server is still down, two-phase commit protocol is violated and the integrity of the transaction is not guaranteed.

  If the System Administrator waits until commit server is back up to kill the infected process, **probe** executes automatically when the System Administrator attempts to kill the process. **probe** communicates with the commit server and determines whether the participating server should commit the transaction locally. The integrity of the transaction is guaranteed.

Program note 5

The application has entered the *committed* phase of the two-phase commit transaction. This means that any **probe** process querying the commit server will be told to commit the transaction locally. After this point, the application does not need to concern itself with aborting the transaction.

Program note 6

If the above **dbsqlexec** to server1 fails because the application's link to the server has gone down, server1 will use **probe** to communicate with the commit server. **probe** will find that the transaction is committed in the commit server and will tell server1 to commit locally.

If **probe** cannot communicate with the commit server, server1 will infect the user task process in Adaptive Server. If the System Administrator kills the infected process before communication with the commit server is reestablished, the transaction will be rolled back, thus violating two-phase protocol and leaving the database in an inconsistent state. If possible, the System Administrator should always wait until communication with the commit server is reestablished before killing the infected process.

If the **dbsqlexec** to server1 fails because server1 has gone down, the local transaction will remain in limbo until server 1 is restored. As part of the recovery process, server1 will use **probe** to communicate with the commit server. **probe** will find that the transaction is committed in the commit server and will tell server1 to commit locally.

If **probe** cannot communicate with the commit server, server1 will mark the database as suspect. After communication with the commit sever is reestablished, the suspect database should be re-recovered.

Program note 7

If the above **dbsqlexec** to server2 fails because the application's link to the server has gone down, server2 will use **probe** to communicate with the commit server. **probe** will find that the transaction is committed in the commit server and will tell server2 to commit locally.

If **probe** cannot communicate with the commit server, server2 will infect the user task process in Adaptive Server. If the System Administrator kills the infected process before communication with the commit server is reestablished, the transaction will be rolled back, thus violating two-phase protocol and leaving the database in an inconsistent state. If possible, the System Administrator should always wait until communication with the commit server is reestablished before killing the infected process.

If the **dbsqlexec** to server2 fails because server2 has gone down, the local transaction will remain in limbo until server 2 is restored. As part of the recovery process, server2 will use **probe** to communicate with the commit server. **probe** will find that the transaction is committed in the commit server and will tell server2 to commit locally.

If **probe** cannot communicate with the commit server, server2 will mark the database as suspect. After communication with the commit sever is reestablished, the suspect database should be re-recovered.

Program note 8

close_commit marks the transaction as complete in the **spt_committab** table on the commit server. If **close_commit** fails, the transaction is not marked as complete. No actual harm is done by this, but the System Administrator may choose to manually update **spt_committab** in this case.

abort_xact

Description — Mark a distributed transaction as being aborted.

Syntax — RETCODE abort_xact(connect, commid)

```
DBPROCESS    *connect;
DBINT        commid;
```

Parameters

connect
A pointer to the DBPROCESS used to communicate with the commit service.

commid
The *commid* used to identify the transaction to the commit service.

Return value — SUCCEED or FAIL.

Usage

- This routine informs the commit service that the status of a distributed transaction should be changed from “begin” to “abort.”

See also — commit_xact, remove_xact, scan_xact, start_xact, stat_xact

build_xact_string

Description

Build a name for a distributed transaction.

Syntax

```
void build_xact_string(xact_name, service_name,
                commid, result)

 char        *xact_name;
 char        *service_name;
 DBINT       commid;
 char        *result;
```

Parameters

xact_name

The application or user name for the transaction. This name gets encoded in the name string but is not used by the commit service or Adaptive Server. It serves to identify the transaction for debugging purposes.

service_name

The name that will be used by Adaptive Server to contact the commit service, should it be necessary to recover the transaction. If *service_name* is NULL, the name DSCOMMIT is used.

service_name must correspond to name of the interfaces file entry for the commit service. If *service_name* is NULL, the interfaces file must contain an entry for DSCOMMIT.

commid

The number used by the commit service to identify the transaction. *commid* is the number returned by the call to **start_xact**.

result

Address of buffer where the string should be built. The space must be allocated by the caller.

Return value

None.

Usage

- This routine builds a name string for use in the SQL **begin transaction** and **prepare transaction** of an Adaptive Server transaction. If Adaptive Server has to recover the transaction, it uses information encoded in the name to determine which commit service to contact and which transaction in that service to inquire about. The application should issue a SQL **begin transaction** using the string built by **build_xact_string**.
- The string built by **build_xact_string** must be large enough to hold the ASCII representation of *commid*, *xact_name*, *service_name*, two additional characters, and a null terminator.

See also

commit_xact, start_xact

close_commit

Description
End a connection with the commit service.

Syntax
```
void close_commit(connect)

 DBPROCESS     *connect;
```

Parameters
connect
A pointer to the DBPROCESS structure that was originally returned by **open_commit**.

Return value
None.

Usage
- This routine calls **dbclose** to end a connection with the commit service. A call to **close_commit** should be made when the application is through with the commit service, to free resources.

See also
dbclose

commit_xact

Description
Mark a distributed transaction as being committed.

Syntax
```
RETCODE commit_xact(connect, commid)

 DBPROCESS     *connect;
 DBINT         commid;
```

Parameters
connect
A pointer to the DBPROCESS used to communicate with the commit service.

commid
The *commid* used to identify the transaction to the commit service.

Return value
SUCCEED or FAIL.

If **commit_xact** fails, you *must* roll back the transaction.

Usage
- This routine informs the commit service that the status of a distributed transaction should be changed from “begin” to “commit.”

See also
abort_xact, remove_xact, scan_xact, start_xact, stat_xact

open_commit

Description
Establish a connection with the commit service.

Syntax
```
DBPROCESS *open_commit(login, servername)

LOGINREC    *login;
char        *servername;
```

Parameters

login
This is a LOGINREC containing information about the user initiating the session, such as login name, password, and options desired. The LOGINREC must have been obtained from a prior call to the DB-Library routine **dblogin**. The caller may wish to initialize fields in the LOGINREC. See the reference page for dblogin for more details.

servername
The name of the commit service; for example, DSCOMMIT_SALESNET. If *servername* is NULL, the name DSCOMMIT is used. The name cannot contain a period (.) or a colon (:).

Return value
A pointer to a DBPROCESS structure to be used in subsequent commit service calls, or NULL if the open failed.

Usage

- This routine calls **dbopen** to establish a connection with the commit service. A call to **open_commit** must precede any calls to other commit service routines, such as **start_xact**, **commit_xact**, **abort_xact**, **remove_xact**, and **scan_xact**. A session with the commit service is closed by calling **close_commit**.
- This routine returns a DBPROCESS structure, which is used to communicate with the commit service. The DBPROCESS must be dedicated to its role with the commit service and should not be used otherwise in the distributed transaction.

See also
dblogin, dbopen

remove_xact

Description
Decrement the count of sites still active in the distributed transaction.

Syntax
```
RETCODE remove_xact(connect, commid, n)

DBPROCESS    *connect;
```

```
DBINT          commid;
int            n;
```

Parameters

connect
A pointer to the DBPROCESS used to communicate with the commit service.

commid
The *commid* used to identify the transaction to the commit service.

n
The number of sites to remove from the transaction.

Return value

SUCCEED or FAIL.

Usage

- The commit service keeps a count of the number of sites participating in a distributed transaction. This routine informs the commit service that one or more sites has done a local commit or abort on the transaction and is hence no longer participating. The commit service removes the sites from the transaction by decrementing the count of sites.
- The transaction record is deleted entirely if the count drops to 0.

See also

abort_xact, commit_xact, scan_xact, start_xact, stat_xact

scan_xact

Description

Print commit service record for distributed transactions.

Syntax

```
RETCODE scan_xact(connect, commid)

DBPROCESS    *connect;
DBINT        commid;
```

Parameters

connect
A pointer to the DBPROCESS used to communicate with the commit service.

commid
The *commid* used to identify the transaction to the commit service. If *commid* is -1, all commit service records are displayed.

Return value

SUCCEED or FAIL.

Usage

- This routine displays the commit service record for a specific distributed transaction, or for all distributed transactions known to the commit service.

See also abort_xact, commit_xact, remove_xact, start_xact, stat_xact

start_xact

Description Start a distributed transaction using the commit service.

Syntax

```
DBINT start_xact(connect, application_name, xact_name,
        site_count)

DBPROCESS    *connect;
char         *application_name;
char         *xact_name;
int          site_count;
```

Parameters

connect
A pointer to the DBPROCESS used to communicate with the commit service.

application_name
The name of the application. This name can be anything the application developer chooses. It will appear in the table maintained by the commit service but is not used by the commit service or the Adaptive Server recovery system.

xact_name
The name of the transaction. This name will appear in the table maintained by the commit service and must be supplied as part of the transaction name string built by **build_xact_string**. The name cannot contain a period (.) or a colon (:).

site_count
The number of sites participating in the transaction.

Return value An integer called the *commid*. This number is used to identify the transaction in subsequent calls to the commit service. In case of error, this routine will return 0.

Usage

- This routine records the start of a distributed transaction with the commit service. A record is placed in the commit service containing the *commid*, which is a number that caller subsequently uses to identify the transaction to the commit service.

See also abort_xact, build_xact_string, commit_xact, remove_xact, scan_xact, stat_xact

stat_xact

Description

Return the current status of a distributed transaction.

Syntax

int stat_xact(connect, commid)

```
DBPROCESS    *connect;
DBINT        commid;
```

Parameters

connect

A pointer to the DBPROCESS used to communicate with the commit service.

commid

The *commid* is used to identify the transaction to the commit service. If *commid* is -1, all commit service records are displayed.

Return value

A character code: “a” (abort), “b” (begin), “c” (commit), “u” (unknown), or -1 (request failed).

Usage

- This routine returns the transaction status for the specified distributed transaction.

See also

abort_xact, commit_xact, remove_xact, scan_xact, start_xact

APPENDIX A

Cursors

This appendix introduces the DB-Library cursor.

Topic	Page
Cursor overview	471
Sensitivity to change	473
DB-Library cursor functions	476
Holding locks	476
Stored procedures used by DB-Library cursors	477

Cursor overview

Because relational databases are oriented toward sets, no concept of next row exists, meaning that you cannot operate on an individual row in a set. Cursor functionality solves this problem by letting a result set be processed one row at a time, similar to the way you read and update a file on a disk. A DB-Library cursor indicates the current position in a result set, just as the cursor on your screen indicates the current position in a block of text.

DB-Library cursors are *client-side* cursors. This means that they do not correspond to an Adaptive Server cursor, but emulate a cursor that appears to the user to be in the server. The DB-Library cursor transparently does keyset management, row positioning, and concurrency control entirely on the client side.

DB-Library cursor capability

The DB-Library cursor routines offer the following capabilities, with certain limitations:

- Forward and backward scrolling (depending on how the keyset is defined during **dbcursoropen**)

- Direct access by position in the result set
- Positioned updates (even if the result set was defined with **order by**)
- Sensitivity adjustments to changes made by other users
- Concurrency control through several options

Differences between DB-Library cursors and browse mode

Cursors let the user scroll through and update a result set with fewer restrictions than browse mode. Although cursors require a unique index, they do not require a timestamp nor a second connection to a database for updates. Also, they do not create a copy of the entire result set. The following table summarizes these differences:

Table A-1: Cursors and browse mode

Item	Cursors	Browse mode
Row timestamps	Not required	Required
Multiple connections for updates	Unnecessary	Necessary
Table usage	Use original tables	Uses a copy of tables

Differences between DB-Library and Client-Library cursors

A DB-Library cursor does not correspond to an actual Adaptive Server cursor. Instead, at the time the cursor is declared with **dbcursoropen**, DB-Library fetches keysets from Adaptive Server "under the covers." It then builds qualifiers based on the keys for the current row and sends them to Adaptive Server. The server parses the query and returns a result set. When **dbcursorfetch** is called to retrieve more data, the DB-Library cursor may have to do additional selects. In addition, Adaptive Server may have to parse the query each time **dbcursorfetch** is called.

A Client-Library cursor corresponds to an actual cursor in Adaptive Server. It is sometimes referred to, therefore, as a *native* cursor. A new TDS protocol allows Client-Library to interact with the server to manage the cursor.

A Client-Library cursor is faster than a DB-Library cursor because it does not have to send SQL commands to the server, which causes multiple re-parsing of the query. But because the result set remains on the server side, it cannot offer the same options for concurrency control as a DB-Library cursor.

The following table summarizes these and additional differences between the two cursors:

Table A-2: Differences between DB-Library cursors and Client-Library cursors

DB-Library cursor	Client-Library cursor
Cursor row position is defined by the client.	Cursor row position is defined by the server.
Can define optimistic concurrency control (allows dirty reads).	Cannot define optimistic concurrency control (does not allow dirty reads).
Can fetch backward (if CUR_KEYSET or CUR_DYNAMIC is specified for *scrollopt* during **dbcursoropen**).	Can only fetch forward.
More memory may be required if you query very large row sizes, unless you specify a smaller number of rows in the fetch buffer during **dbcursoropen**.	More memory is not required, regardless of how large the row sizes are.
You cannot access an Open Server application unless the application installs the required DB-Library stored procedures.	You can access a version 10.0 (or later) Open Server application that is coded to support cursors.
Slower performance.	Faster performance.

Sensitivity to change

Three broad categories identify cursors according to their sensitivity to change:

- *Static* – Values, order, and membership in the result set do not change while the cursor is open.
- *Keyset-Driven* – Values can change, but order and membership in the result set remain fixed at *open time* (the moment the cursor is opened).
- *Dynamic* – Values, order, and membership in the result set can all change.

Static cursor

In a static cursor, neither the cursor owner nor any other user can change the result set while the cursor is open. Values, membership, and order remain fixed until the cursor is closed. You can either take a snapshot of the result set (which begins to diverge from the snapshot as updates are made), or you can lock the entire result set to prevent updates.

It is not necessary for cursor routines to support static cursors directly. You can achieve static behavior through one of the following methods:

- Take a snapshot copy of the result set (with **select...into**) and then call **dbcursoropen** against the snapshot (temporary table).
- Lock the result set by calling **dbcursoropen** with the **holdlock** keyword in a **select** statement. This method, however, significantly reduces concurrency.

Keyset-driven cursor

In a keyset-driven cursor, the order and the membership of rows in the result set are fixed at open time, but changes to values may be made by the cursor owner. Committed changes made by other users are visible. If a change affects a row's order, or results in a row no longer qualifying for membership, the row does not disappear or move unless the cursor is closed and reopened. If the cursor remains open, deleted rows, when accessed, return a special error code that says they are missing. Updating the key also causes the rows to be "missing."

Inserted data does not appear, but changes to existing data do appear when the buffer is refreshed. With or without **order by**, the user can access rows by either *relative* or *absolute position*.

To access a row by relative position, move the cursor relative to its current position. For example, if the cursor is on row three and you want to access row eight, tell the cursor to jump five rows relative to its current position. The cursor jumps five rows to row eight.

To access a row by absolute position, tell the cursor the number of the row you want to access. For example, if the cursor is on row three and you want to access row eight, tell the cursor to jump to row eight.

Dynamic cursor

In a dynamic cursor, uncommitted changes made by the cursor owner and committed changes made by anyone become visible the next time the user scrolls. Changes include inserts and deletes as well as changes in order and membership. (Deleted rows do not leave holes.) The user can access rows by relative (but not absolute) position in the result set. Dynamic cursors cannot use an **order by** clause.

Concurrency control

Cursors control, through several options, *concurrent access*. Concurrent access occurs when more than one user accesses and updates the same data at the same time. During concurrent access, data can become unreliable without some kind of control. To activate the particular concurrency control desired, specify one of the following options when you open a cursor:

Table A-3: Concurrency control options

Option	Result
CUR_READONLY	Updates are not permitted.
CUR_LOCKCC	The set of rows currently in the client buffer is locked when they are fetched inside a user-initiated transaction. No other user can update or read these rows. Updates issued by the cursor owner are guaranteed to succeed. No locks are held unless the application first issues **begin transaction**. Locks are held until the application issues a **commit transaction**. Locks are not automatically released when the next fetch is executed.
CUR_OPTCC and CUR_OPTCCVAL-	Rows currently in the buffer are not locked, and other users can update or read them freely.

To detect collisions between updates issued by the cursor owner and those issued by other users, cursors save and compare timestamps or column values. Therefore, if you specify either of the optimistic concurrency control options (CUR_OPTCC or CUR_OPTCCVAL) your updates may fail because of collisions with other updates. You may want to design the application to refresh the buffer and then retry updates that fail.

The two optimistic concurrency control options differ in the way they detect collisions:

Table A-4: Detecting concurrency collisions

Option	Method of Detection
CUR_OPTCC	Optimistic concurrency control based on timestamp values. Compares timestamps if available; otherwise, saves and compares the value of all non-text, non-image columns in the table with their previous values.
CUR_OPTCCVAL	Optimistic concurrency control based on values. Compares selected values whether or not a timestamp is available.

DB-Library cursor functions

The following list summarizes the DB-Library cursor routines:

Routine	Description
dbcursoropen	Declares and opens the cursor, specifies the size of the fetch buffer and defines the keyset, and sets the concurrency control option.
dbcursorinfo	Returns the number of columns and the number of rows in the open cursor.
dbcursorcolinfo	Returns column information for the specified column number in the open cursor.
dbcursorbind	Associates program variables with columns.
dbcursorfetch	Scrolls the fetch buffer.
dbcursor	Updates, deletes, inserts, and refreshes the rows in the fetch buffer.
dbcursorclose	Closes the cursor.

For details about an individual routine, see its reference page.

Holding locks

To retain the flexibility of the Adaptive Server transaction model, cursors do not automatically issue **begin transaction** or **commit transaction**. The duration of locks acquired during cursor operations is entirely under the control of the application. In other words, an application that uses CUR_LOCKCC on either the **dbcursoropen** or **dbcursor** routine must also issue **begin transaction** for the locking to have any effect.

To hold the lock on the currently buffered rows when CUR_LOCKCC is used on **dbcursoropen**, the application must issue **commit transaction** and **begin transaction** before each **dbcursorfetch** that scrolls the local buffer (except for the very first **dbcursorfetch**, which should be preceded only by **begin transaction**).

To use the short-duration locking feature, issue **begin transaction** before locking the row to be updated with the CUR_LOCKCC option of **dbcursor**. If each update is independent, issue **commit transaction** after each update. If multiple updates to the same screenful of data depend on each other, issue **commit transaction** when the screen is scrolled.

For repeatable-read consistency, specify **holdlock** in the **select** statement in **dbcursoropen**, and issue **begin transaction** before the first **dbcursorfetch**. Locks are obtained as the data is fetched and are retained until the application issues **commit transaction** or **rollback transaction**.

Although you can close and reopen a repeatable-read cursor, you can get the same effect with FETCH_FIRST.

Other combinations are possible as well. The important thing to remember is that locks are not held unless **begin transaction** is in effect. Locks acquired while **begin transaction** is in effect are held until a **commit transaction** or **rollback transaction** is issued.

Stored procedures used by DB-Library cursors

DB-Library's cursor routines call the Adaptive Server's catalog stored procedures to find out table formats and identify key columns.

For more information about the catalog stored procedures, see the *Adaptive Server Enterprise Reference Manual*.

A

abort_xact 464
aggregate operators
  returning for a compute column 70
application names
  setting in LOGINREC 317
applications
  DB-Library/C 7, 12
  DB-Library/C example 7
  gateway 252, 304
applications, compiling and linking. See Open Client/Server Programmer's Supplement xii
arithmetic exceptions 402, 403

B

batches
  command. See Command batches 86
bcp
  BCPLABELED option 443
  binding data 418
  changing allowable number of errors 433
  changing default data formats 433
  changing first row to copy 433
  changing last row to copy 433
  changing number of rows to copy 433
  changing program variable data address 431
  changing program variable data length 430
  character set translations for 446
  copying multiple files 444, 448
  default data formats 439
  enabling 446
  ending bulk copy from program variables 435
  executing 436
  host file format 423, 426
  initializing 438
  overriding default data formats 423, 426, 432
  reading format definitions 444
  routines for 36
  saving preceding rows in SQL Server 417
  and Secure SQL Server 443
  sending data from program variables 445
  sending text/image values 440
  setting LOGINREC for 446
  setting number of columns in host file 432
  setting options for 443
  specifying host file format 423, 426
  writing format definitions to a file 448
bcp_batch 417, 418
bcp_bind 418, 422
bcp_colfmt 423, 426
bcp_colfmt_ps 426, 430
bcp_collen 430, 431
bcp_colptr 431, 432
bcp_columns 432, 433
  and bcp_bind 420
bcp_control 433, 435
bcp_done 435, 436
  and bcp_bind 422
bcp_exec 436, 437
bcp_getl 437, 438
bcp_init 438, 440
bcp_moretext 440, 443
bcp_options 443
bcp_readfmt 444
bcp_sendrow 445
BCP_SETL 446
bcp_setxlate 446, 447
bcp_writefmt 448
binary data
  reading page of 248
  writing page of to the server 377
bind result column to program variable 74, 80
browse mode 26, 28
  and DBPROCESS 27
  determining number of tables involved 362
  determining whether regular column source is updatable 96

identifying browsable tables 360
buffers
command. See Command buffers 6
determining size for results 340
placing query results header in 341
row. See Row buffers 6
build_xact_string 465
bulk copy 413, 416
bulk copy processing. See bcp 36
bylist 86
returning 85

C

chained transactions 403
character set
returning client 158
returning server 306
setting 317
setting default 310
translation 90
character set translations
freeing tables 154
loading tables 176
specifying for bcp 446
for strings 383
tables 155
character strings
and quotation marks 298
translating from one character set to another 383
characters
getting from command buffer 158
Client/server
architecture 1, 2
Client-Library
definition of 4
Clients
types of clients 2
close_commit 466
columns
compute. See Compute columns 51
regular. See Regular columns 74
returning id of in order by clause 235
returning number in order by clause 228
command batches 15
canceling current 86
determining whether more results to process 212
sending to the server 344, 351
setting results for next command 276
and switching databases 89
verifying correctness of 346
command buffers 6, 15
adding text to 94, 150
checking for Transact-SQL constructs 163
clearing 155
copying portions of 356
getting characters 158
limiting clearing 404
and message handling 215
returning character length of 358
setting no clear option 95
commands
canceling entire batch 87
determining whether it can return rows 95
determining whether it returned rows 290
determining whether more to process 212
getting stored procedures status number 286
processing 15, 17
returning number of current 116
returning number of rows affected by 115
setting up results for next 276
commit_xact 466
comparing
datetime values 130
compiling and linking. See Open Client/Server Programmer's Supplement xii
compute clauses
returning number in results 228
compute columns
associating with indicator variables 73
binding to program variables 55, 62
getting data 51
order returned 52
returning data length of 54
returning maximum data length of 69
returning number in row 226
returning select-list id 68
returning server datatype for 71
returning type of row aggregate 70
returning user-defined datatypes for 72
summing or averaging 52

compute rows 17
determining 218
getting data for a column 52
reading next 218
returning bylist for 85
returning number of columns 226
creating a notification procedure 220, 223
creating a registered procedure 220, 223
CS-Library
definition of 4
cursor
binding 119
closing 121
fetching against 123
opening 125
retrieving column information for 122
retrieving information about 125
updating 117

D

data
getting user-allocated 168
high-speed loading of. See bcp 36
reading binary 248
reading server (UNIX and AOS/VS) 171
saving user-allocated 333
writing binary 377
writing to the server (UNIX and AOS/VS) 173
databases
determining whether changed 89
multi-user updates 26
reading pages 248
returning name of current 217
updating 26, 28
updating on multiple servers 449
using specified 373
datatypes
binding compute columns to 55, 62
binding regular columns to 74, 80
compute columns 71
conversions supported 106, 112
conversions supported by dbaltbind 57
conversions supported by dbaltbind_ps 63
conversions supported by dbbind 75
conversions supported by dbbind_ps 80
converting 104, 109
converting to same 107, 114
DB-Library 408, 411
DB-Library/C 12
determining supported conversions 376
getting precision and scale for regular column 102
returning for compute columns 71, 72
returning for regular column 100
server 106, 111
server, list of 408
SQL Server 12
user-defined, for compute columns 72
user-defined, for regular columns 103
date formats
input 403
DATEFIRST option 402
DATEFORMAT option of set command 403
dates
converting parts to character strings 136
converting to character format 132
converting values into usable format 134
determining month name in specified language 210
parts of 132
returning name of day in specified language 143
returning order for specified language 139
returning parts as numeric values 140
symbols recognized by DB-Library 132
datetime routines 34, 35
datetime values
comparing 130
days
returning name of in specified language 143
db12hour 49, 51
dbadata 51, 53
as alternate to dbaltbind 60, 67
dbadlen 54, 55
dbaltbind 55, 61
as alternate to dbadata 52
dbaltbind_ps 62, 67
dbaltcolid 68
dbaltlen 69
dbaltop 70
dbalttype 71, 72
dbaltutype 72, 73

dbanullbind 73, 74
DBARITHABORT option 402
DBARITHIGNORE option 403
DBAUTH option 403
dbbind 74, 79
  as alternate to dbdata 130
dbbind_ps 80, 84
DBBUFFER option 403
  and DBFIRSTROW 154
  and dbgetrow 167
  and DBLASTROW 176
  and reading result rows 219
dbbufsize 84, 85
dbbylist 85, 86
dbcancel 86, 87
dbcanquery 88
DBCHAINXACTS option 403
dbchange 89
dbcharsetconv 90
dbclose 90, 91
dbclrbuf 91, 92
dbclropt 92, 93
dbcmd 94, 95
DBCMDROW 95
dbcolbrowse 96
dbcollen 97, 98
dbcolname 98, 99
  and returning bylist 85
dbcolsource 99, 100
dbcoltype 100, 102
dbcoltypeinfo 102, 103
dbcolutype 103, 104
dbconvert 104, 108
dbconvert_ps 109, 114
DBCOUNT 115
DBCURCMD 116
DBCURROW 116, 117
dbcursor 117, 119
dbcursorbind 119, 121
dbcursorclose 121
dbcursorcolinfo 121, 122
dbcursorfetch 122, 125
dbcursorinfo 125
dbcursoropen 125, 128
dbdata 129, 130
  as alternate to dbbind 79, 84
dbdate4cmp 130, 131
dbdate4zero 131
dbdatechar 132, 133
dbdatecmp 133, 134
dbdatecrack 134, 136
dbdatename 136, 139
dbdateorder 139, 140
dbdatepart 140, 141
DBDATEREC structure 134
DBDATETIME structure 135
  converting date parts to character strings 136
  converting integer component to character format 132
  converting values into usable format 134
  returning parts as numeric values 140
dbdatezero 141, 142
dbdatlen 142, 143
dbdayname 143, 144
DBDEAD 144, 145
dberrhandle 145, 149
dbexit 149, 150
dbfcmd 150, 153
DBFIPSFLAG option 404
DBFIRSTROW 153, 154
  and dbgetrow 167
dbfree_xlate 154, 155
dbfreebuf 155, 156
dbfreequal 156
dbfreesort 156, 157
dbgetchar 158
dbgetcharset 158, 159
dbgetloginfo 159, 161
dbgetlusername 161, 162
dbgetmaxprocs 162, 163
dbgetnatlang 163
dbgetoff 163, 165
dbgetpacket 165, 166
dbgetrow 166, 168
DBGETTIME 168
dbgetuserdata 168, 169
dbhasretstat 169, 171
dbinit 171
  and dbexit 150
DBIORDESC (UNIX and AOS/VS) 171, 172
DBIOWDESC (UNIX and AOS/VS) 173
DBISAVAIL 174

DBISOLATION option 404
dbisopt 174, 175
DBLASTROW 175, 176
DB-Library 4
  determining version in use 375
  initializing 171
dbload_xlate 176, 177
  and freeing translation tables 155
dbloadsort 177, 178
dblogin 178, 180
dbloginfree 180
dbmny4add 180, 181
dbmny4cmp 181, 182
dbmny4copy 182, 183
dbmny4divide 183, 184
dbmny4minus 184, 185
dbmny4mul 185, 186
dbmny4sub 186, 187
dbmny4zero 187, 188
dbmnyadd 188, 189
dbmnycmp 189, 190
dbmnycopy 190, 191
dbmnydec 191, 192
dbmnydivide 192, 193
dbmnydown 193, 194
dbmnyinc 195
dbmnyinit 195, 197
dbmnymaxneg 197, 198
dbmnymaxpos 198
dbmnyminus 199
dbmnymul 200
dbmnyndigit 201, 207
dbmnyscale 207, 208
dbmnysub 209
dbmnyzero 210
dbmonthname 210, 211
DBMORECMDS 212
dbmoretext 212, 213
dbmsghandle 213, 217
  and dberrhandle 148
dbname 217, 218
DBNATLANG option 404
dbnextrow 218, 220
  and DBROWS 290
  and DBROWTYPE 291
DBNOAUTOFREE option 404
  and dbfcmd 152
  and dbfreebuf 156
DBNOCOUNT option 405
  and DBCOUNT 115
DBNOEXEC option 405
dbnpcreate 220, 222
dbnpdefine 223, 224
dbnullbind 225
dbnumalts 226
dbnumcols 226, 227
dbnumcompute 228
DBNUMORDERS 228, 229
dbnumrets 229, 230
DBOFFSET option 405
  and dbgetoff 164
dbopen 231, 234
  getting a LOGINREC 178
  setting login response time 323
dbordercol 235
DBPARSEONLY option 405
dbpoll 235, 240
DBPRCOLSEP option
  and dbspr1row 339
dbprhead 241
DBPRLINELEN option 405
DBPRLINESEP option 405
DBPROCESS structure 6, 7
  allocating 231
  closing a 90
  closing all 149
  de-allocating a 90
  de-allocating all 149
  determining current limit available 162
  determining whether available 174
  determining whether dead 144
  getting client character set from 158
  getting national language from 163
  getting server process id 337
  getting user-allocated data 168
  initializing 231
  marking available 307
  multiple 6
  saving user-allocated data 333
  setting maximum number available 327
  sharing single 174
  and two-phase commit service 451

DBPRPAD option 406
and dbspr1row 339
dbprrow 241, 242
dbprtype 242, 243
dbqual 244, 247
freeing allocated memory 156
DBRBUF (UNIX and AOS/VS) 247, 248
dbreadpage 248, 249
dbreadtext 249, 251
dbrecftos 251, 252
dbrecvpassthru 252, 254
dbregdrop 254, 255
dbregexec 256, 258
dbreghandle 258, 262
dbreginit 262, 263
dbreglist 264
dbregnowatch 265, 266
dbregparam 266, 270
dbregwatch 270, 275
dbregwatchlist 275, 276
dbresults 276, 279
dbretdata 279, 283
dbretlen 283, 284
dbretname 285, 286
dbretstatus 286, 288
dbrettype 288, 290
DBROWCOUNT option 406
DBROWS 290
DBROWTYPE 291
dbrpcinit 291, 293
dbrpcparam 293, 295
and return parameter values 229
dbrpcsend 295, 296
dbrpwclr 296, 297
dbrpwset 297, 298
dbsafestr 298, 299
dbsechandle 300, 303
dbsendpassthru 304, 306
dbservcharset 306
dbsetavail 307
dbsetbusy 307, 310
dbsetdefcharset 310
dbsetdeflang 310, 311
dbsetidle 311, 312
dbsetifile 312, 313
dbsetinterrupt 313, 317
and canceling result rows 88
DBSETLAPP 317
DBSETLCHARSET 317, 318
DBSETLENCRYPT 318, 319
DBSETLHOST 319, 320
DBSETLNATLANG 320, 321
dbsetloginfo 321, 323
dbsetlogintime 323, 324
DBSETLPACKET 324, 325
DBSETLPWD 325, 326
DBSETLUSER 326, 327
dbsetmaxprocs 327
dbsetnull 327, 329
dbsetopt 329, 331
dbsetrow 331, 332
dbsettime 333
dbsetuserdata 333, 336
dbsetversion 336, 337
DBSHOWPLAN option 406
dbspid 337, 338
dbspr1row 338, 340
and dbspr1rowlen 340
dbspr1rowlen 340, 341
dbsprhead 341, 342
and dbspr1rowlen 340
and dbsprline 343
dbsprline 343, 344
and dbspr1rowlen 340
dbsqlexec 344, 346
dbsqlok 346, 351
dbsqlsend 351, 352
DBSTAT option 406
DBSTORPROCID option 406
dbstrbuild 352, 354
dbstrcmp 354, 356
and dbstrsort 360
dbstrcpy 356, 358
dbstrlen 358, 359
dbstrsort 359, 360
dbtabbrowse 360, 361
dbtabcount 361, 362
dbtabname 362, 363
and dbtabcount 362
dbtabsource 363, 364
DBTDS 365
DBTEXTLIMIT option 407

dbtextsize 365, 366
DBTEXTSIZE option 407
dbtsnewlen 366, 367
dbtsnewval 367, 368
dbtsput 368, 369
dbtxptr 369, 370
dbtxtimestamp 371
dbtxtsnewval 372
dbtxtsput 372, 373
dbuse 373, 374
dbvarylen 374, 375
dbversion 375
dbwillconvert 376, 377
dbwritepage 377, 378
dbwritetext 378, 383
dbxlate 383, 385
deadlock
  handling 169, 334
debugging
  and dbprhead 241
  and dbprrow 241
  and dbspr1row 338
  and dbspr1rowlen 340
  and dbsprhead 342
  and dbsprline 343
  recording SQL text sent to the server 251
  for two-phase commit service 465
decimal datatype
  getting precision and scale for regular column 102
default character set
  setting for an application 310
default language
  setting for an application 310
defining a notification procedure 223
defining a registered procedure 223
distributed transactions. See Two-phase commit service 449
dropping a registered procedure 254

E

embedded SQL
  comparing Client-Library to 5
encrypted passwords 318
encryption handler
  installing 300
error handling 22
  and converting datatypes 105, 111
  and DBDEAD 144
  installing a user-function 145, 149
  list of errors 386
  translating messages from one language to another 352
  uninstalling handler 148
error severity values 10
errors 386, 402
  DB-Library 386
example programs
  online example programs xii
executing a registered procedure 256, 262
exit values 9

F

file descriptors (UNIX and AOS/VS)
  access to 171, 173
file names, for libraries. See Open Client/Server Programmer's Supplement xii
files
  header 9
functions
  user-supplied to handle interrupts 313
  user-supplied, indicating DB-Library is finished reading from the server 311
  user-supplied, indicating server access 307

G

gateway applications 33, 252, 304
getting the client character set 158
getting the national language 163
getting the server character set 306

H

handler
  error 22

message 22
notification 258
header files 9
help
Technical Support xiv
host names
setting in LOGINREC 319

I

image values
bulk copying parts 440
bytes left of 365
limiting size of 407
reading parts of 249
and text pointers 370
and text timestamps 371
updating 212, 378
include files 9
input streams
and checking for unread bytes in network buffer (UNIX and AOS/VS) 247
responding to multiple (UNIX and AOS/VS) 172
utilizing multiple (UNIX and AOS/VS) 173, 352
interfaces file
and dbopen 231
specifying name and location 312
interrupt handling 313

L

languages
getting name from DBPROCESS 163
setting default 310
setting name in LOGINREC 320
setting national 404
line length
specifying for rows 405
listing registered procedures 264
listing requested registered procedure notifications 275
logging into the server 231
login record. See LOGINREC structure 7
LOGINREC structure 7
adding remote passwords 297
allocating 178
clearing all remote passwords 296
freeing 180
packet size field 165, 324
setting application name in 317
setting client character set in 317
setting for bcp 446
setting host name in 319
setting password in 325
setting user language name in 320
setting username in 326
logins, secure 300

M

message handling 22
and dberrhandle 148
and dbreadpage 249
and deadlock 334
installing a user function 213, 217
uninstalling handler 216
money routines 34, 35
months
determining name in specified language 210
multiple input streams 235

N

network buffer
polling 235
network buffers
determining whether unread bytes (UNIX and AOS/VS) 247
network connections
closing 90
specifying interfaces file 312
notification handler 258
notification procedure
creating 220, 223
defining 223
notification request
canceling 265
listing 275
notifications

listing registered procedures 275
registered procedure 258
null values
binding 327
default 328
defining 327
numeric datatype
getting precision and scale for regular column 102

O

offsets
setting where to return 405
types of 164
Open Client/Server Programmer's Supplement xii
open_commit 467
options 402, 408
checking status of 174
clearing 92
DB-Library 402
parameter values of 407
setting 329
order by clauses
returning column id in 235
returning number of columns in 228
output streams
utilizing multiple (UNIX and AOS/VS) 173, 352

P

packet size
TDS 165, 324
padding
specifying characters to use 406
parameters
registered procedure 266
passthrough operation 252, 304
passwords
remote, adding 297
remote, clearing 296
setting server 325
polling the network buffer 235
process id
getting 337
processing plan
generating description of 406
programming
DB-Library/C 7, 12

Q

queries
aborting during arithmetic exceptions 402
ignoring arithmetic exceptions 403
not executing 405
quotation marks
and character strings 298

R

registered procedure 31, 33
canceling notification request 265
creating 220, 223
defining 223
dropping 254
example 32
executing 256, 262
handler routine 258
listing currently defined 264
listing requested notifications 275
notifications 235, 256, 258
parameters 266
requesting notifications 270
routines 33
uses of 31
regular columns
associating indicator variables with 225
binding to program variables 74, 80
determining number of in results 226
determining whether data length can vary 374
determining whether source column is updatable with browse mode 96
getting data 129
getting precision and scale with dbcoltypeinfo 102
returning data length of 142
returning datatypes for 100
returning maximum data length of 97

returning name of 98
returning name of source column 99
returning user-defined datatypes for 103
regular rows 17
determining 218
limiting number to return 406
reading next 218
remote procedure calls 30, 292
adding parameters to 293
adding passwords for 297
advantages of 292
clearing passwords 296
determining number of return parameter values 229
determining whether status number was generated 169
getting datatype of return parameter value 288
getting length of return parameter value 283
getting name of return parameter value 285
getting return parameter values 279
getting status number 286
initializing 291
processing 30, 279
signaling end of 295
remove_xact 467, 468
requesting a registered procedure notification 270
result columns
compute. See Compute columns 51
regular. See Regular columns 74
returning name and number of source table 363
result rows 17
buffering 403
canceling 88
compute 17
dropping from buffer 91
placing header in buffer 341
printing 241
printing column headings of 241
processing 17, 22
putting one in buffer 338
reading next 218
regular 17
results
setting up for next query 276
return parameter values 229
determining number of 229
getting 279
getting datatype of 288
getting length of 283
getting parameter name 285
returning TDS packet size 165
returning the client character set 158
returning the national language 163
returning the server character set 306
routines 12, 36
bcp 36
browse mode 27
command processing 15
error handling 22
image handling 28
information retrieval 24
initialization 13
message handling 22
process control 30
registered procedure 33
remote procedure call 30
results processing 17
TDS 33
text handling 28
two-phase commit service 36
row aggregates
returning for a compute column 70
row buffers 167, 332
clearing 91
reading specified rows 166
returning number of first row 153
returning number of last row 175
rows
buffering 403
compute 17
determining type 218
determining whether returned 290
determining whether returned by command 95
dropping from buffer 91
limiting number to return 406
printing 241
printing column headings of 241
reading next 218
reading specified in buffer 166
regular 17
result. See Result rows 17
returning number affected by a command 115
returning number of current 116
returning number of first in buffer 153

returning number of last in buffer 175
returning type of 291
specifying line length 405
specifying separator characters 405
turning counting off 405
updating current in browsable table 244

S

scan_xact 468, 469
secure logins
installing user function for 300
secure SQL Server
and bcp 443
routines for 35
security label handler
installing 302
separator characters
specifying for rows 405
server 87, 168
communicating with 6, 7
converting token values 242
datatypes 106, 111
logging into 231
reading data from (UNIX and AOS/VS) 171
recording SQL text sent to 251
sending text/image values to 212
setting response time 323
setting user passwords 325
types of servers 2
writing data to (UNIX and AOS/VS) 173
servers
multiple 6
setting TDS packet size 324
setting the client character set 317
sort orders 157
comparing two character strings 354
determining order of two character strings 359
freeing 156
loading 177
sprintf function 150
SQL Server
updating among multiple 449
SQL Server Reference Manual xii
SQL text
recording 251
start_xact 469
stat_xact 470
statistics
performance, determining when returned 406
status numbers
for current command 286
determining whether generated 169
stored procedures
calling remotely 30
determining number of return parameter values 229
return parameter values, getting 279
return parameter values, getting datatype of 288
return parameter values, getting length of 283
return parameter values, getting parameter name of 285
returning status number 286
sending ids of 406
and status numbers 169
sybdb.h header file 9, 146, 164
and DB-Library options 402
and error handling 386
syberror.h header file 10, 148
and error severities 386
SYBESMSG
and error handling 148
sybfront.h header file 9
and interrupt handling 314
syntax
checking 405

T

tables
determining names of 362
identifying browsable 360
returning name and number associated with result columns 363
returning name of 362
returning number involved in a select query 361
server work 362
tabular Data Stream
routines 33, 165, 252, 304, 324
tabular Data Stream protocol

determining version being used 365
versions available 365
TDS
determining packet size 165
passthrough operation 252, 304
routines 33
setting packet size 324
TDS buffer
polling 235
TDS packet
receiving 252
sending 304
technical Support xiii
text and image data 165, 324
text pointers 370
returning value of 369
text timestamps 371
putting new value into DBPROCESS 372
returning value of 371
returning value of after update 372
text values
bulk copying parts 440
bytes left of 365
limiting size of 407
reading parts of 249
and text pointers 370
and text timestamps 371
updating 212, 378
text/image data
updating 213
time
amount DB-Library waits for a server response 168
determining when to return status 406
determining whether 12 or 24-hour 49
setting length DB-Library waits for server response 333
setting server login response 323
timestamp columns 26
putting new value in DBPROCESS 368
returning length of after update 366
returning value of after update 367
and updating rows 244
token values
converting to readable strings 242
transactions
distributed. See Two-phase commit service 449
Transact-SQL commands
and DBPROCESS 6
translation tables
freeing 154
loading 176
two-phase commit service 449, 464
building names for recovery purposes 465
closing connections 466
and DBPROCESS 451
debugging 465
decrementing site count 467
diagnostic routines 468, 470
and interfaces file 454
marking transactions as aborted 464
marking transactions as committed 466
opening connections 467
printing record of distributed transactions 468
returning status of a distributed transaction 470
routines for 36
starting a distributed transaction 469
typedefs
DB-Library/C 12
DB-Library/C, list of 409

U

unchained transactions 403
updating databases 26, 28
on multiple servers 449
multi-user situations 26
and text/image data 213
user-defined datatypes
returning for a compute column 72
returning for regular columns 103
usernames
setting 326
user-supplied data
retrieving for a DBPROCESS 168
saving in a DBPROCESS 333
user-supplied functions
calling to handle interrupts 313
indicating DB-Library is finished reading from the server 311
indicating server access 307

V

versions
DB-Library, determining which 375

W

where clauses
for use in updating a browsable table 244